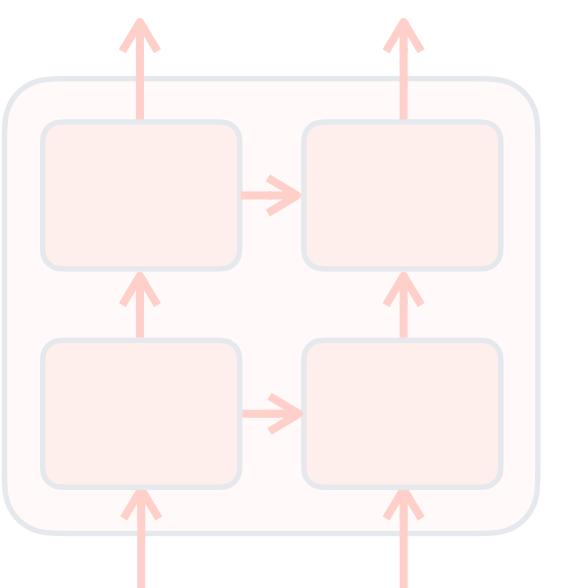


Part II

Advanced Transformers

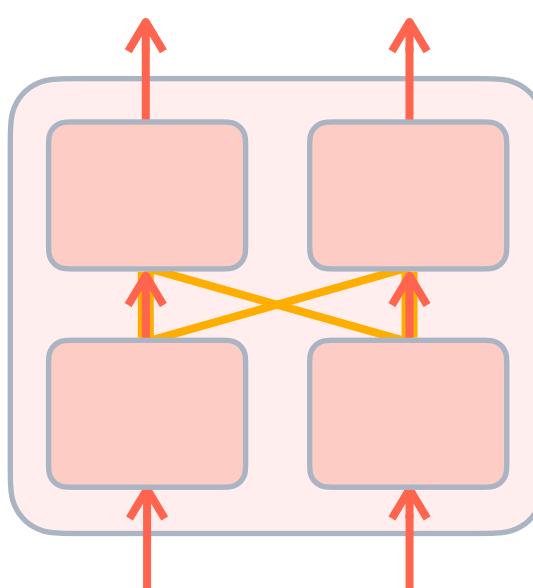
Part I

Recurrent Networks,
Seq2Seq and Attention



Part II

Transformers



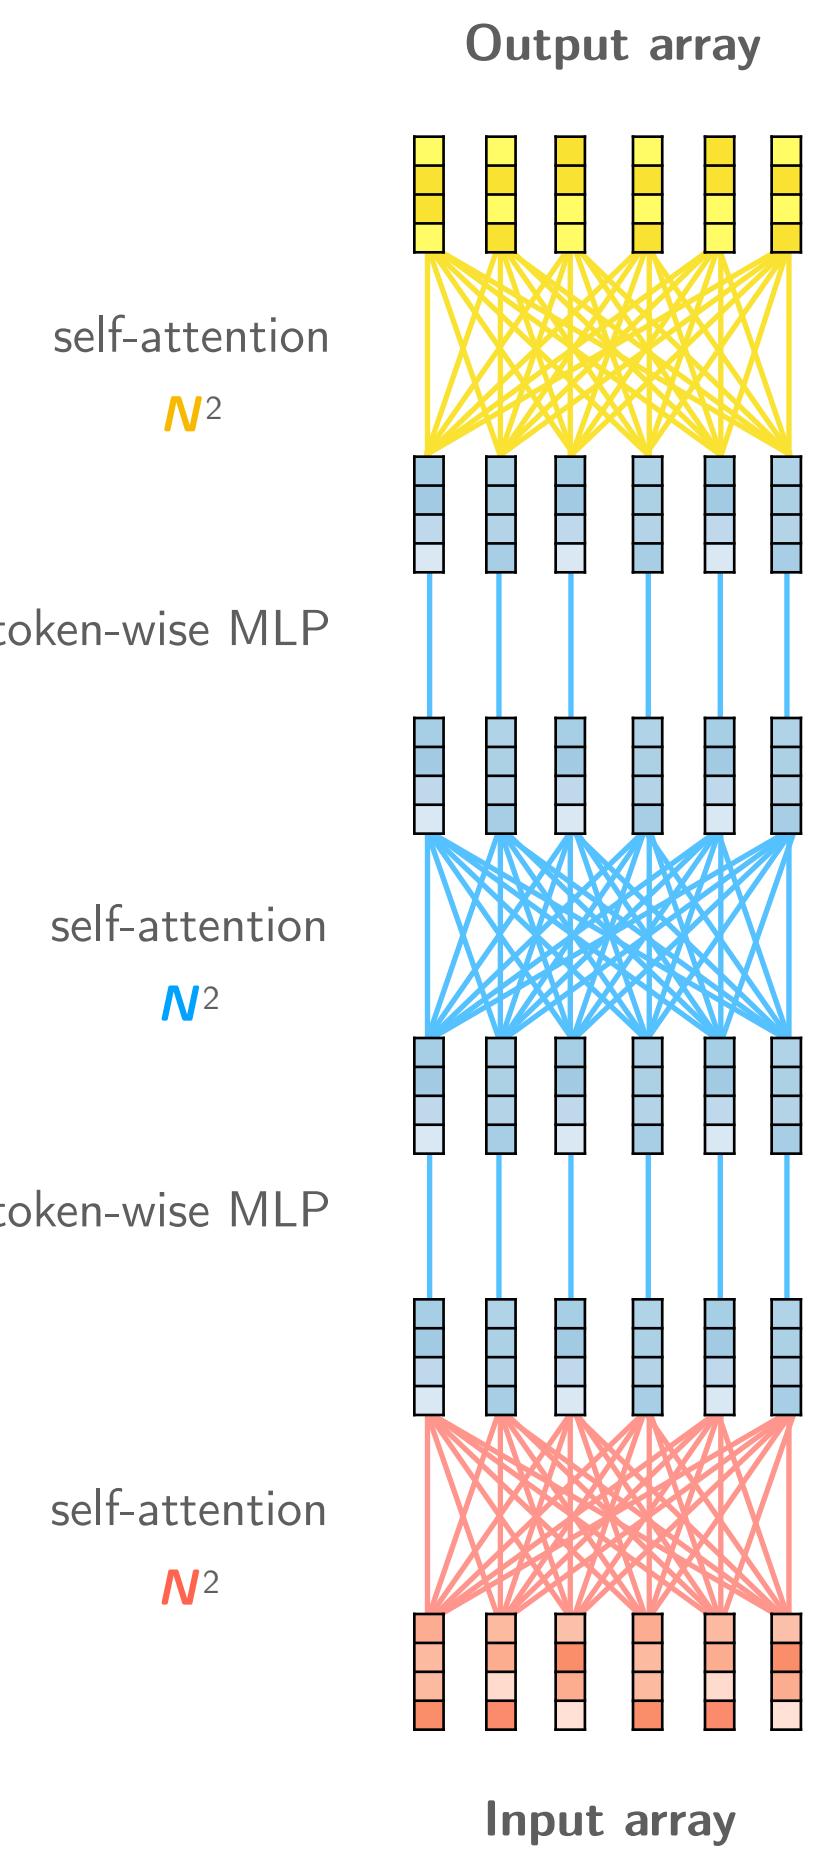
Part III

State Space Models

???

Transformer

as a kind of “token net” *



*using the word “token” a bit loosely

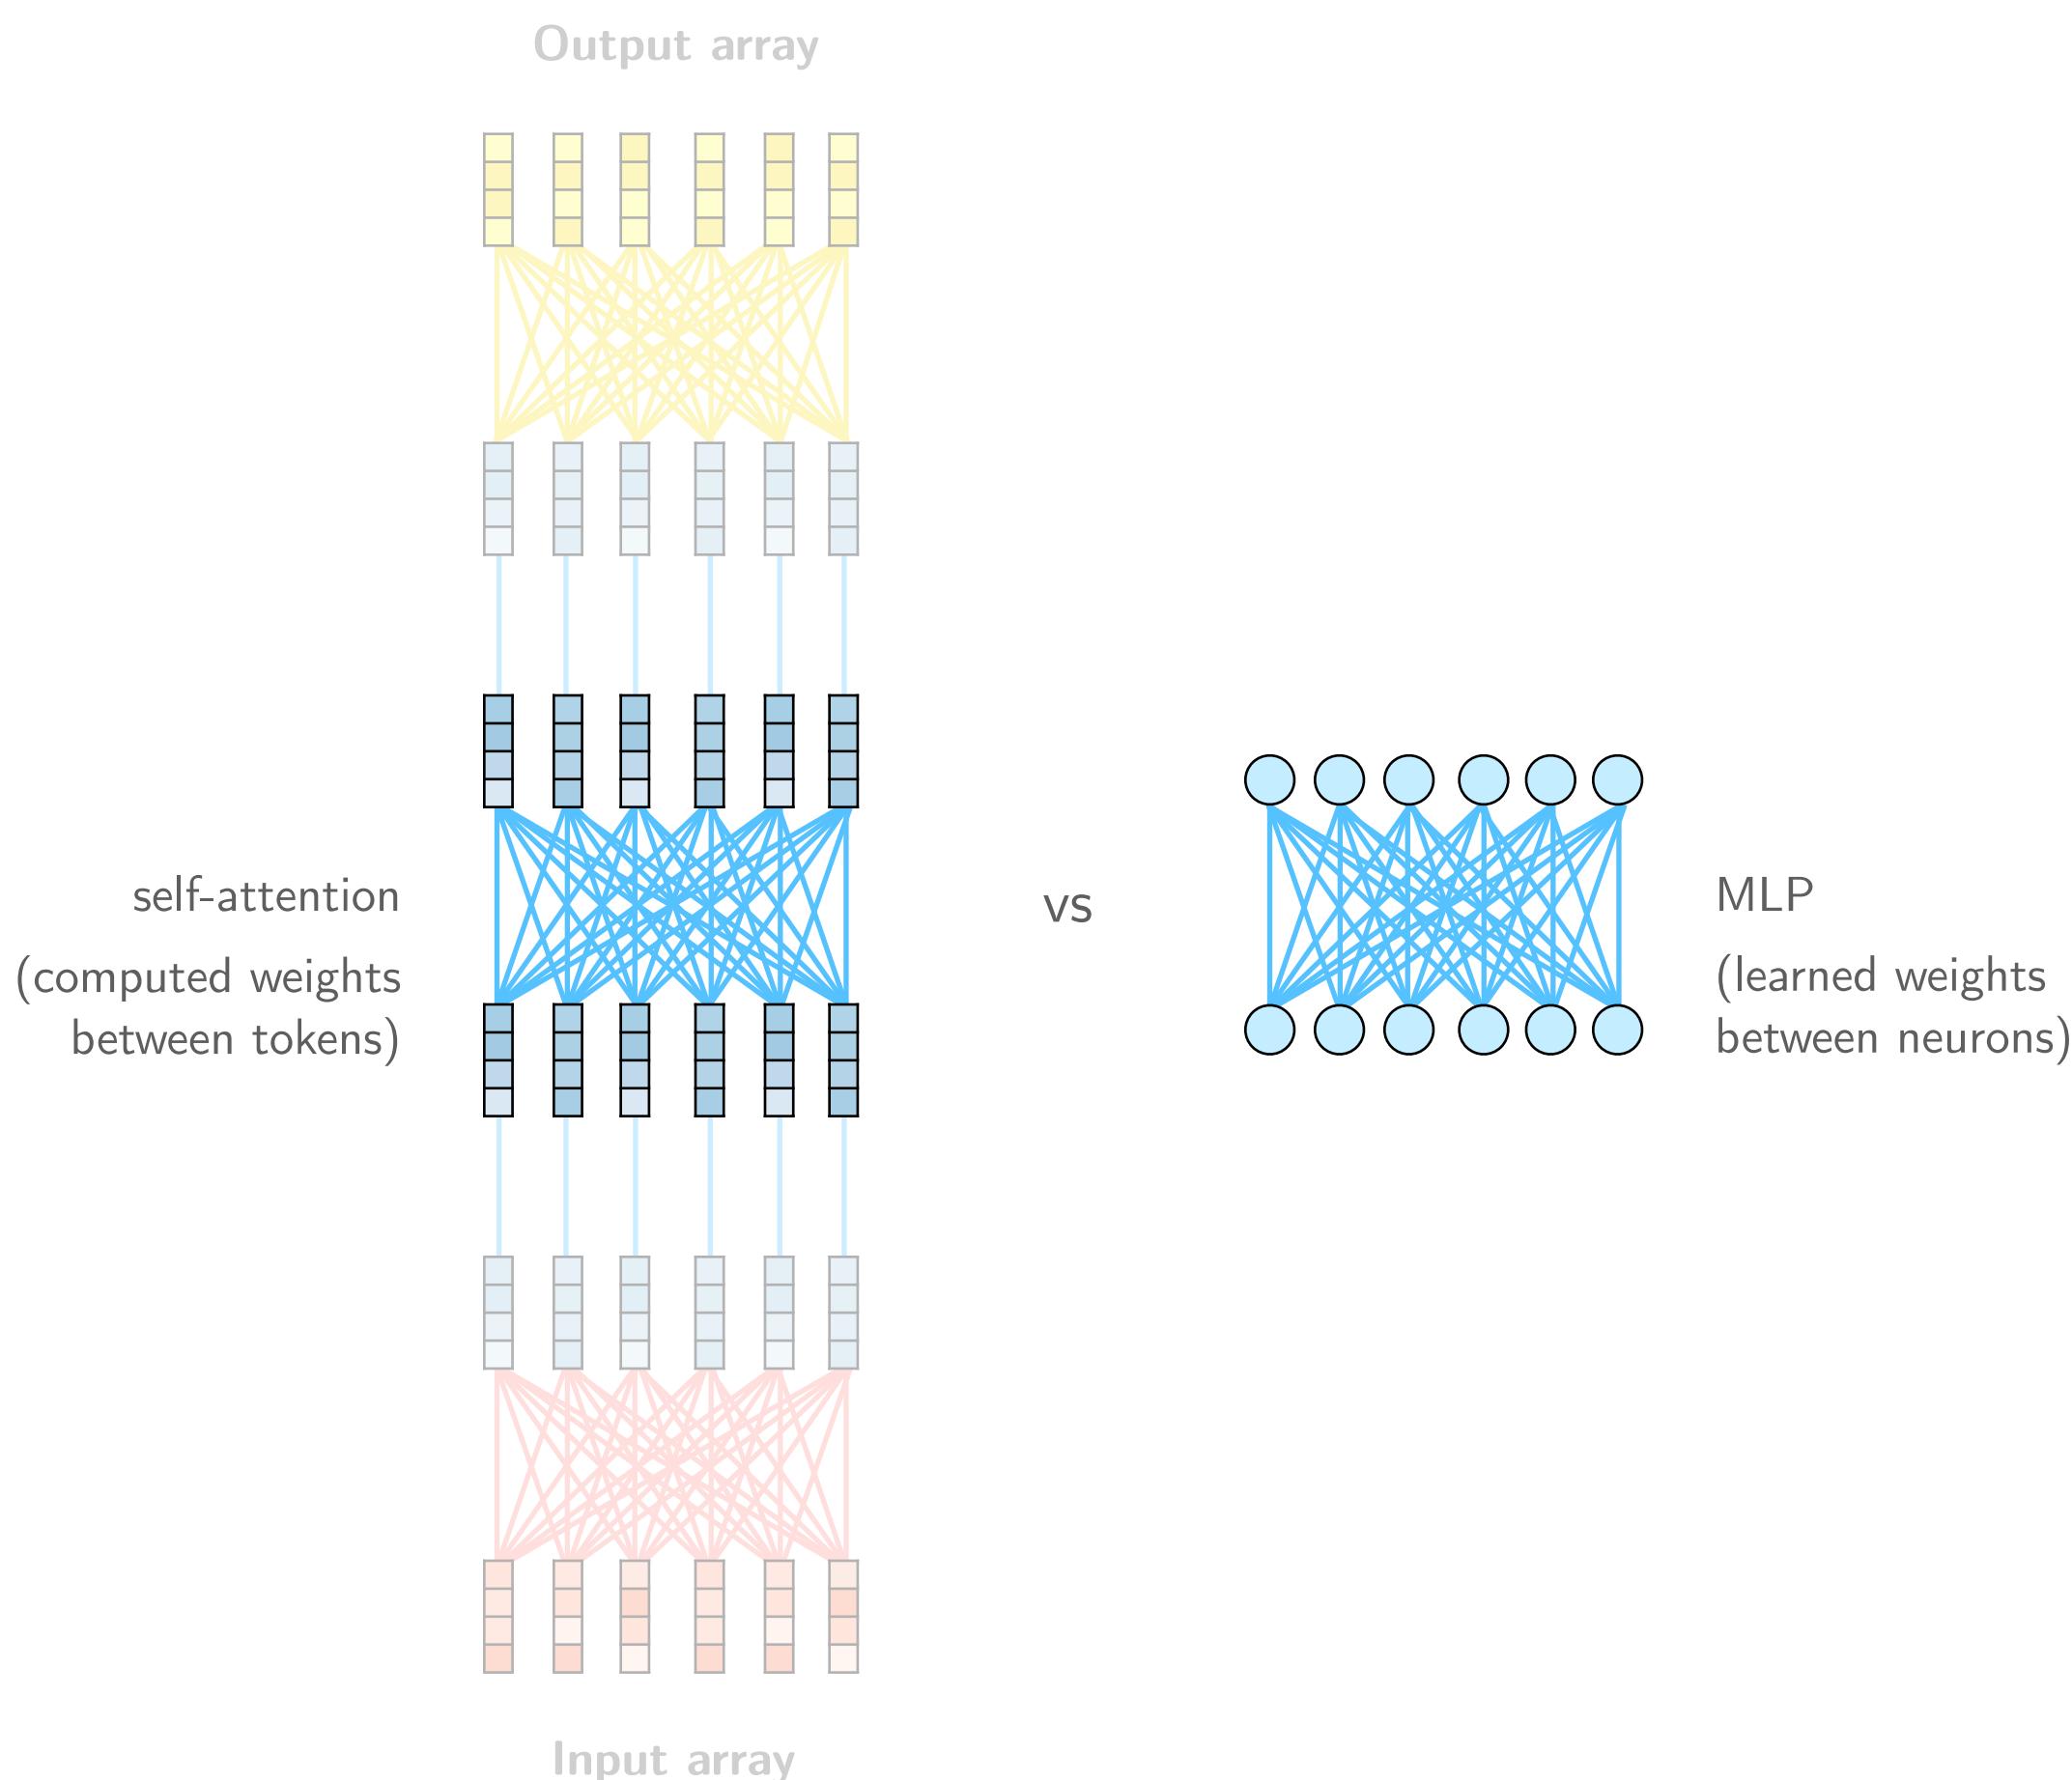
Transformer

as a kind of “token net”

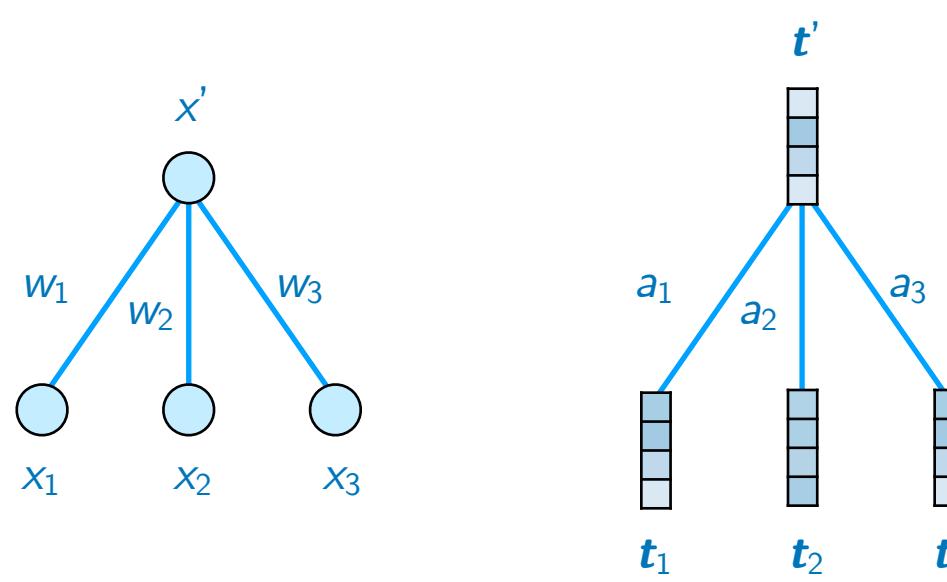
Self-attention is how we “wire up” the tokens.

The weights between neurons are learned,
but the weights between tokens are computed
in the forward pass from the tokens’ vectors.

Arguably, self-attention is the most “natural”
way to wire up tokens.



In both cases, we do a linear combination:



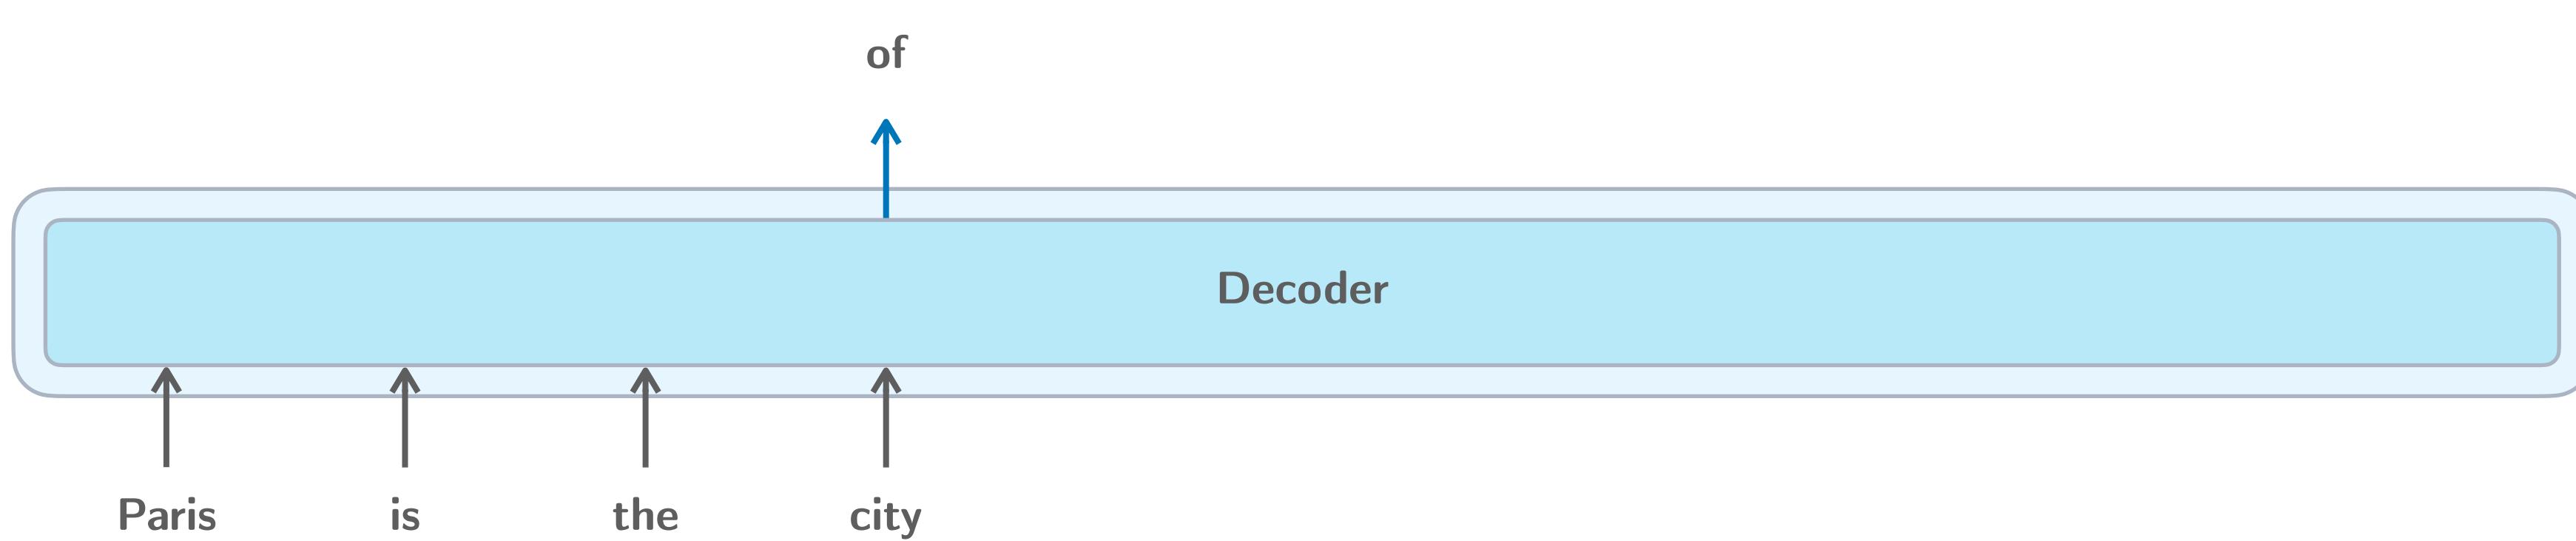
$$\sum_i w_i x_i$$

$$\sum_i a_i t_i$$

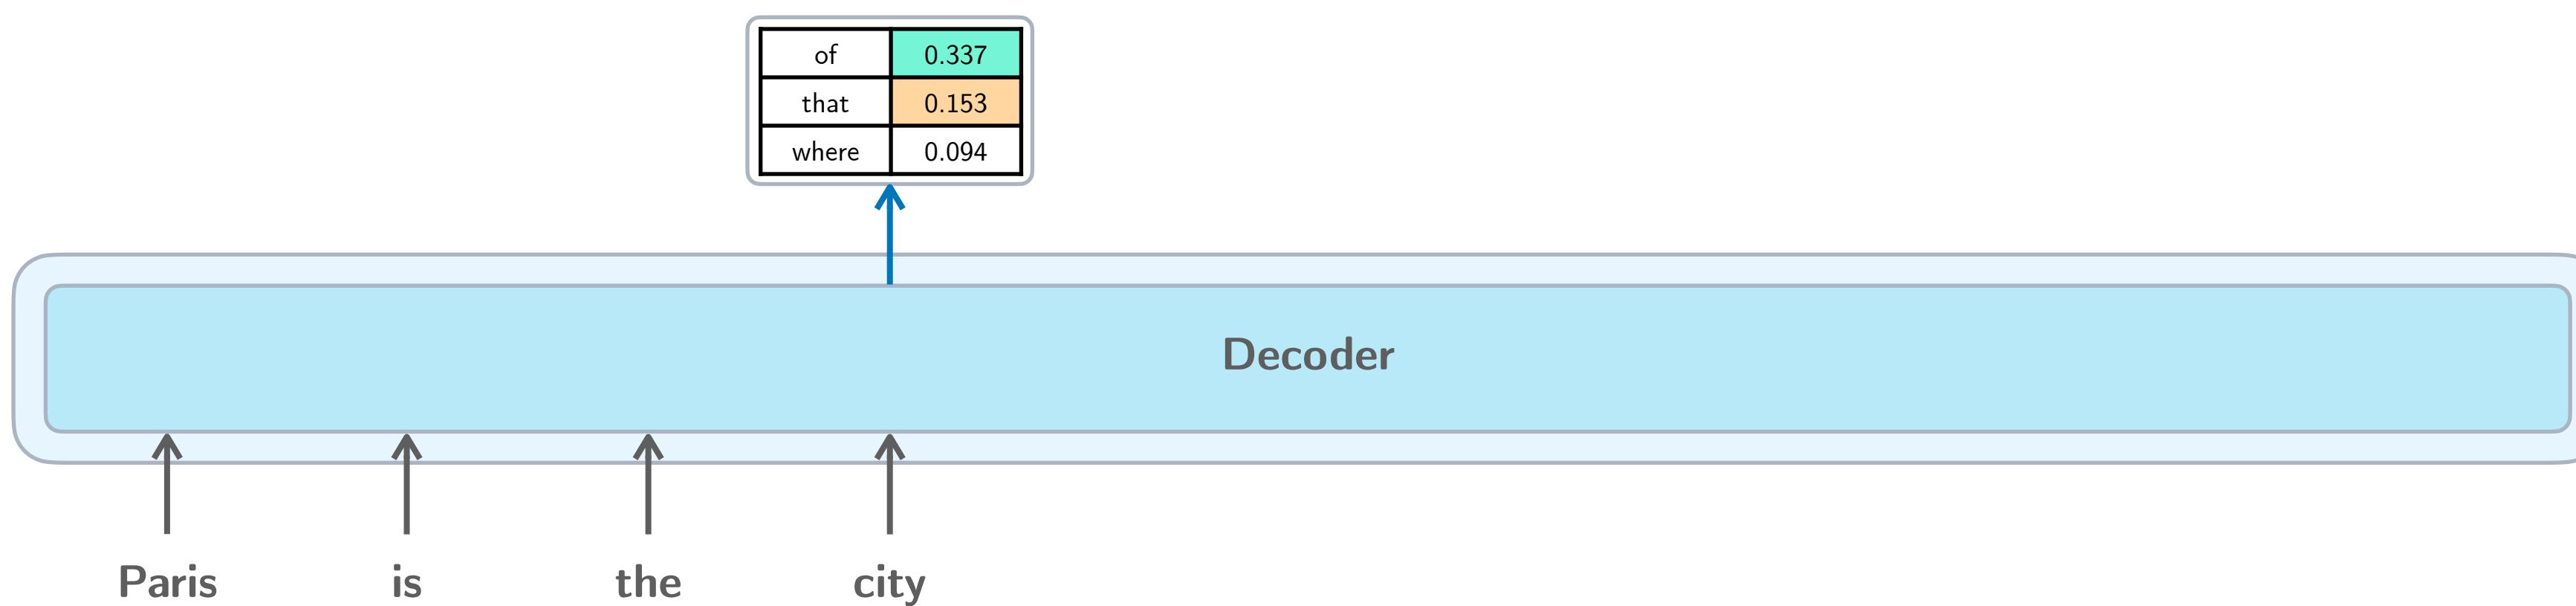


Decoding

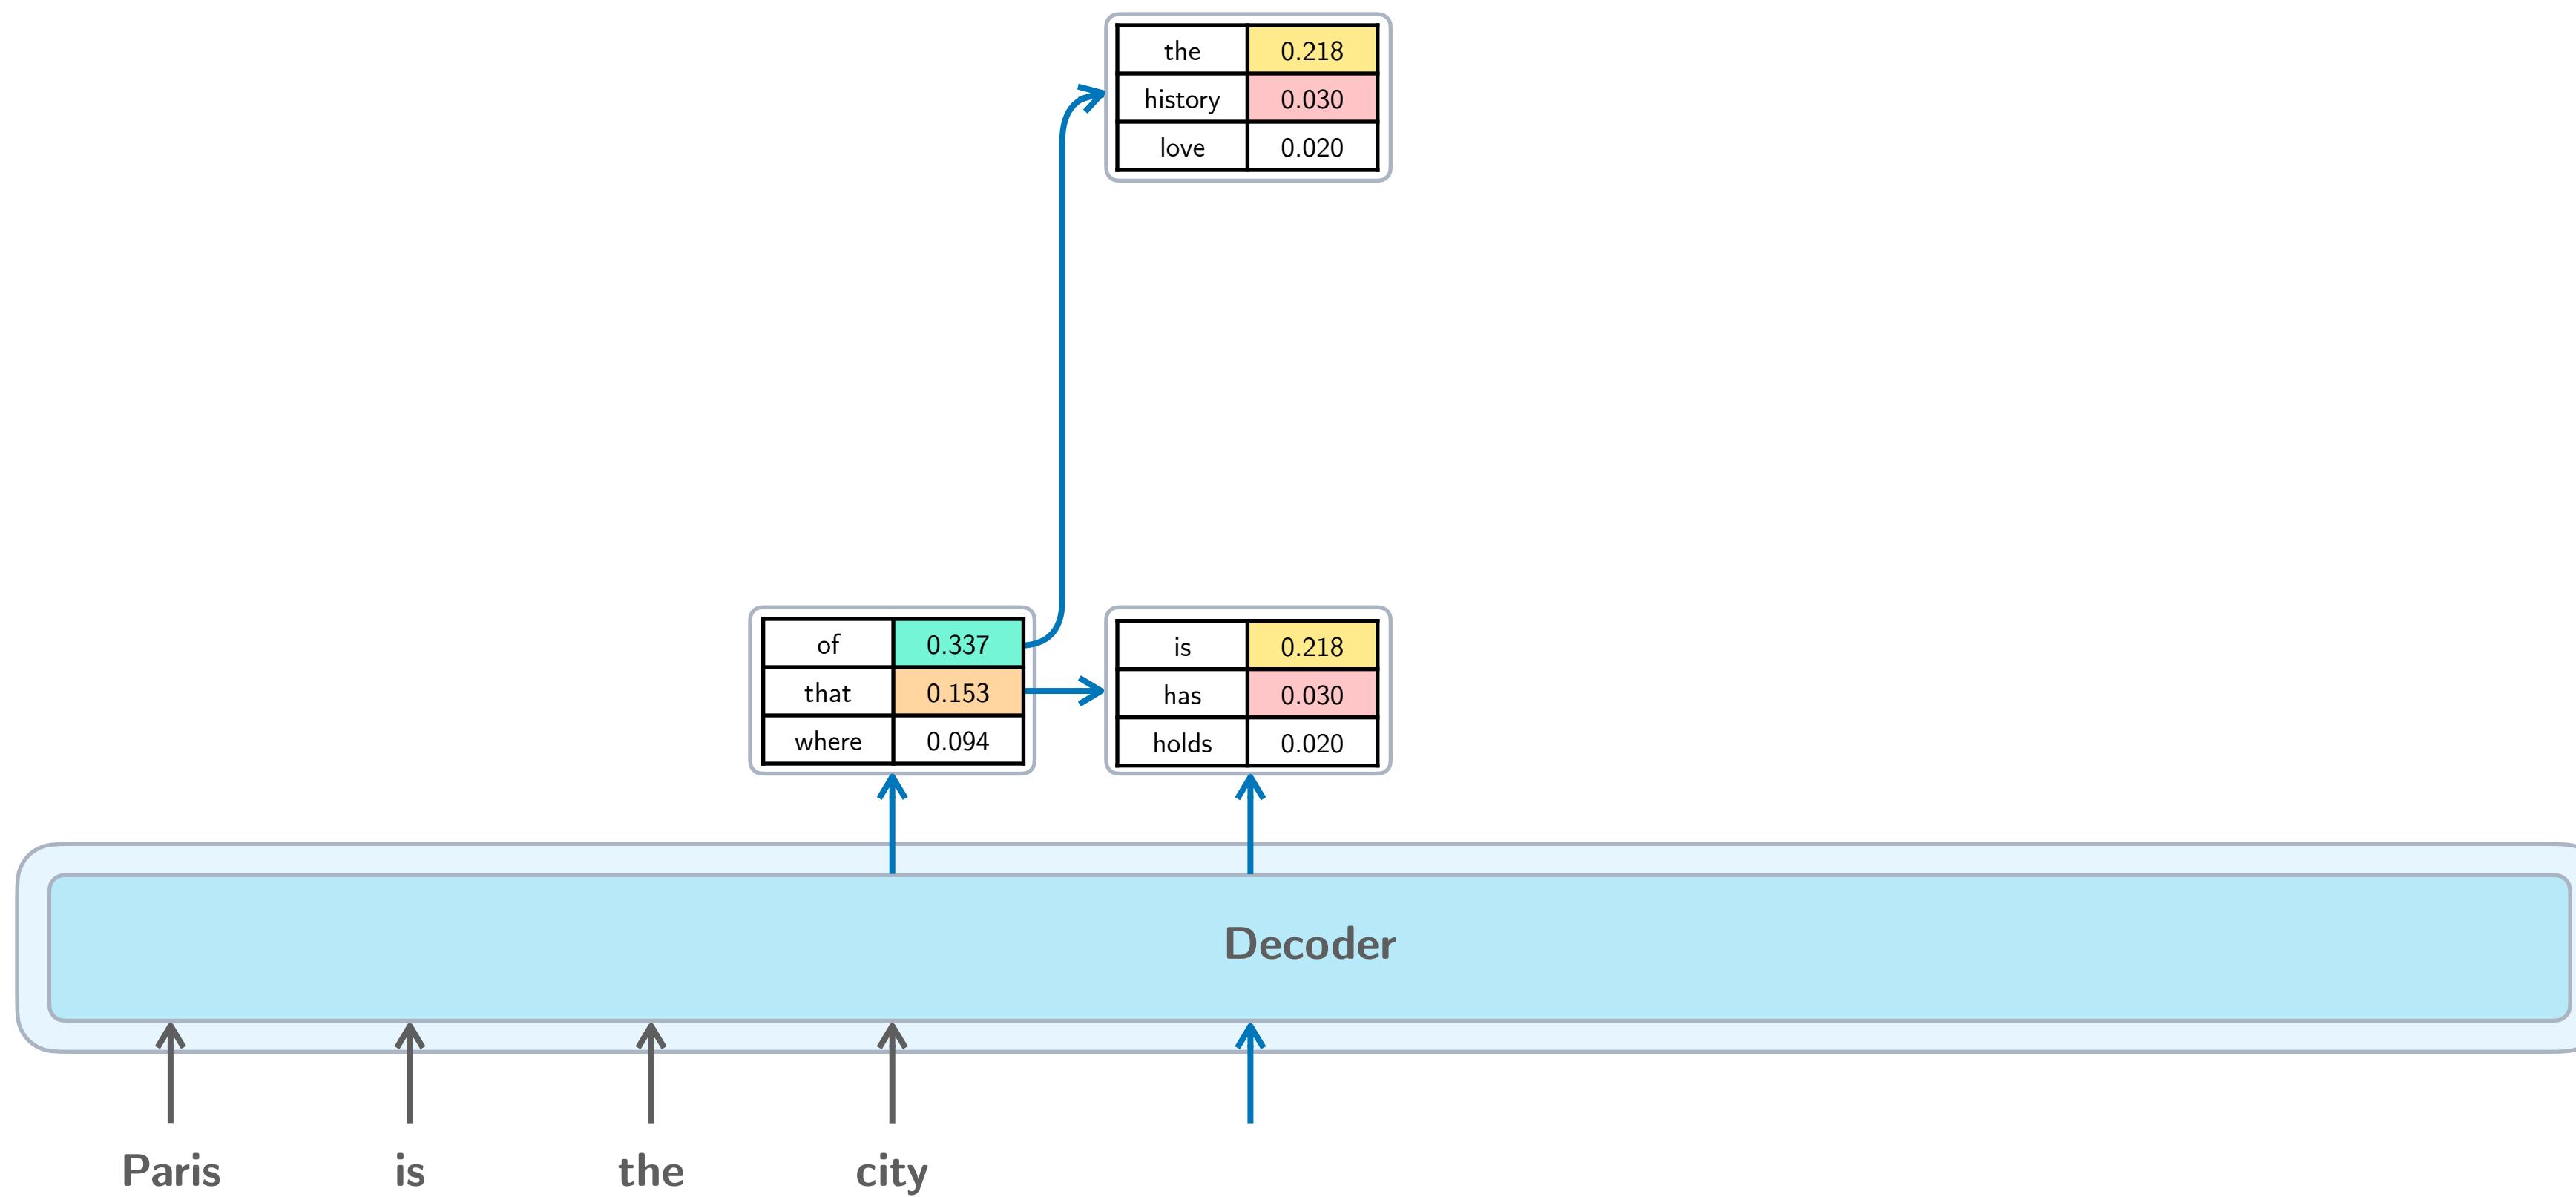
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>



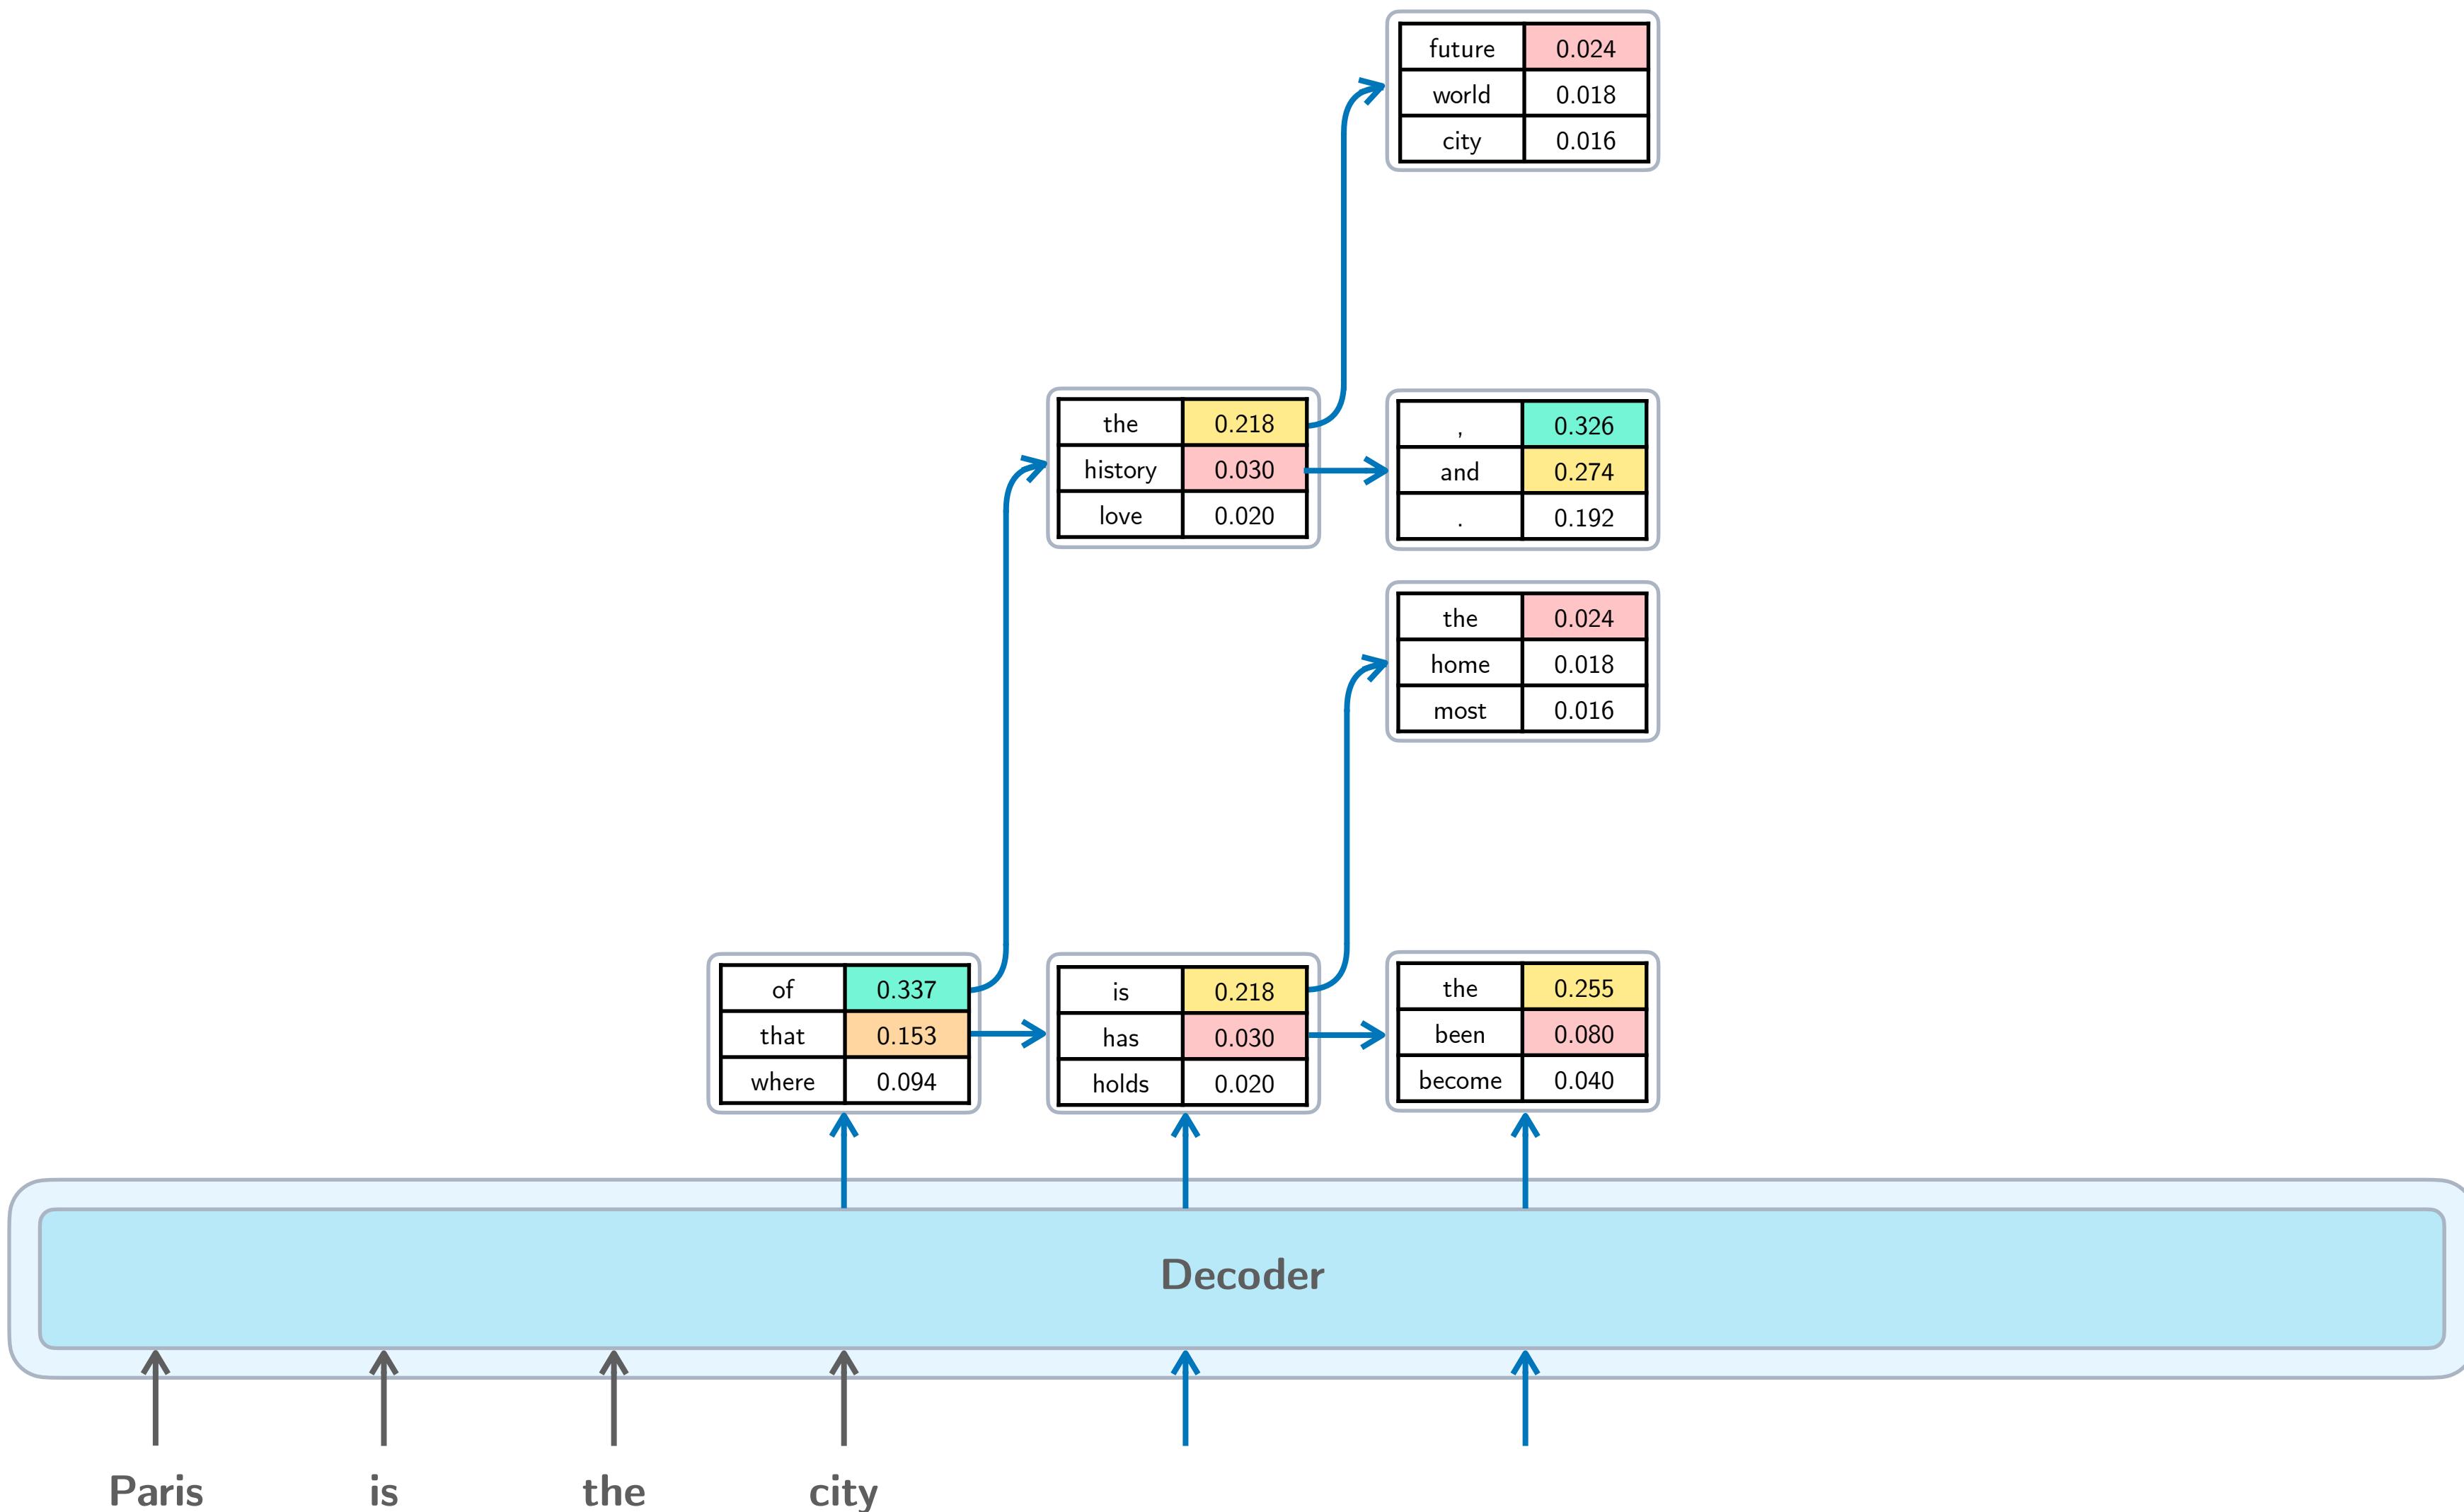
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>



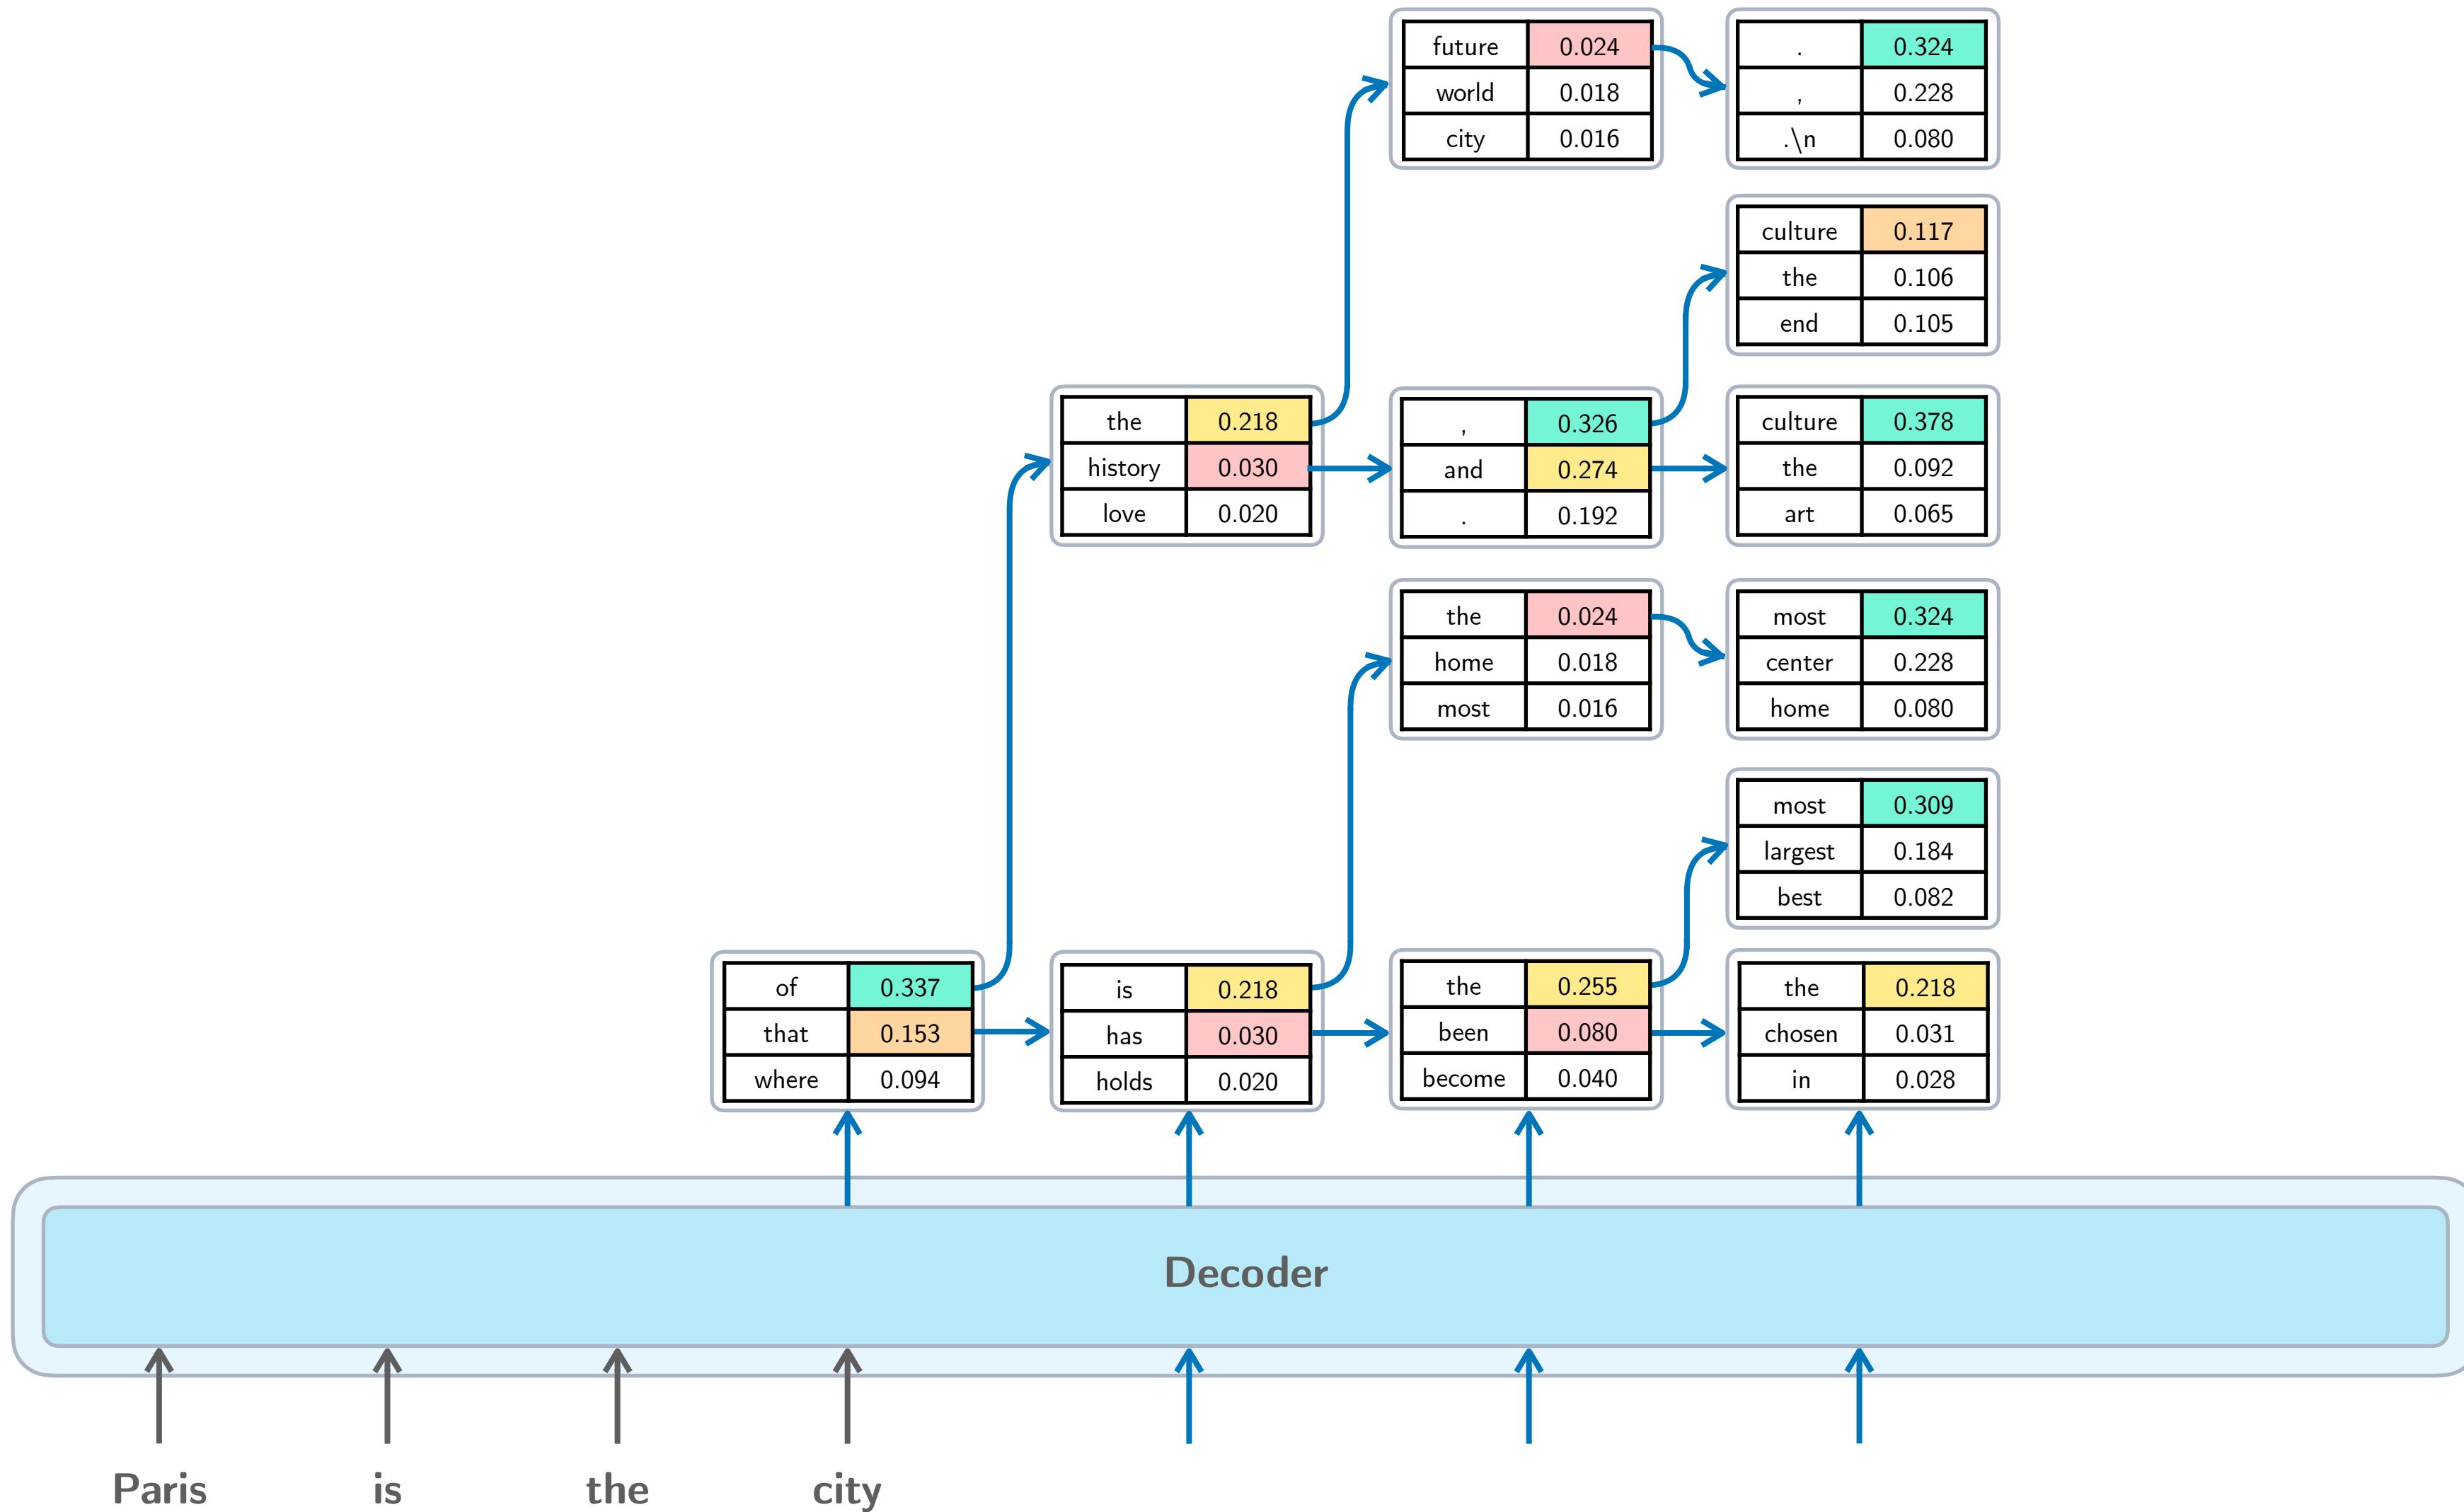
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>



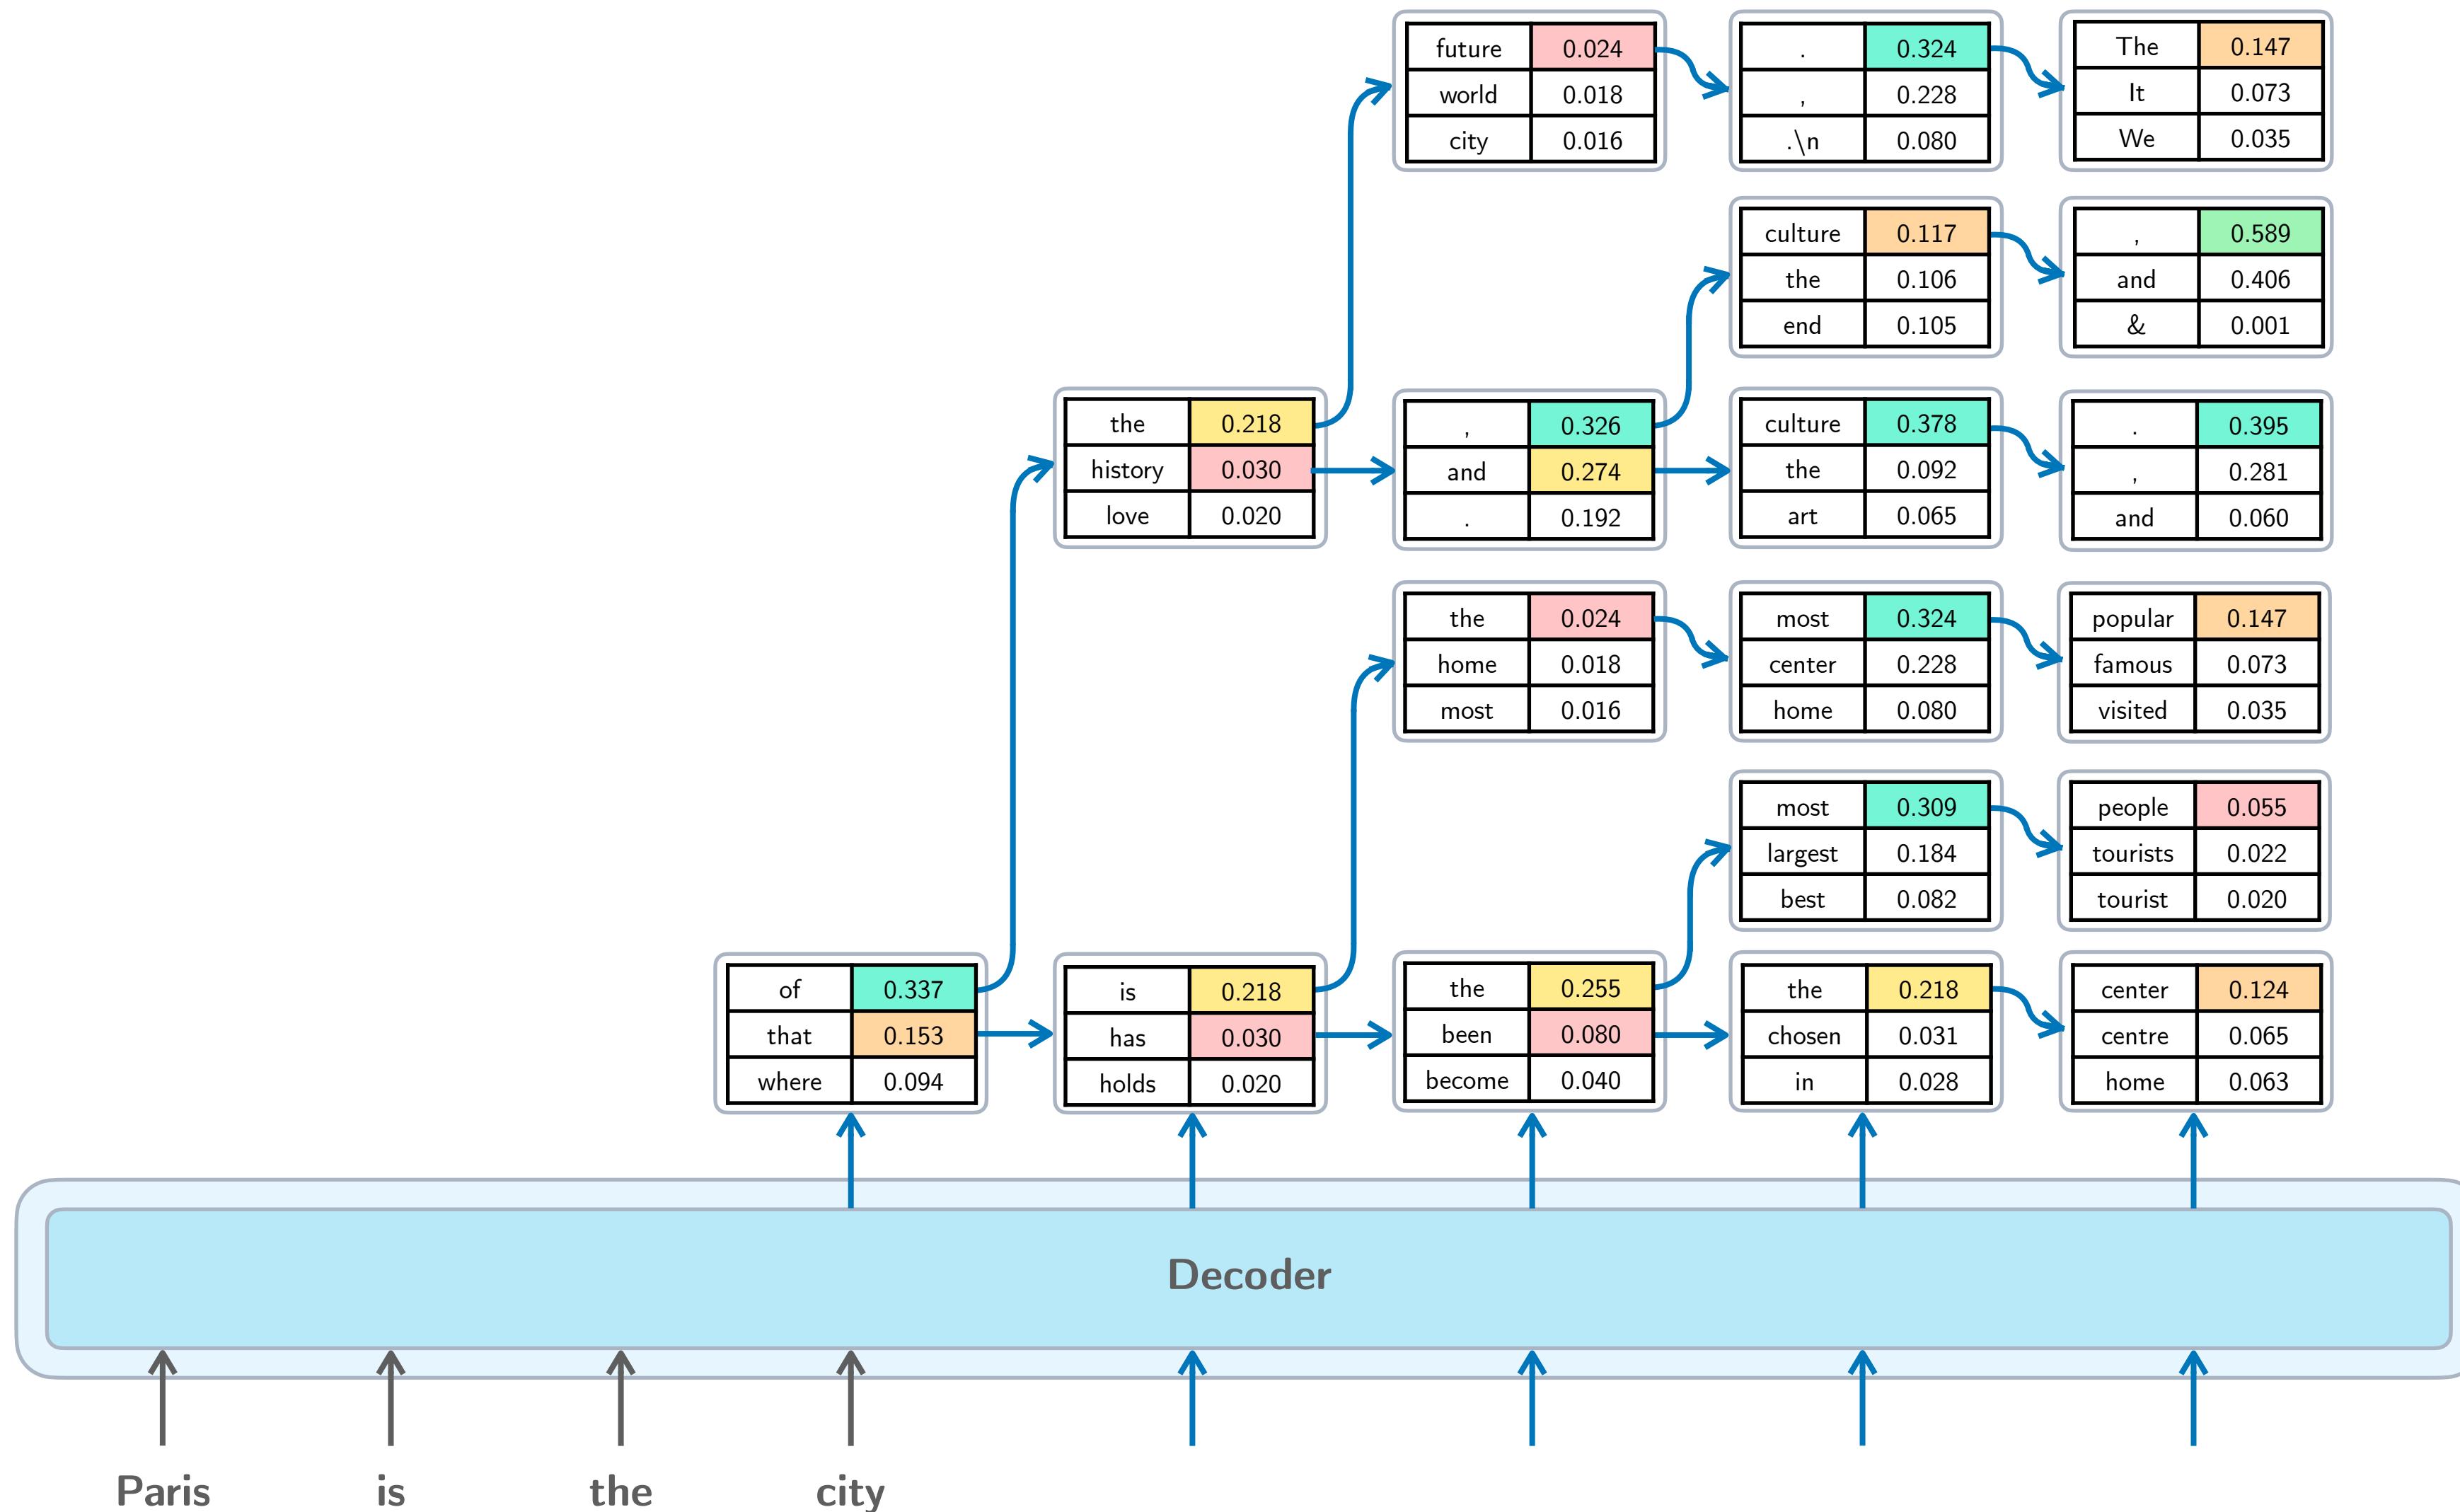
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>



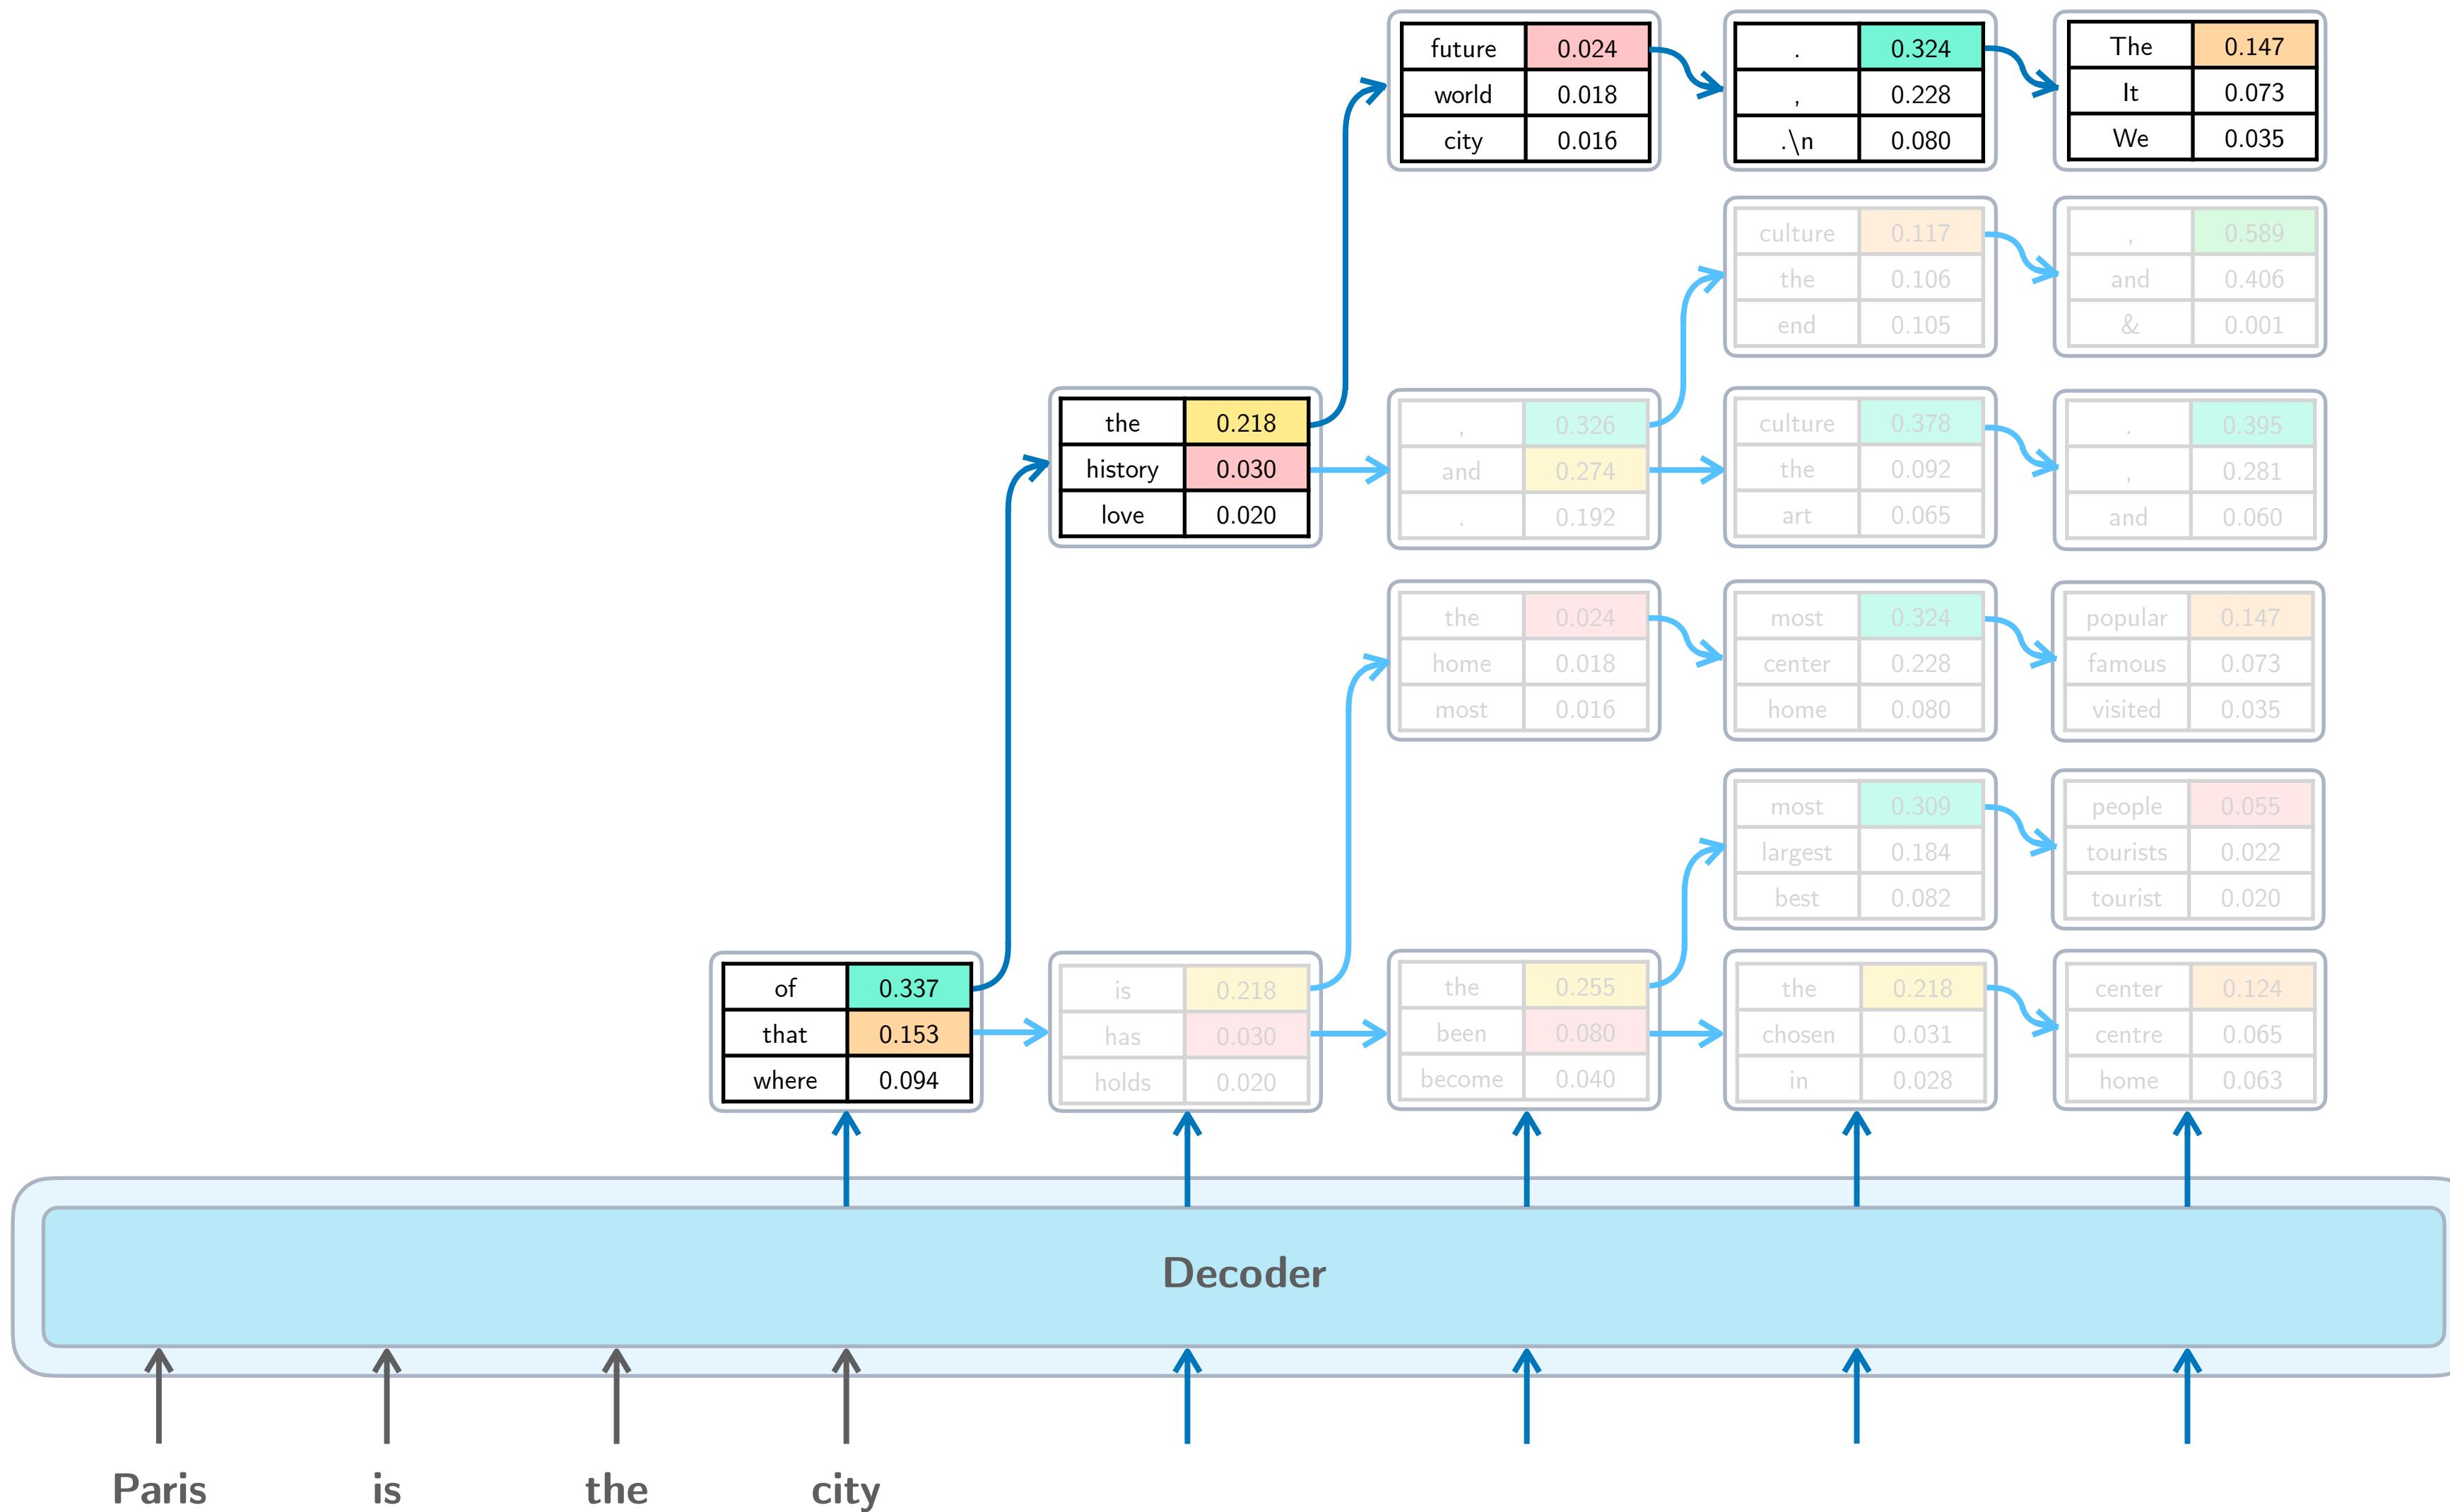
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>



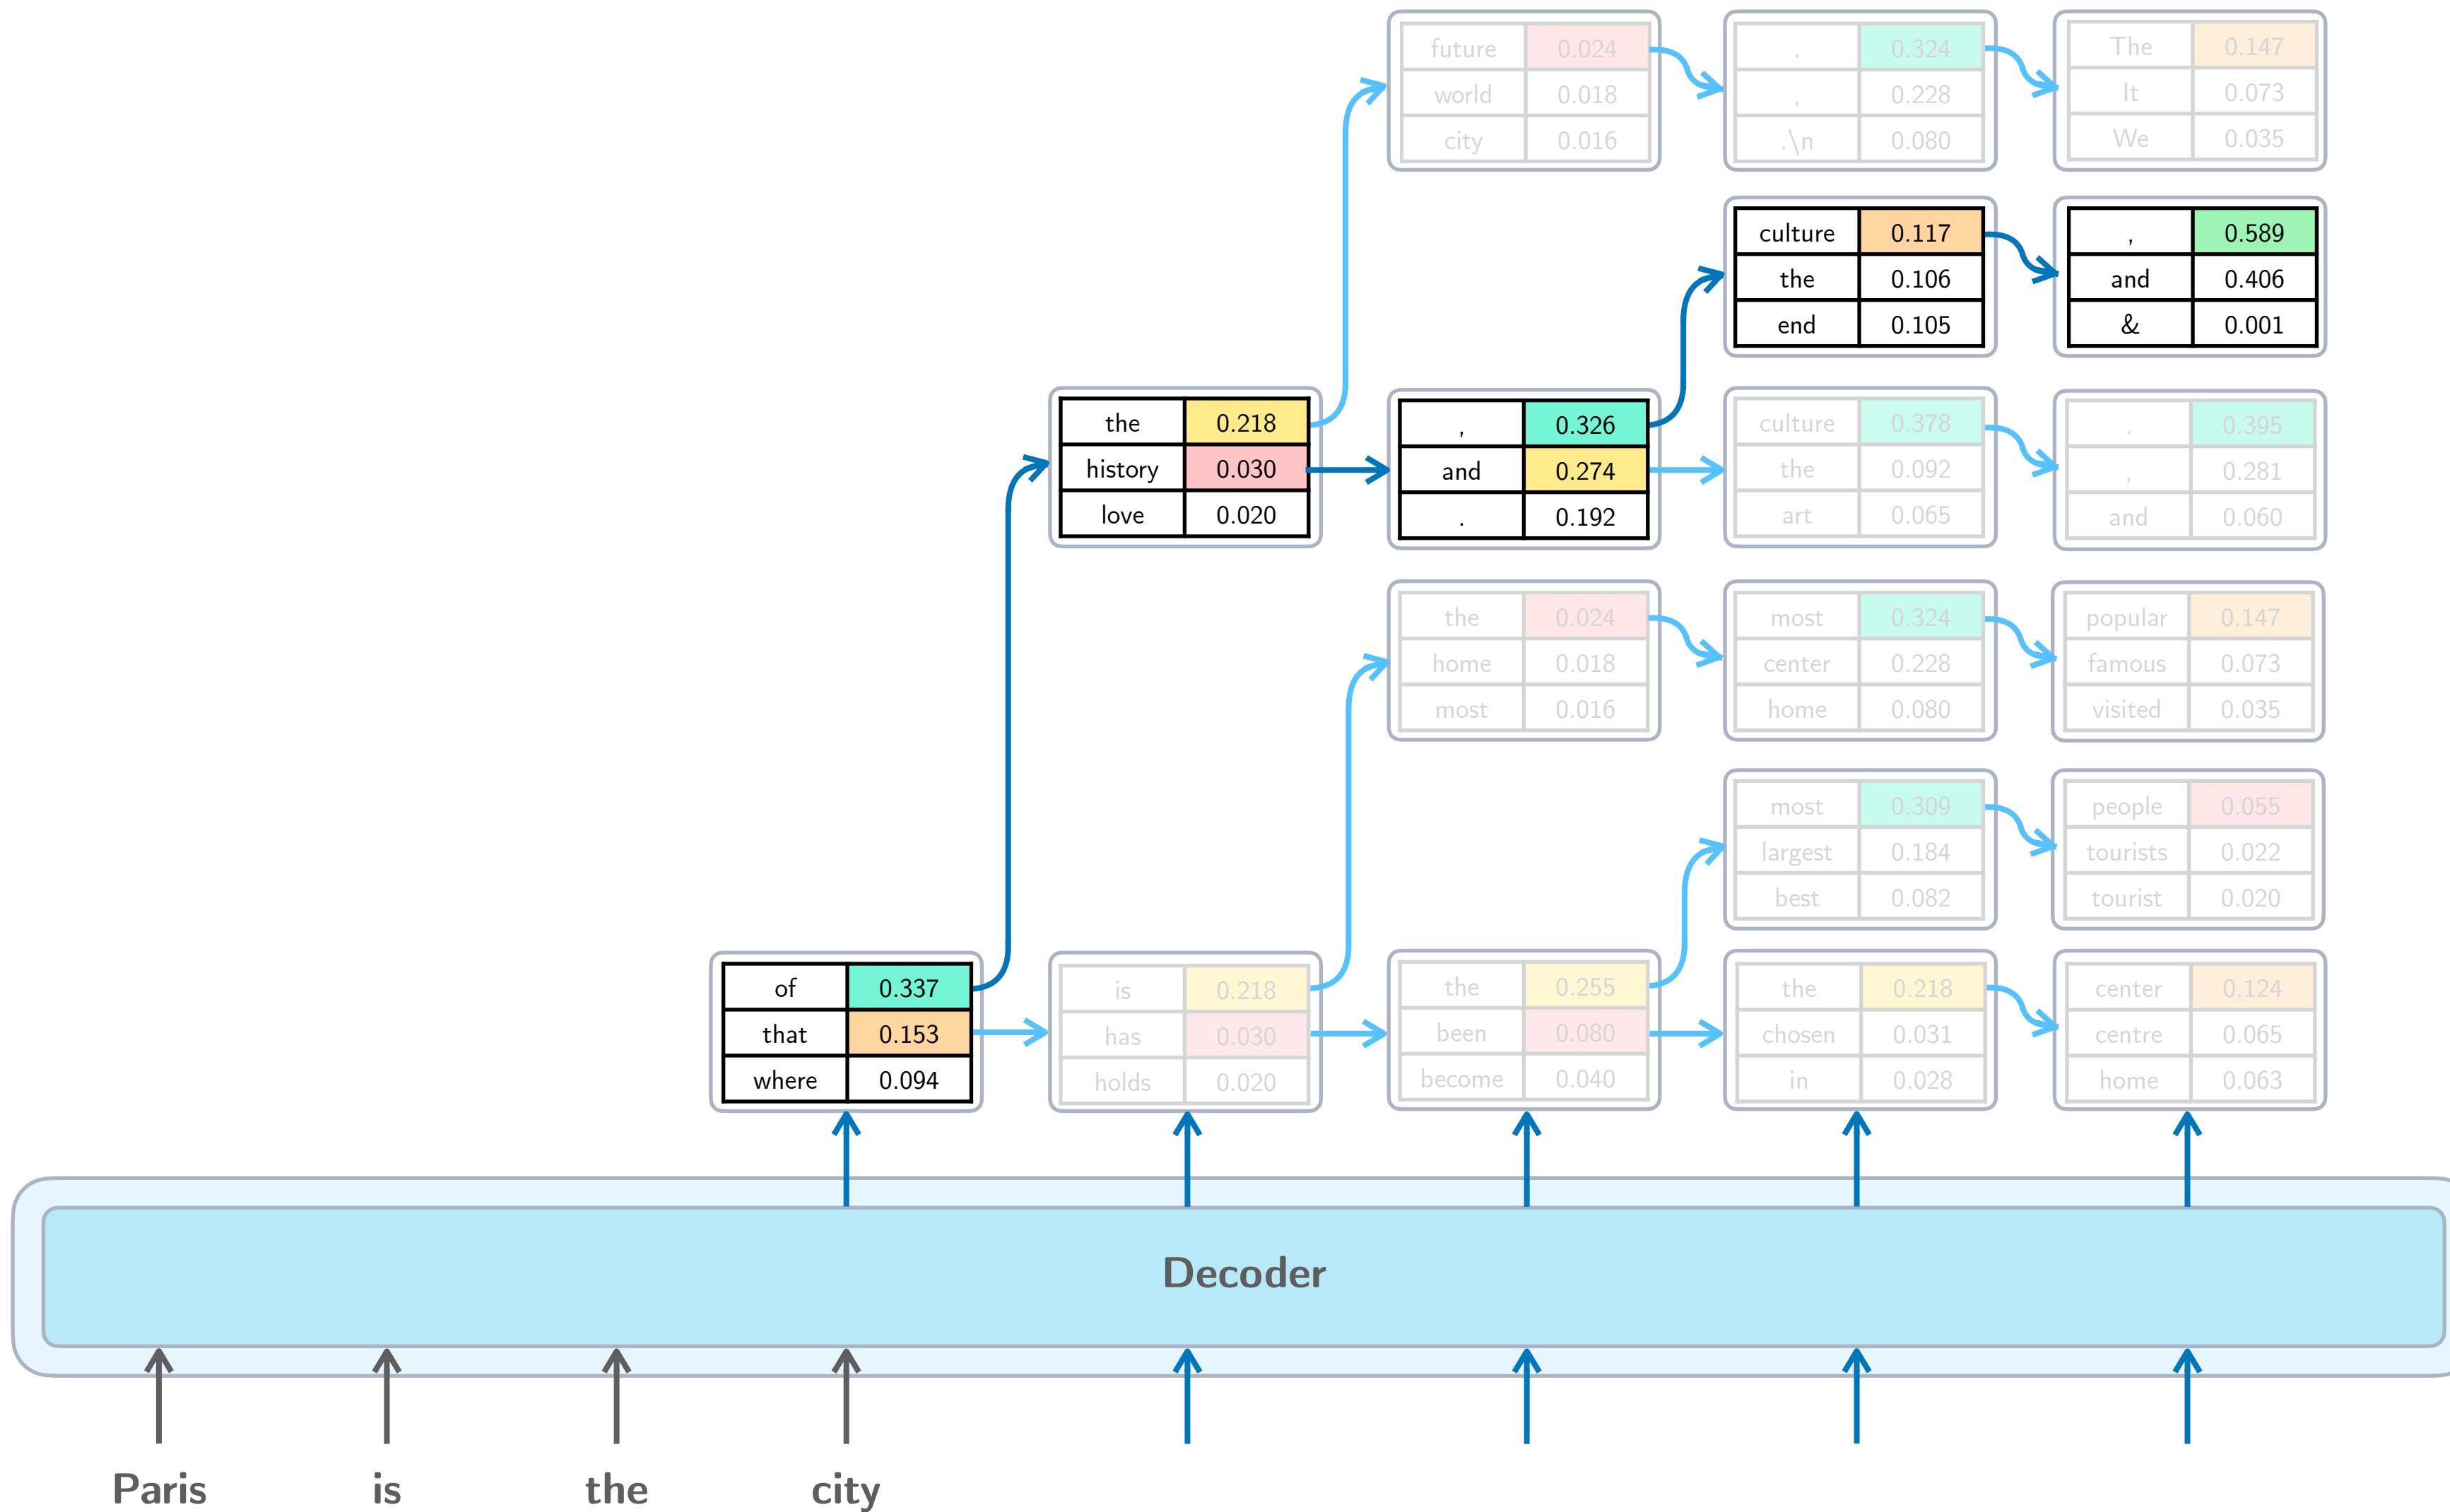
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>



Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>

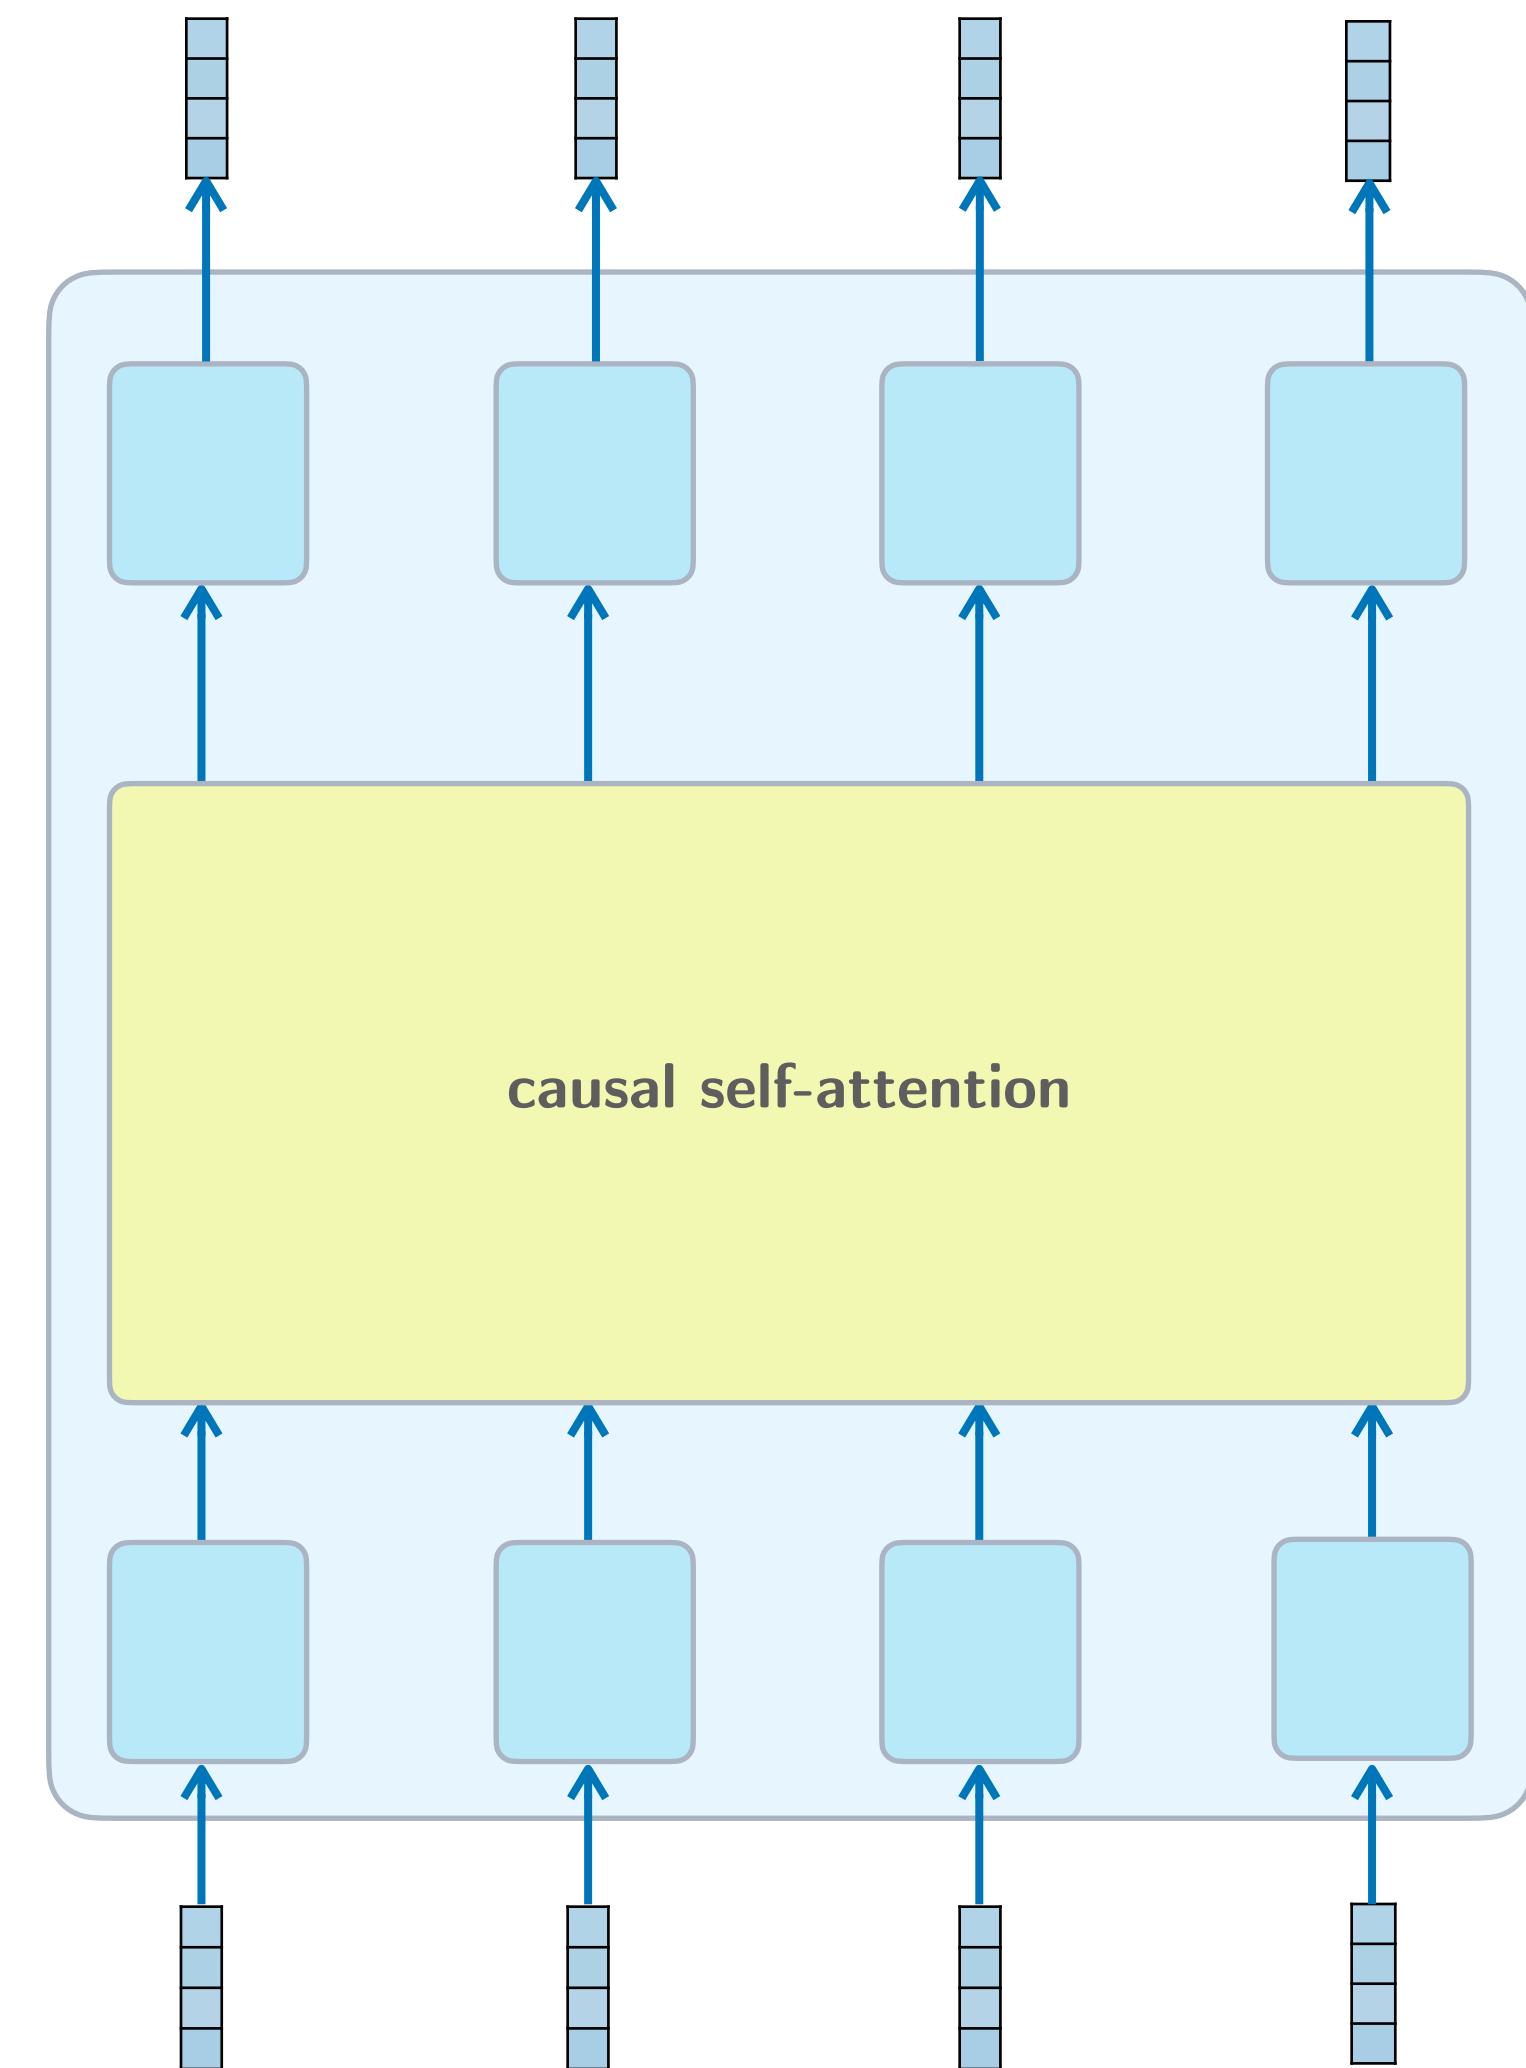


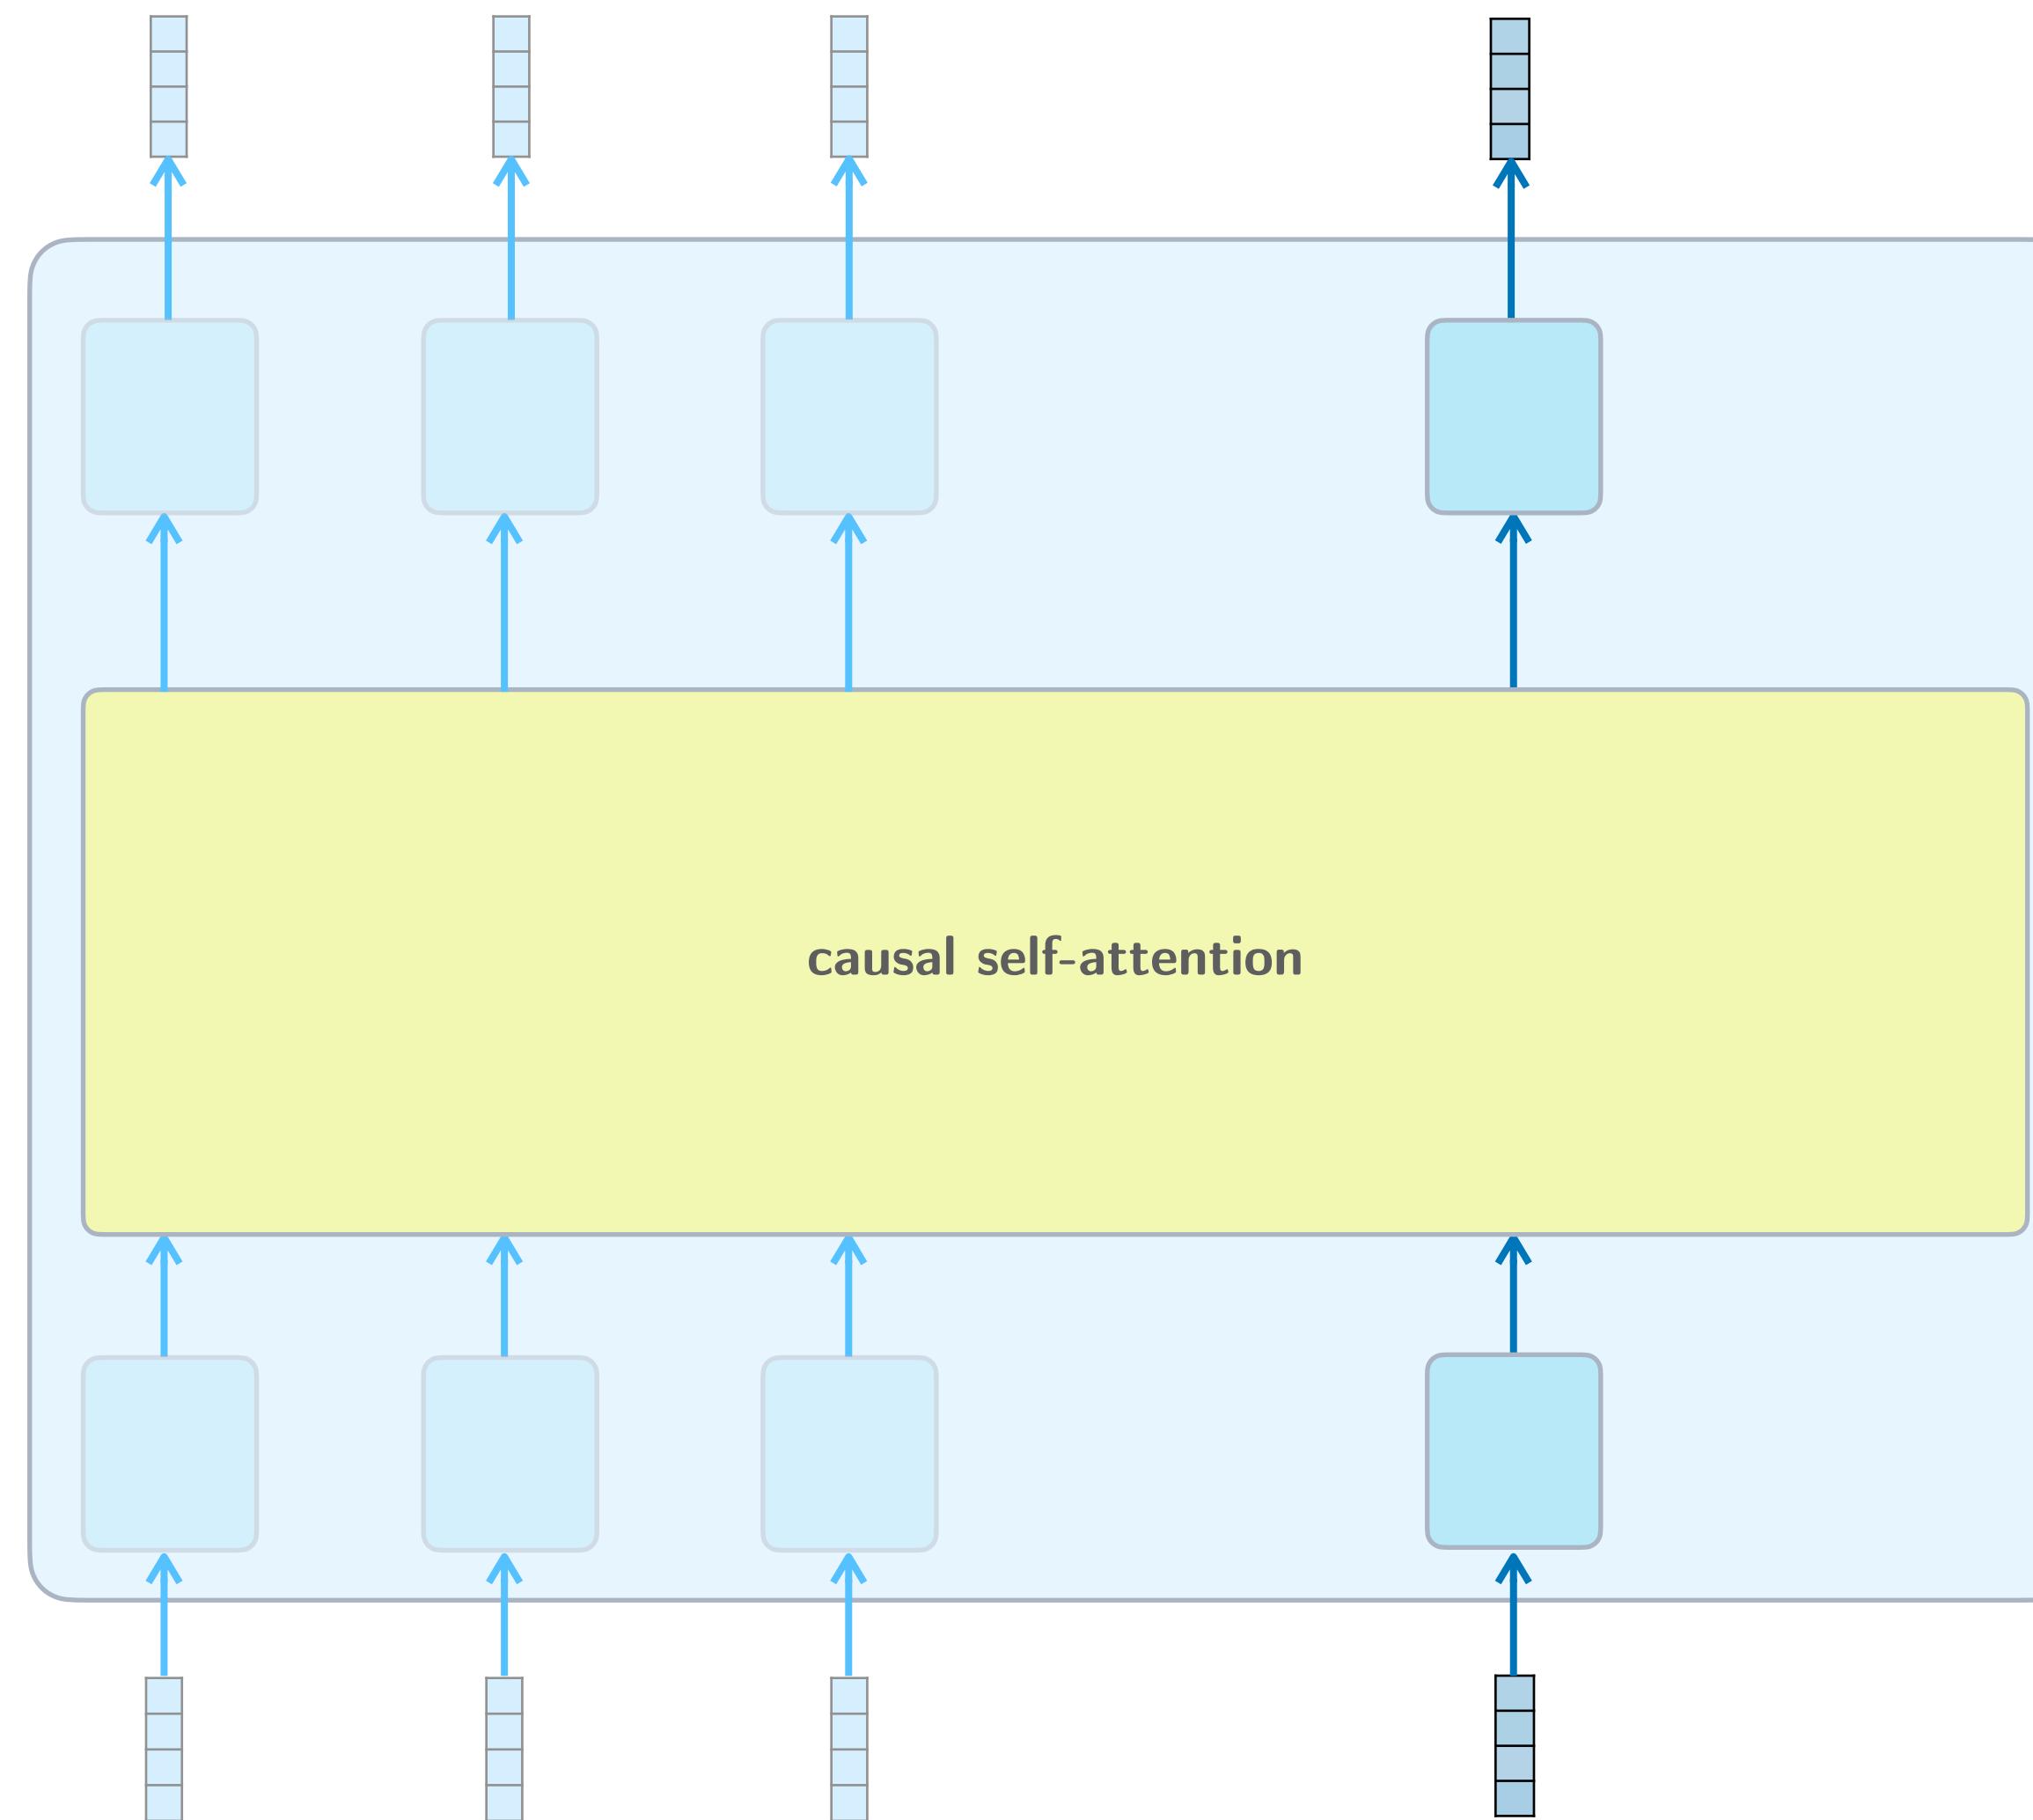
Based on <https://pub.towardsai.net/how-does-an-lm-generate-text-fd9c57781217>

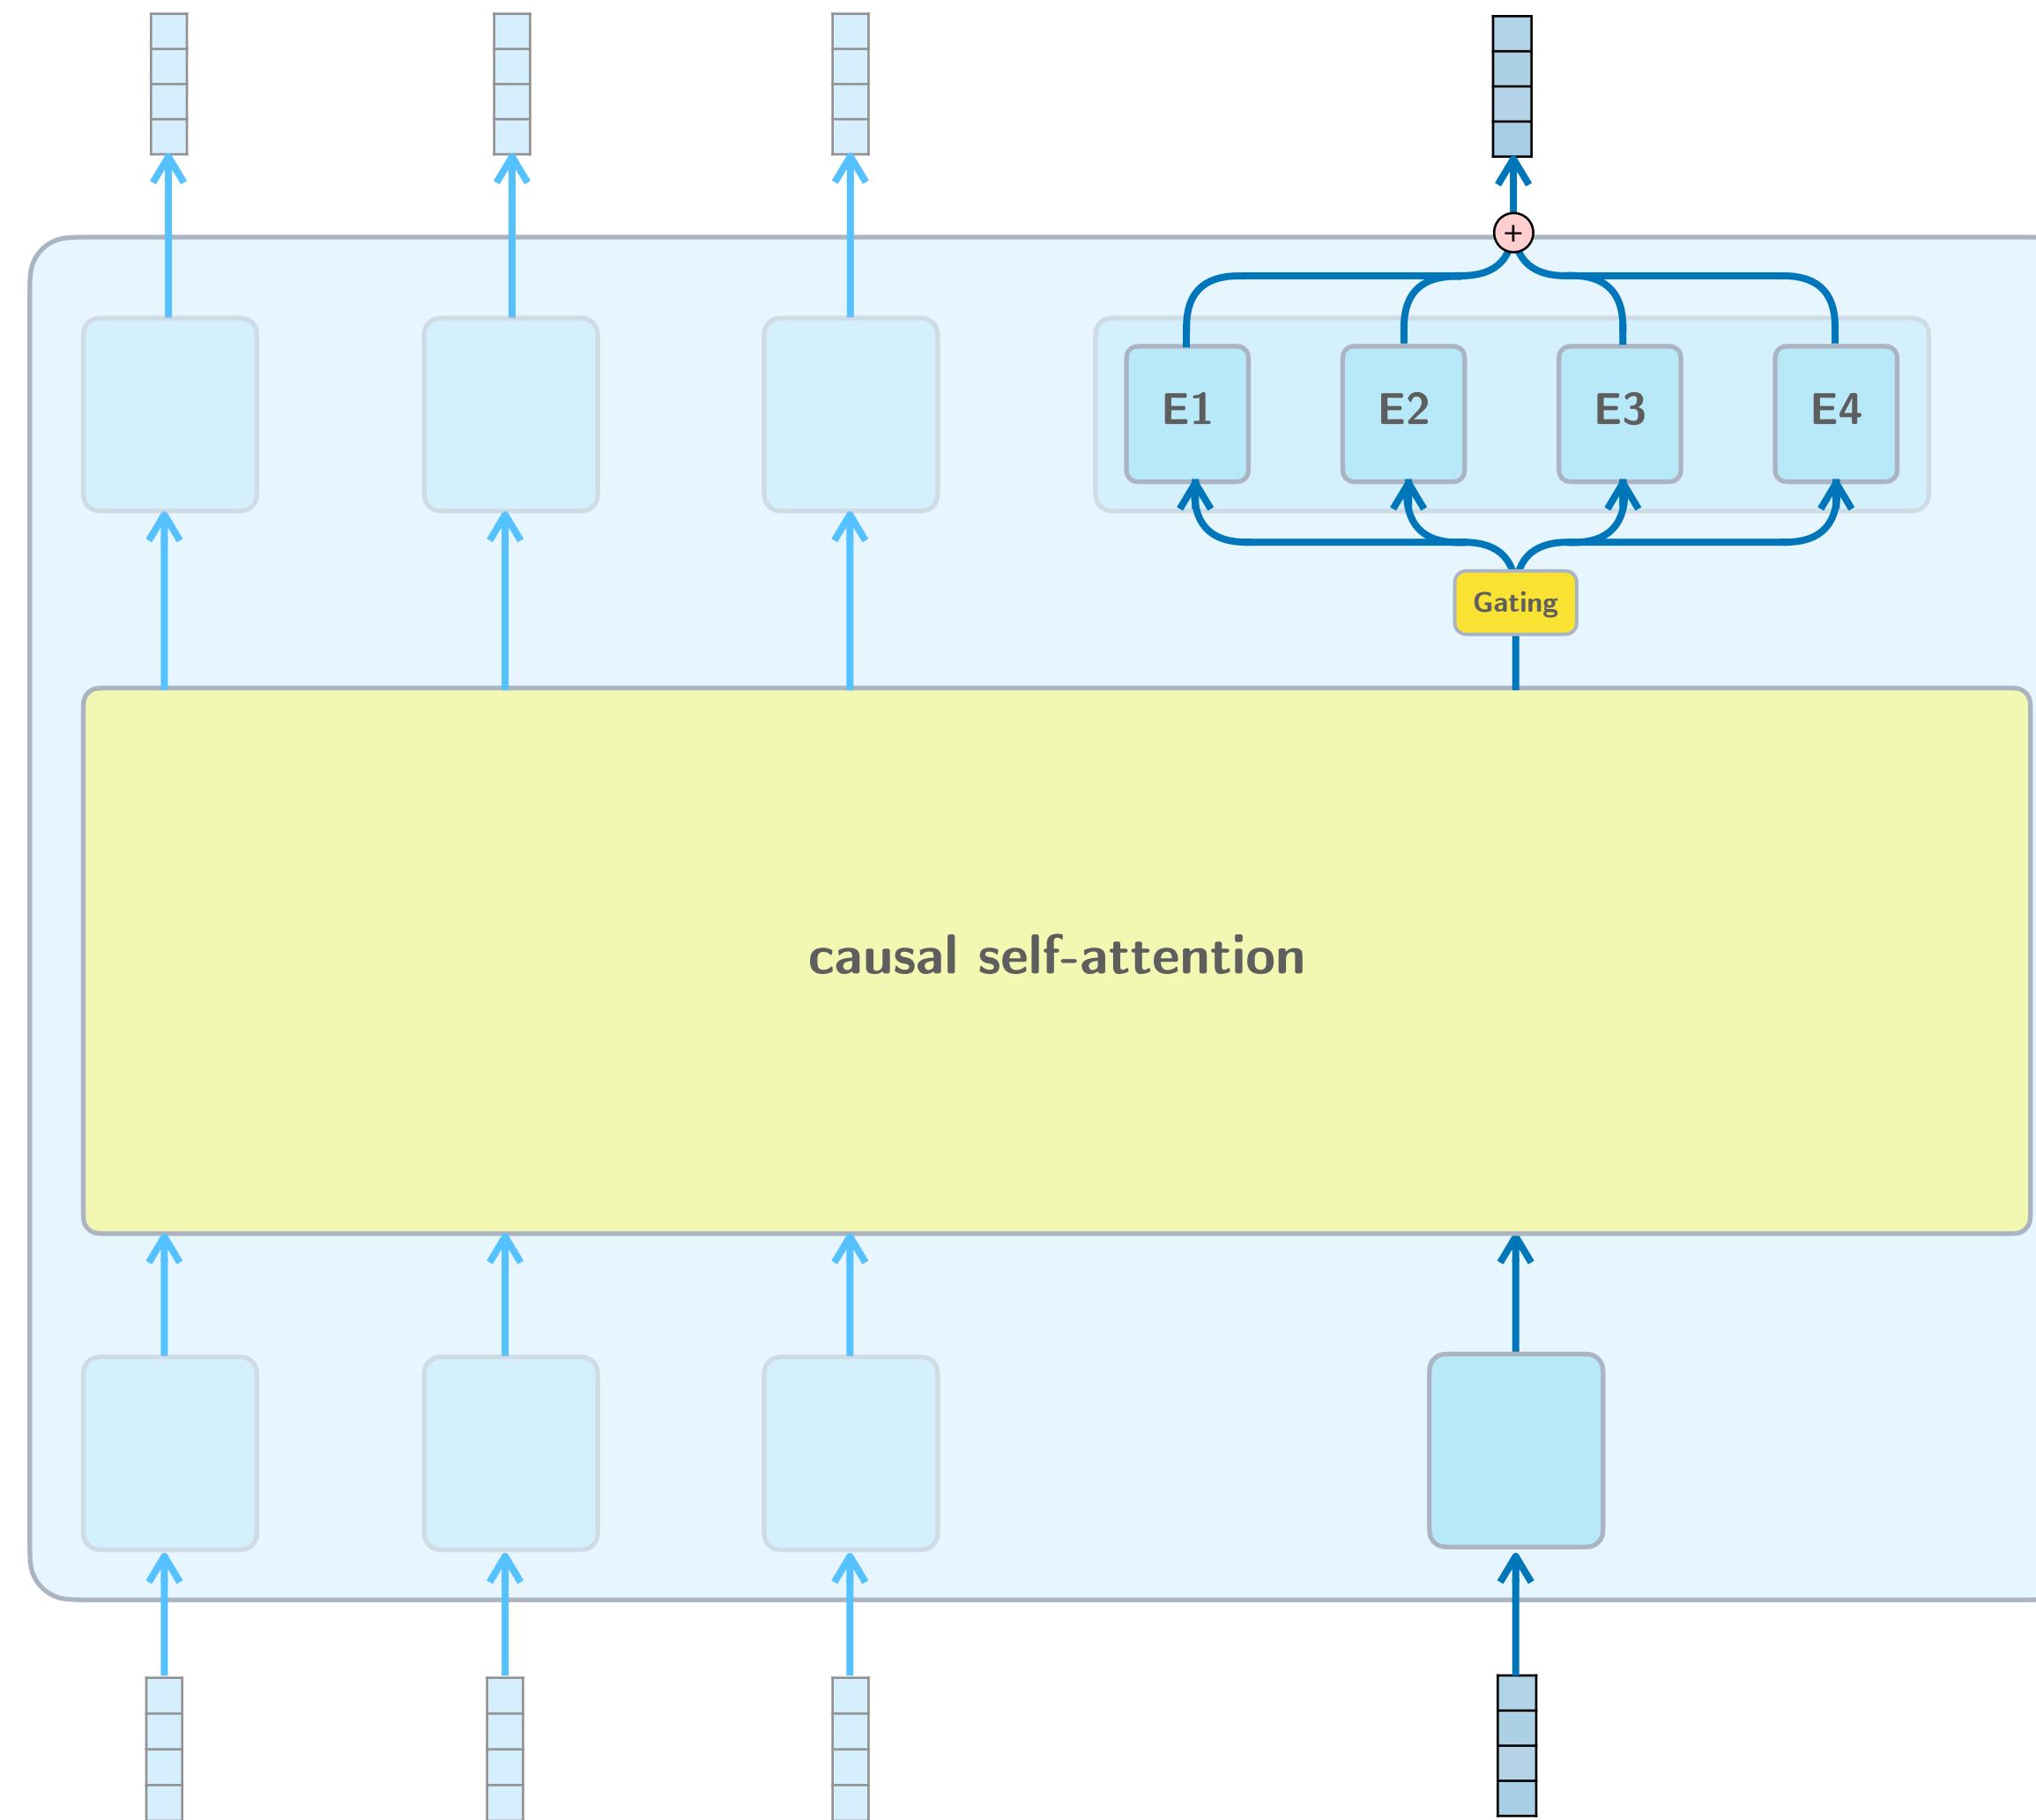




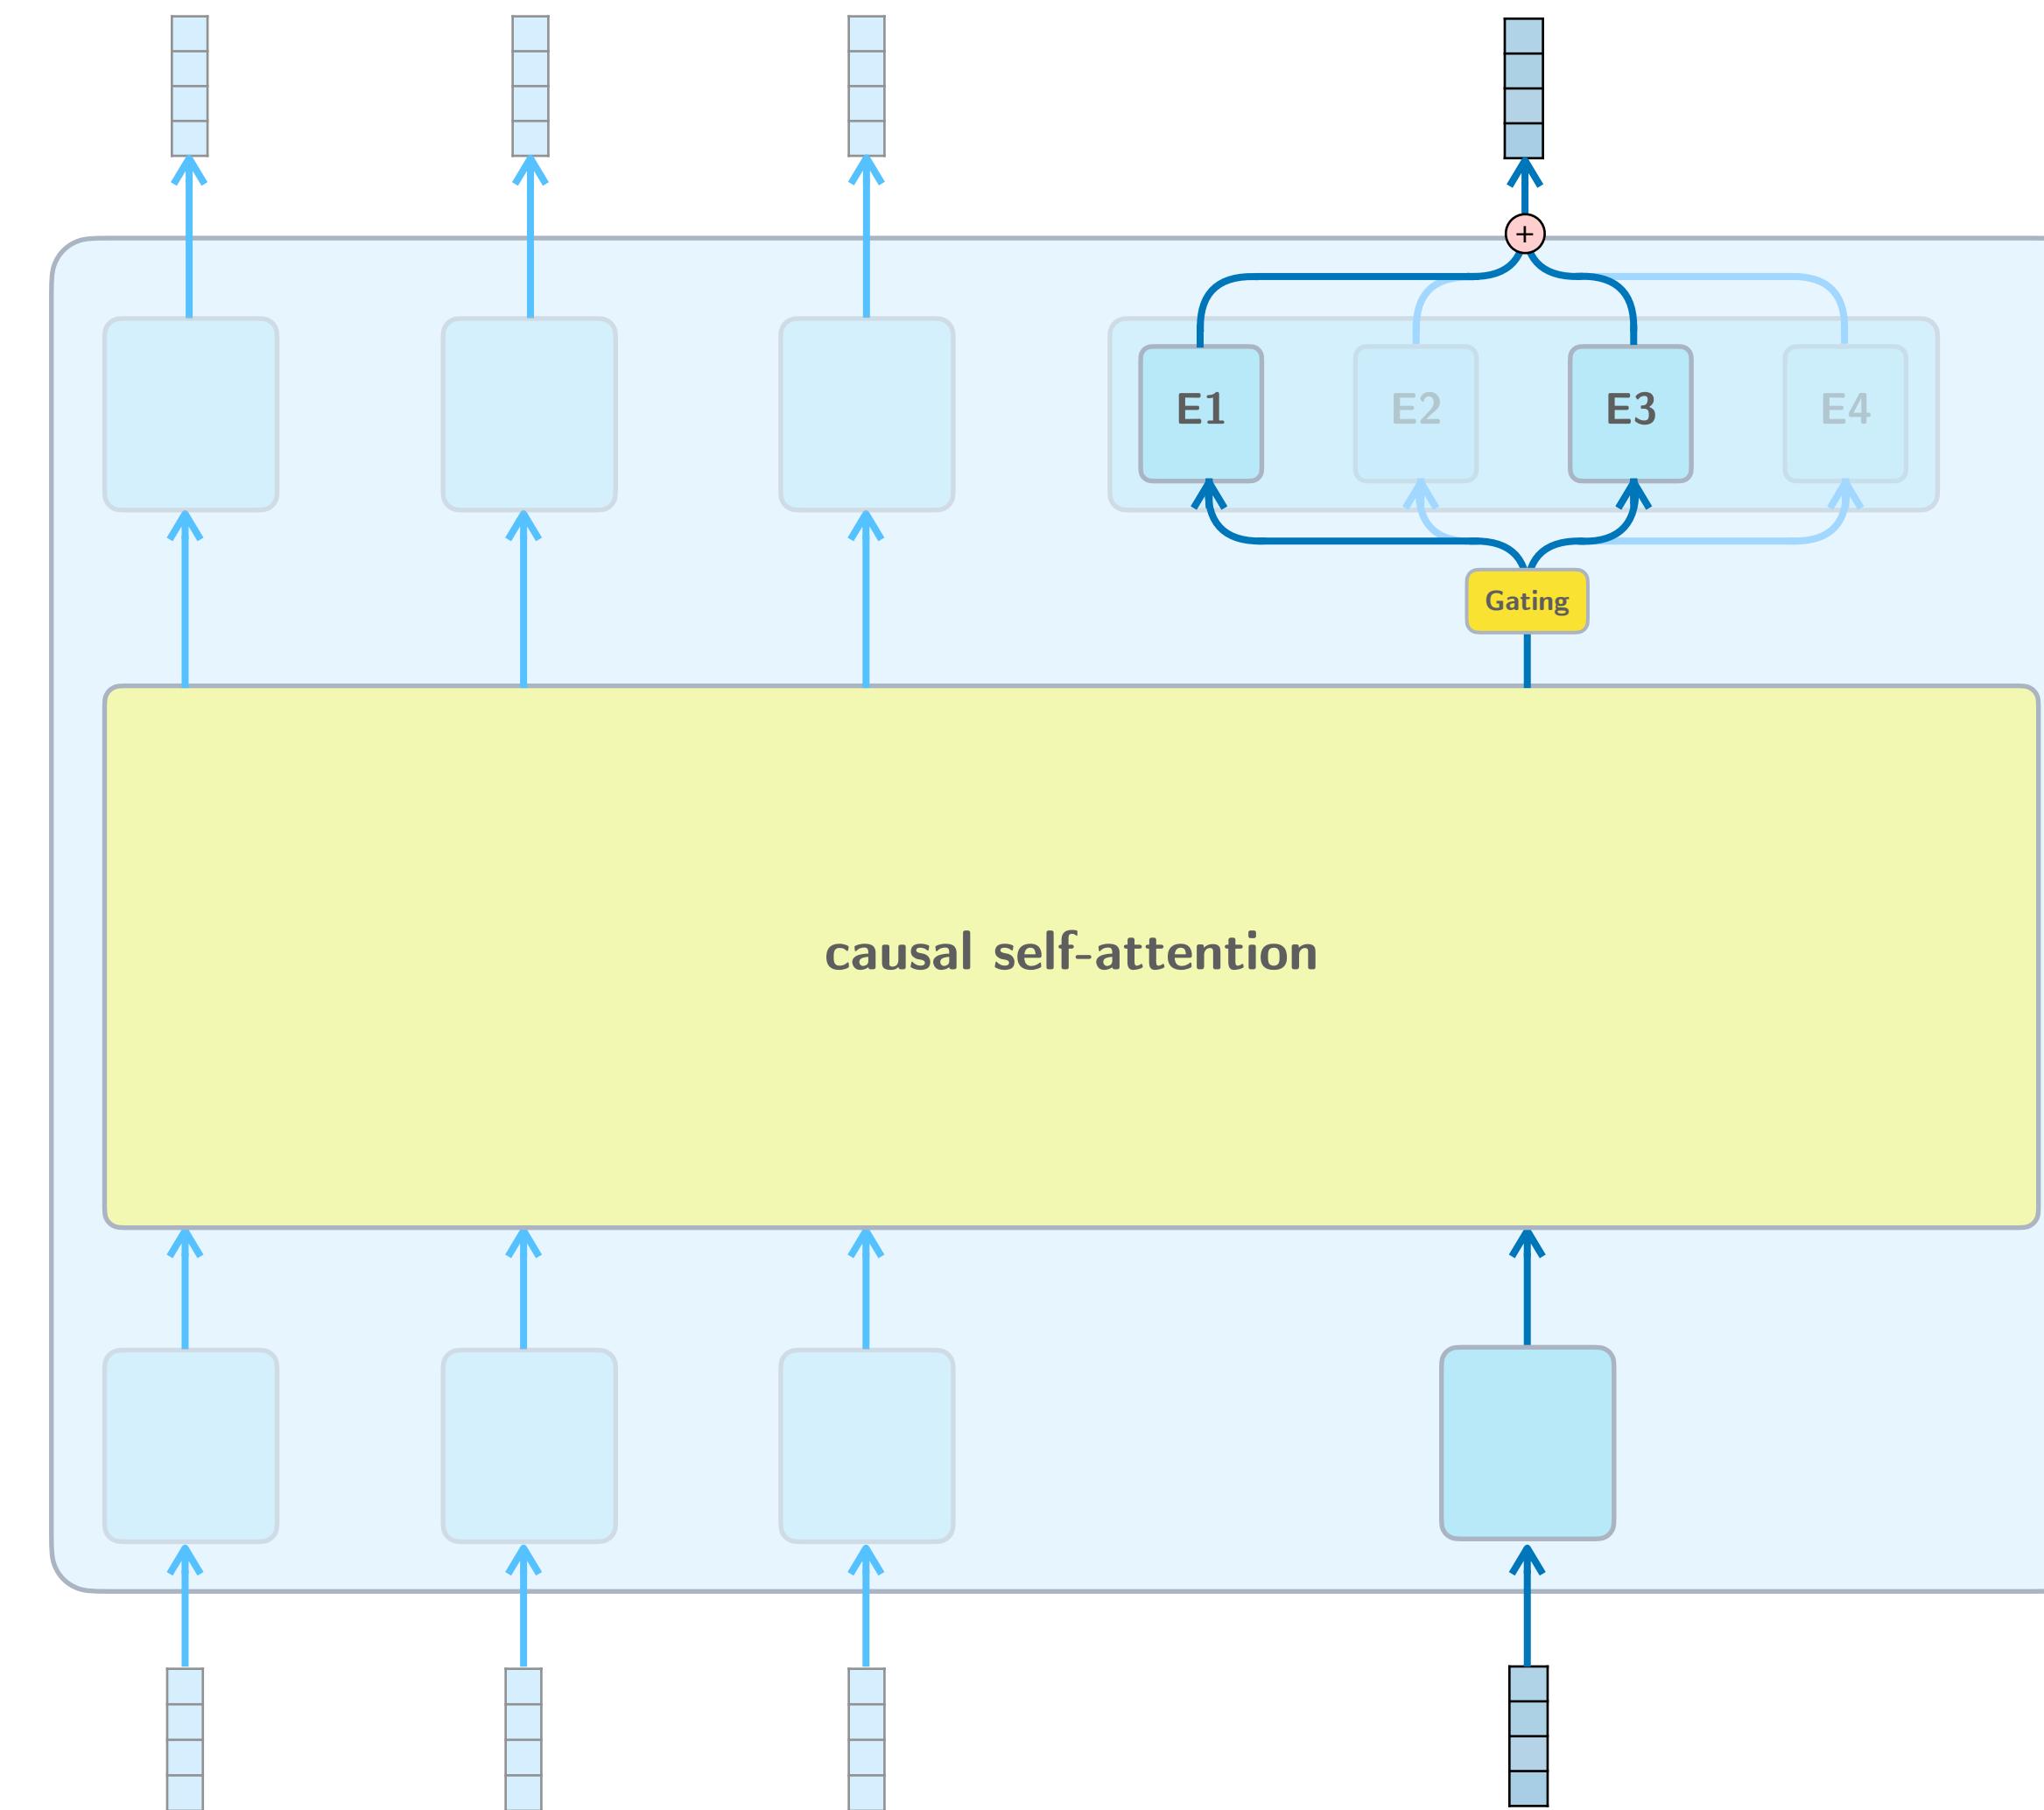
Mixture of Experts



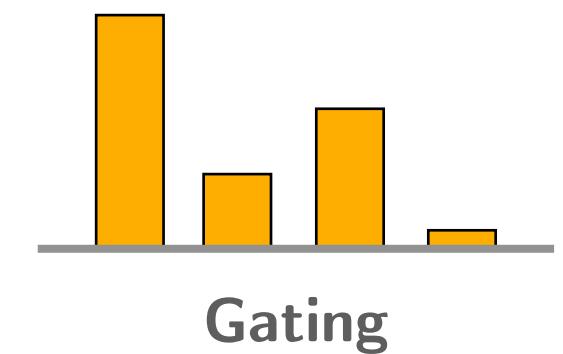


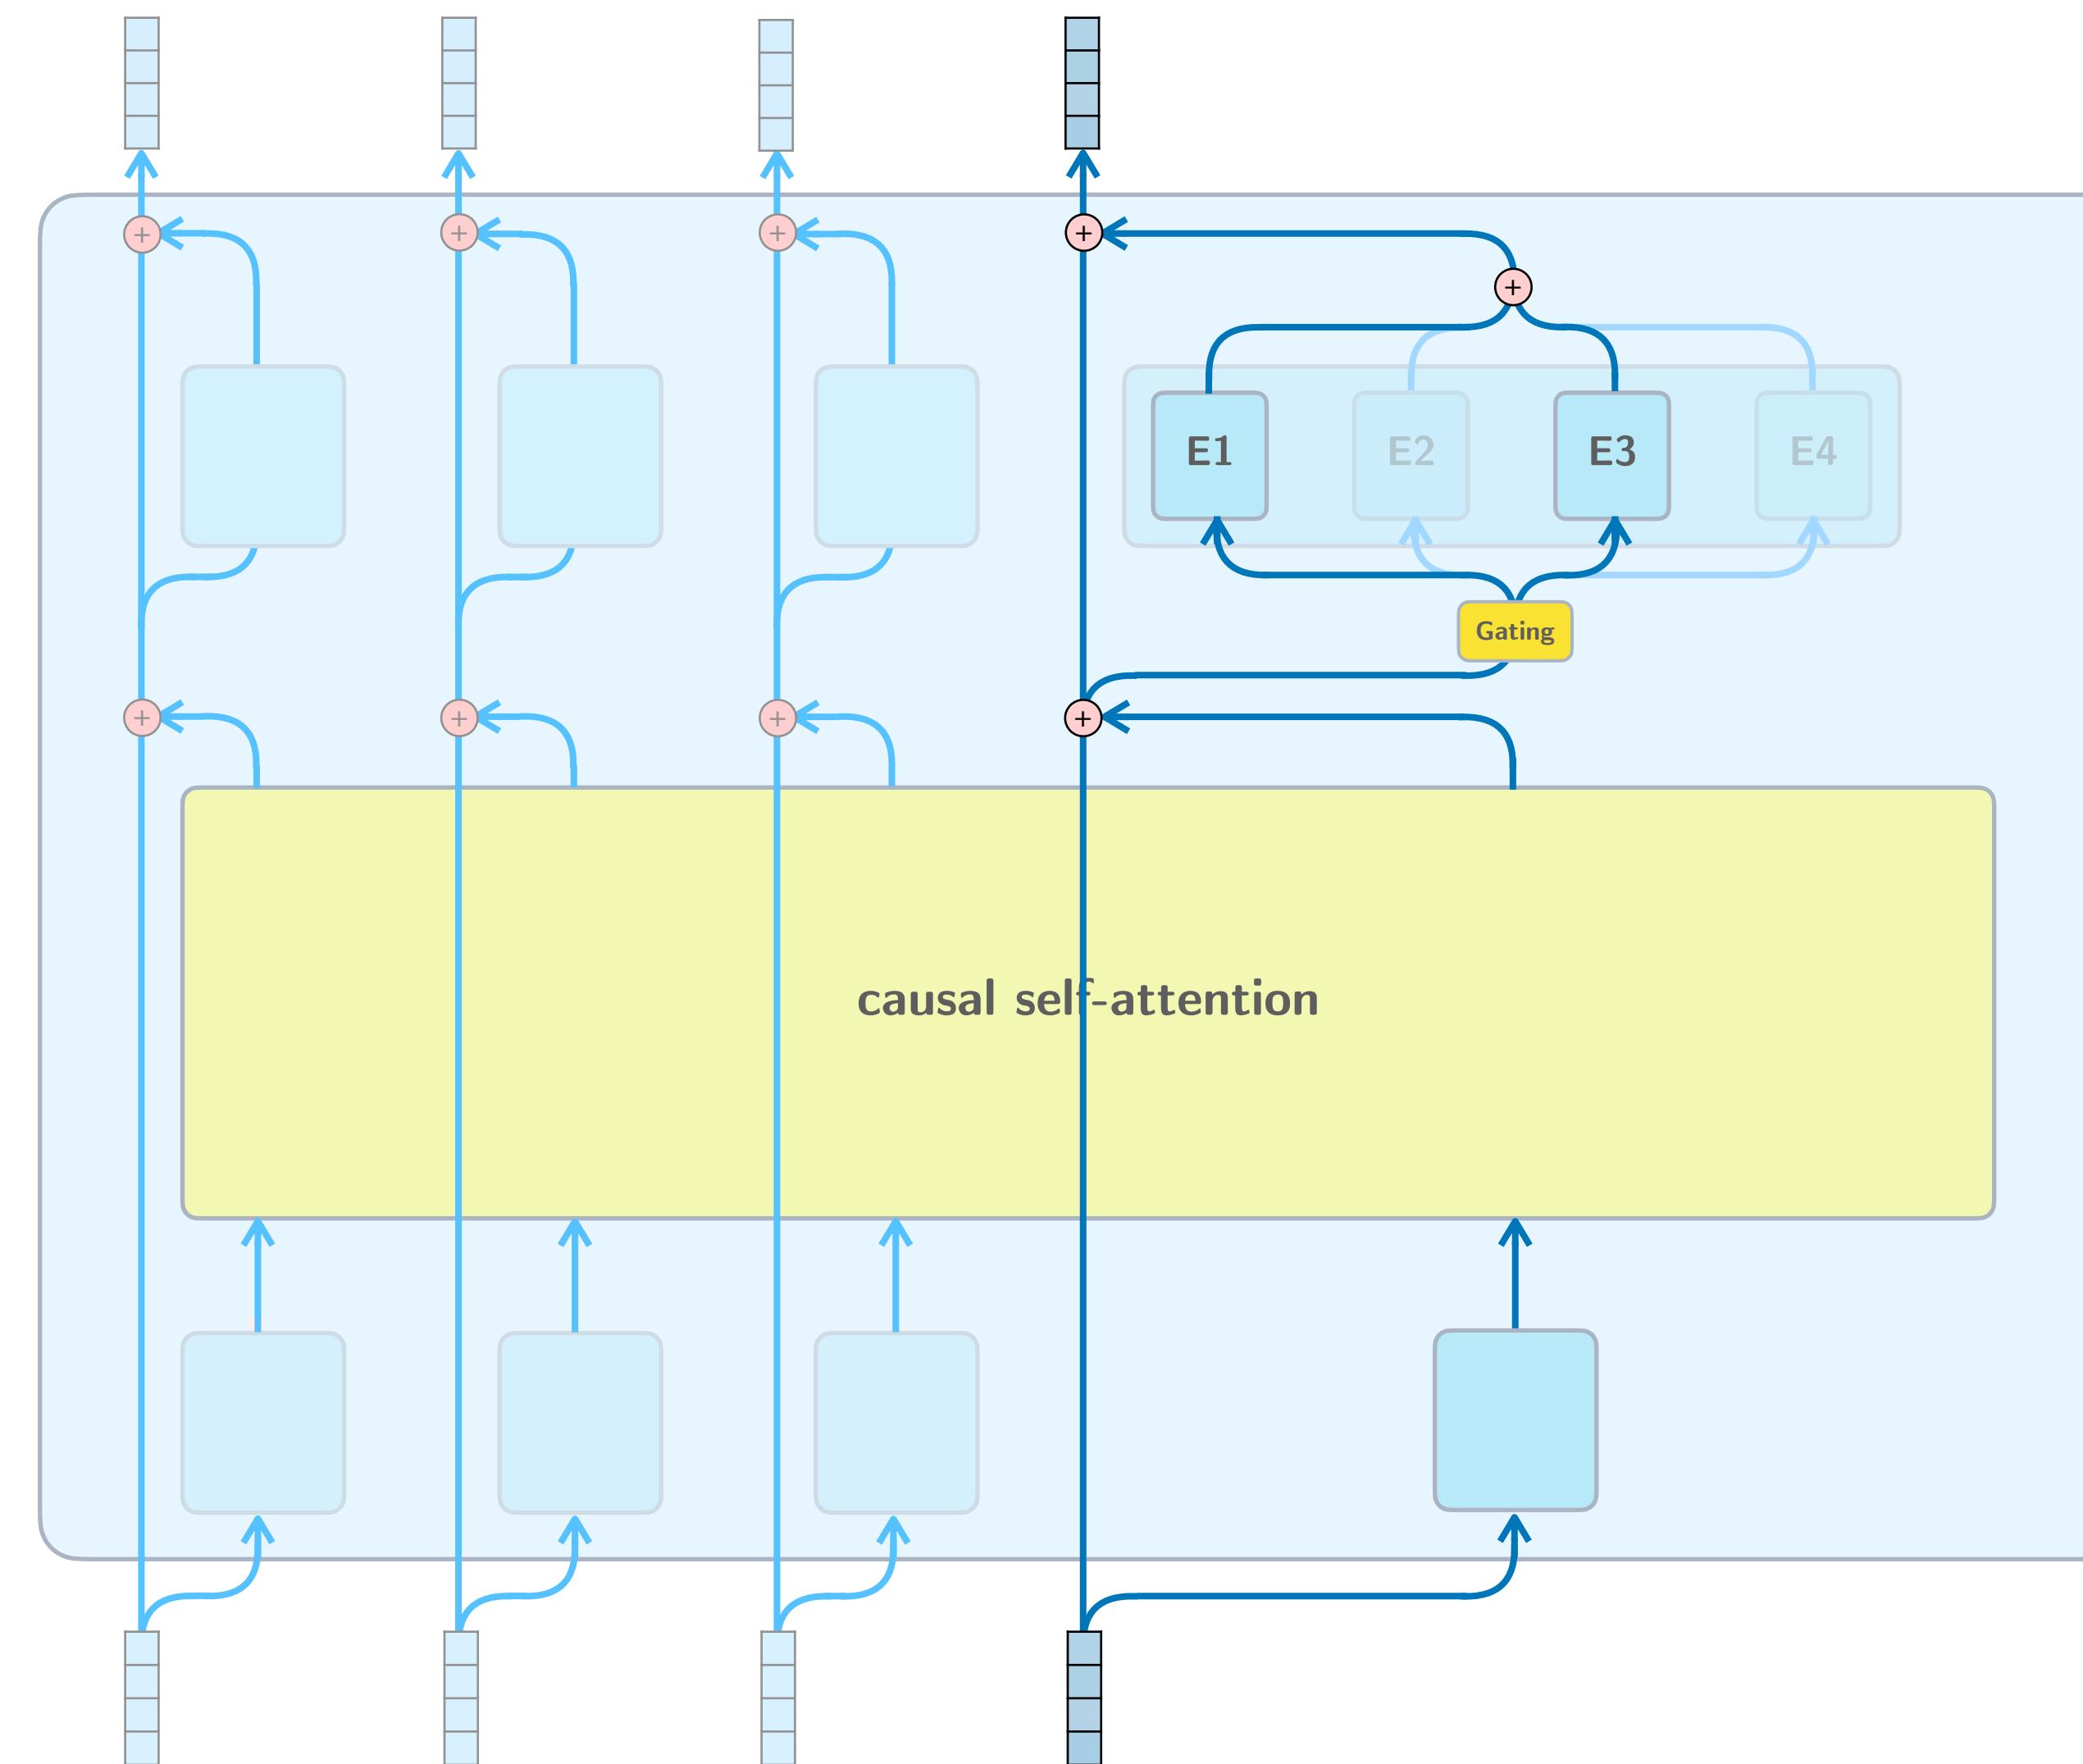


Multiple “expert” MLPs instead of a single MLP.

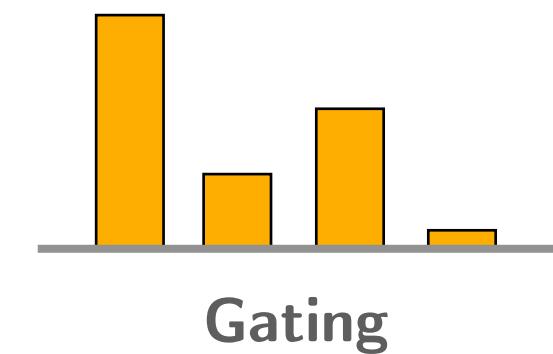


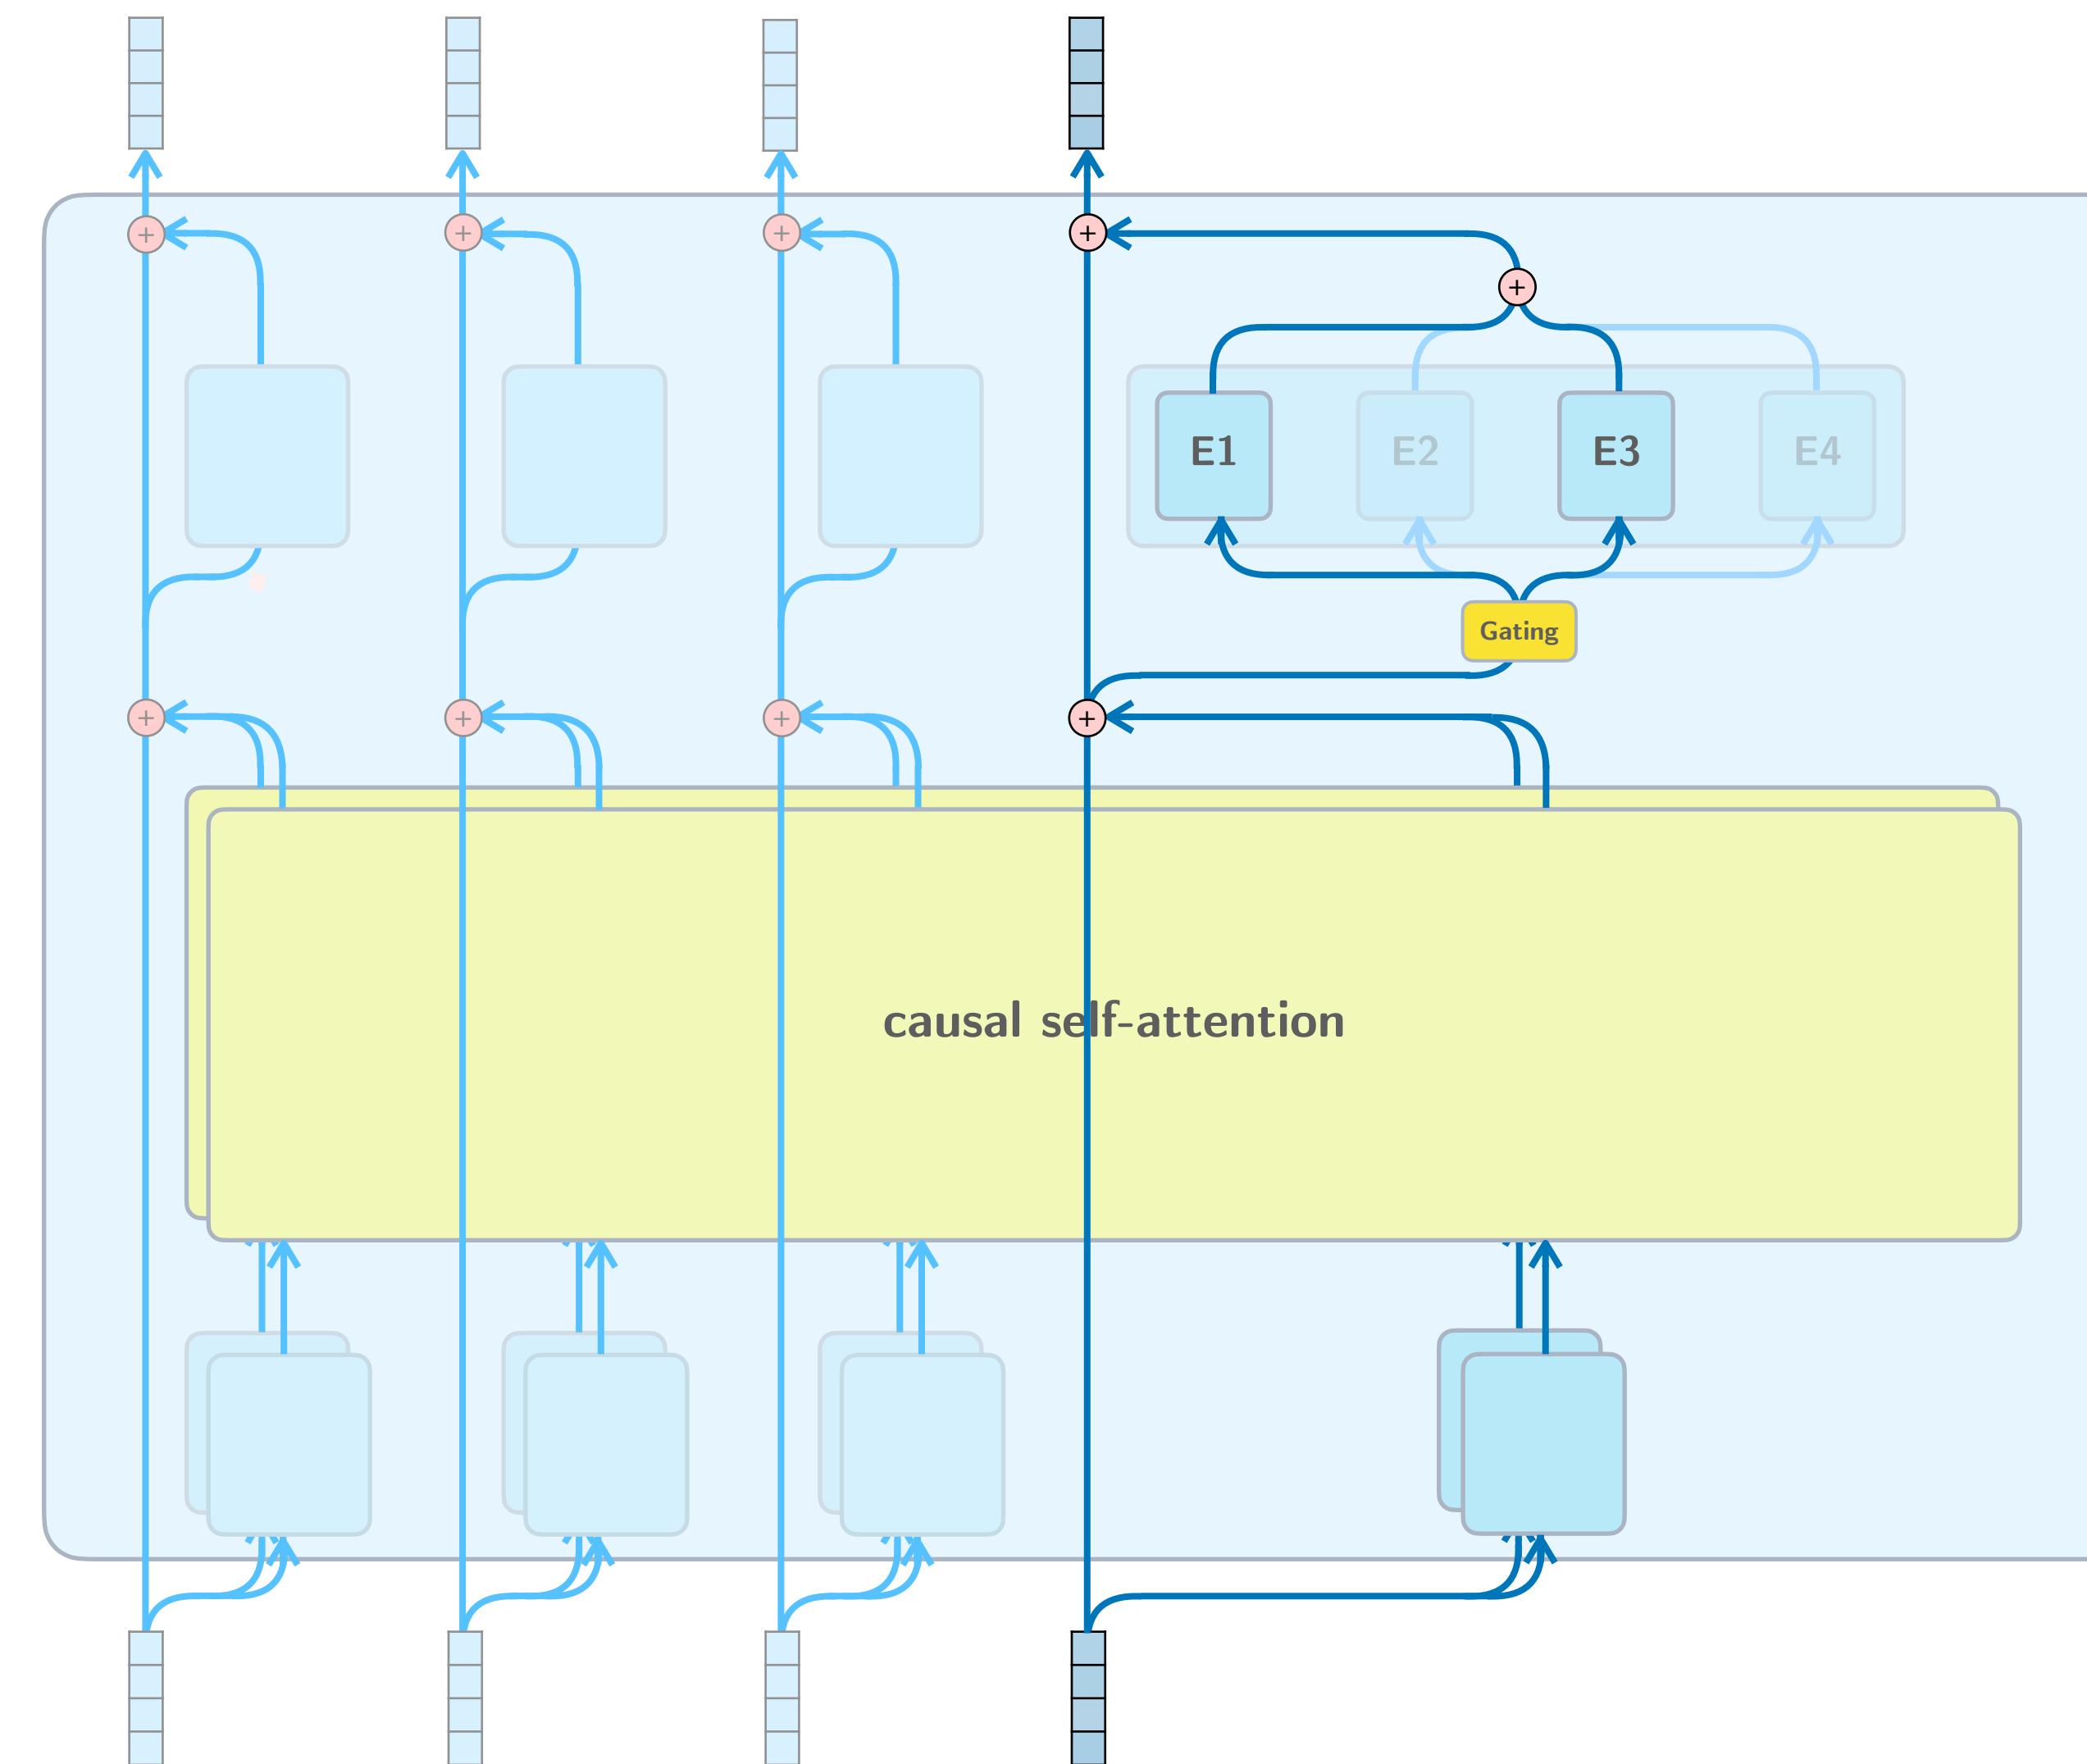
Multiple “expert” MLPs instead of a single MLP.



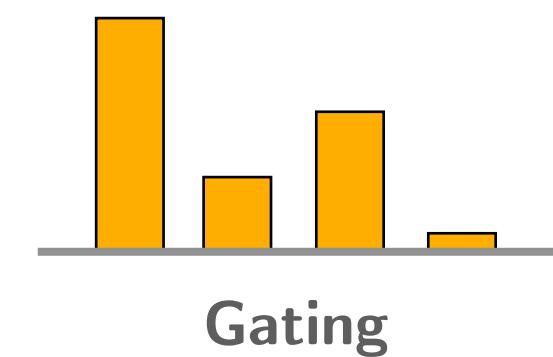


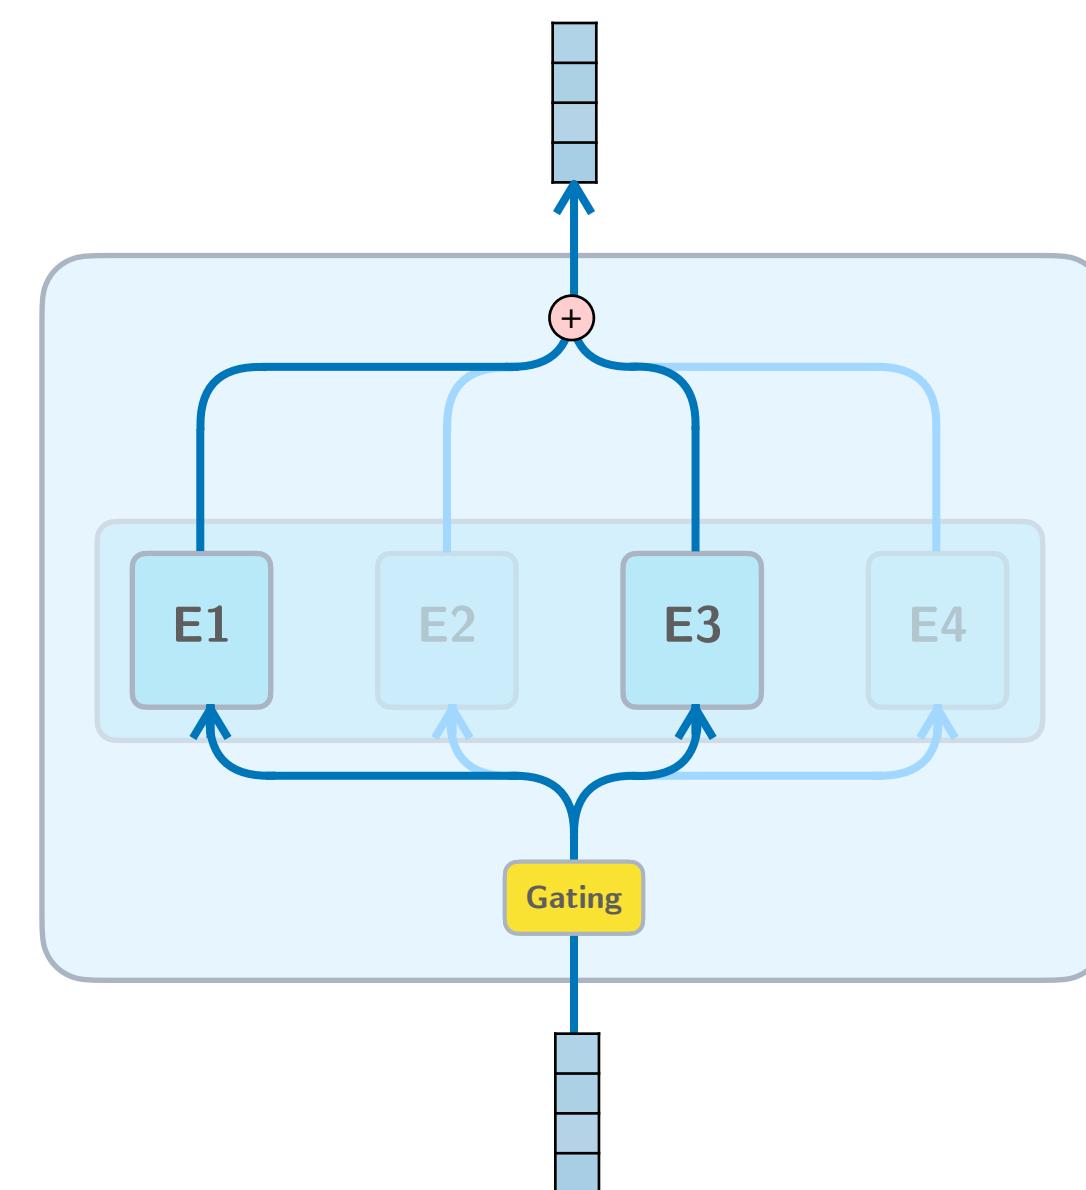
Multiple “expert” MLPs
instead of a single MLP.



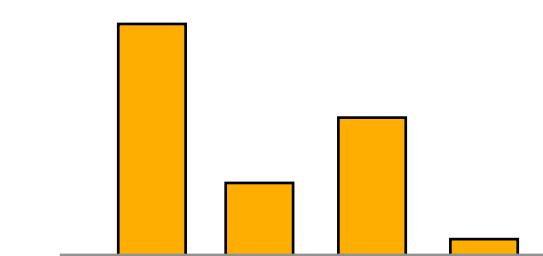


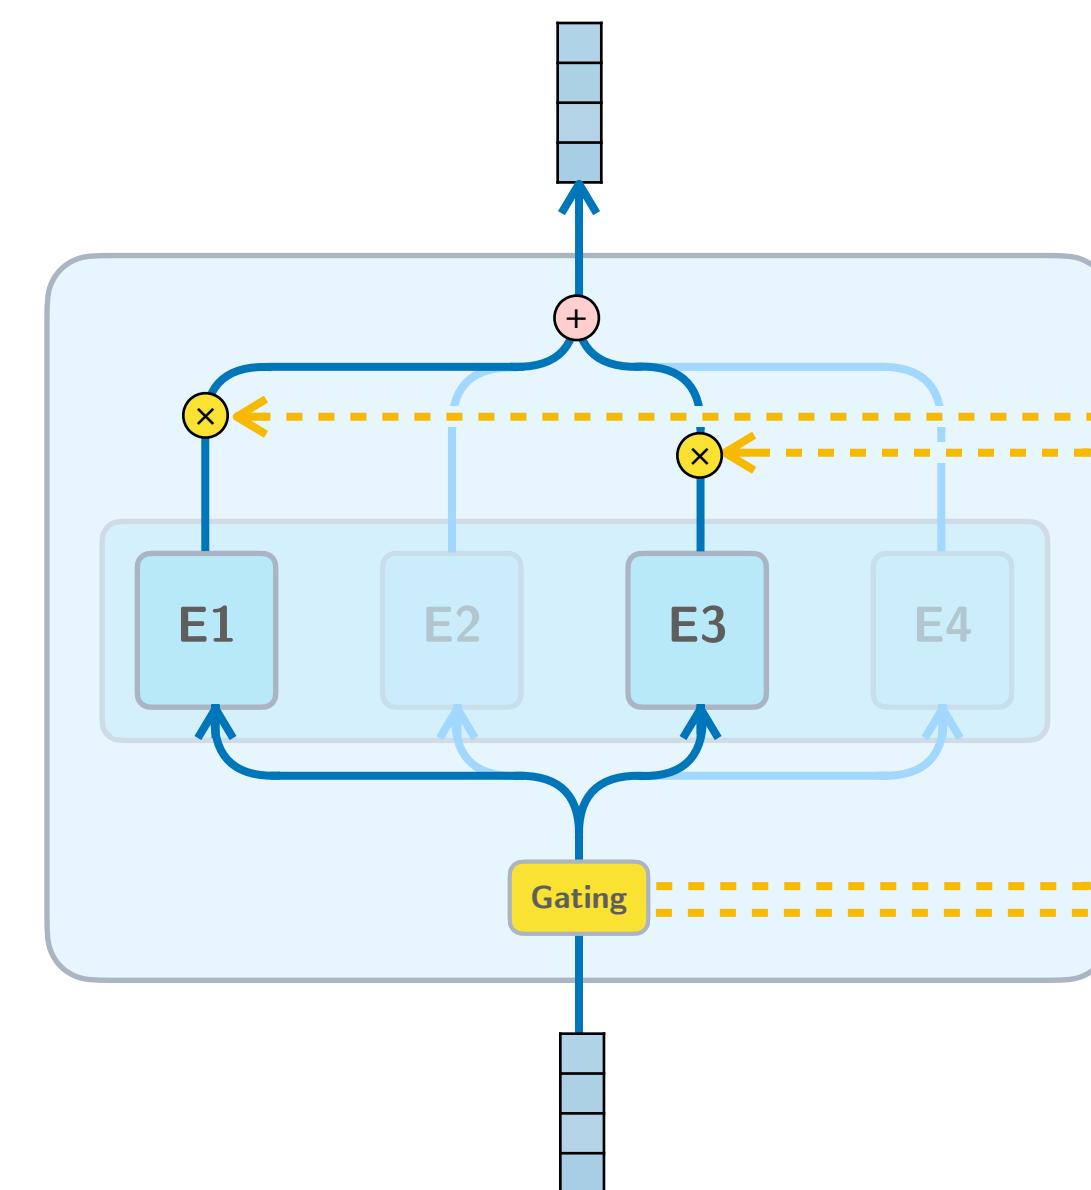
Multiple “expert” MLPs instead of a single MLP.





Multiple “expert” MLPs
instead of a single MLP.

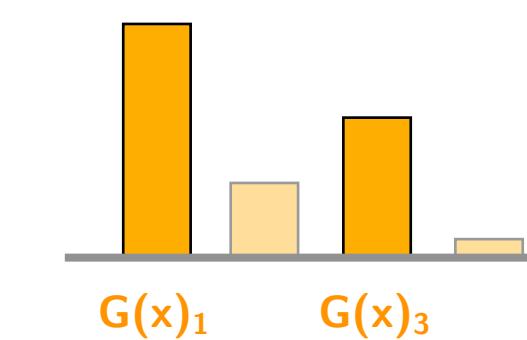




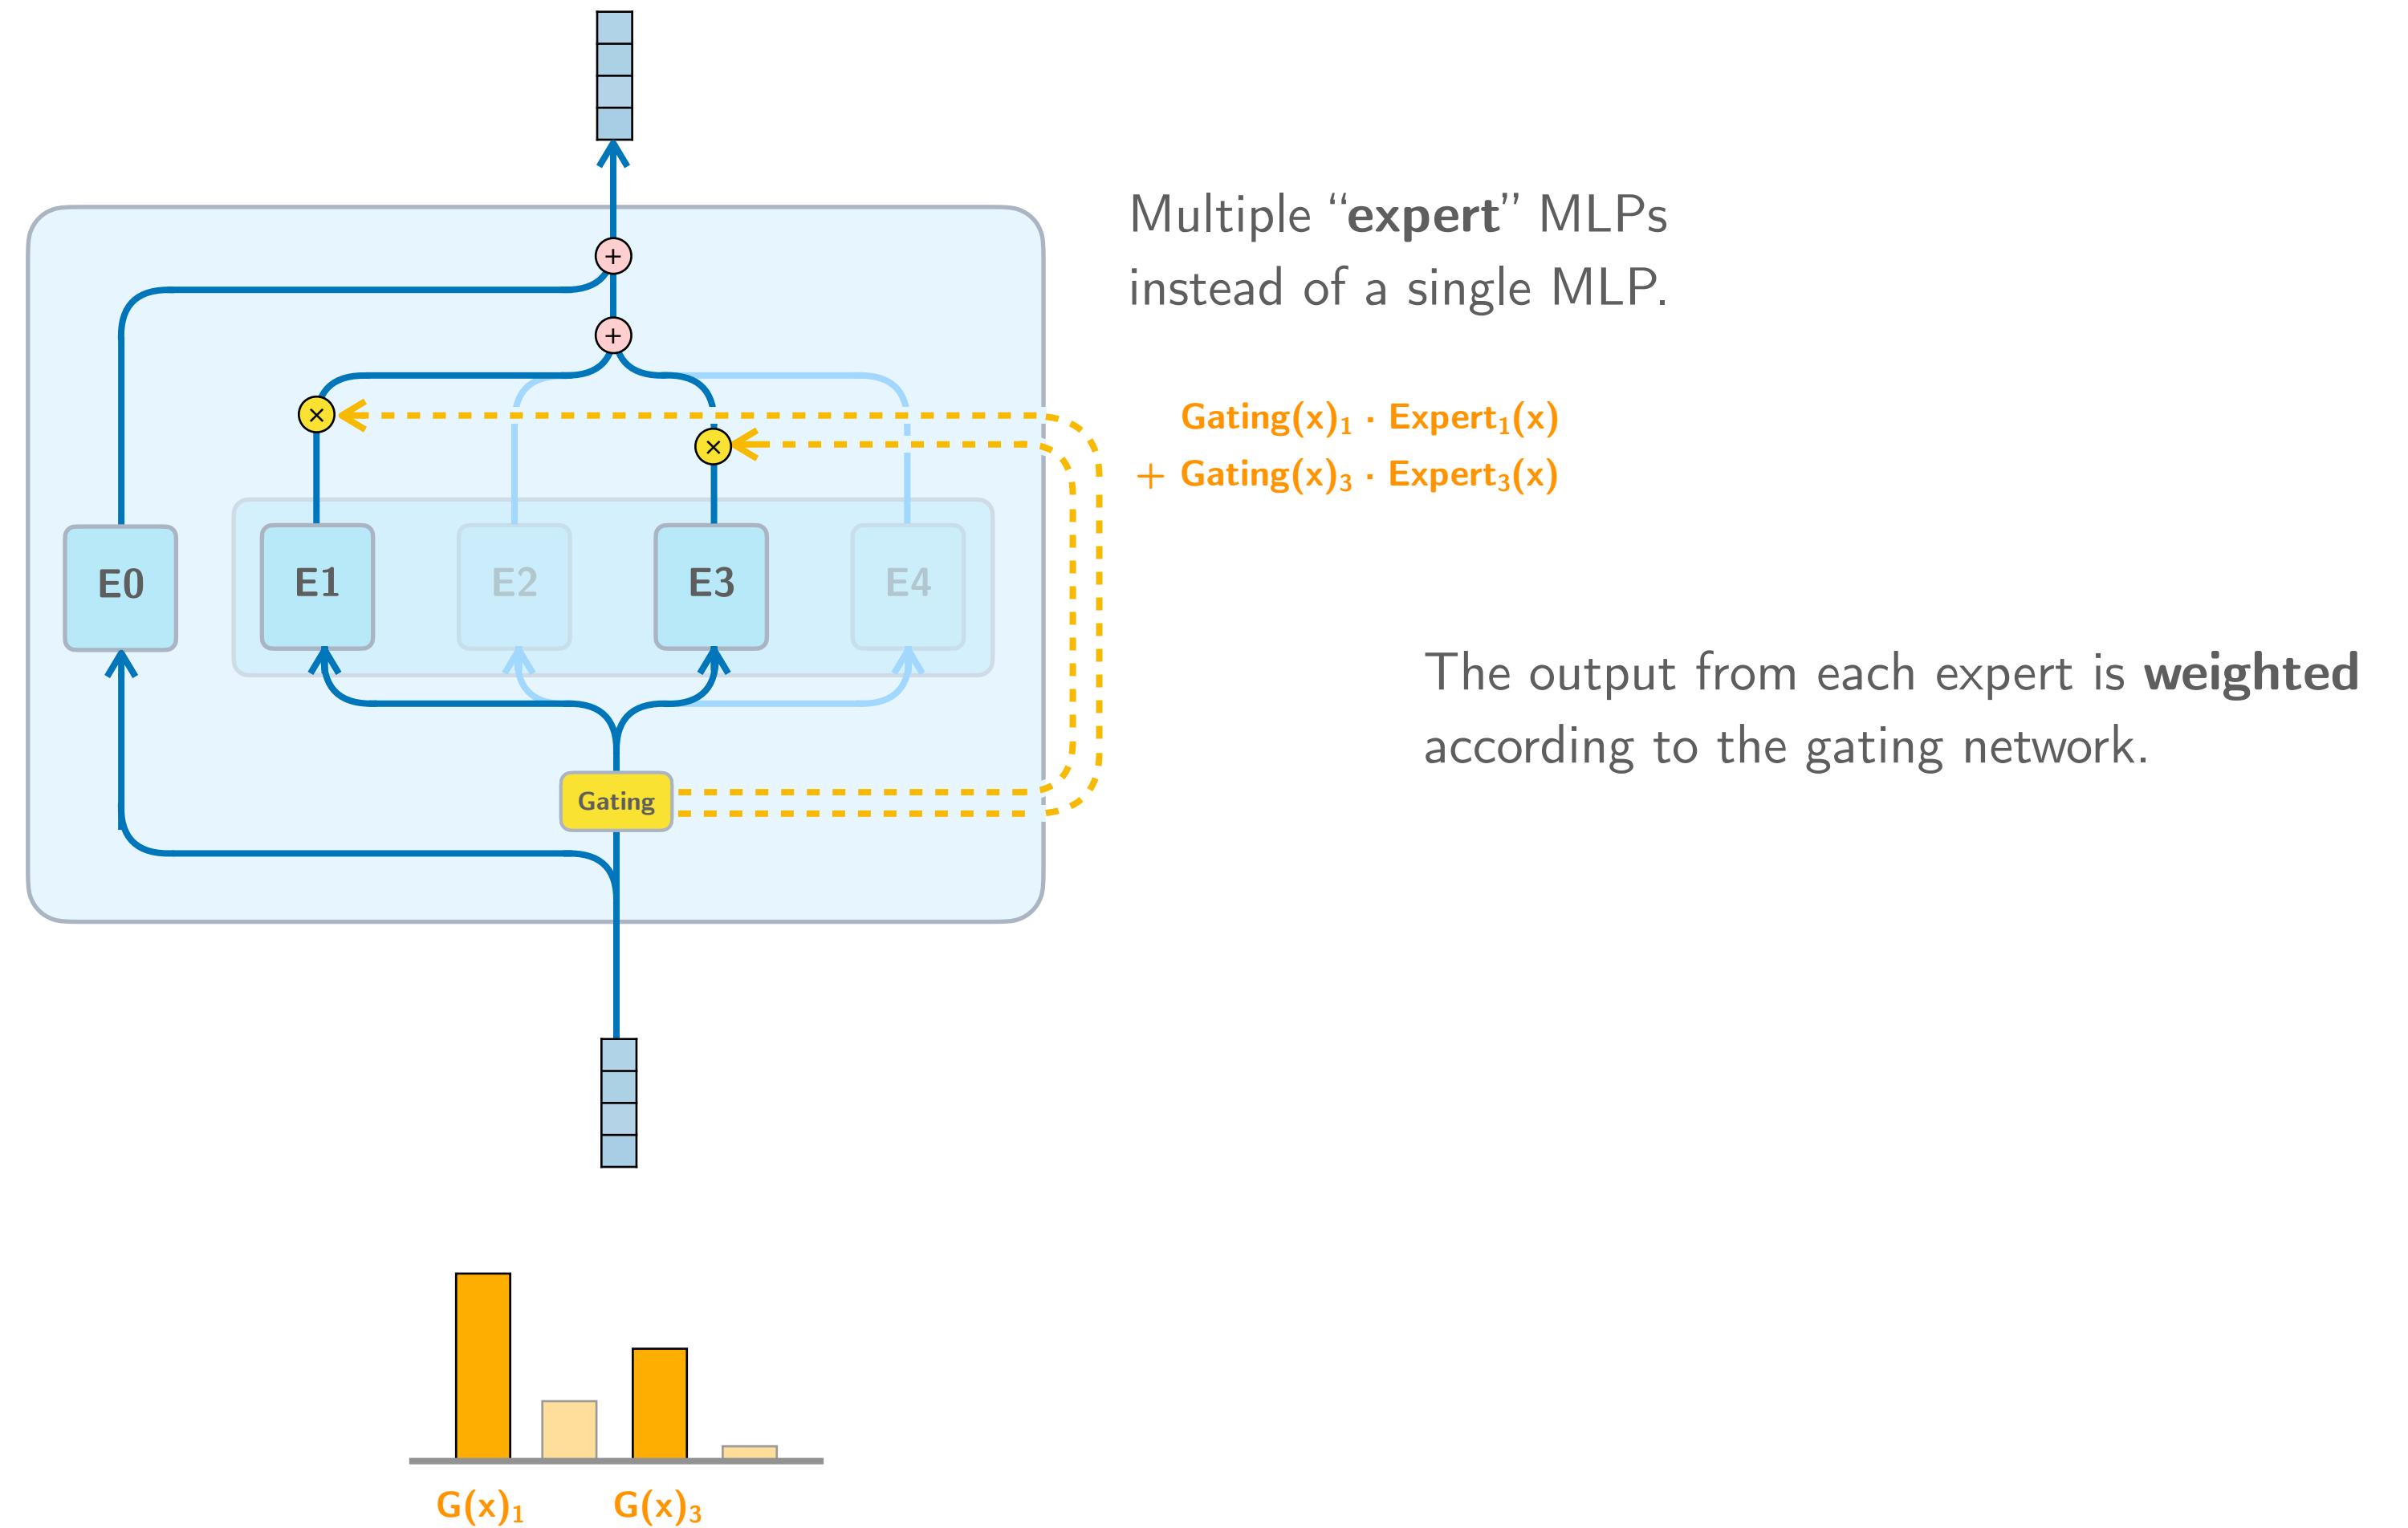
Multiple “expert” MLPs
instead of a single MLP.

$$\begin{aligned} & \text{Gating}(x)_1 \cdot \text{Expert}_1(x) \\ & + \text{Gating}(x)_3 \cdot \text{Expert}_3(x) \end{aligned}$$

The output from each expert is **weighted**
according to the gating network.



Shared experts and routed experts



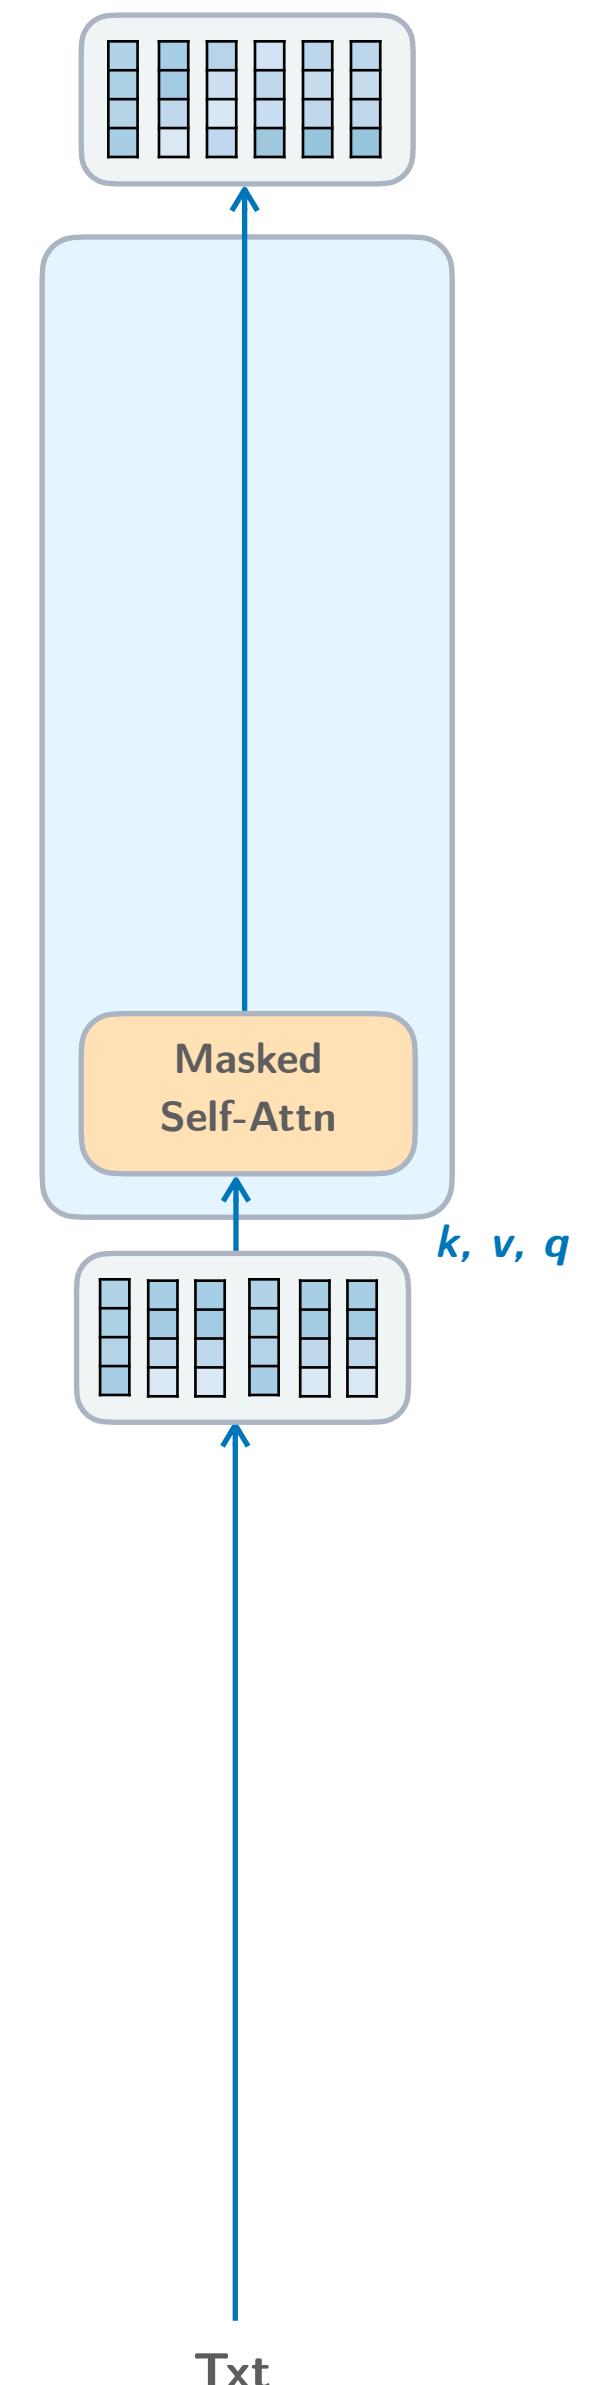


Retrieval and memory

RETRO = Retrieval-Enhanced Transformer

Model does not need to store all knowledge in its parameters.

Fewer parameters with same performance?

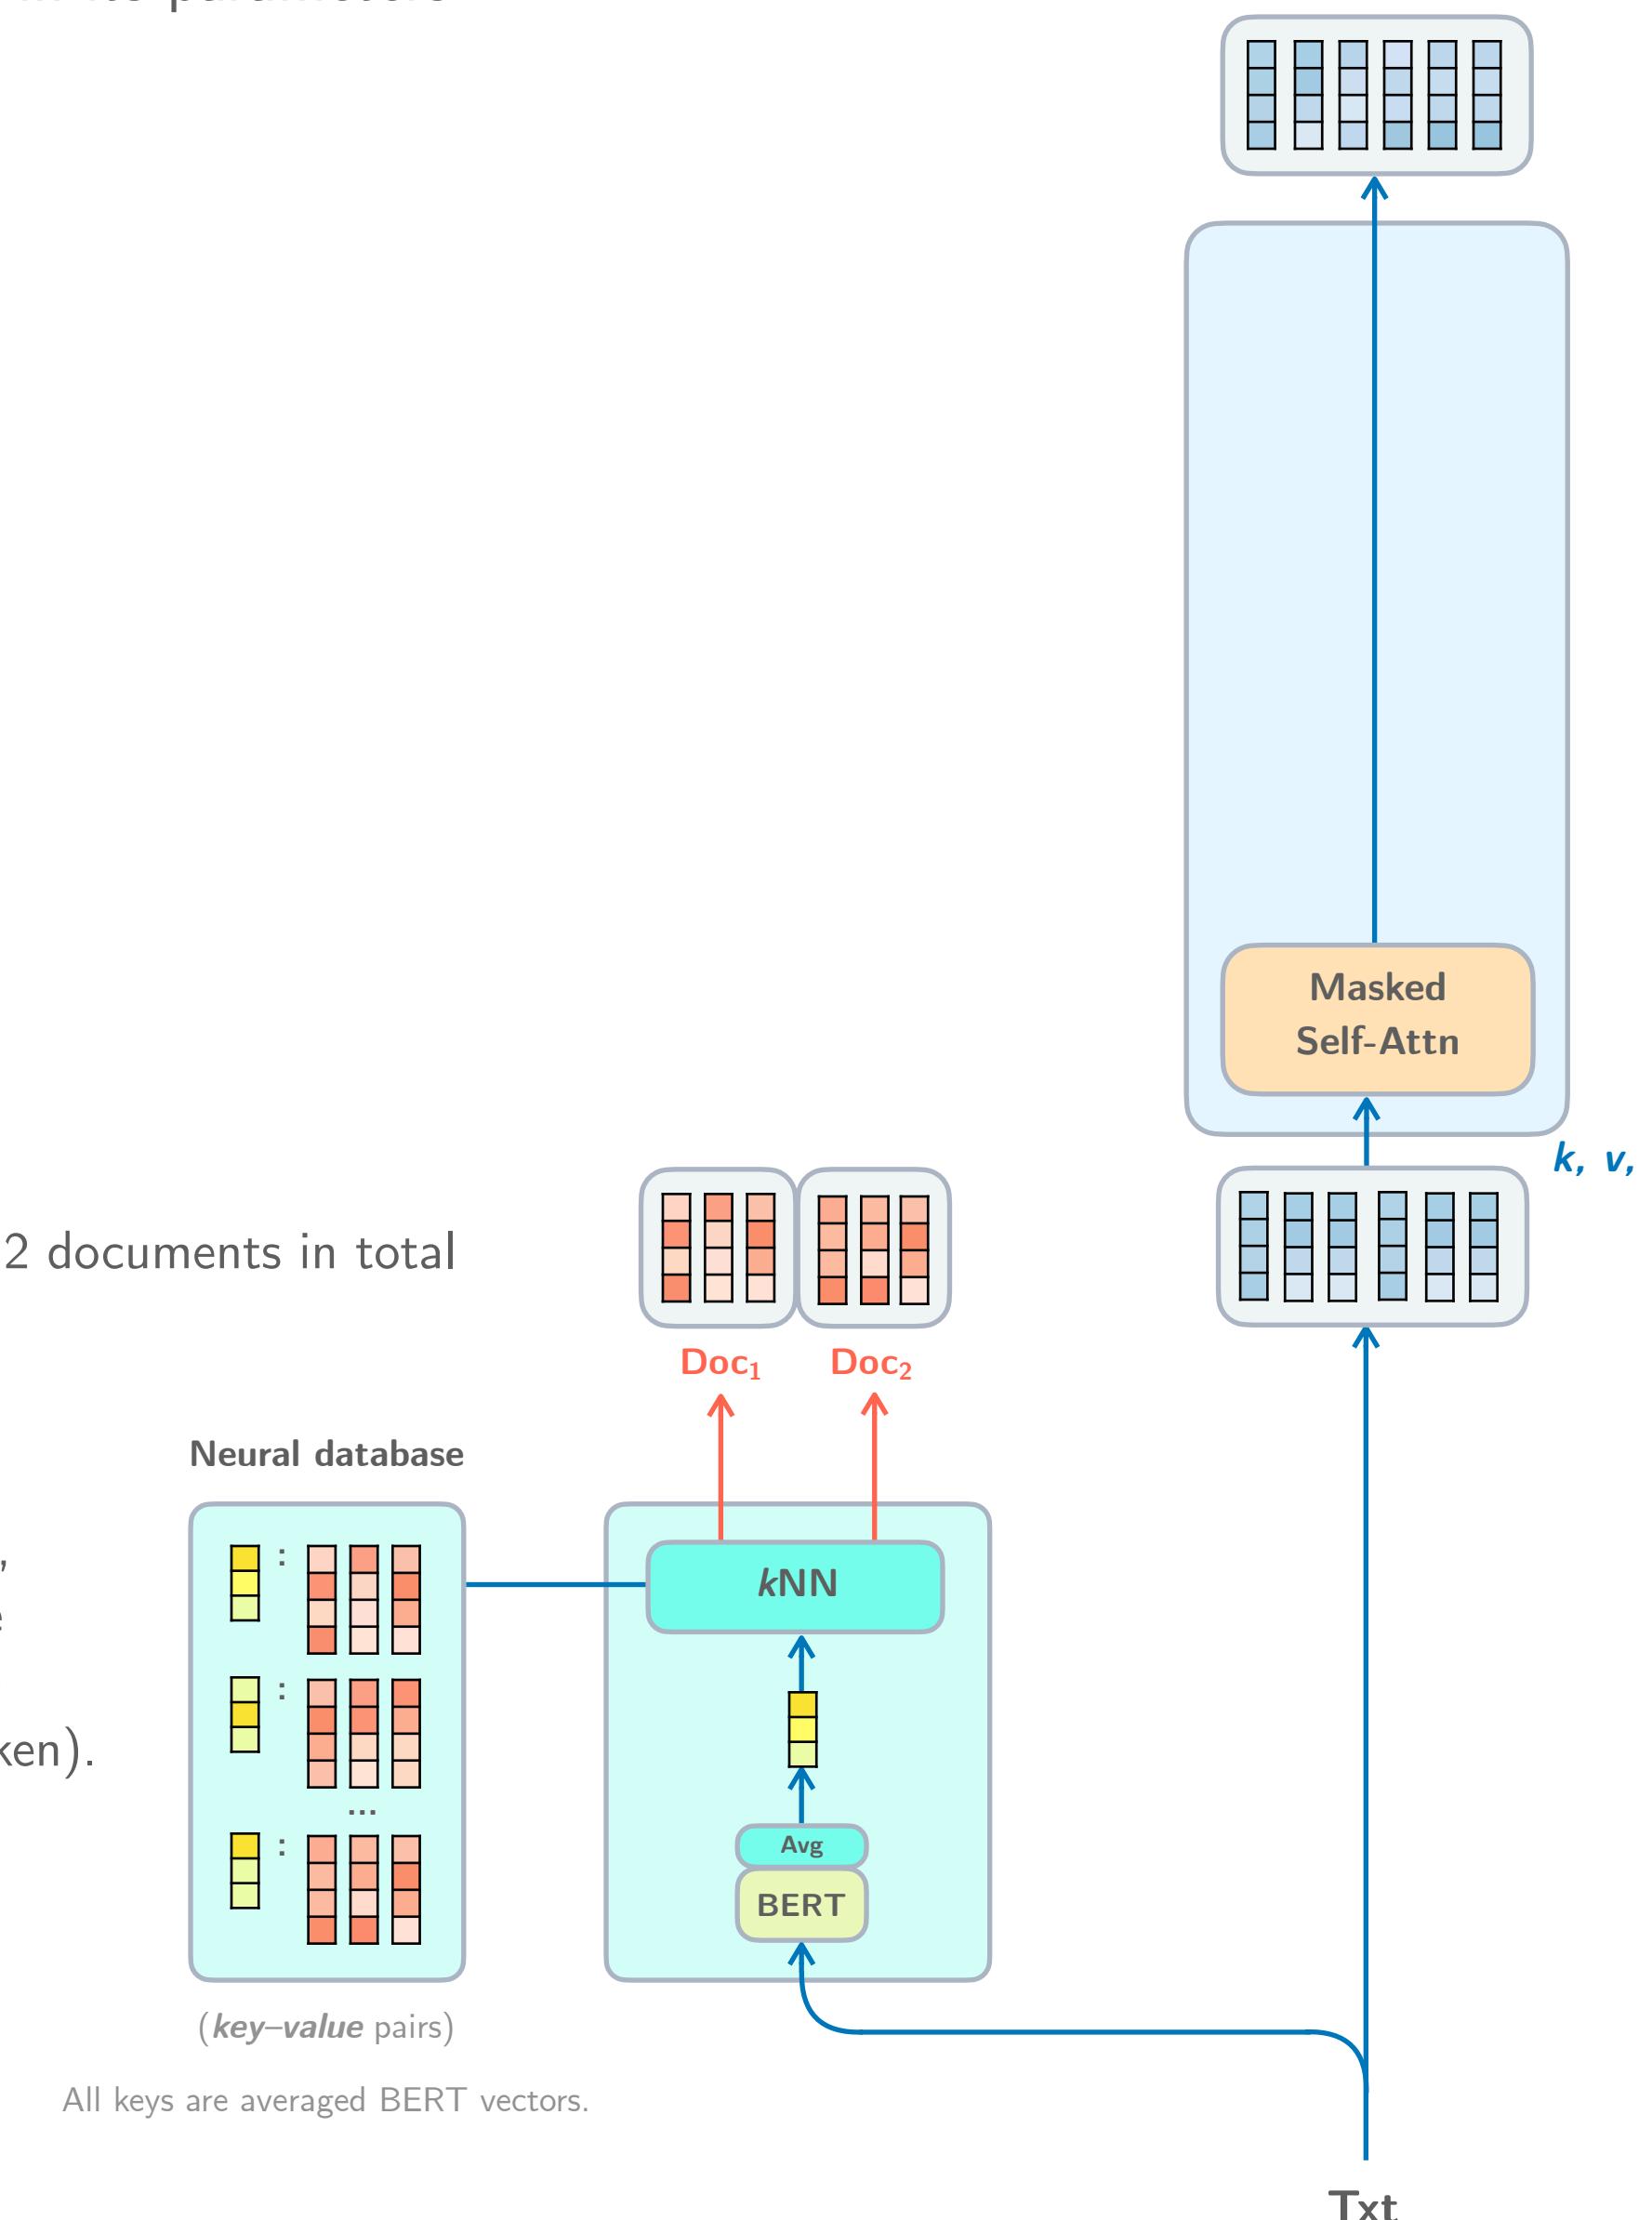


Leaving out Fwd for simplicity

RETRO = Retrieval-Enhanced Transformer

Model does not need to store all knowledge in its parameters.

Fewer parameters with same performance?



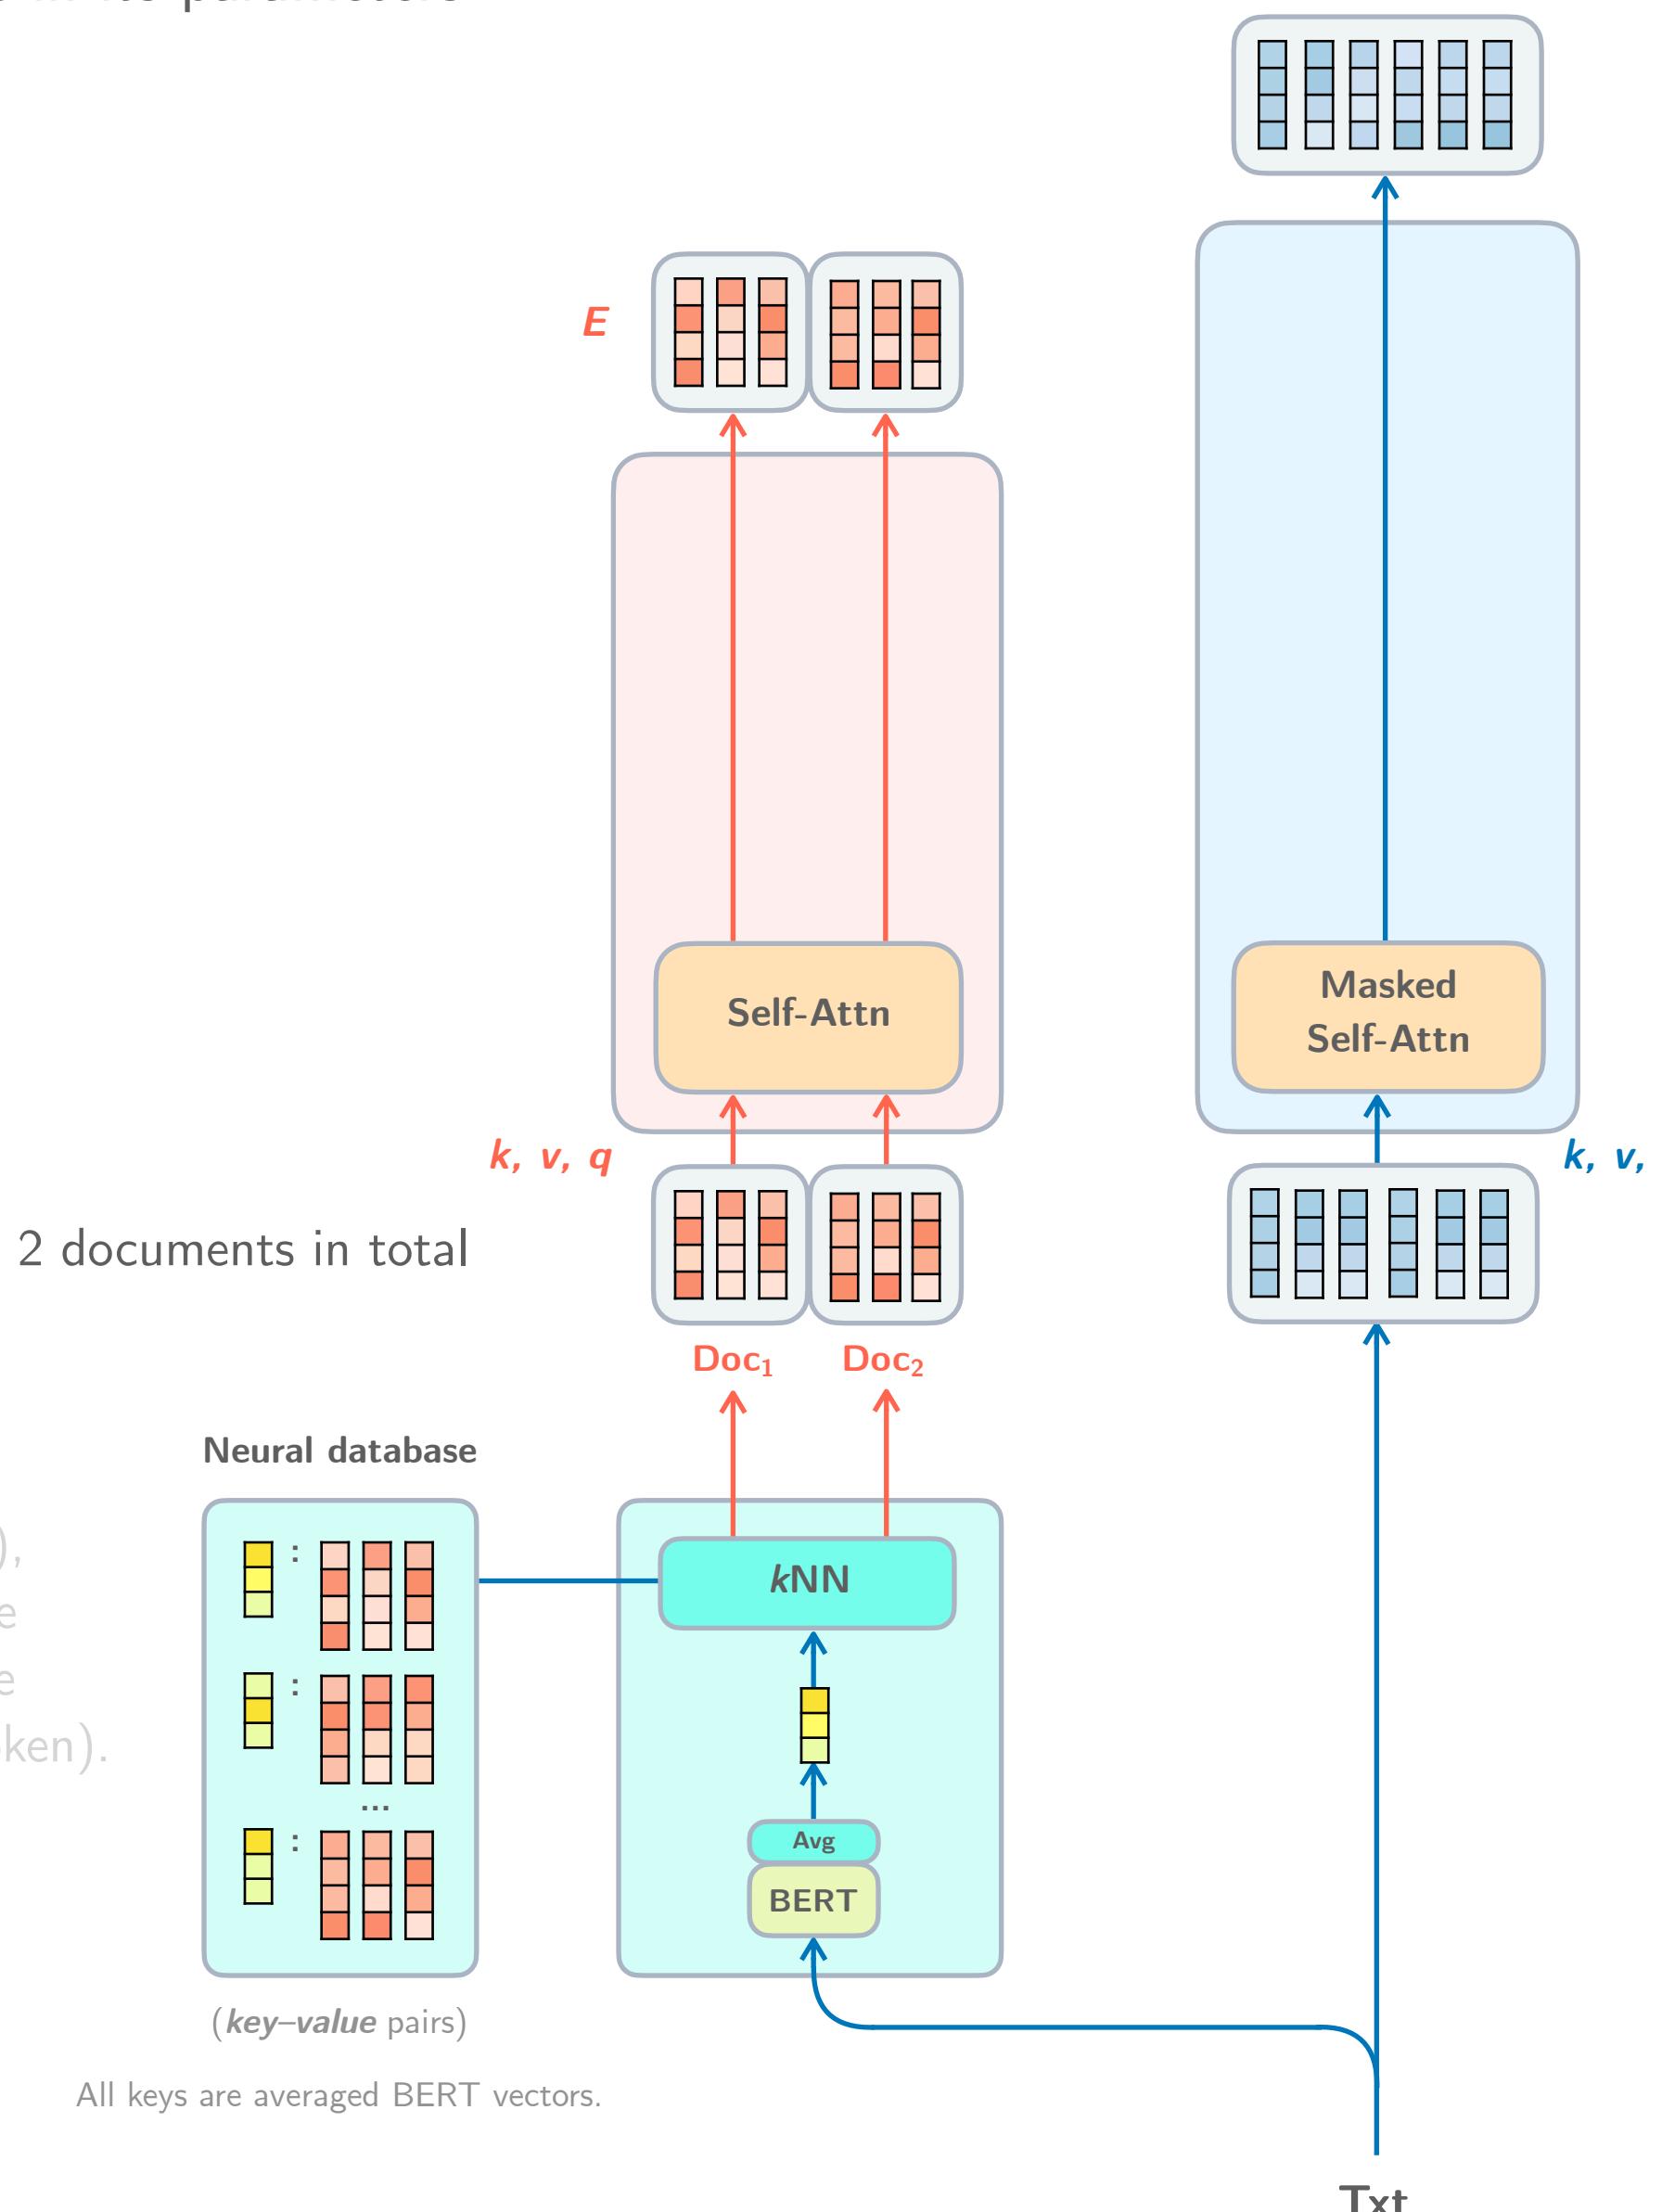
Retrieval itself is not learned (it's not differentiable), but during training, the model can expect to receive as input useful documents and it learns to make use of them (by attending over it to predict the next token).

Leaving out Fwd for simplicity

RETRO = Retrieval-Enhanced Transformer

Model does not need to store all knowledge in its parameters.

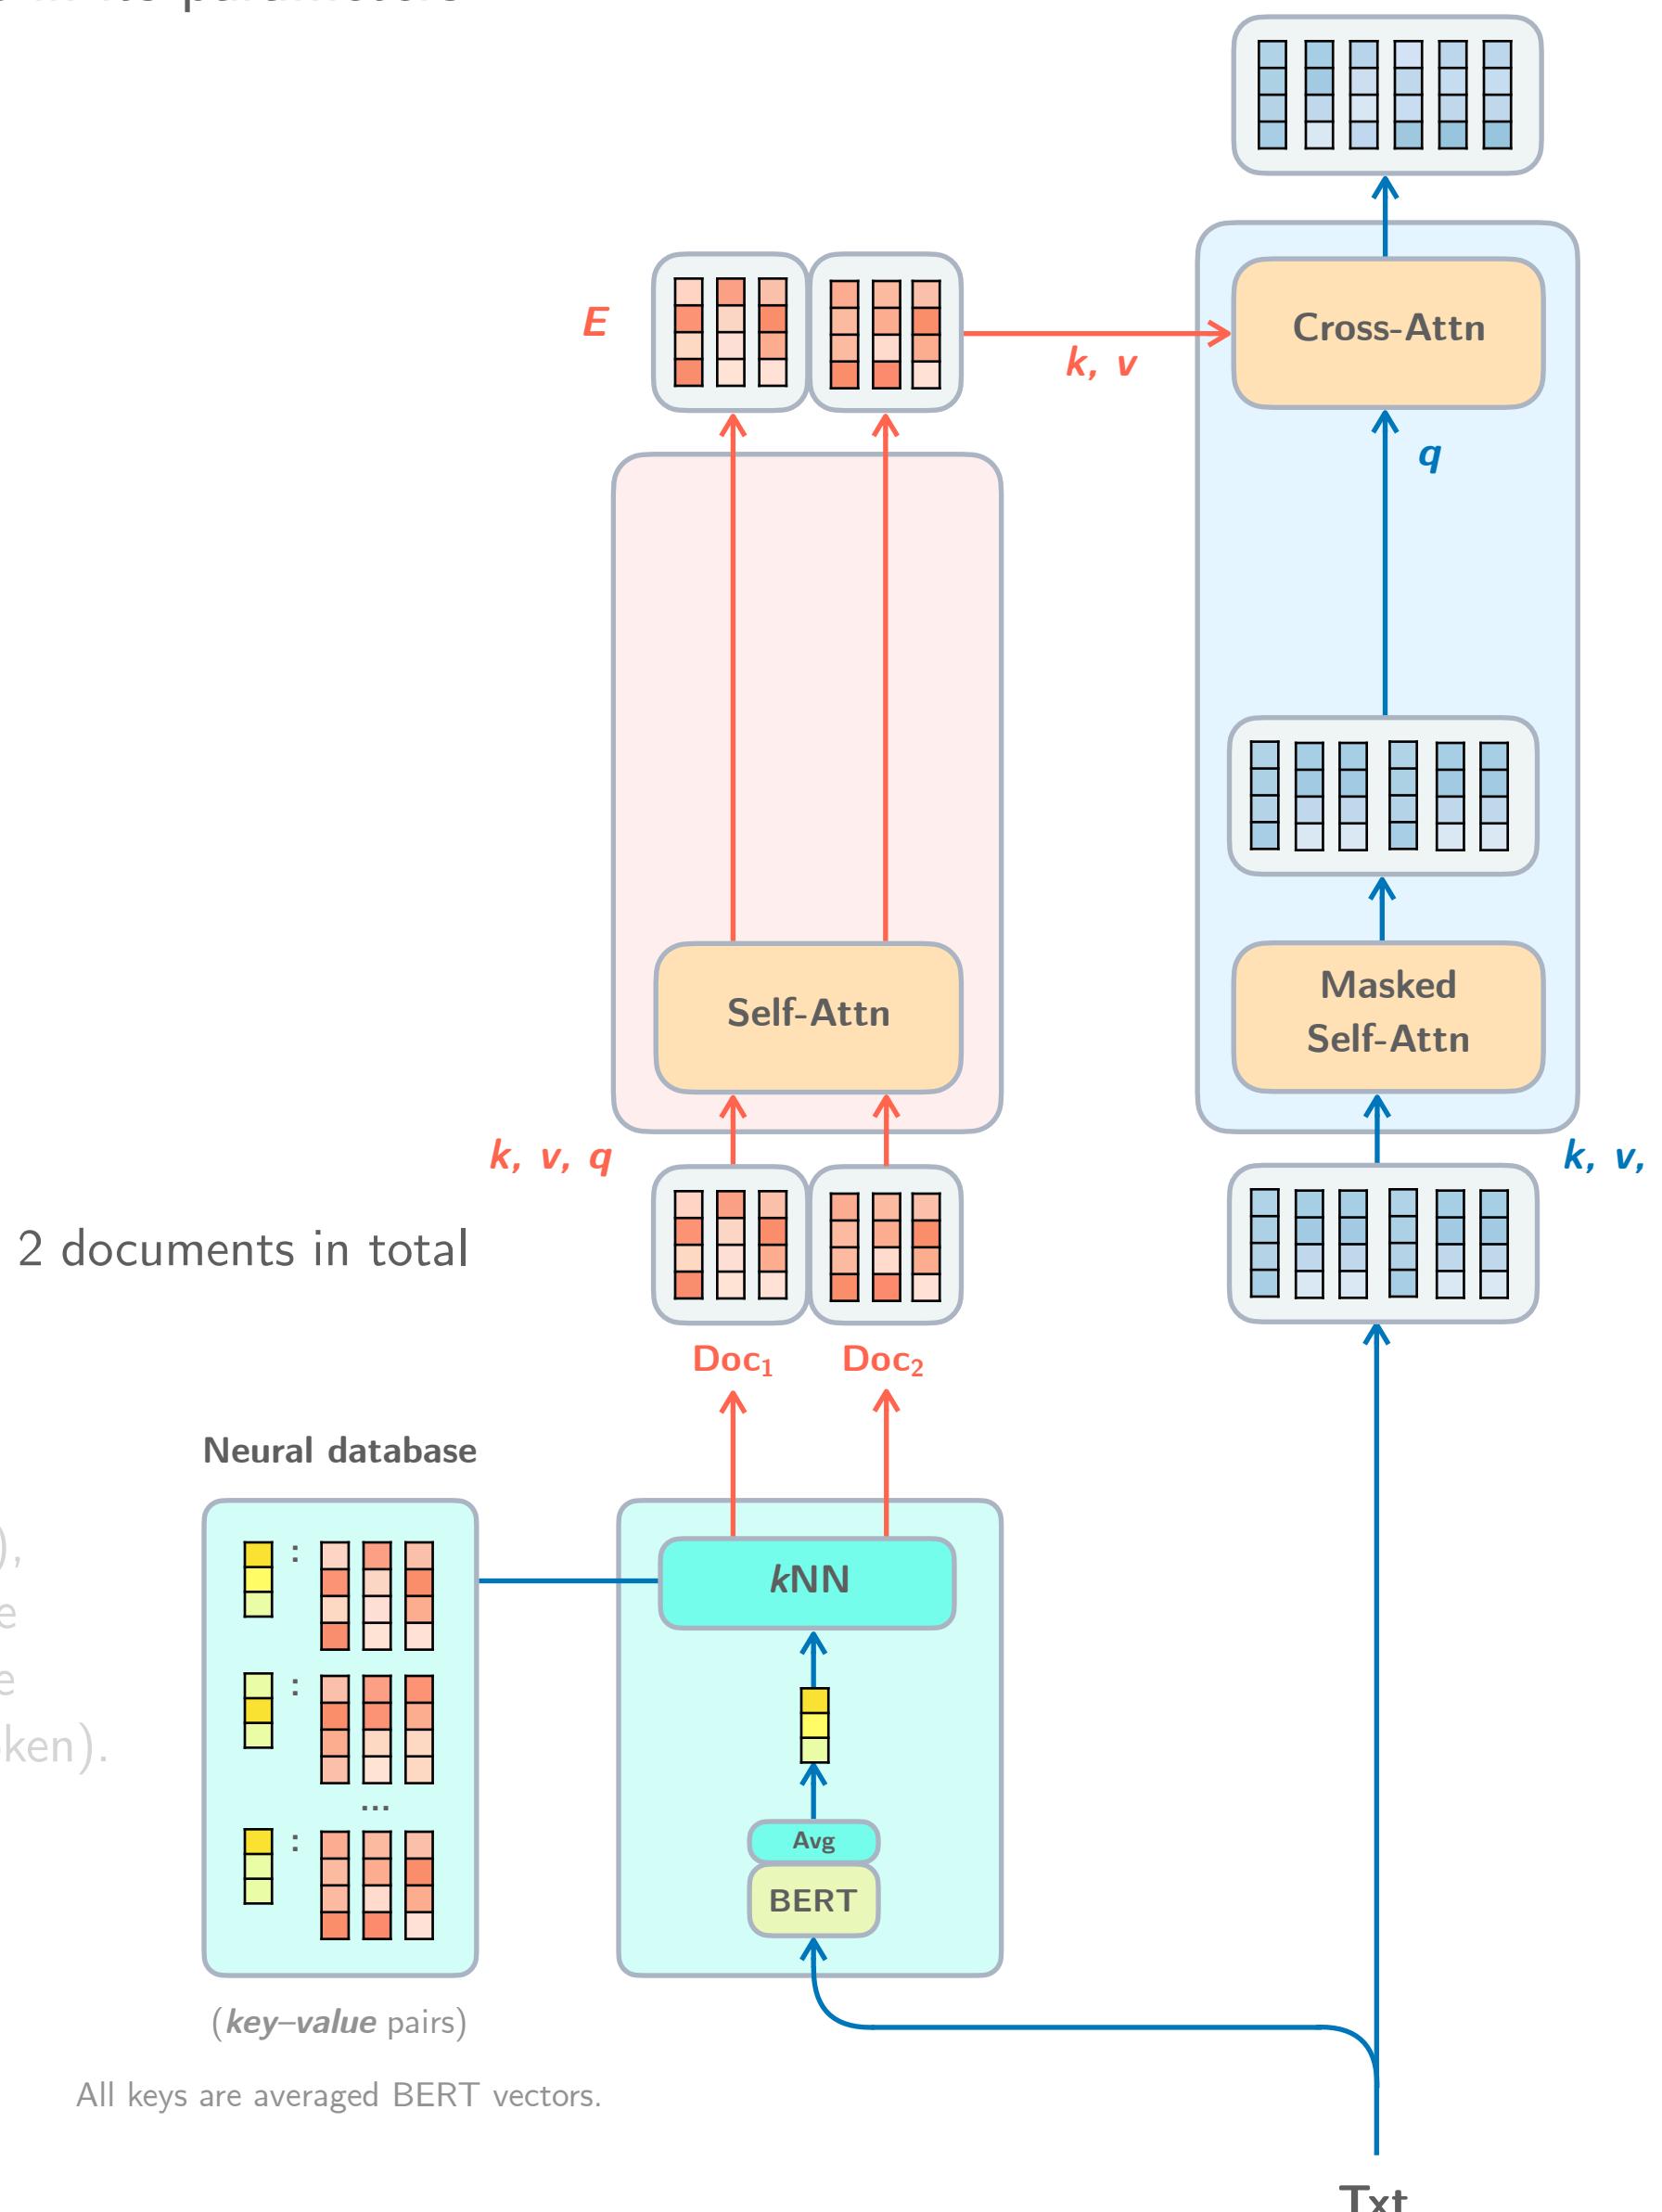
Fewer parameters with same performance?



RETRO = Retrieval-Enhanced Transformer

Model does not need to store all knowledge in its parameters.

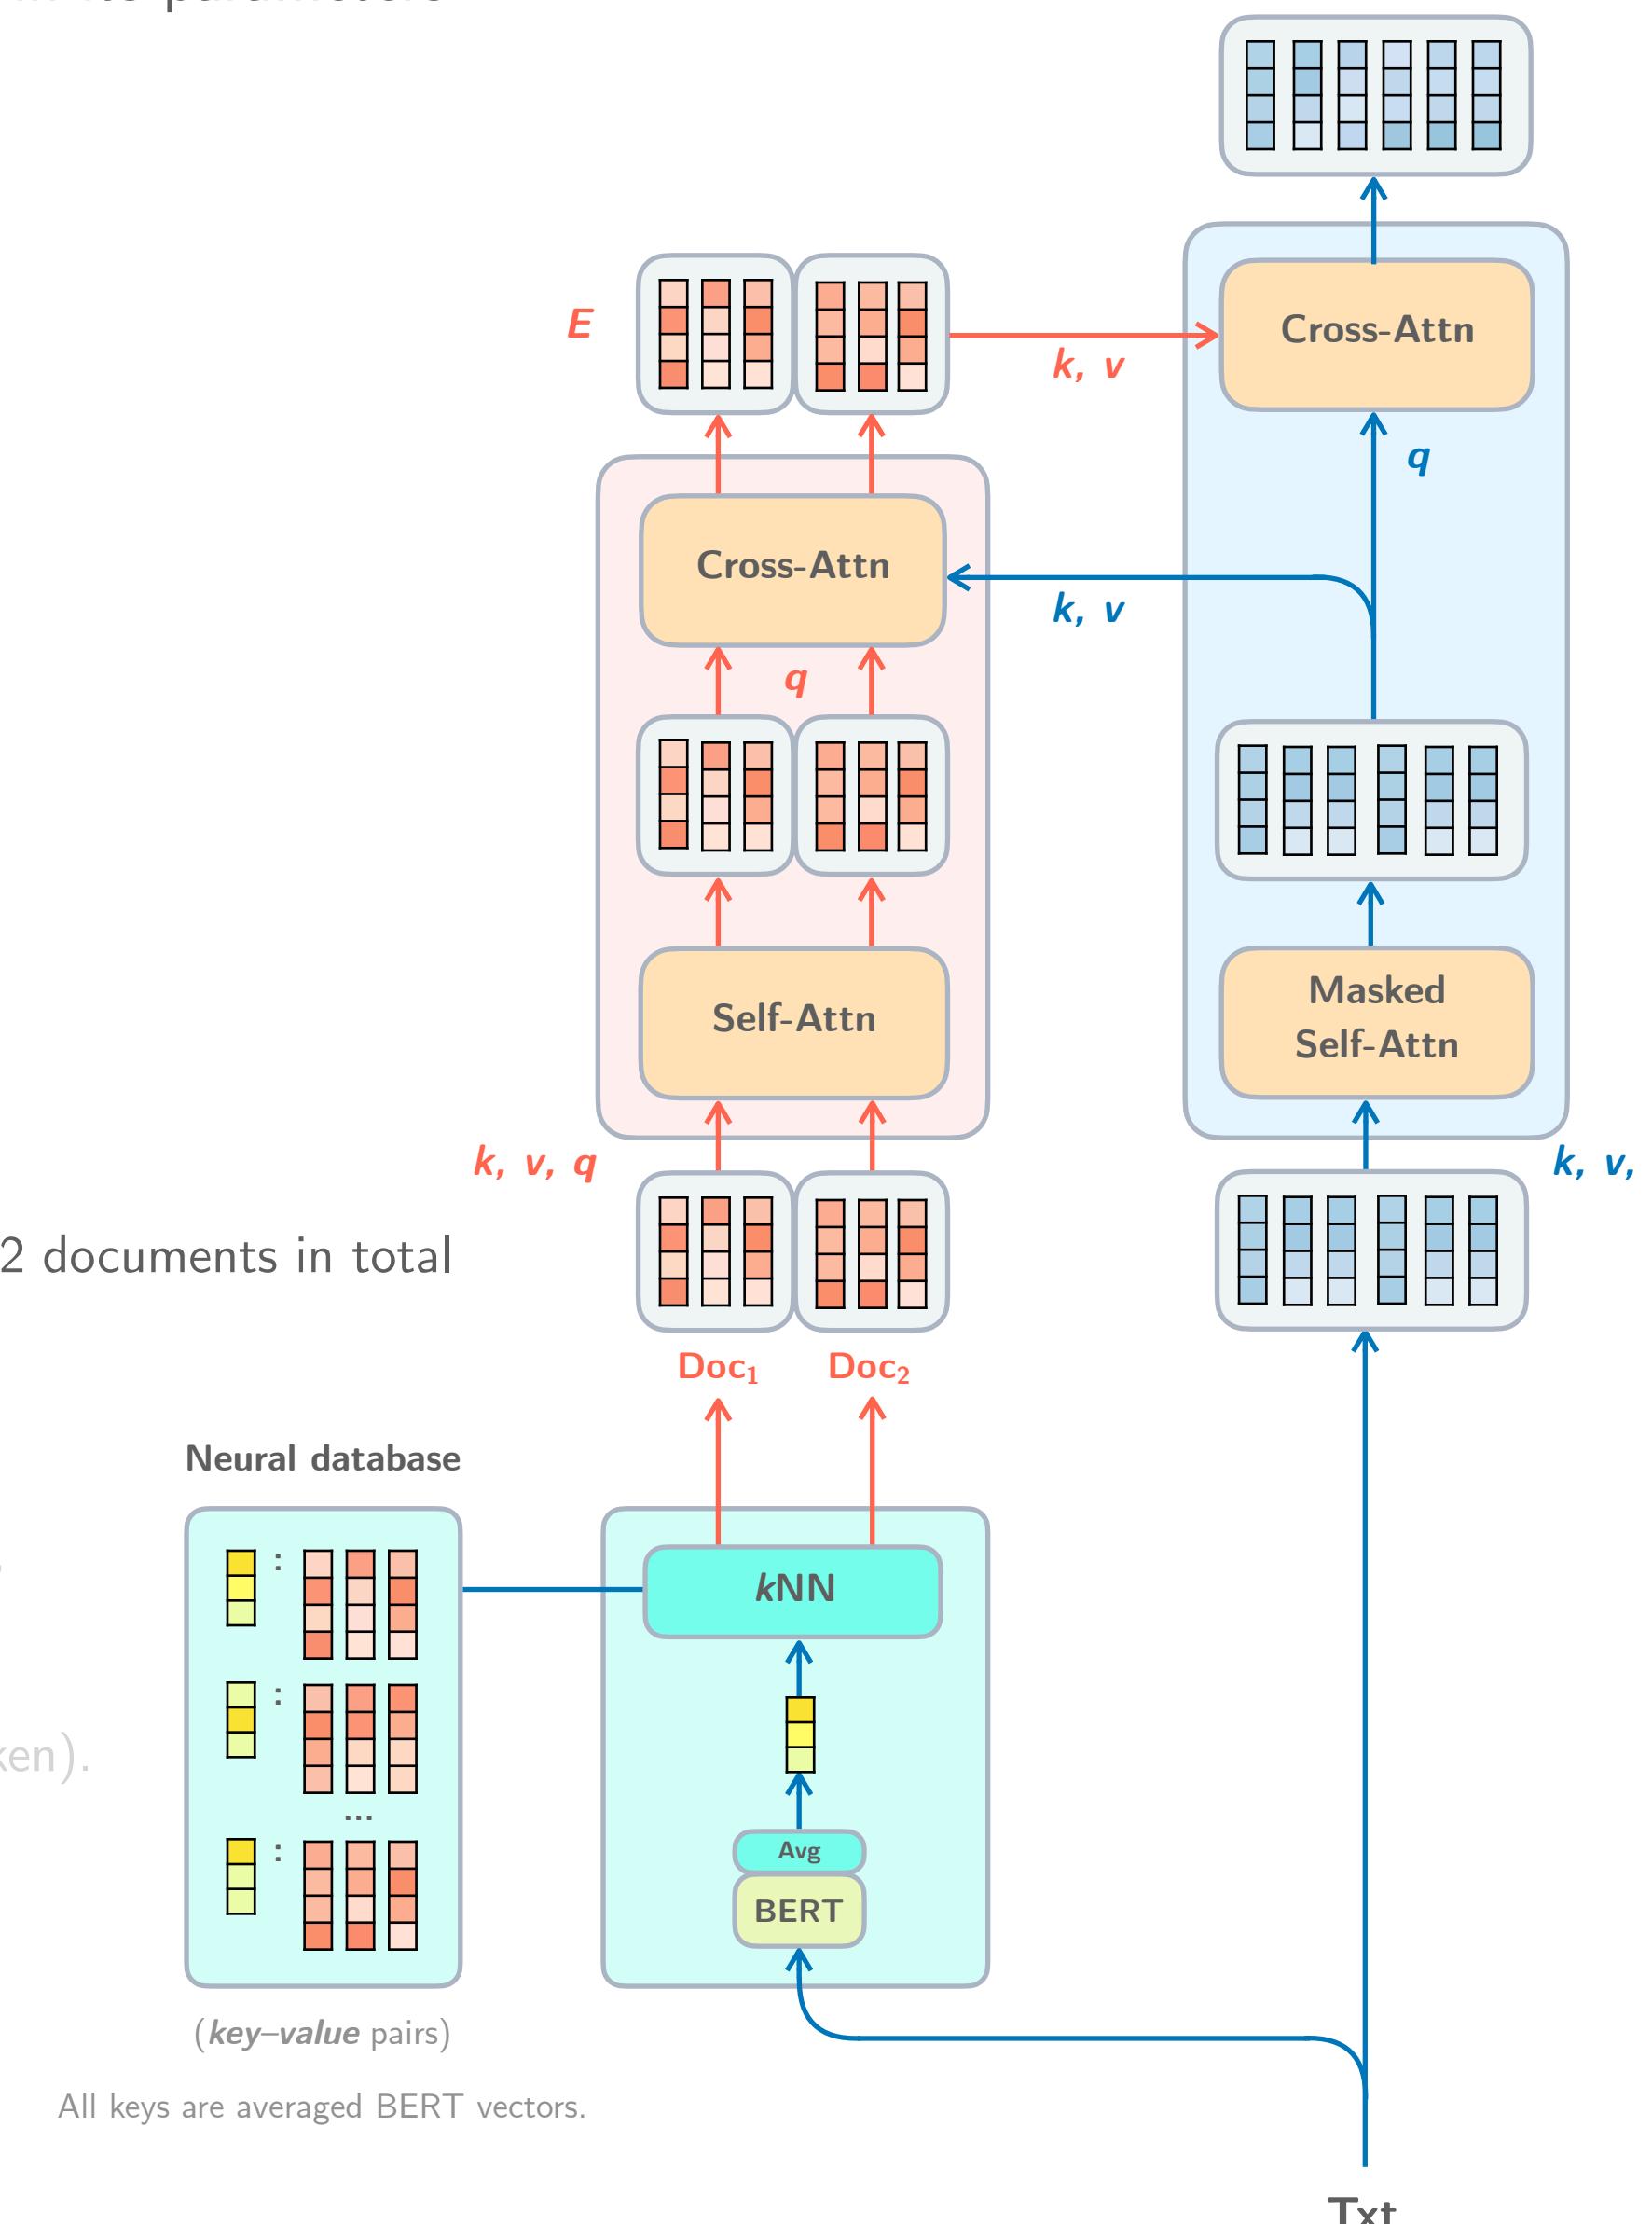
Fewer parameters with same performance?



RETRO = Retrieval-Enhanced Transformer

Model does not need to store all knowledge in its parameters.

Fewer parameters with same performance?



RETRO = Retrieval-Enhanced Transformer

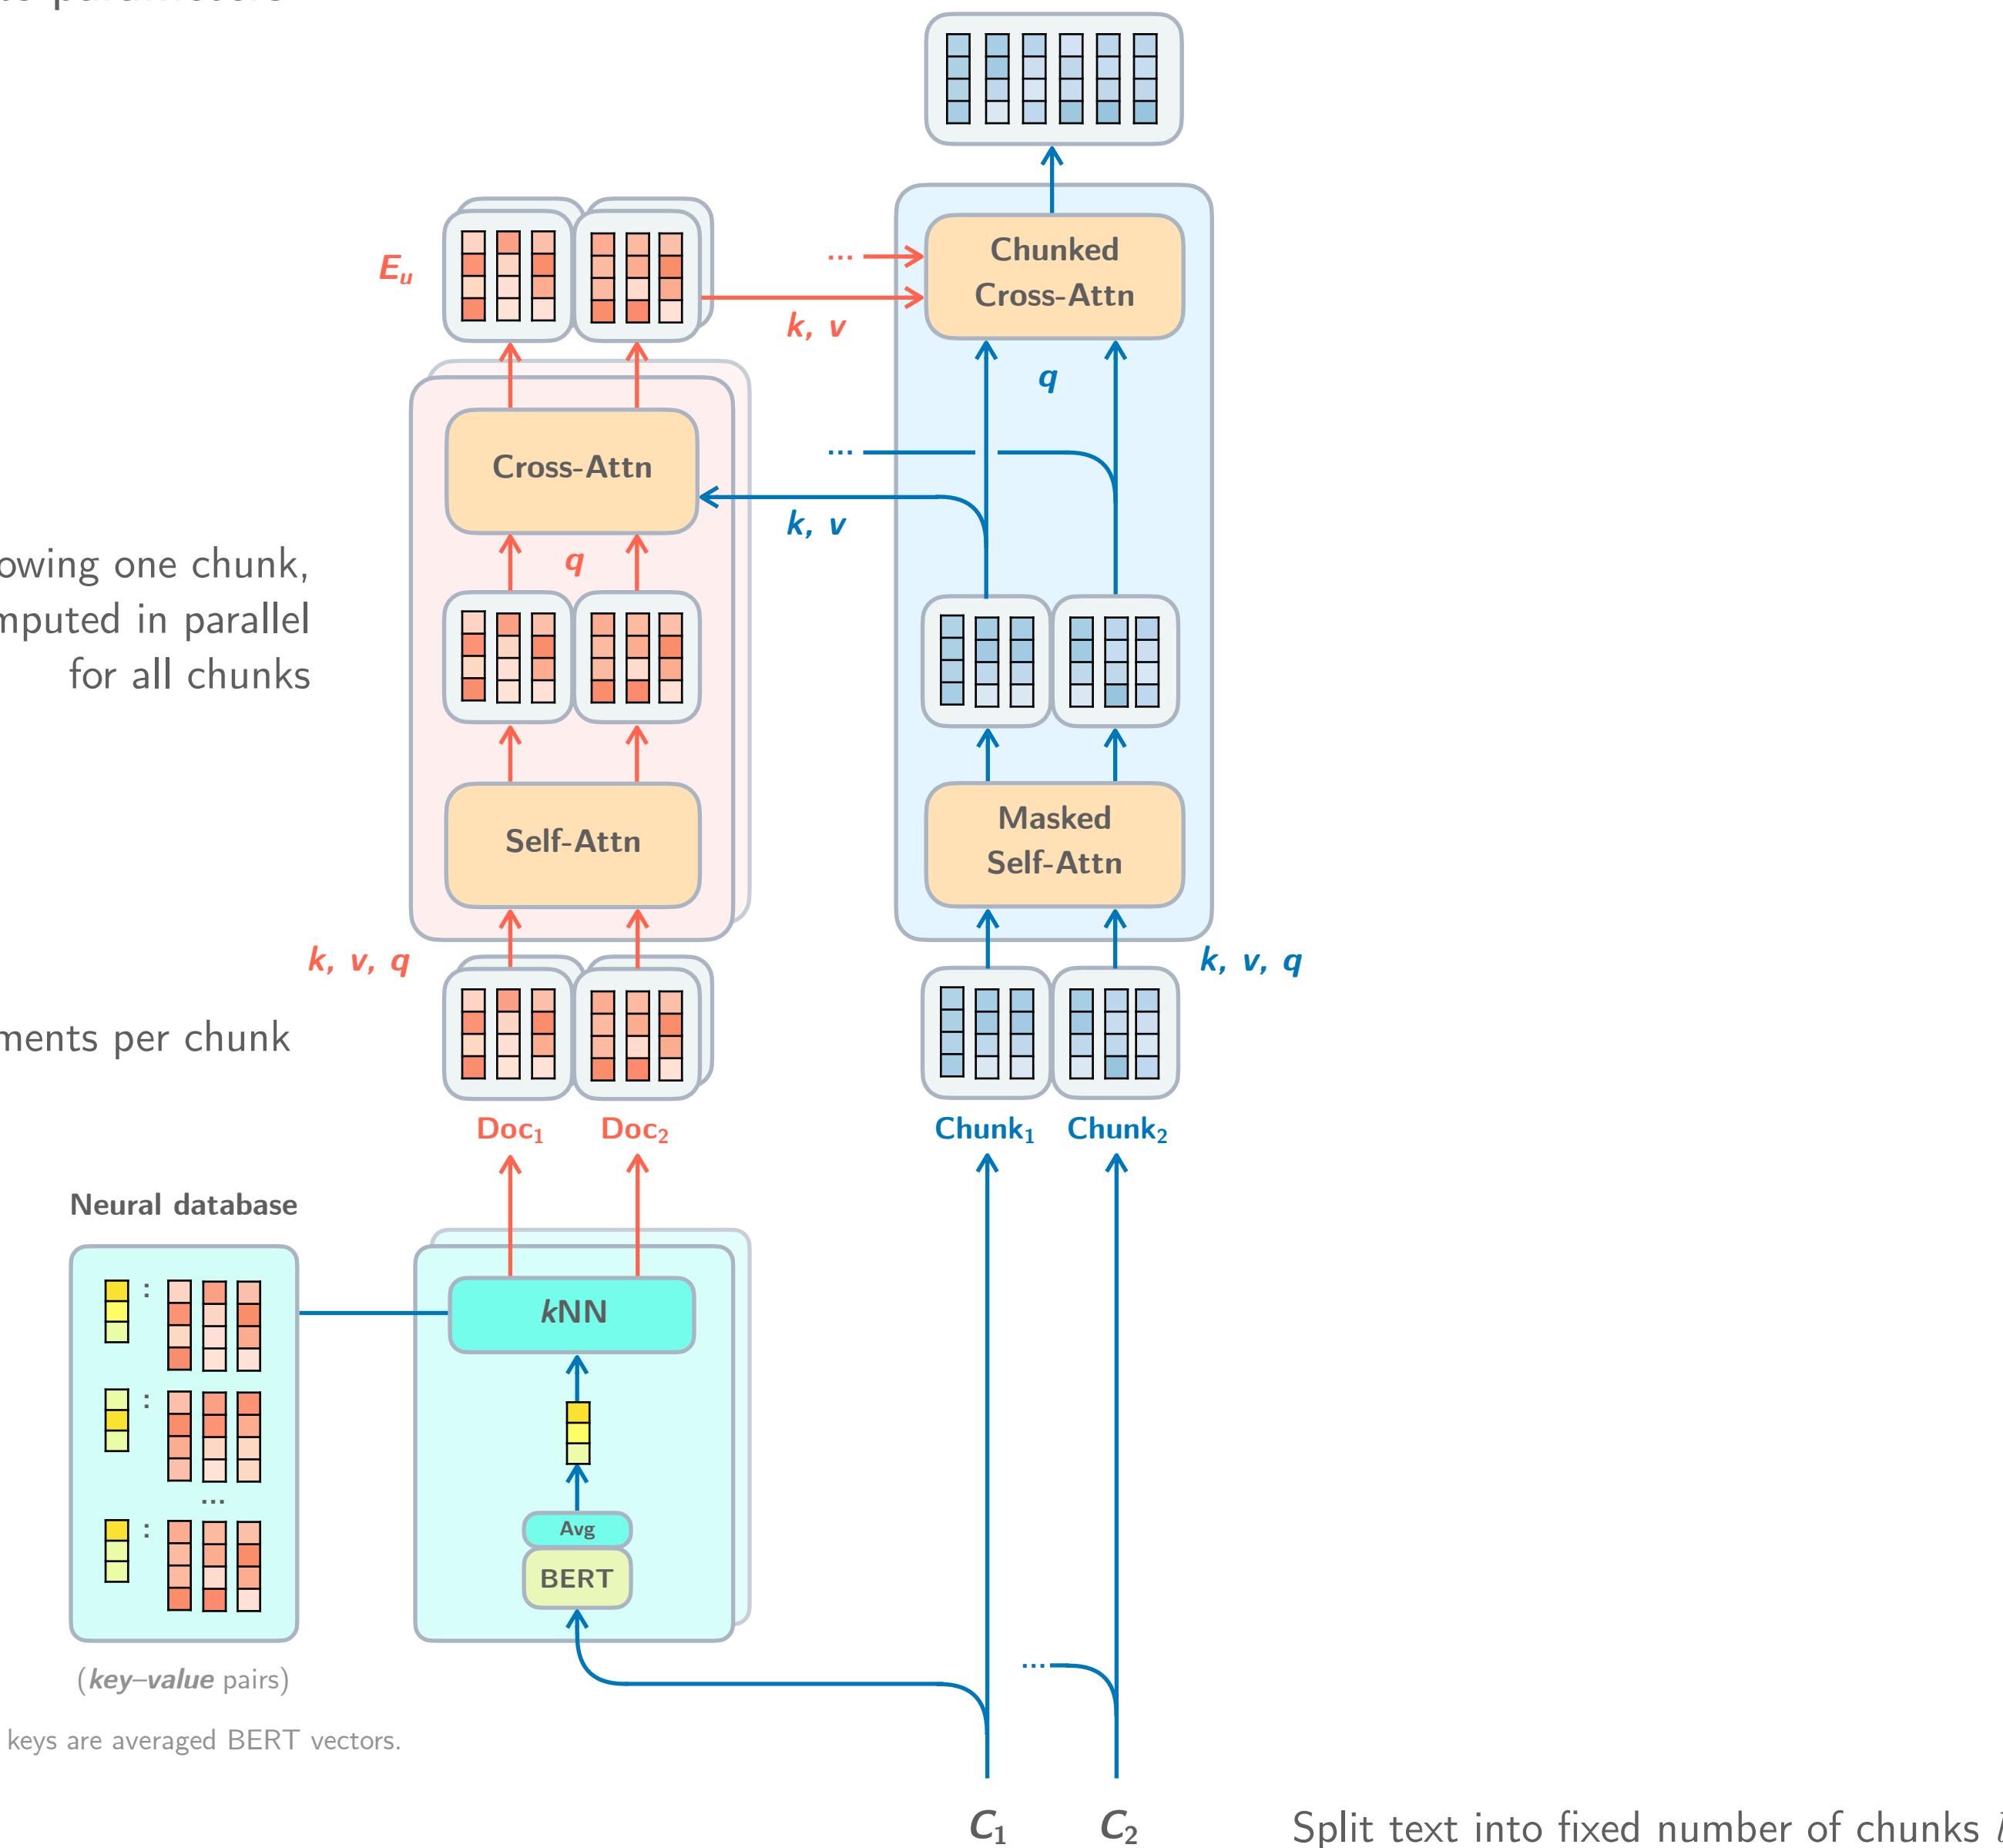
Model does not need to store all knowledge in its parameters.

Fewer parameters with same performance?

Only showing one chunk,
but this is computed in parallel
for all chunks

Here, 2 documents per chunk

Retrieval itself is not learned (it's not differentiable),
but during training, the model can expect to receive
as input useful documents and it learns to make use
of them (by attending over it to predict the next token).



Leaving out Fwd for simplicity

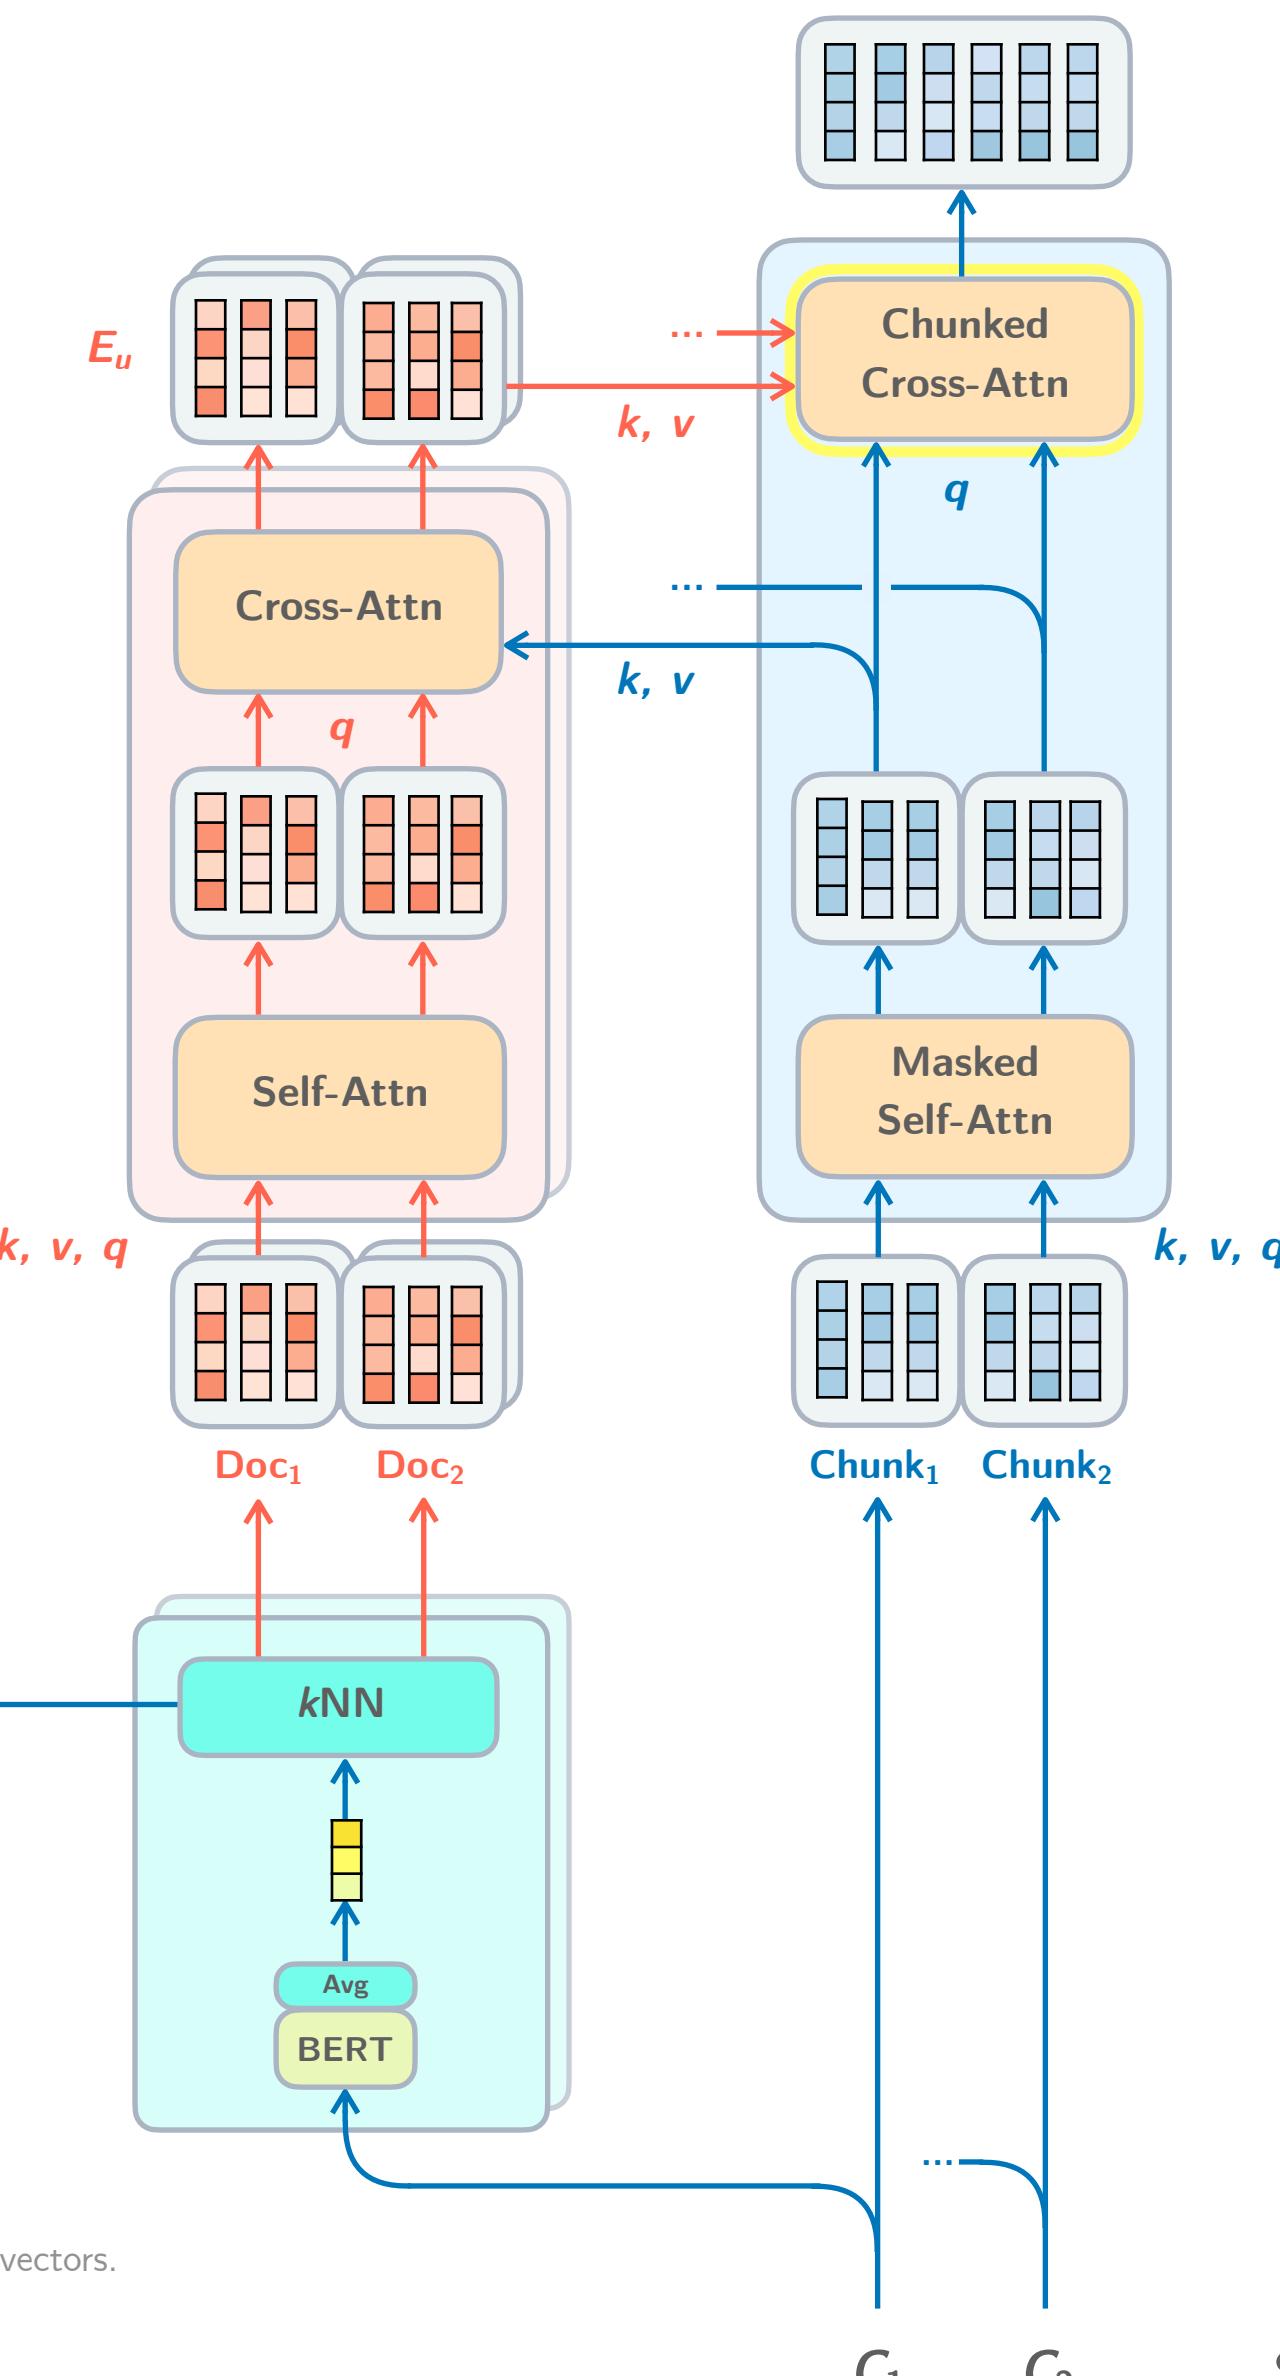
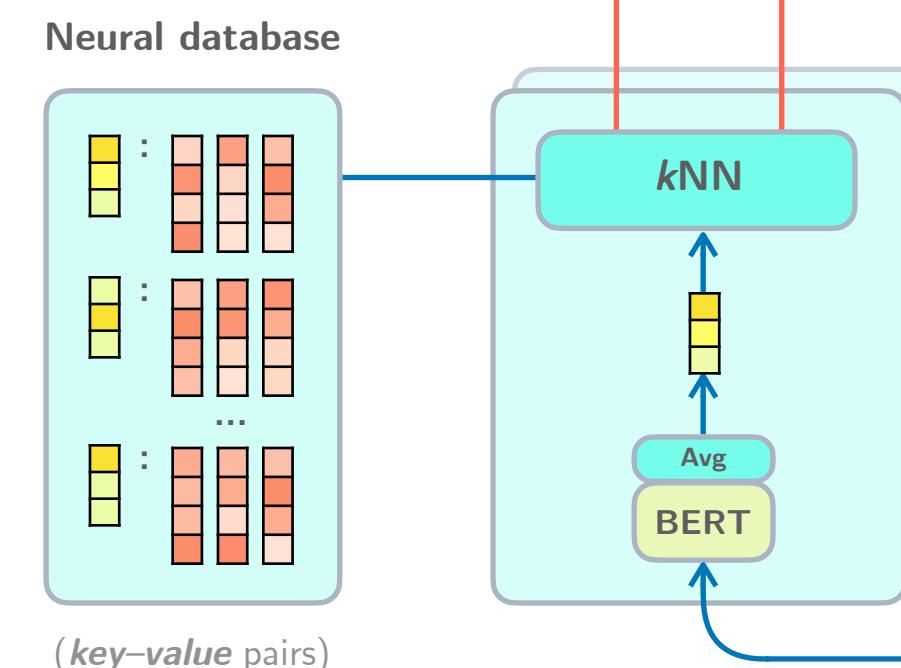
RETRO = Retrieval-Enhanced Transformer

Model does not need to store all knowledge in its parameters.

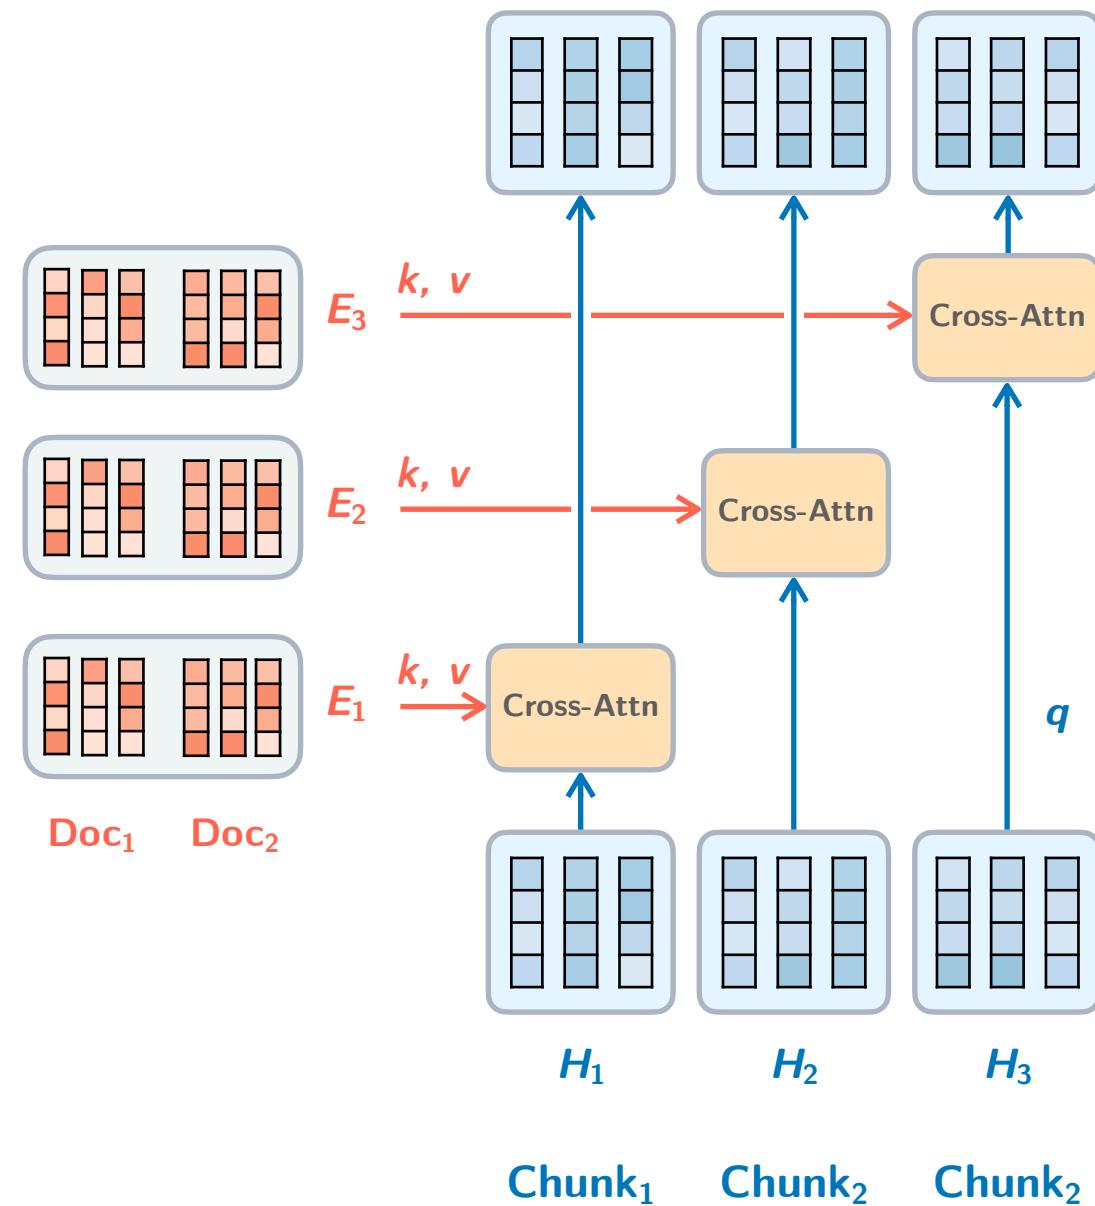
Fewer parameters with same performance?

Only showing one chunk,
but this is computed in parallel
for all chunks

Here, 2 documents per chunk



Chunked cross-attention (simplified)



E_1 are the encoded documents corresponding to chunk 1 (C_1 , and representation H_1)

Retrieval itself is not learned (it's not differentiable),
but during training, the model can expect to receive
as input useful documents and it learns to make use
of them (by attending over it to predict the next token).

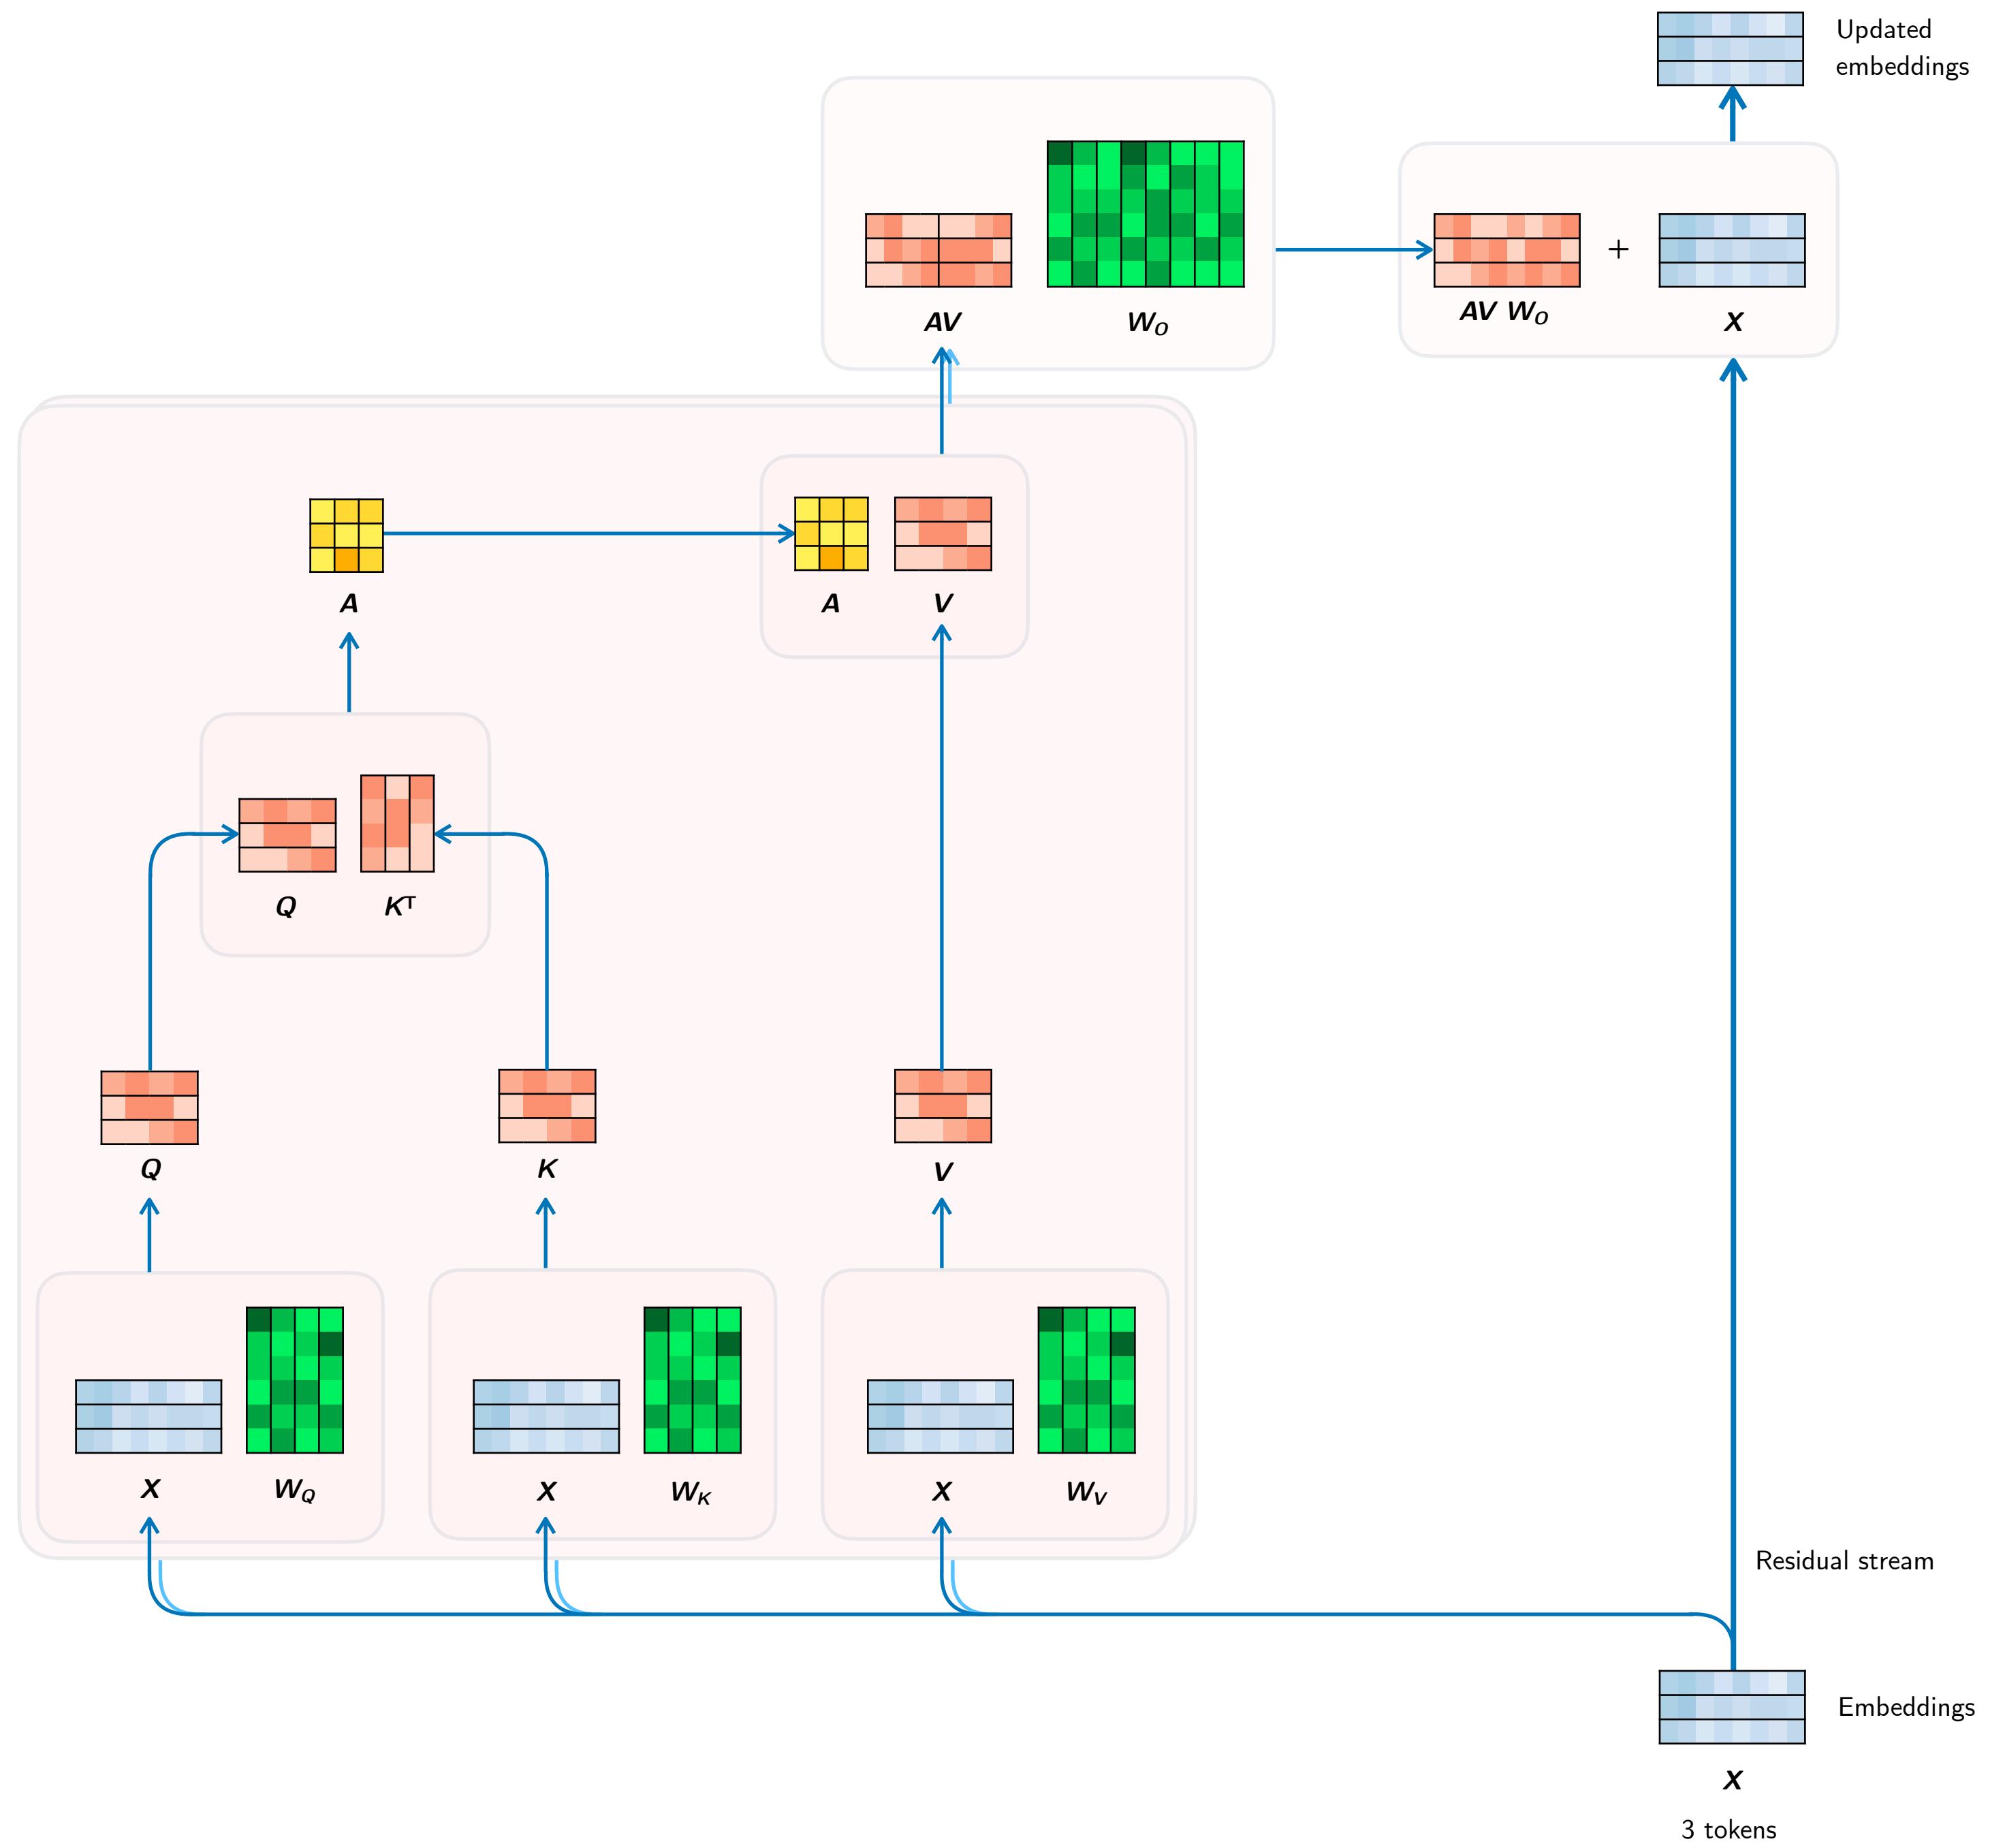
Leaving out Fwd for simplicity

Split text into fixed number of chunks l



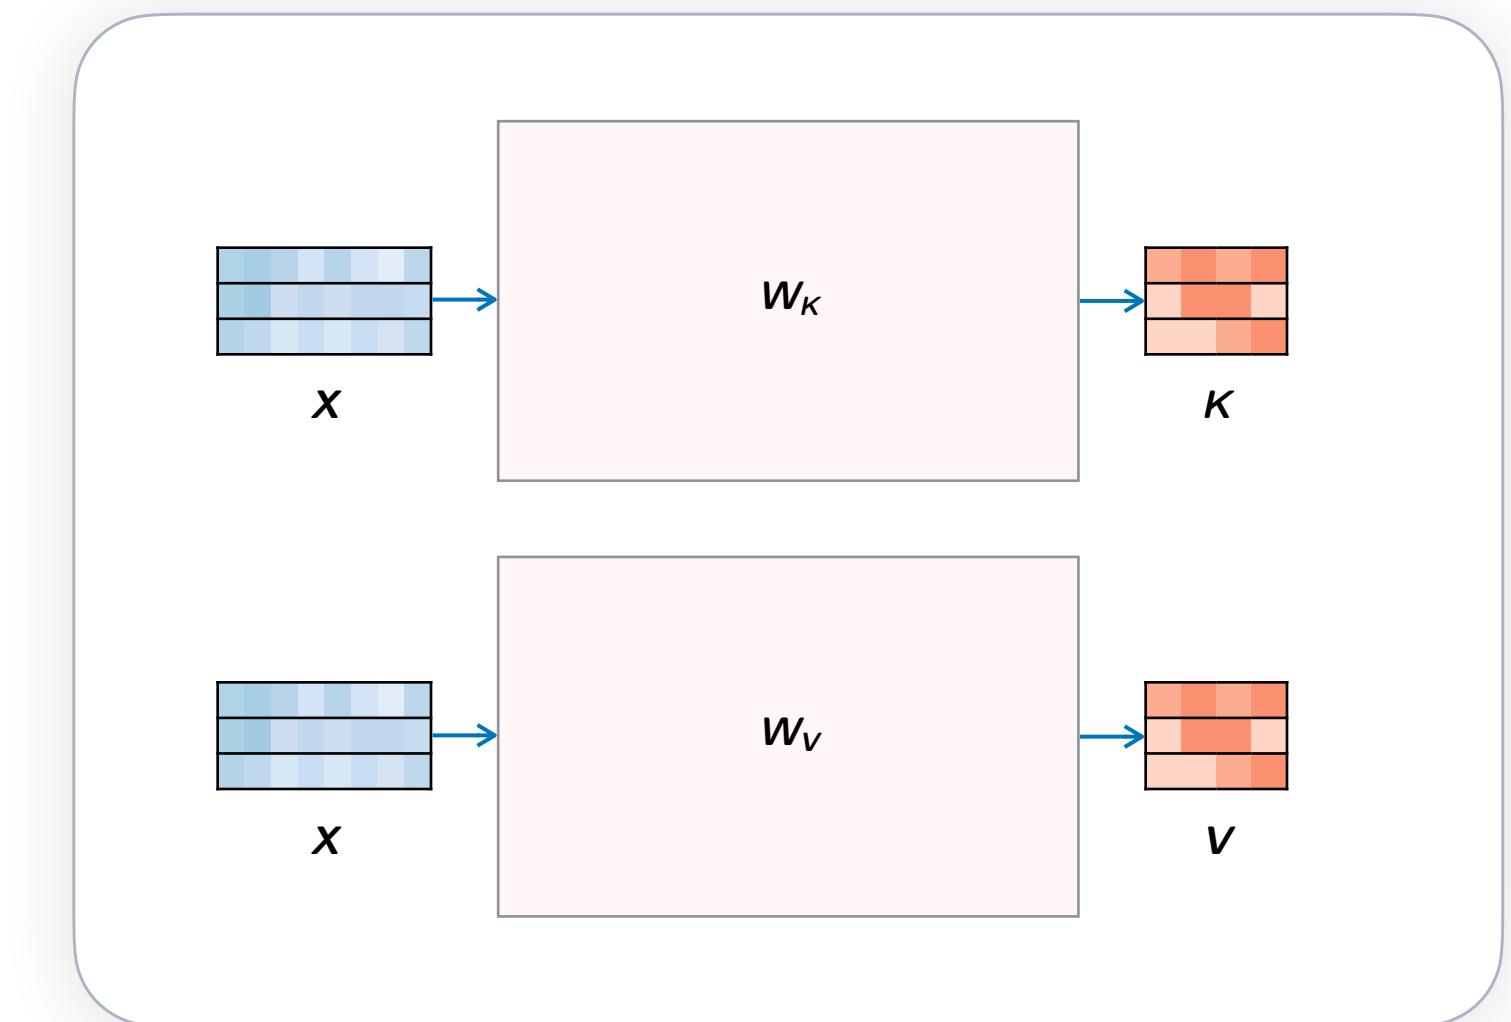
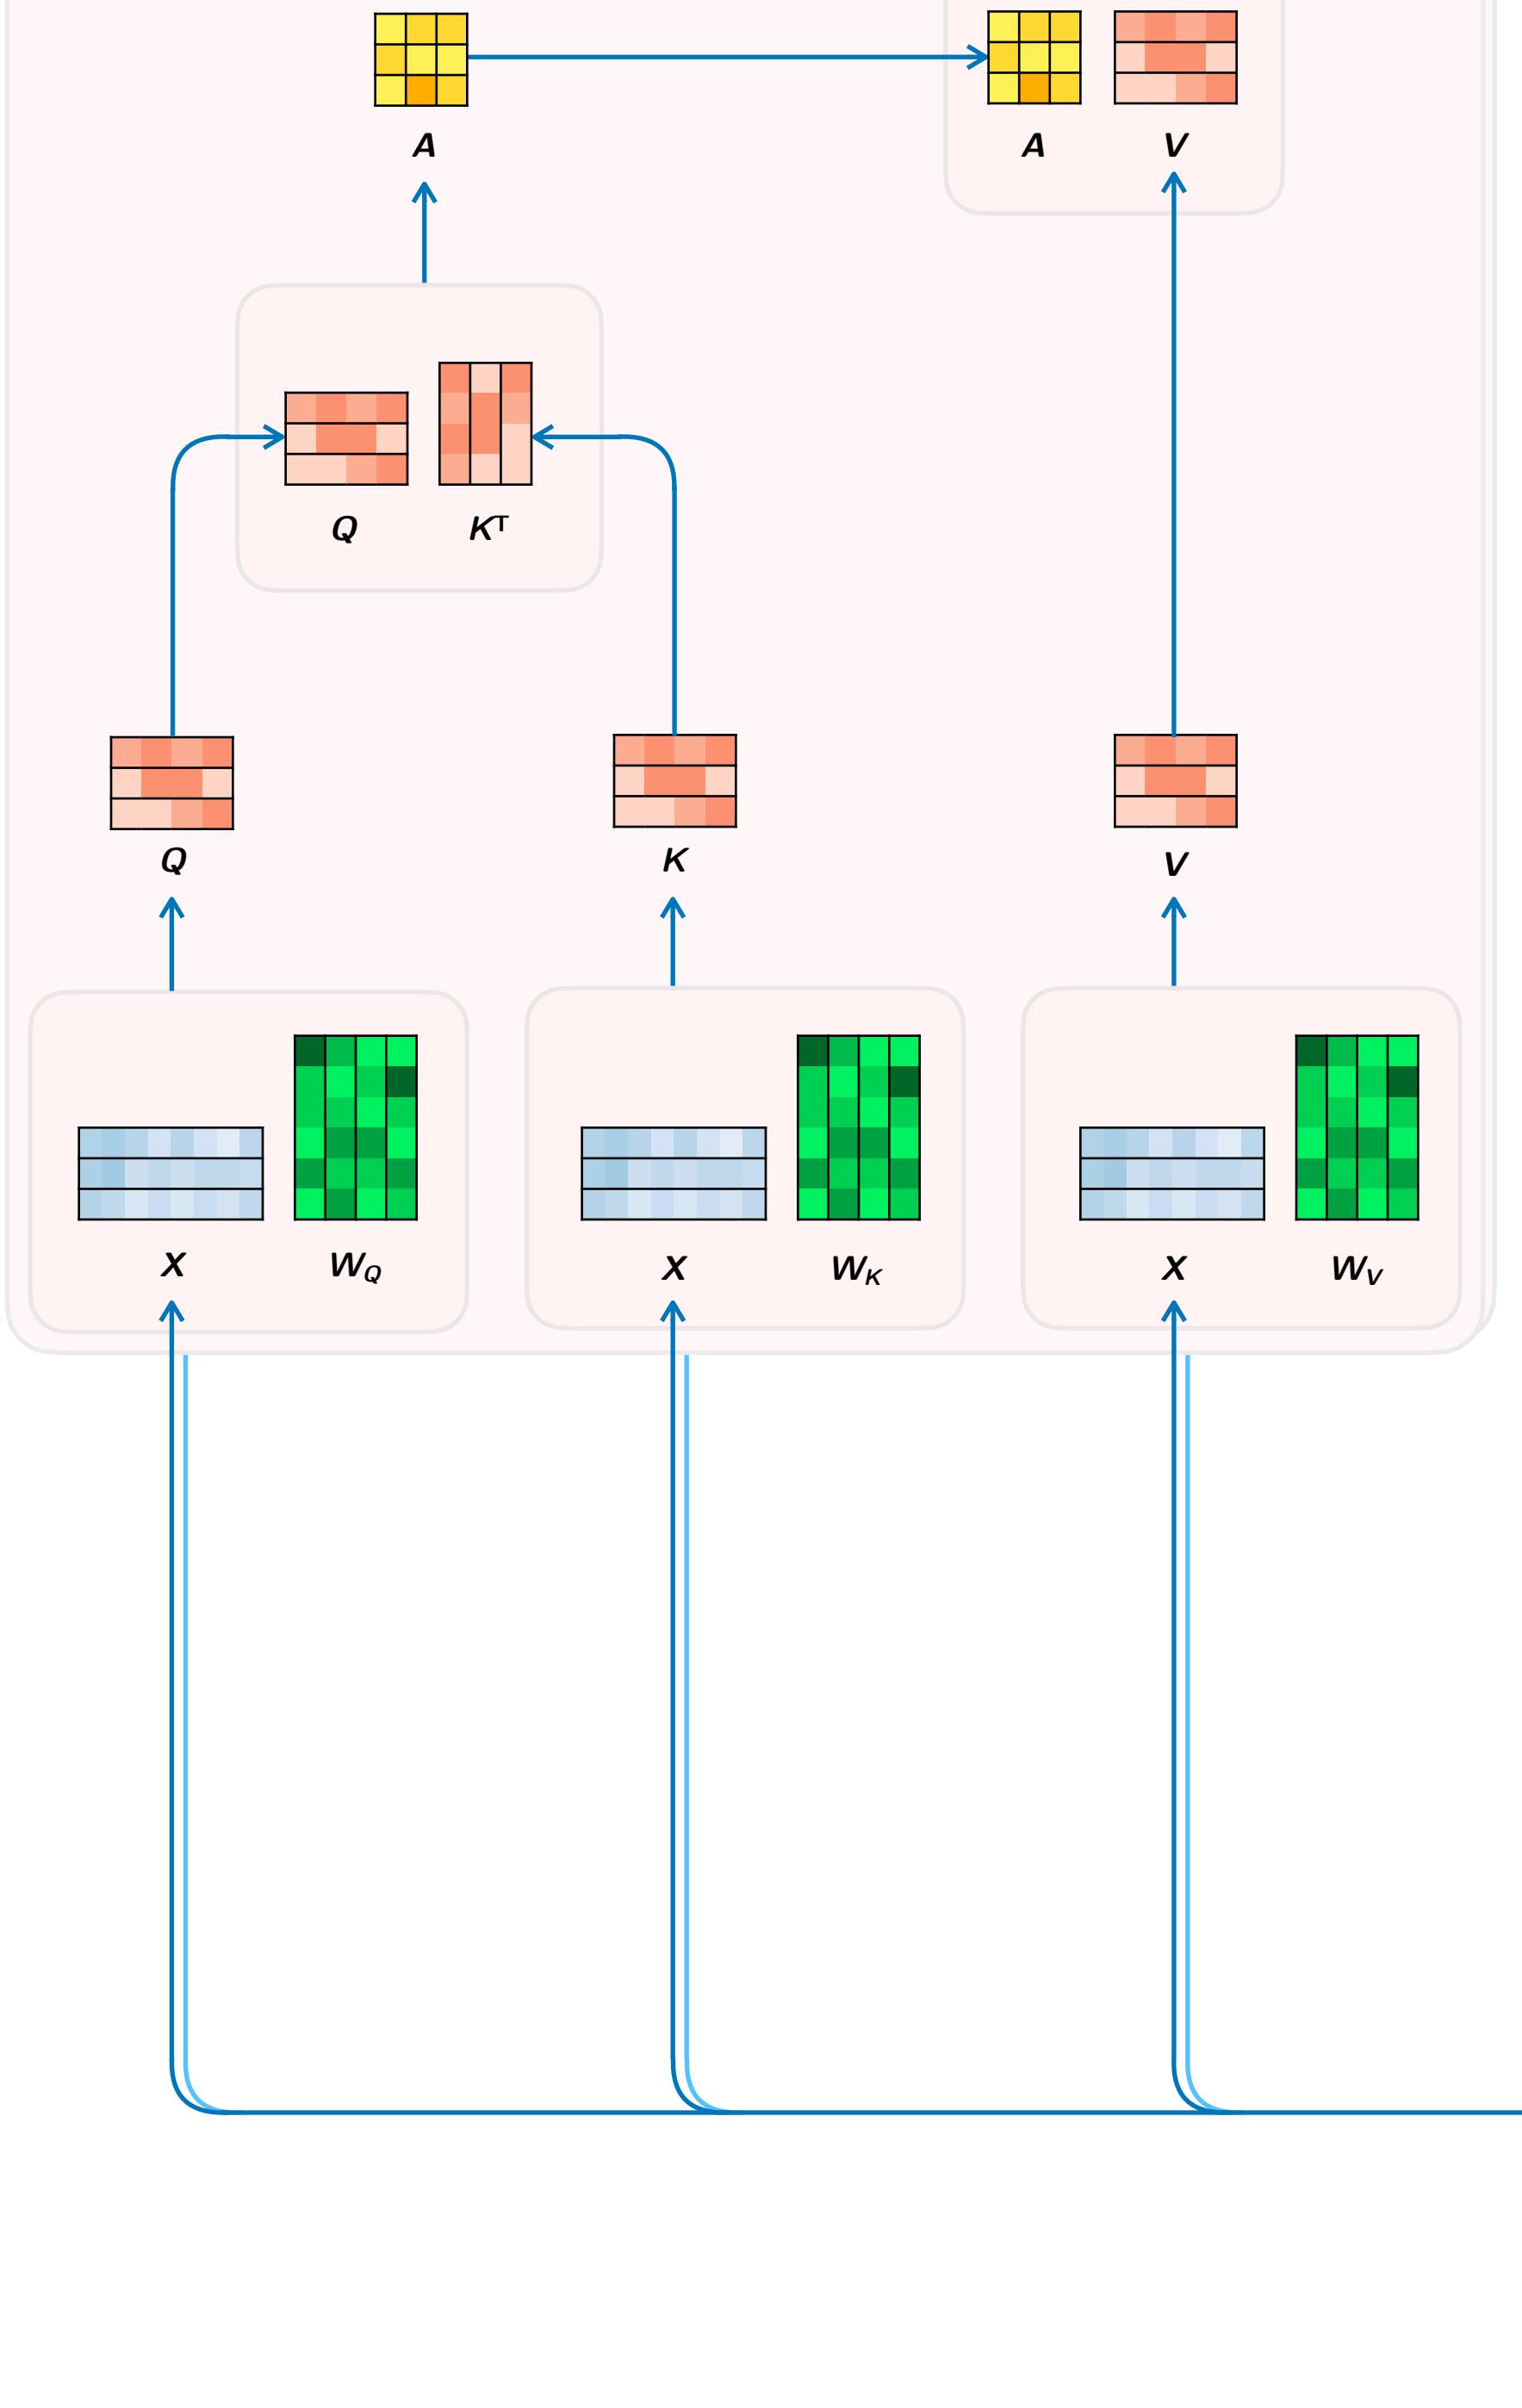
Alternative attention designs

Attention



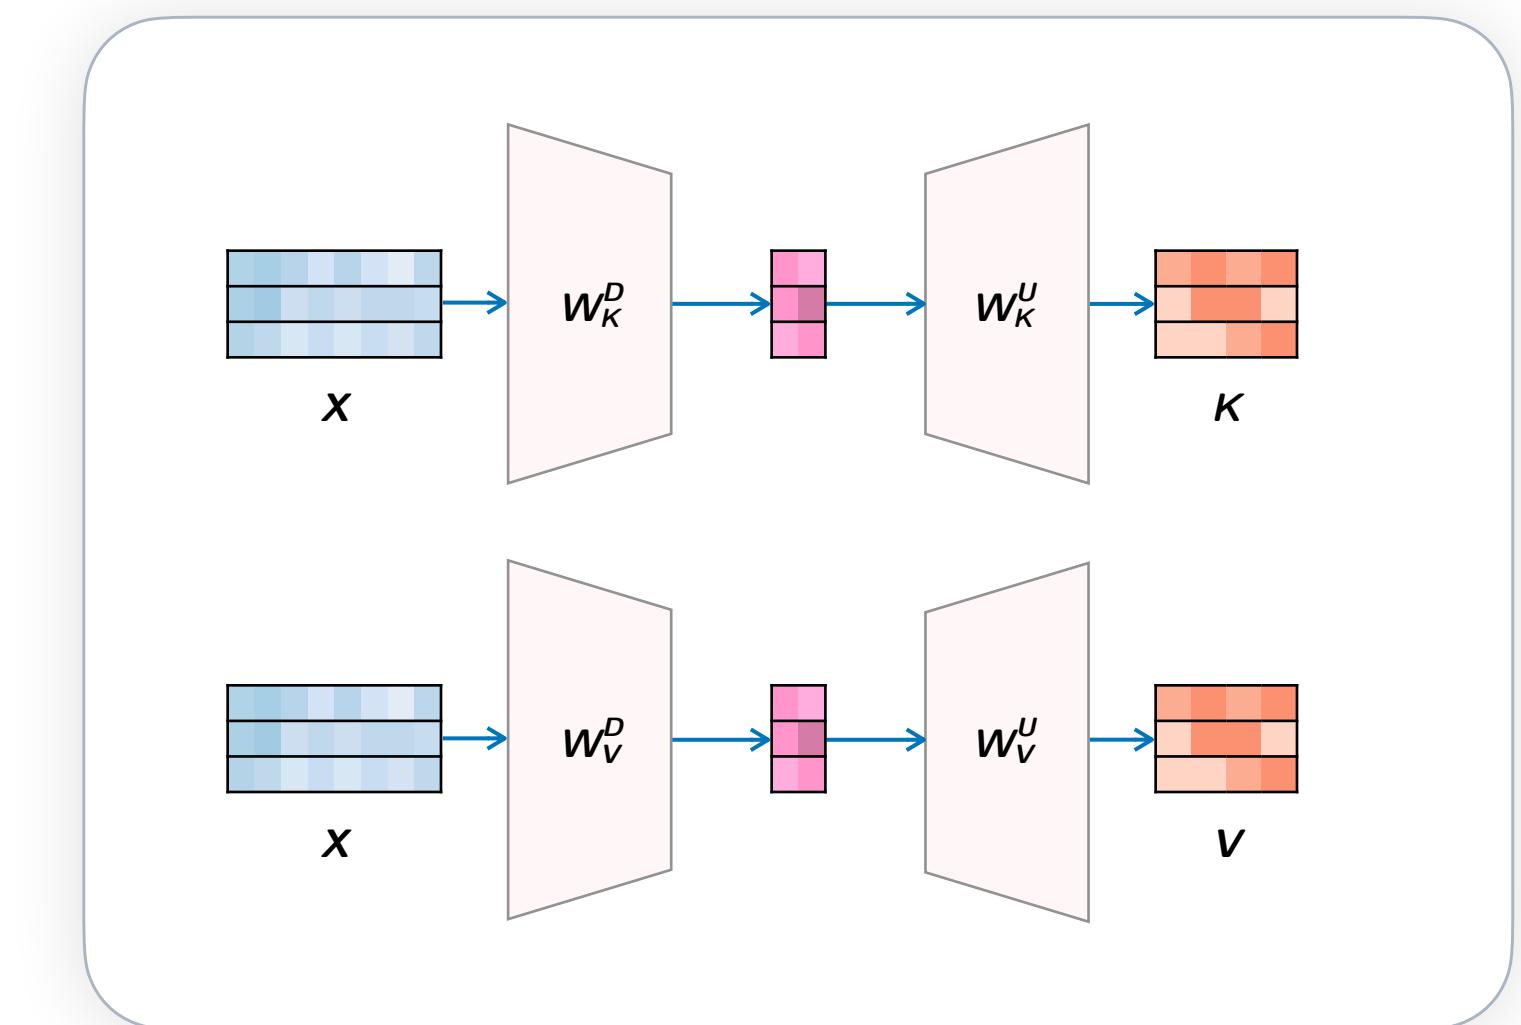
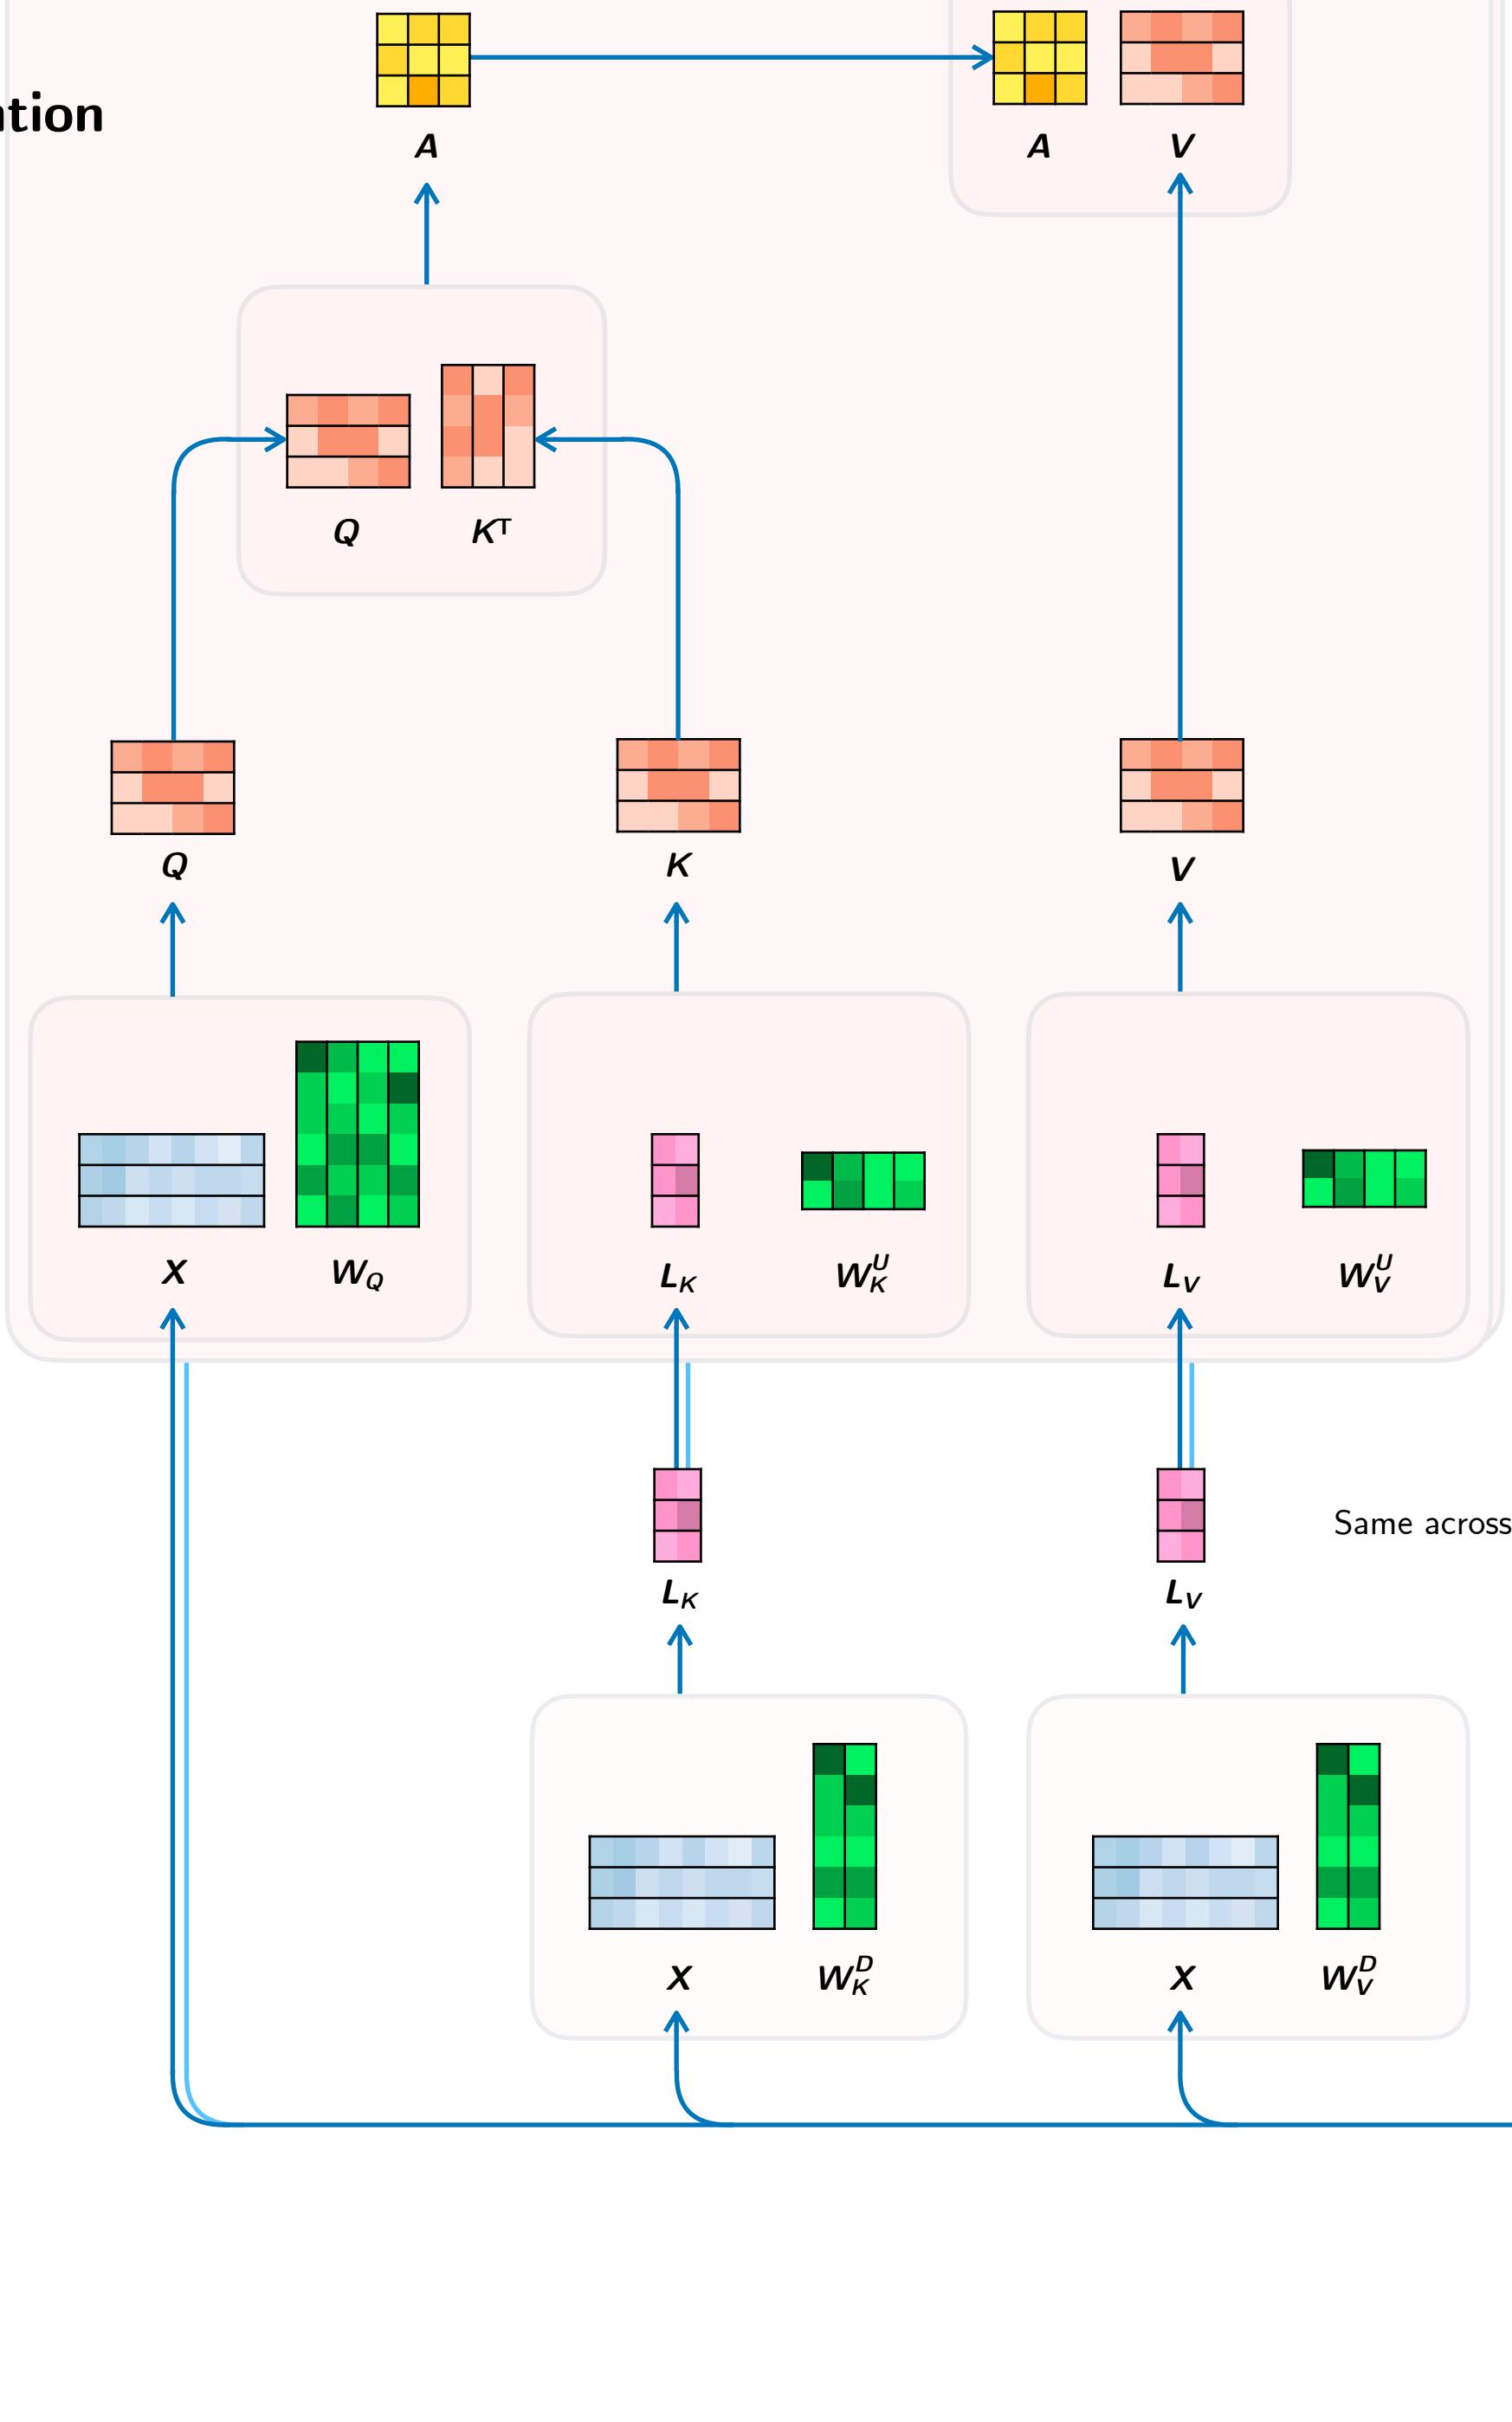
- X has shape $(3, d)$
 - Here: $d=8$
 - W_Q has shape (d, h)
 - We do $Q = X W_Q$, and then Q has shape $(3, h)$ (same with K, V)
 - h is chosen such that $h = d/n_{heads}$ so that we can cleanly concatenate the weighed-value matrices together again to get $(3, n_{heads} * h) = (3, d)$
 - Here: $n_{heads}=2, h=4$

Attention



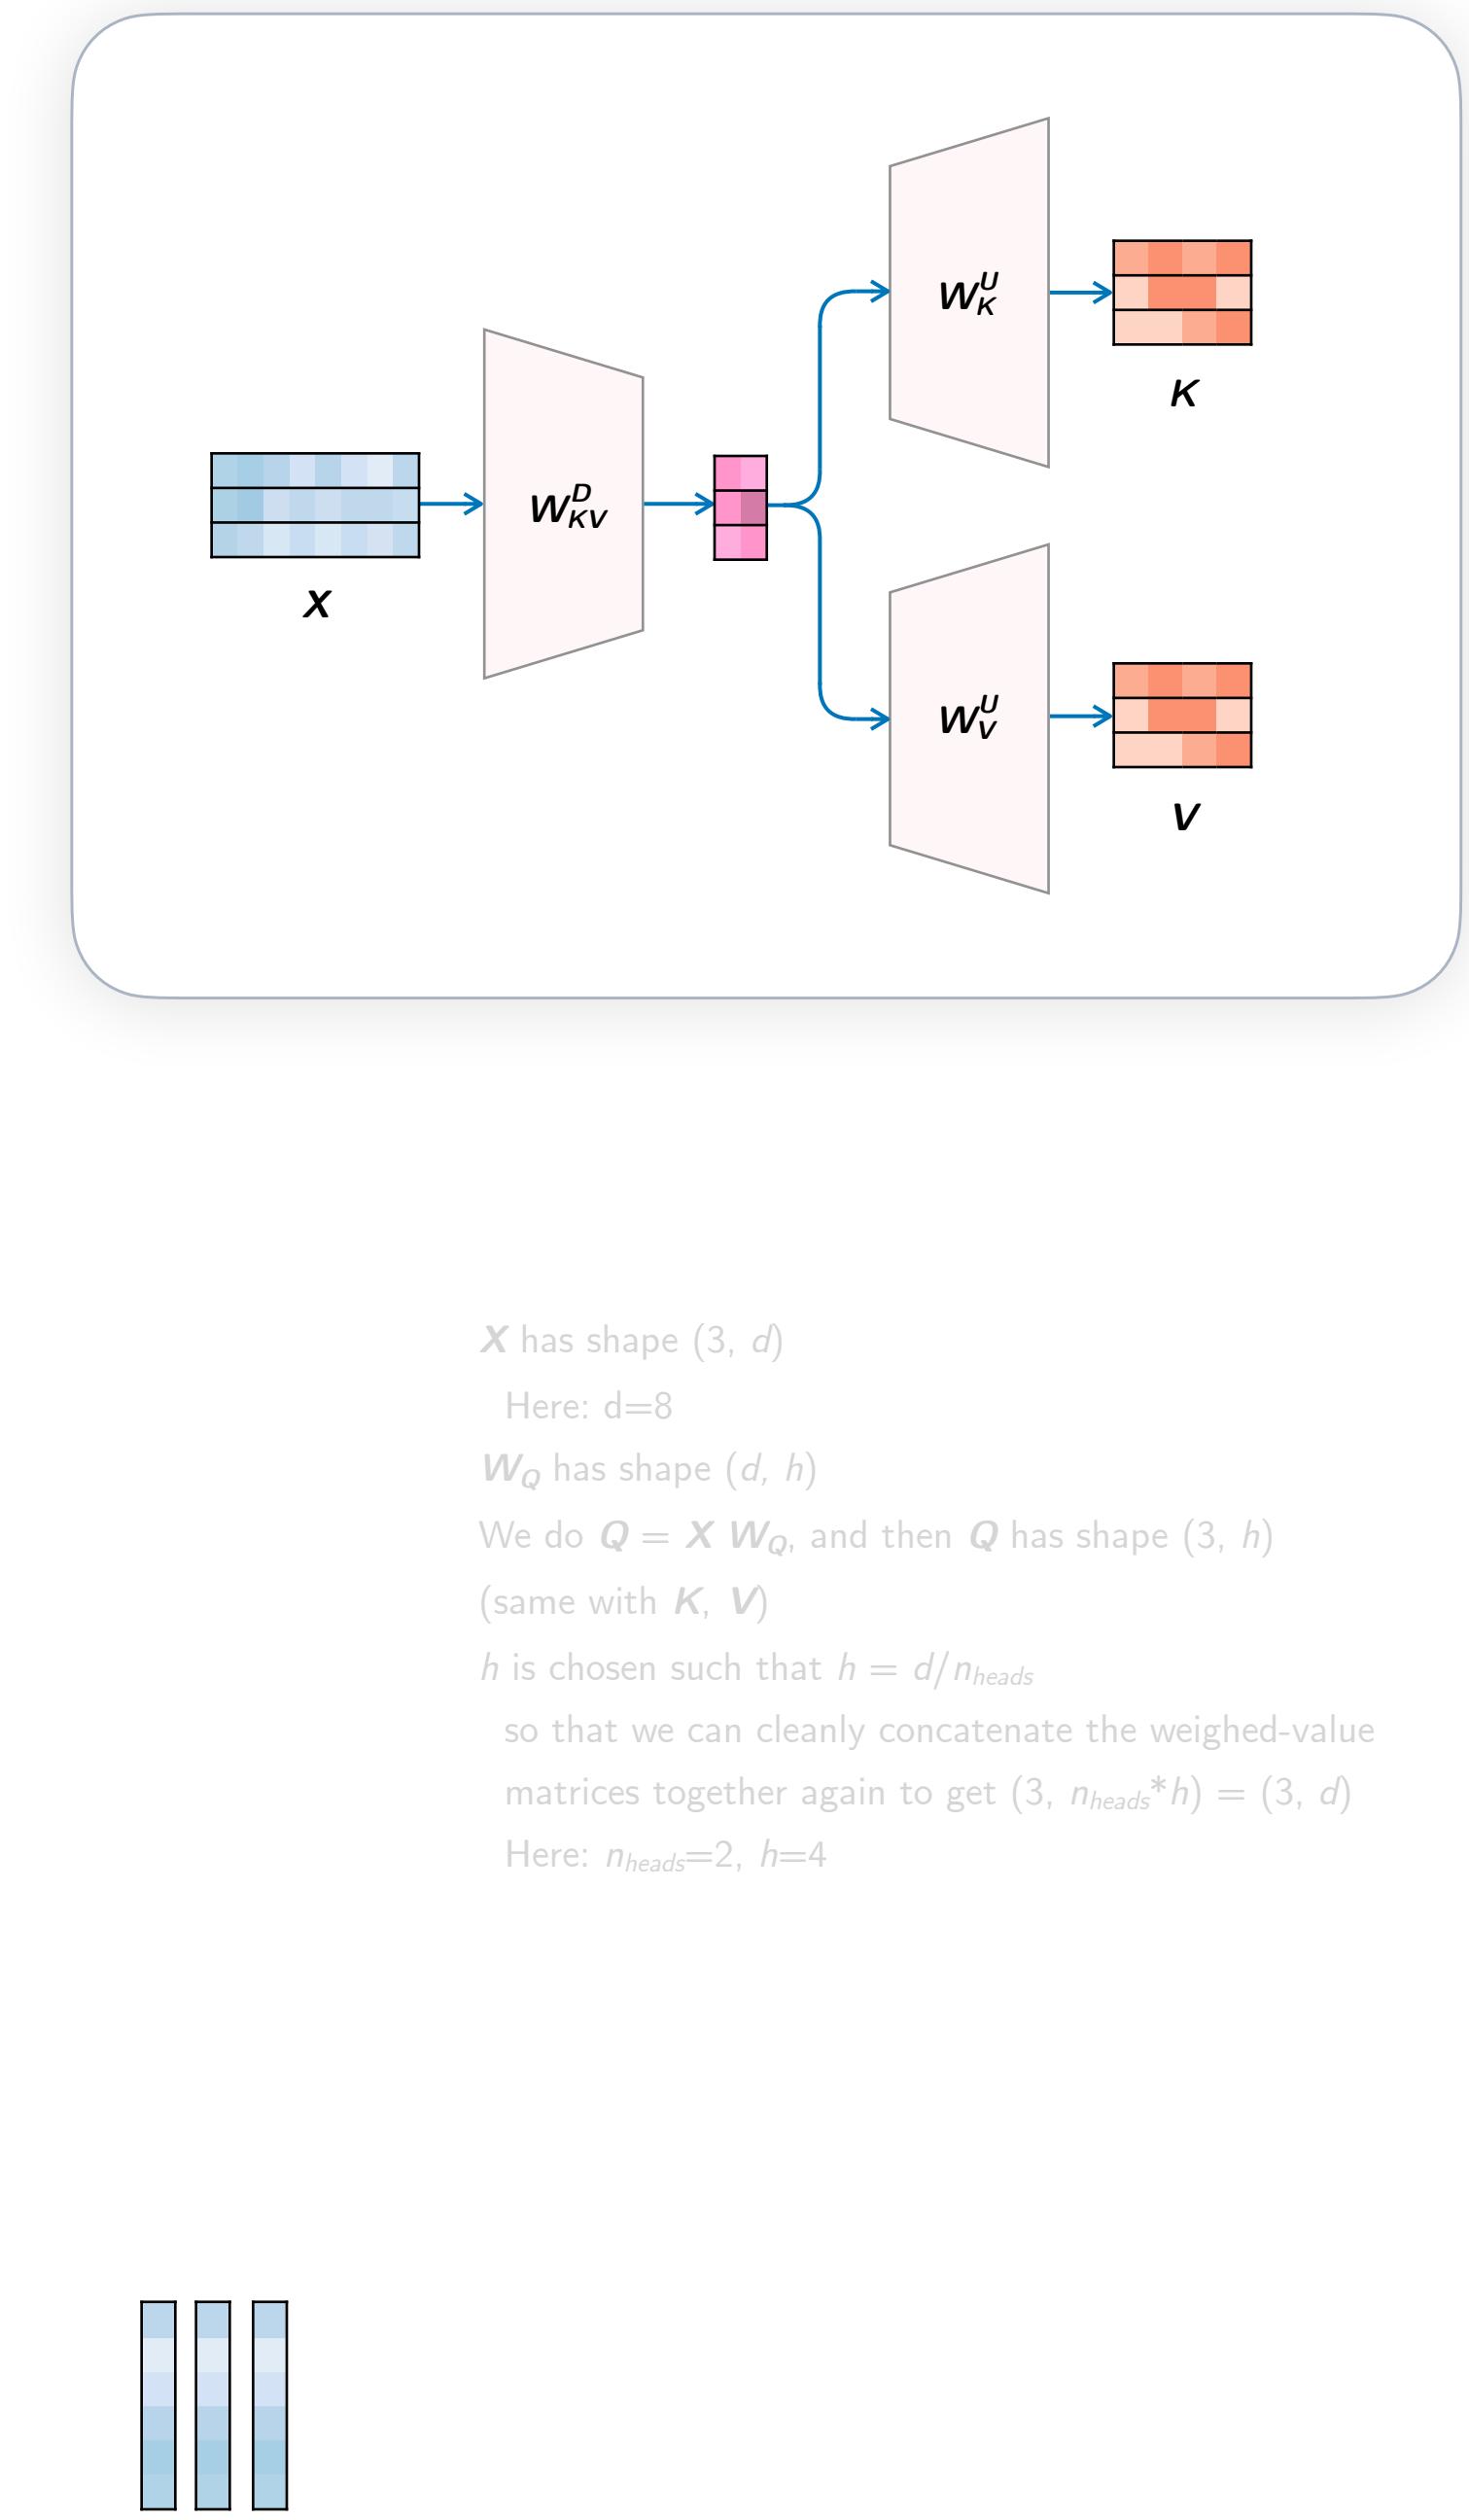
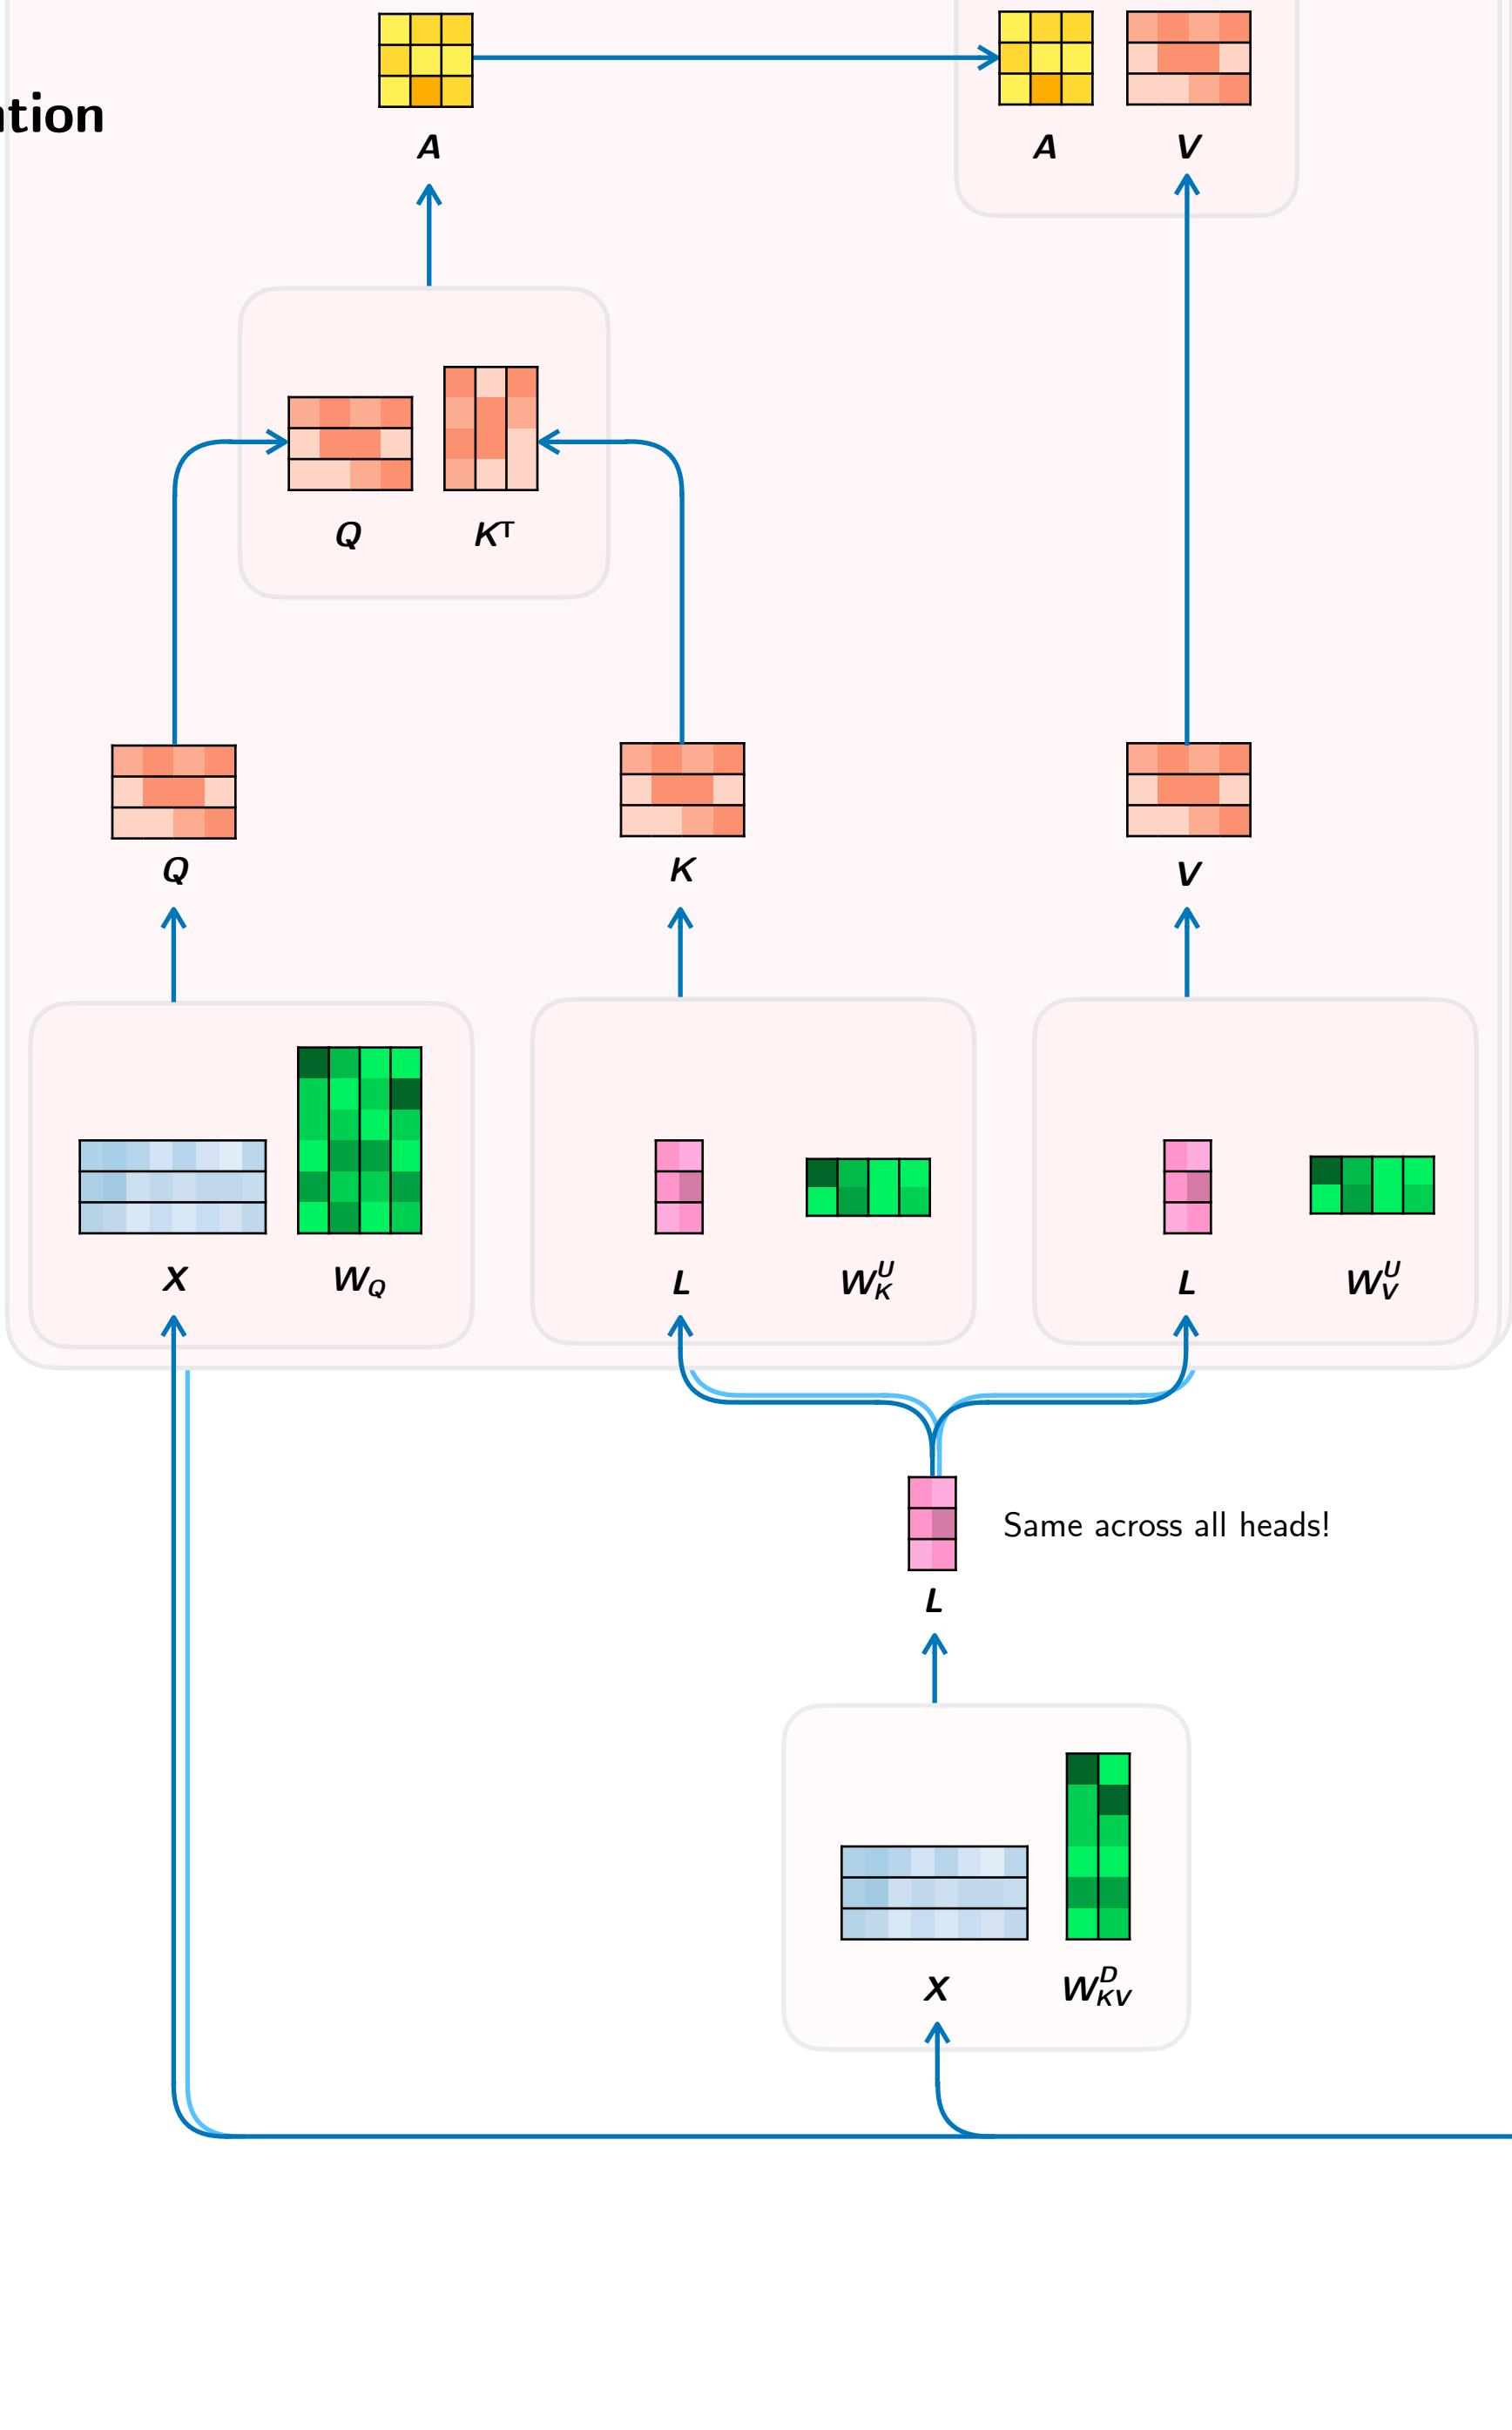
X has shape $(3, d)$
Here: $d=8$
 W_Q has shape (d, h)
We do $Q = X W_Q$, and then Q has shape $(3, h)$
(same with K, V)
 h is chosen such that $h = d/n_{heads}$
so that we can cleanly concatenate the weighed-value
matrices together again to get $(3, n_{heads} \cdot h) = (3, d)$
Here: $n_{heads}=2, h=4$

Latent attention



X has shape $(3, d)$
Here: $d=8$
 W_Q has shape (d, h)
We do $Q = X W_Q$, and then Q has shape $(3, h)$
(same with K, V)
 h is chosen such that $h = d/n_{\text{heads}}$
so that we can cleanly concatenate the weighed-value
matrices together again to get $(3, n_{\text{heads}} \cdot h) = (3, d)$
Here: $n_{\text{heads}}=2, h=4$

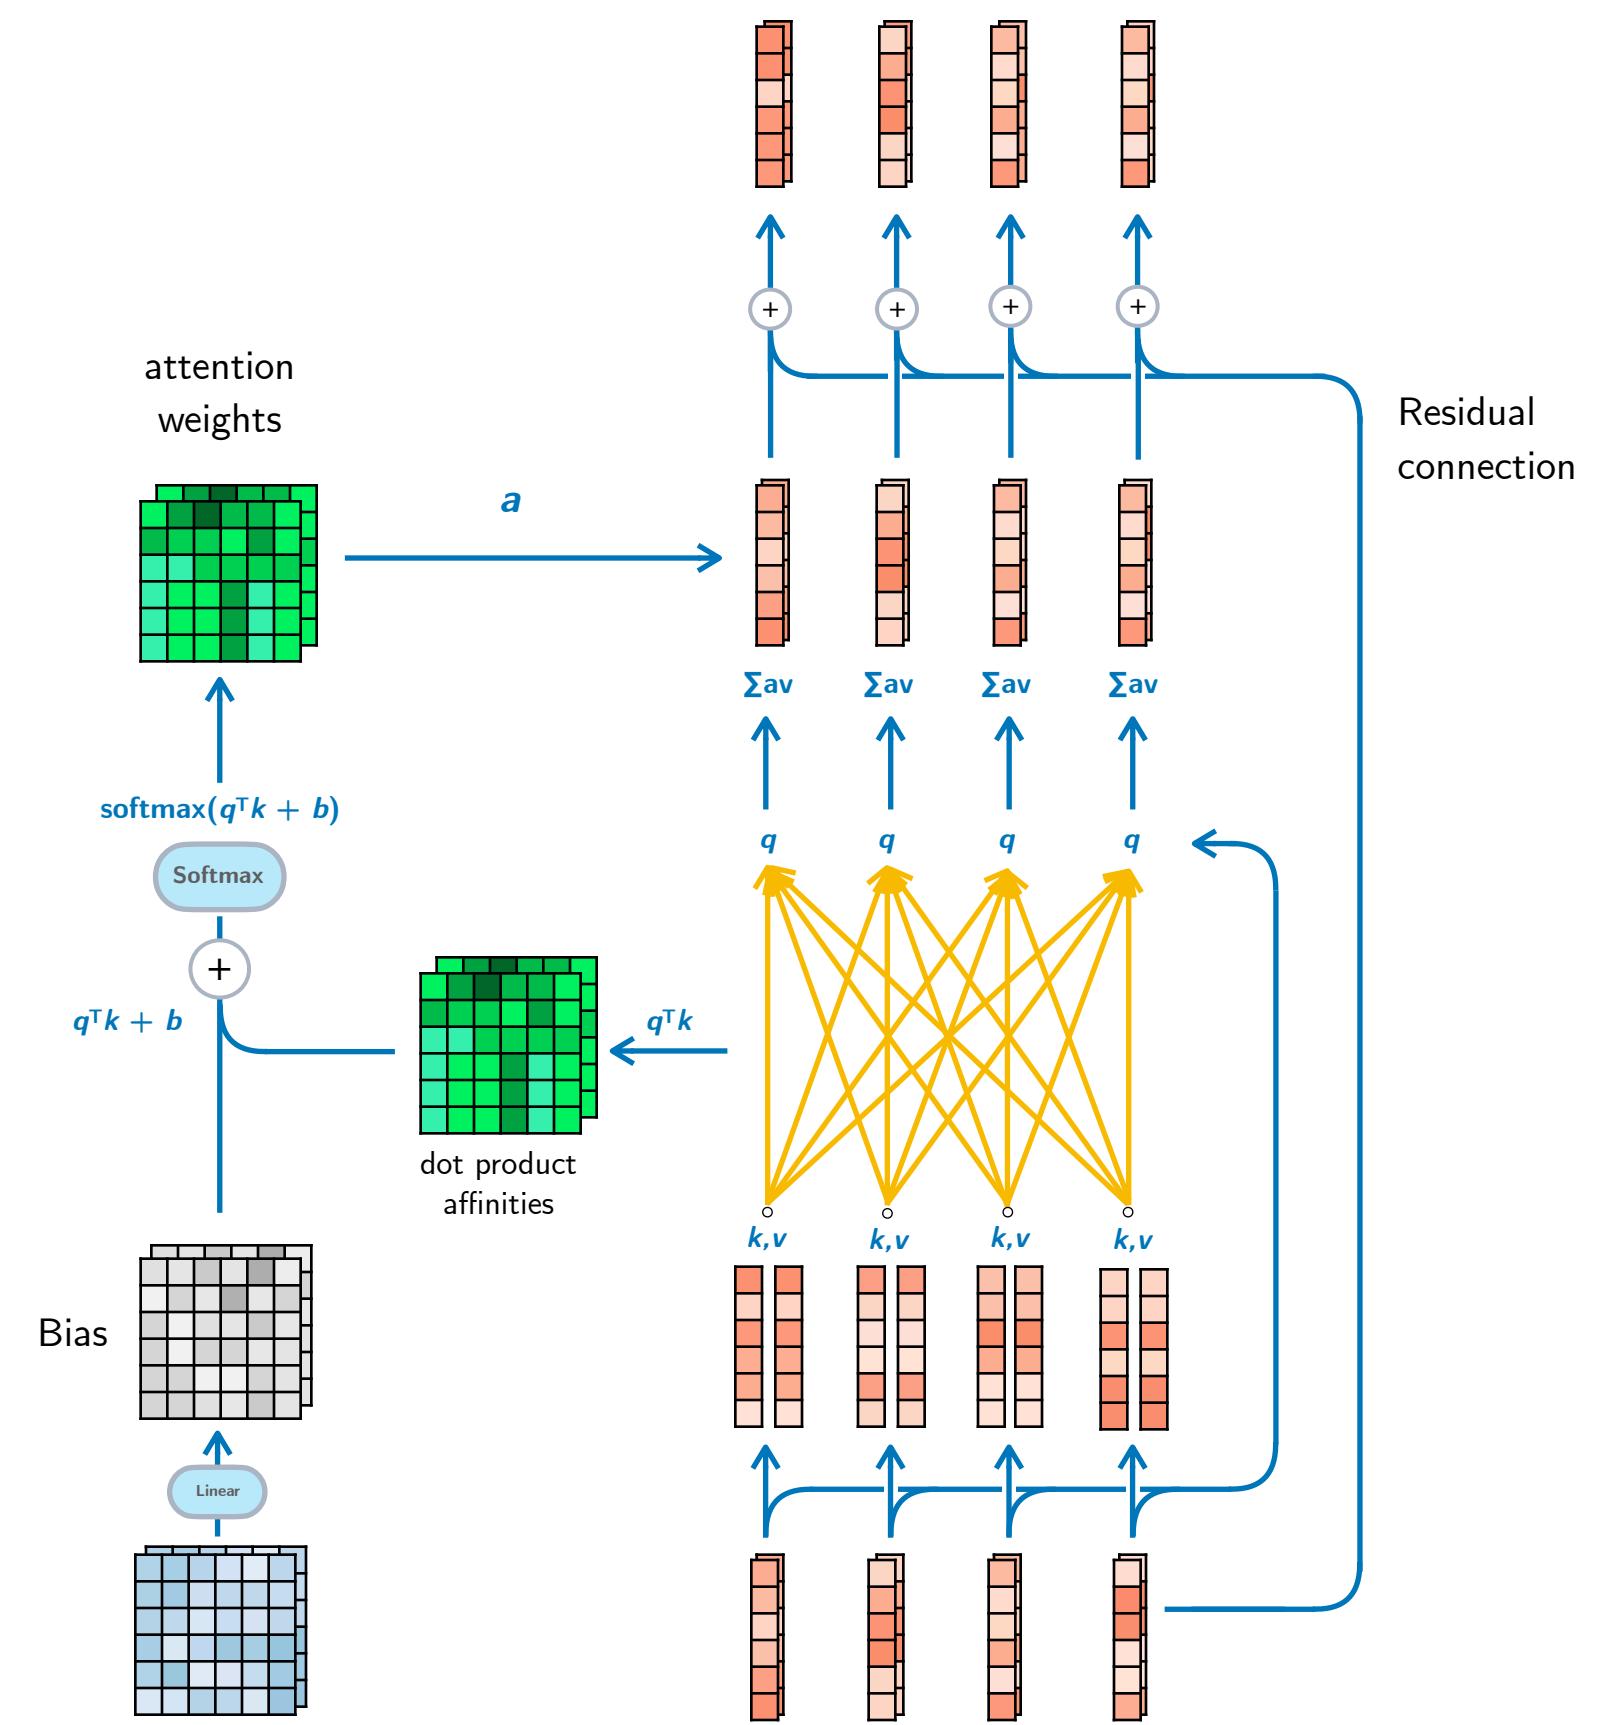
Latent attention



Biased attention

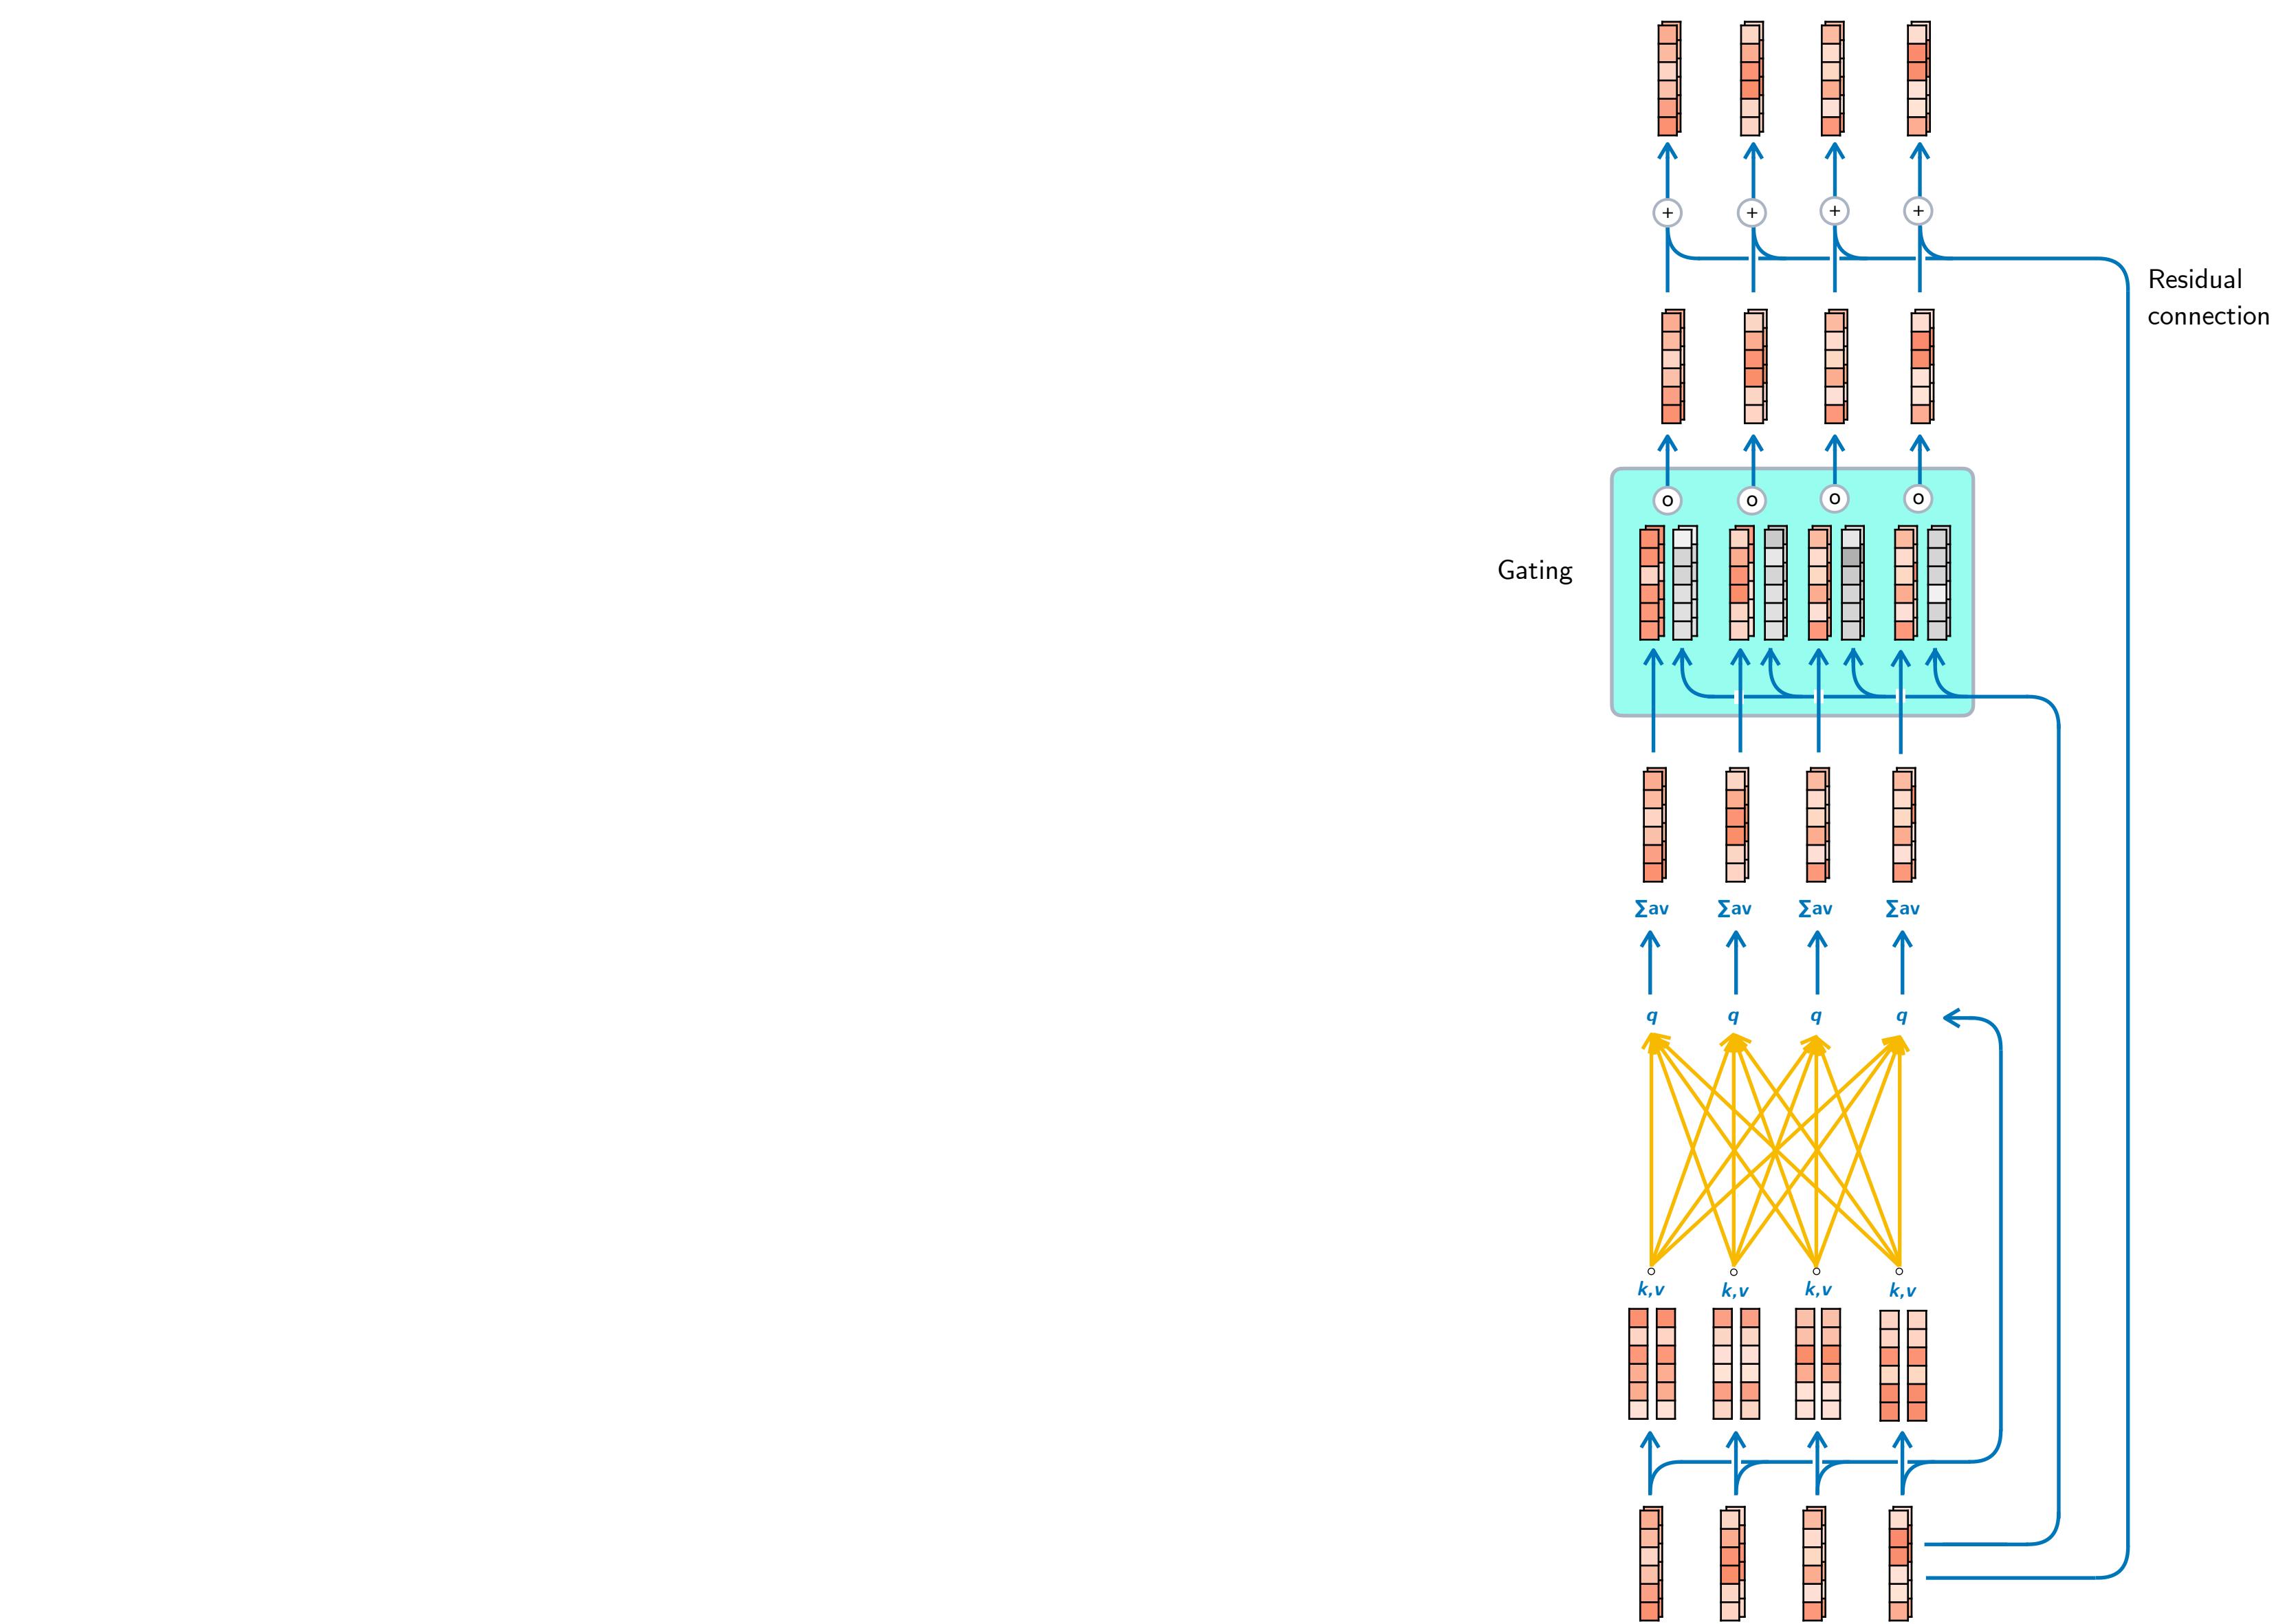
e.g. used in AlphaFold 2

two-lane transformer, where one lane biases the other
by adding a bias tensor to the dot-product similarity scores
alternatively, think of encoder biasing the decoder
or decoder biasing the encoder



Gated attention

e.g. used in AlphaFold 2

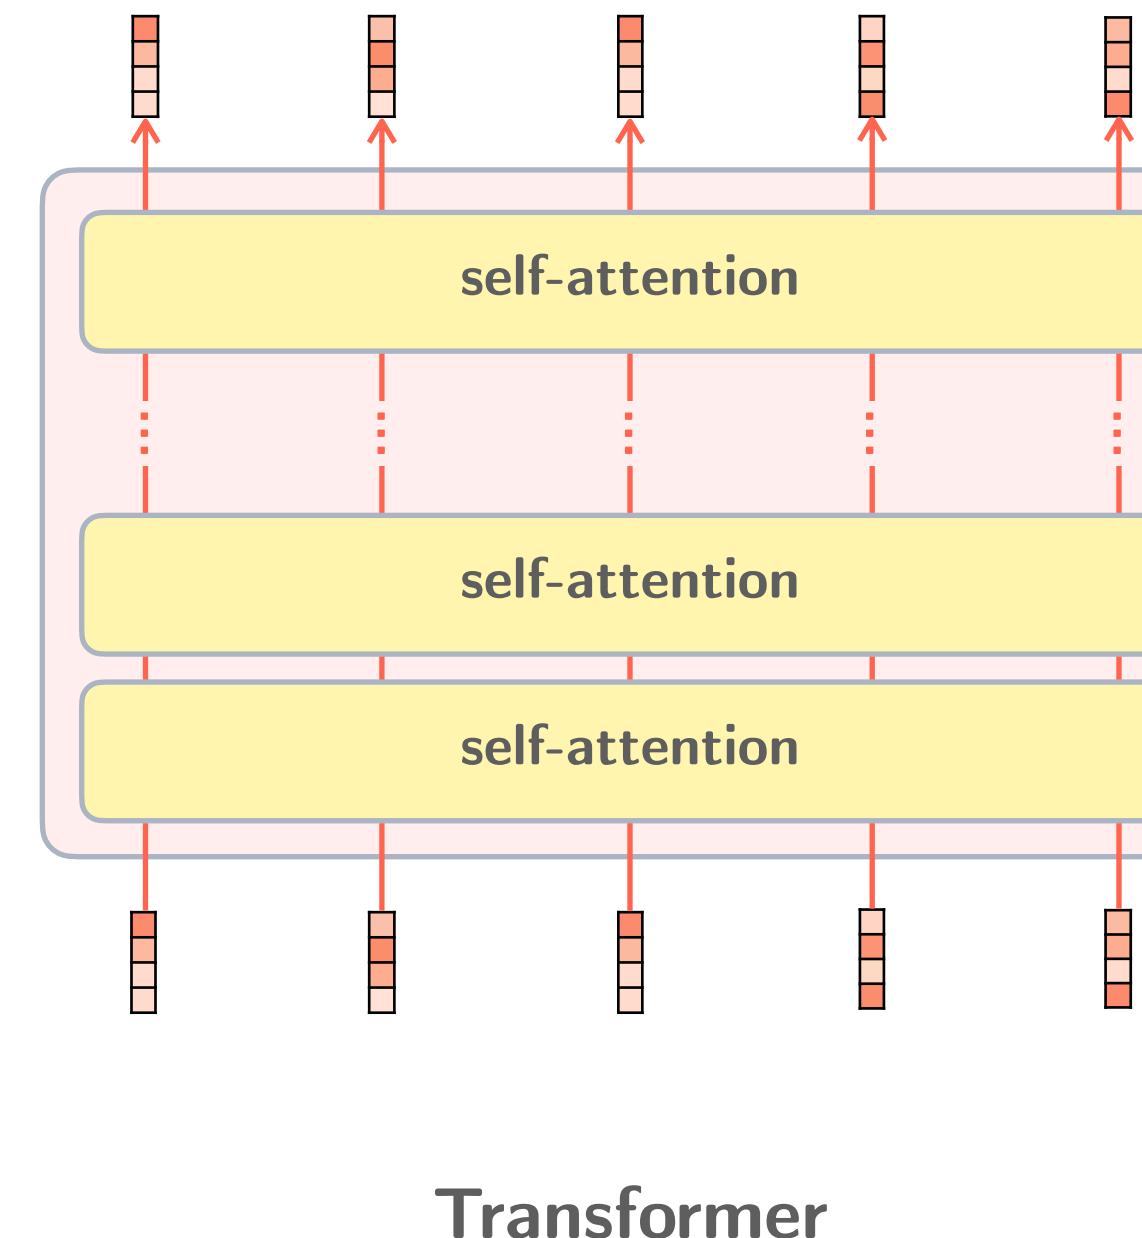


Linear layers to compute keys, values, queries not shown.



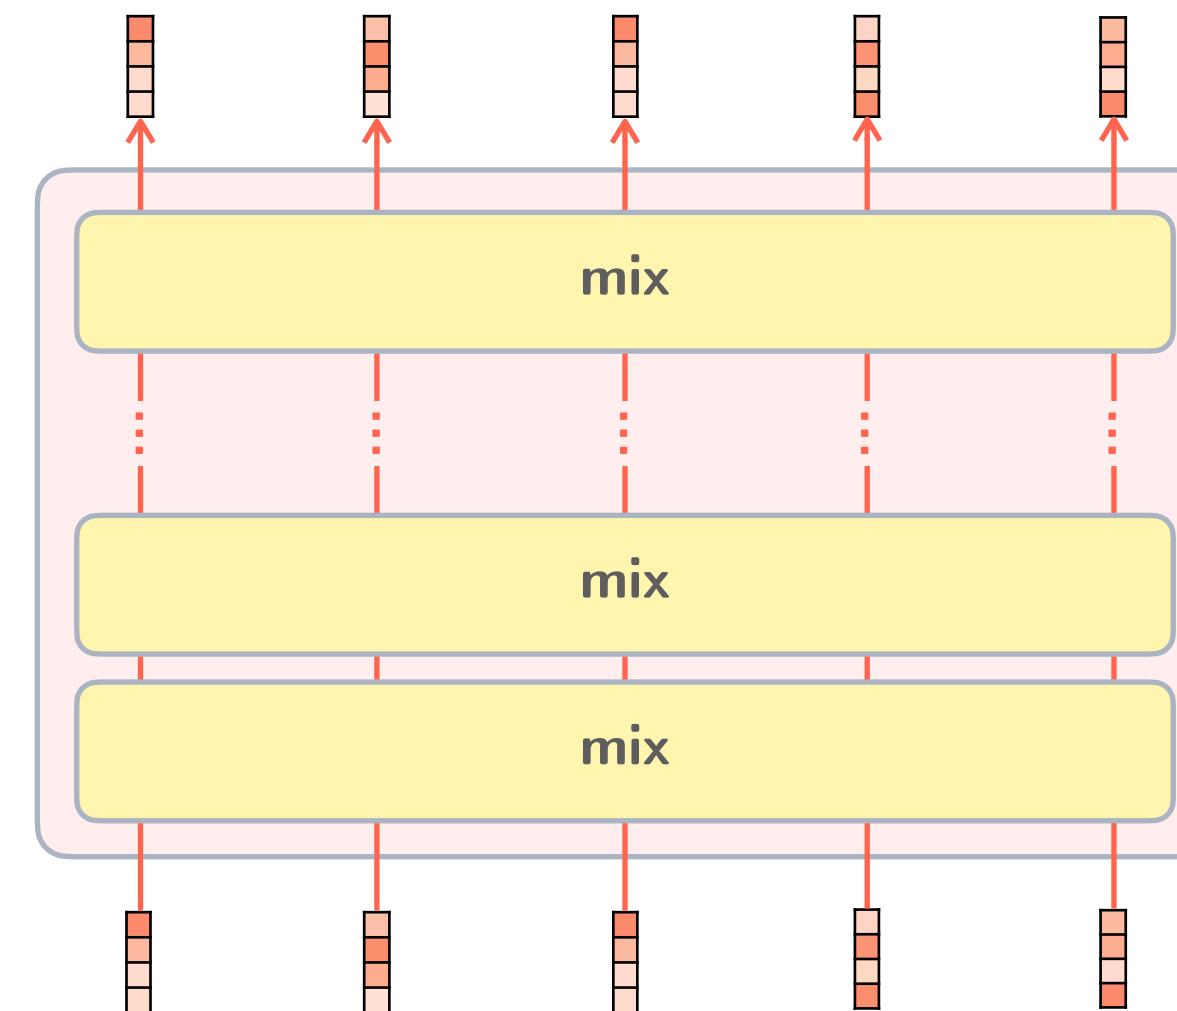
Non-attention transformers

Not showing non-mixing layers for simplicity



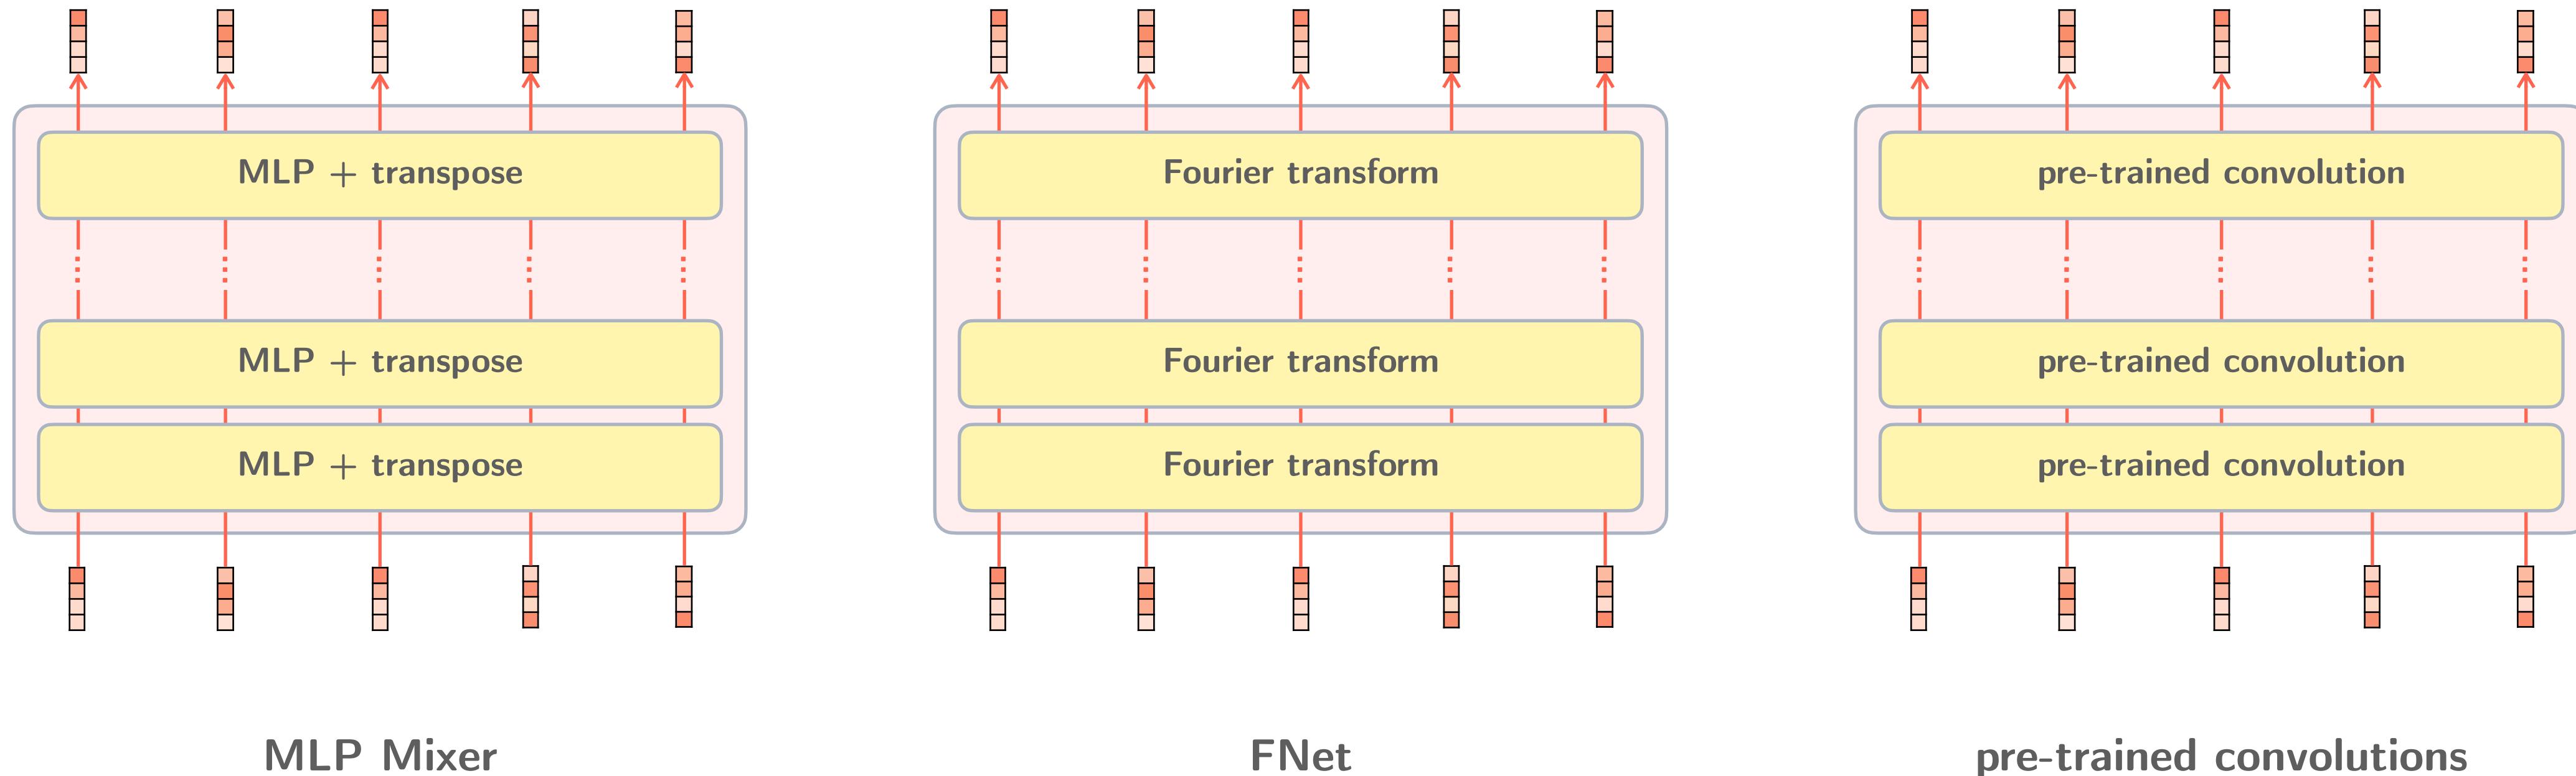
Any non-recurrent sequence-processing neural network where the sequence representation gets refined with each layer by letting individual time steps learn about all the other time steps?

Not showing non-mixing layers for simplicity



“Token network”

Any non-recurrent sequence-processing neural network where the sequence representation gets refined with each layer by letting individual time steps learn about all the other time steps?

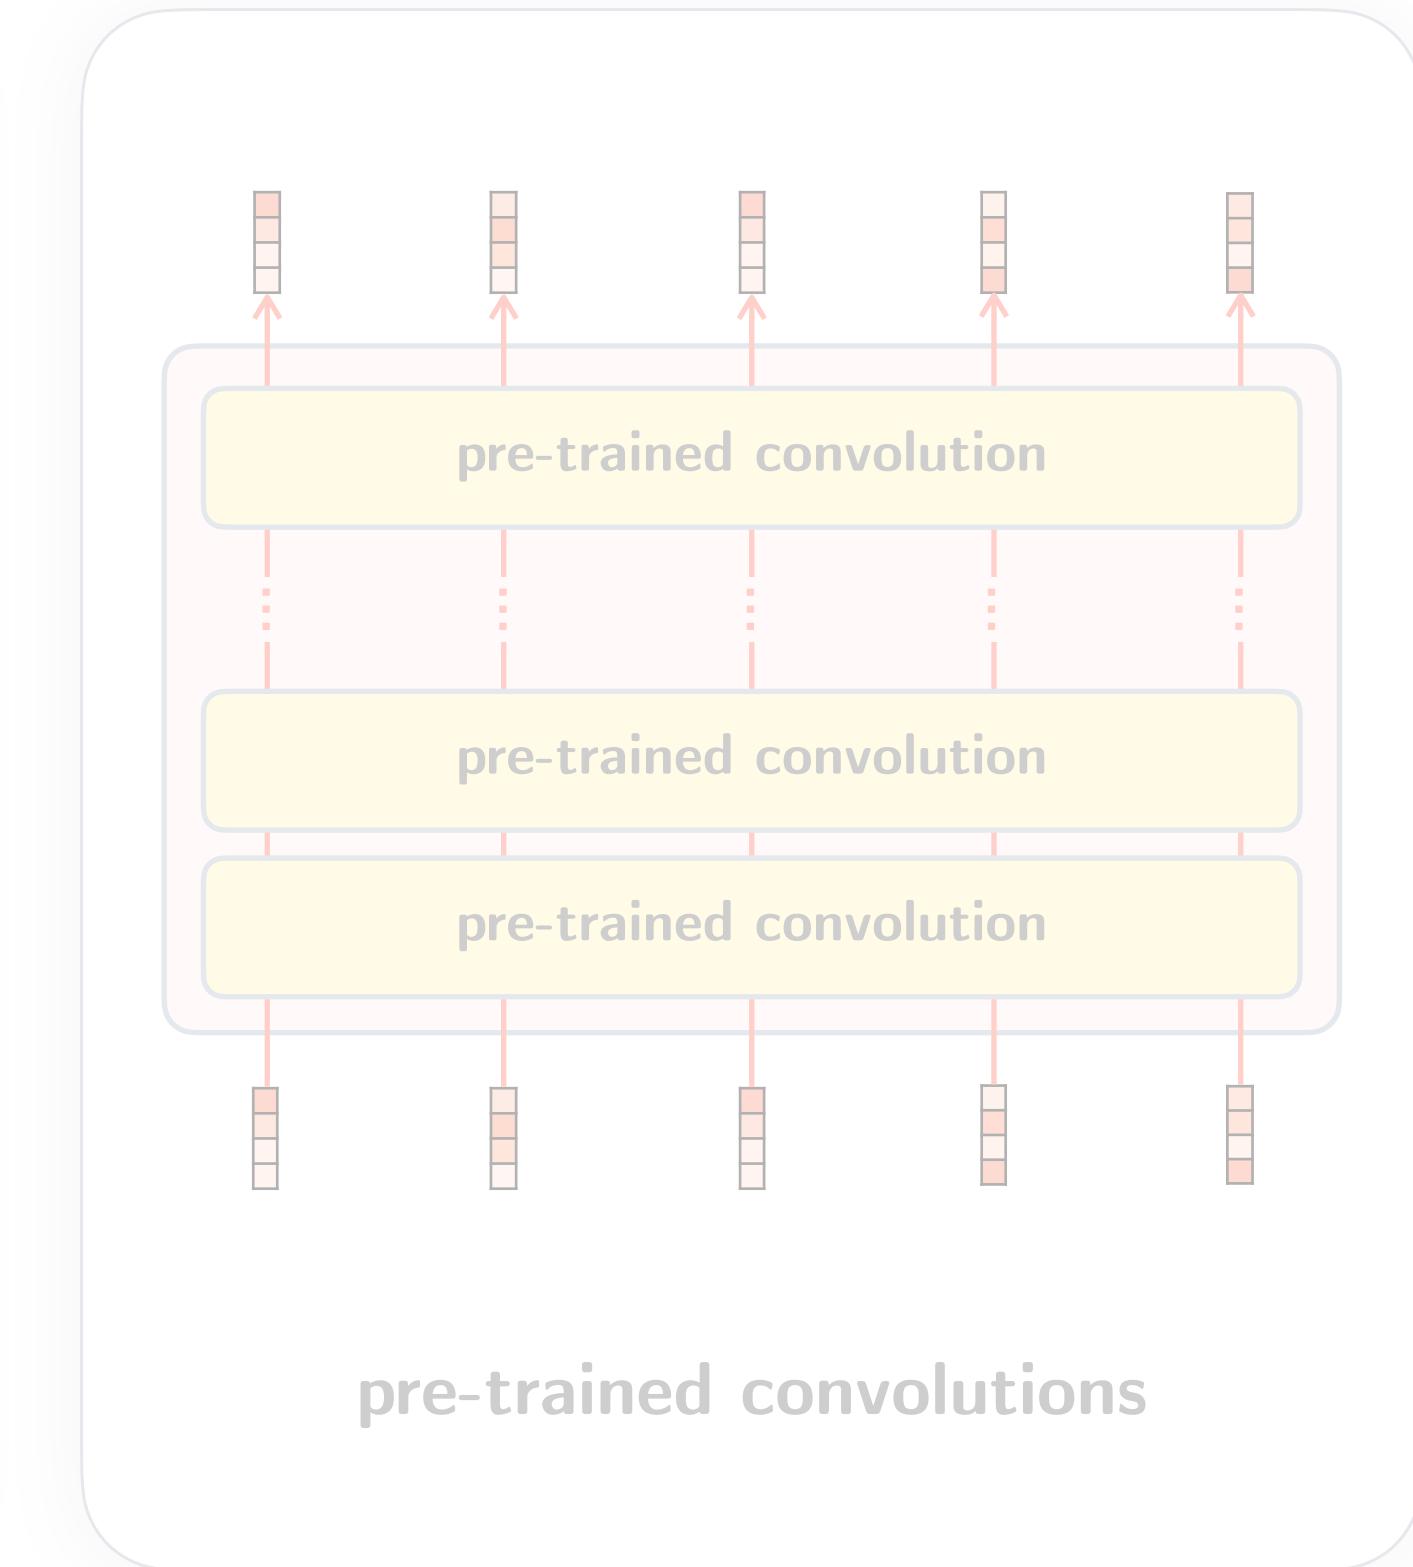
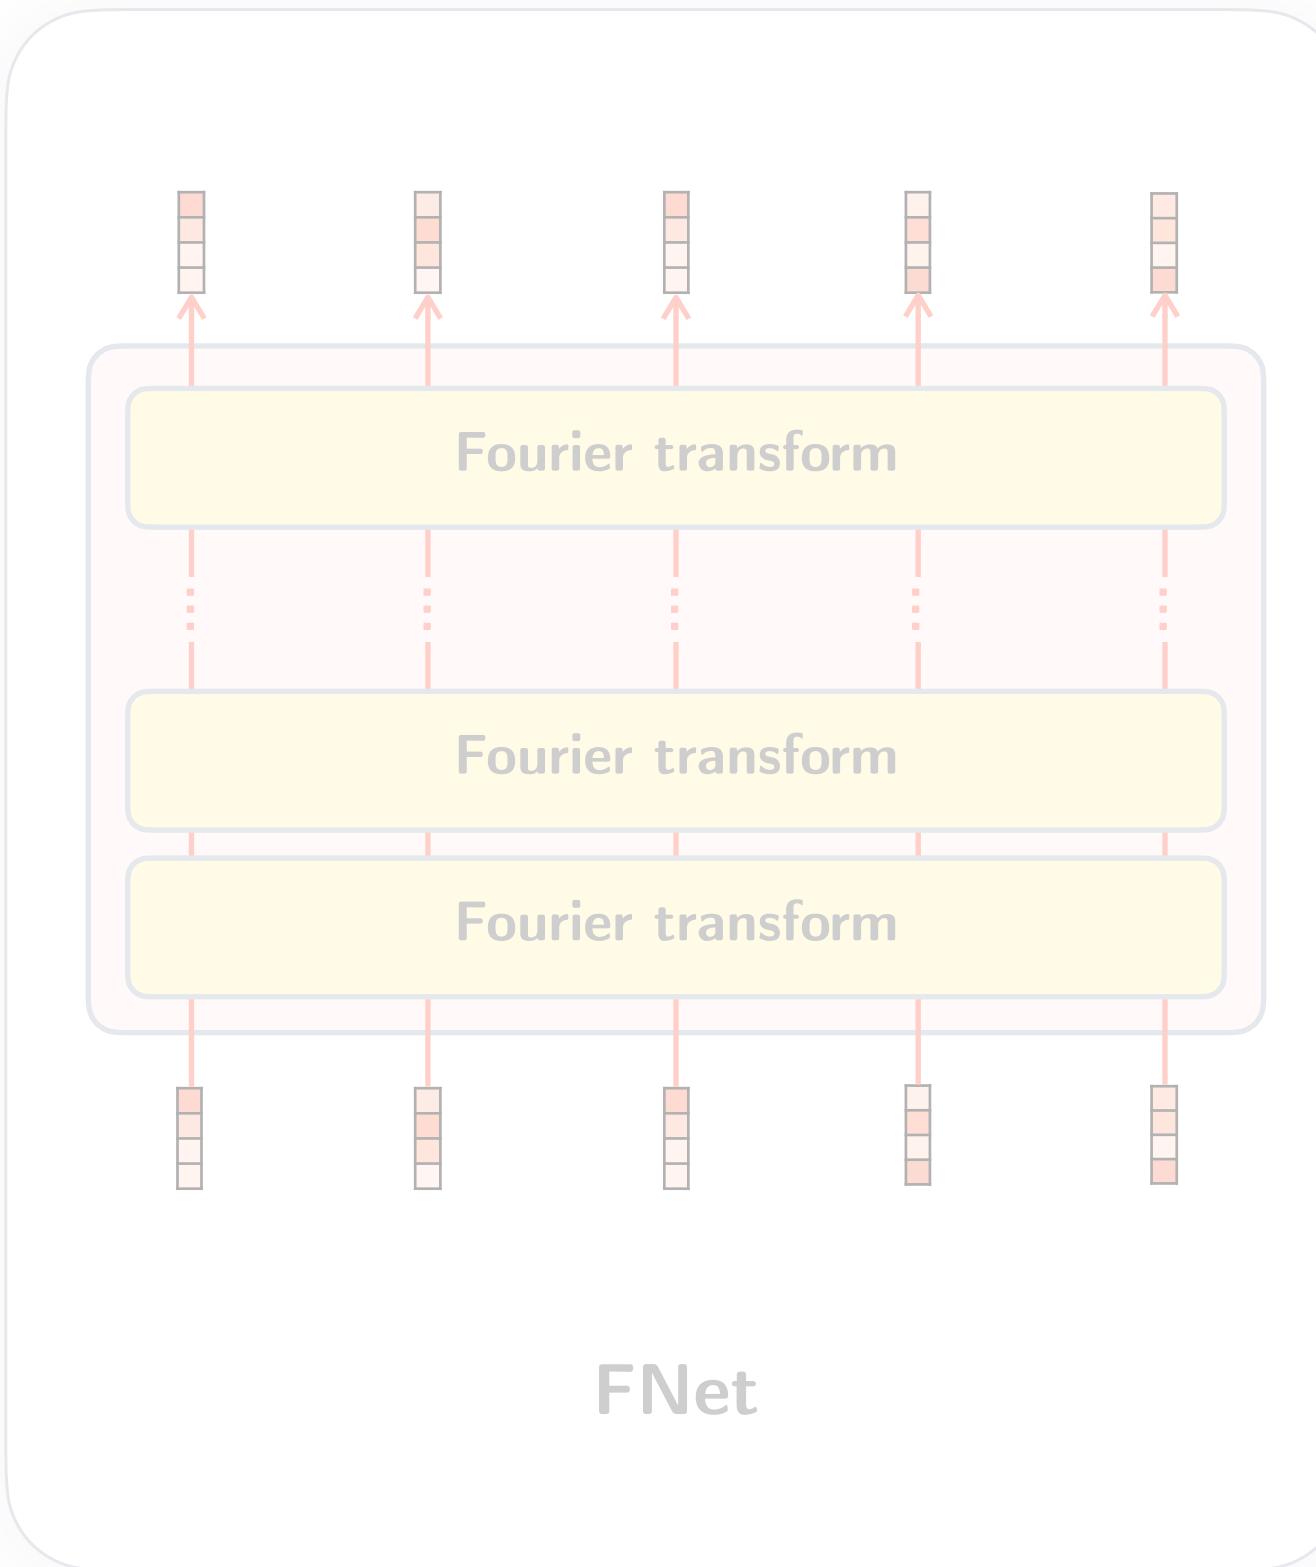
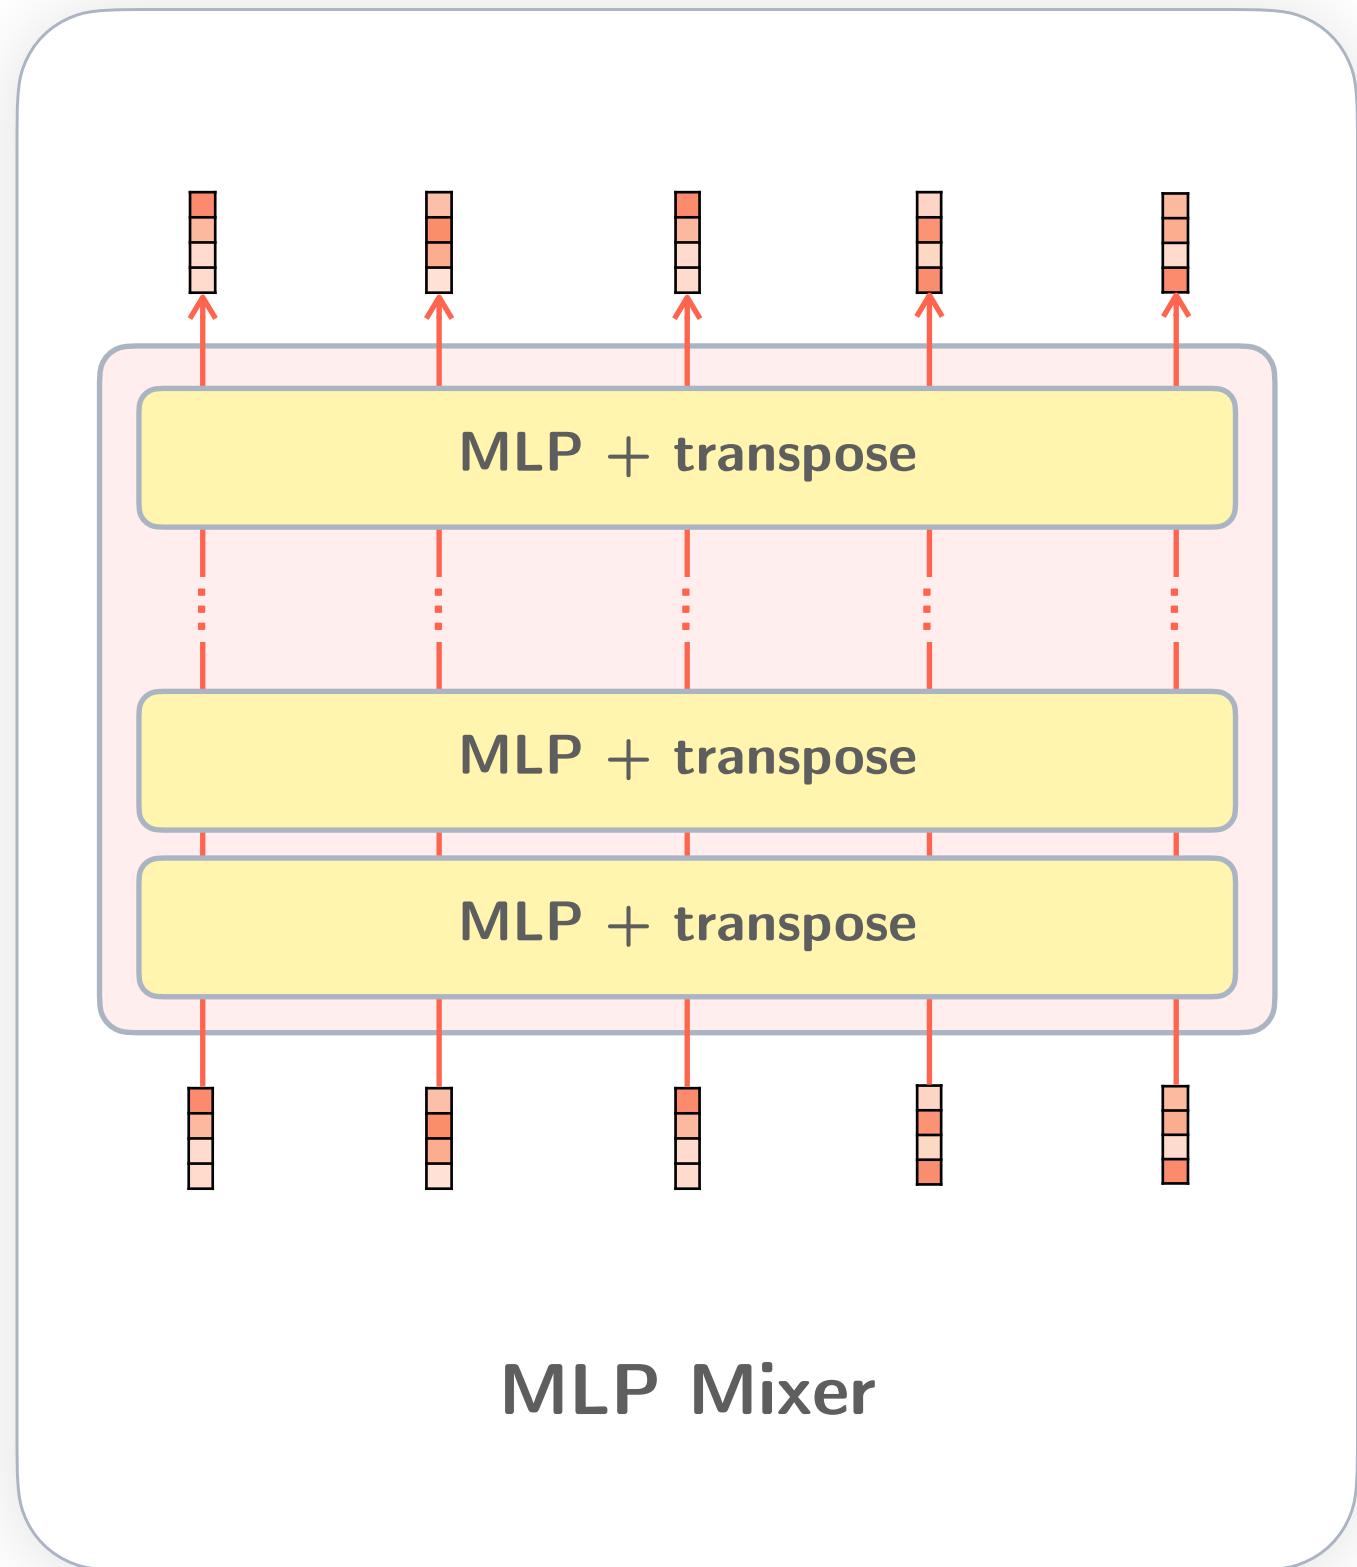


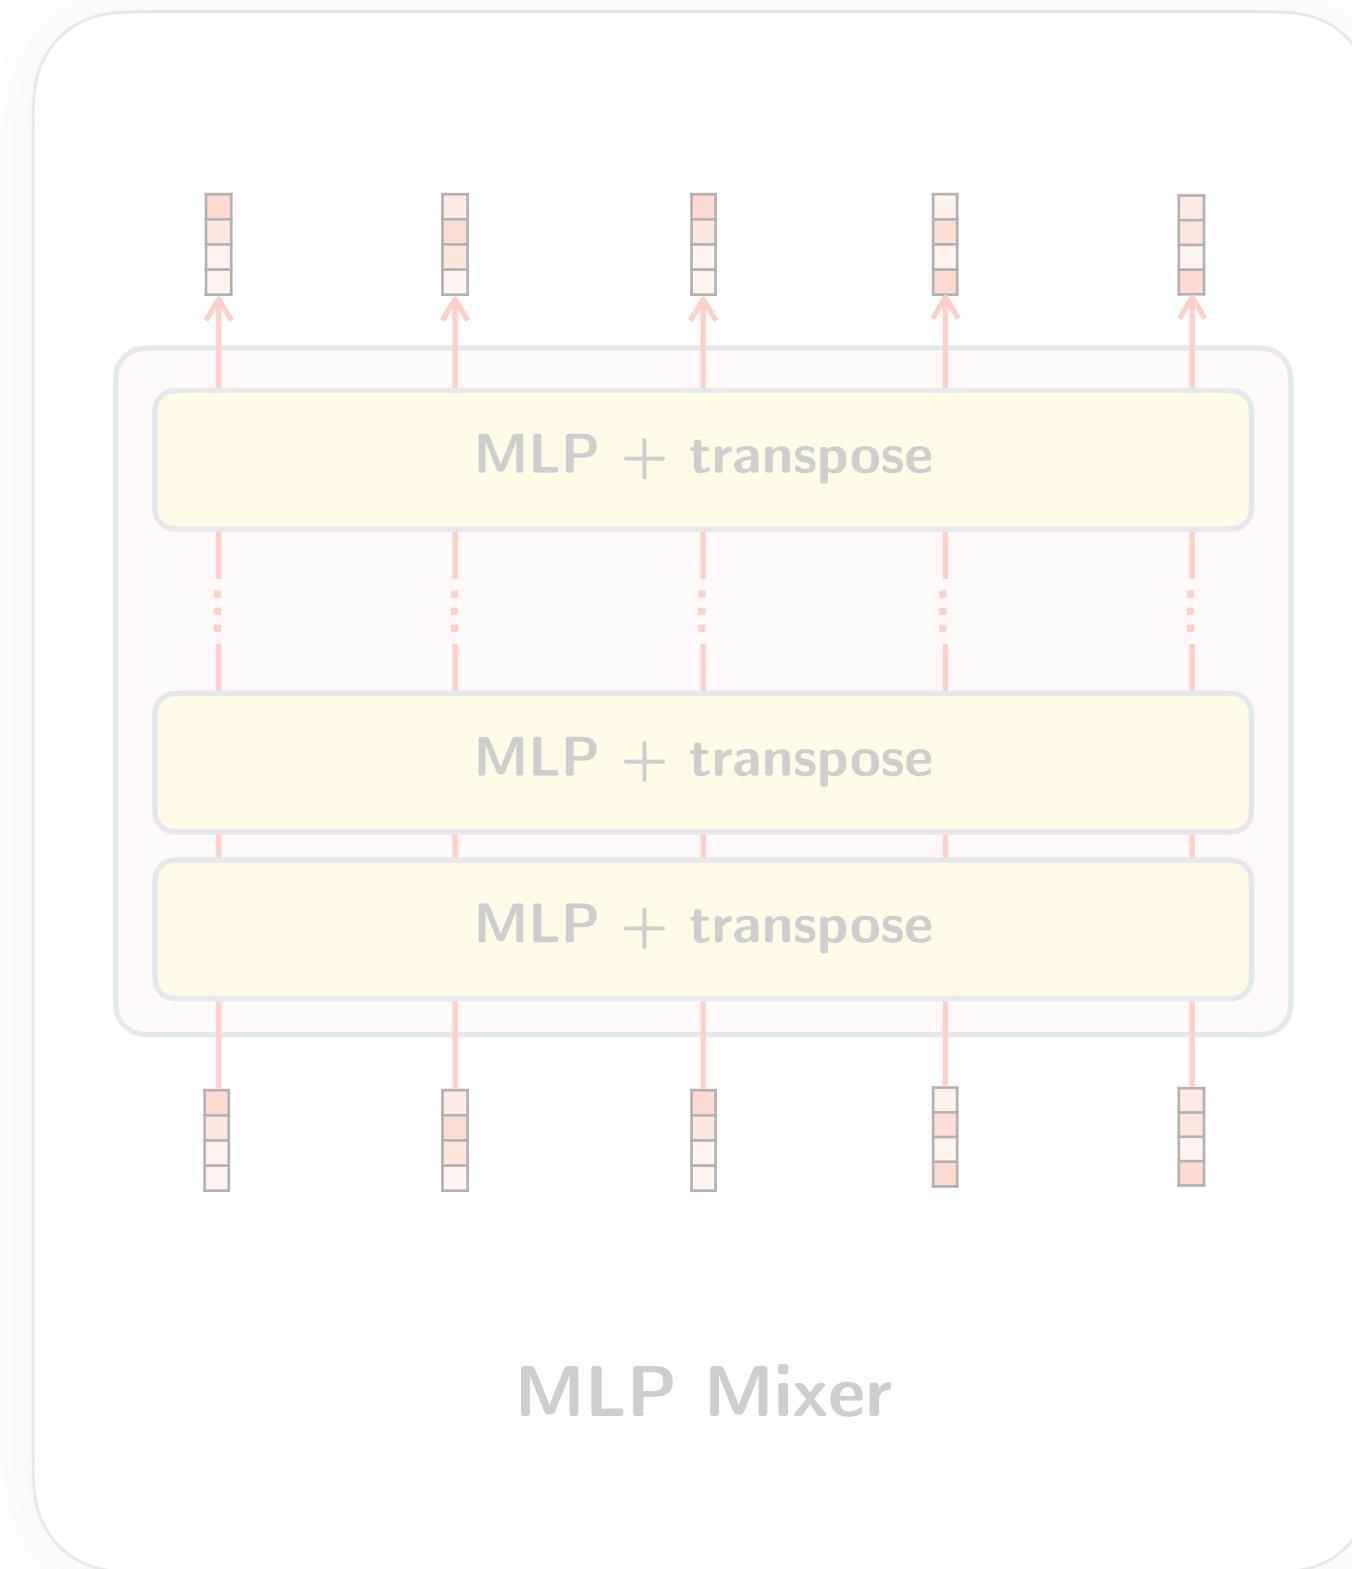
MLP Mixer

FNet

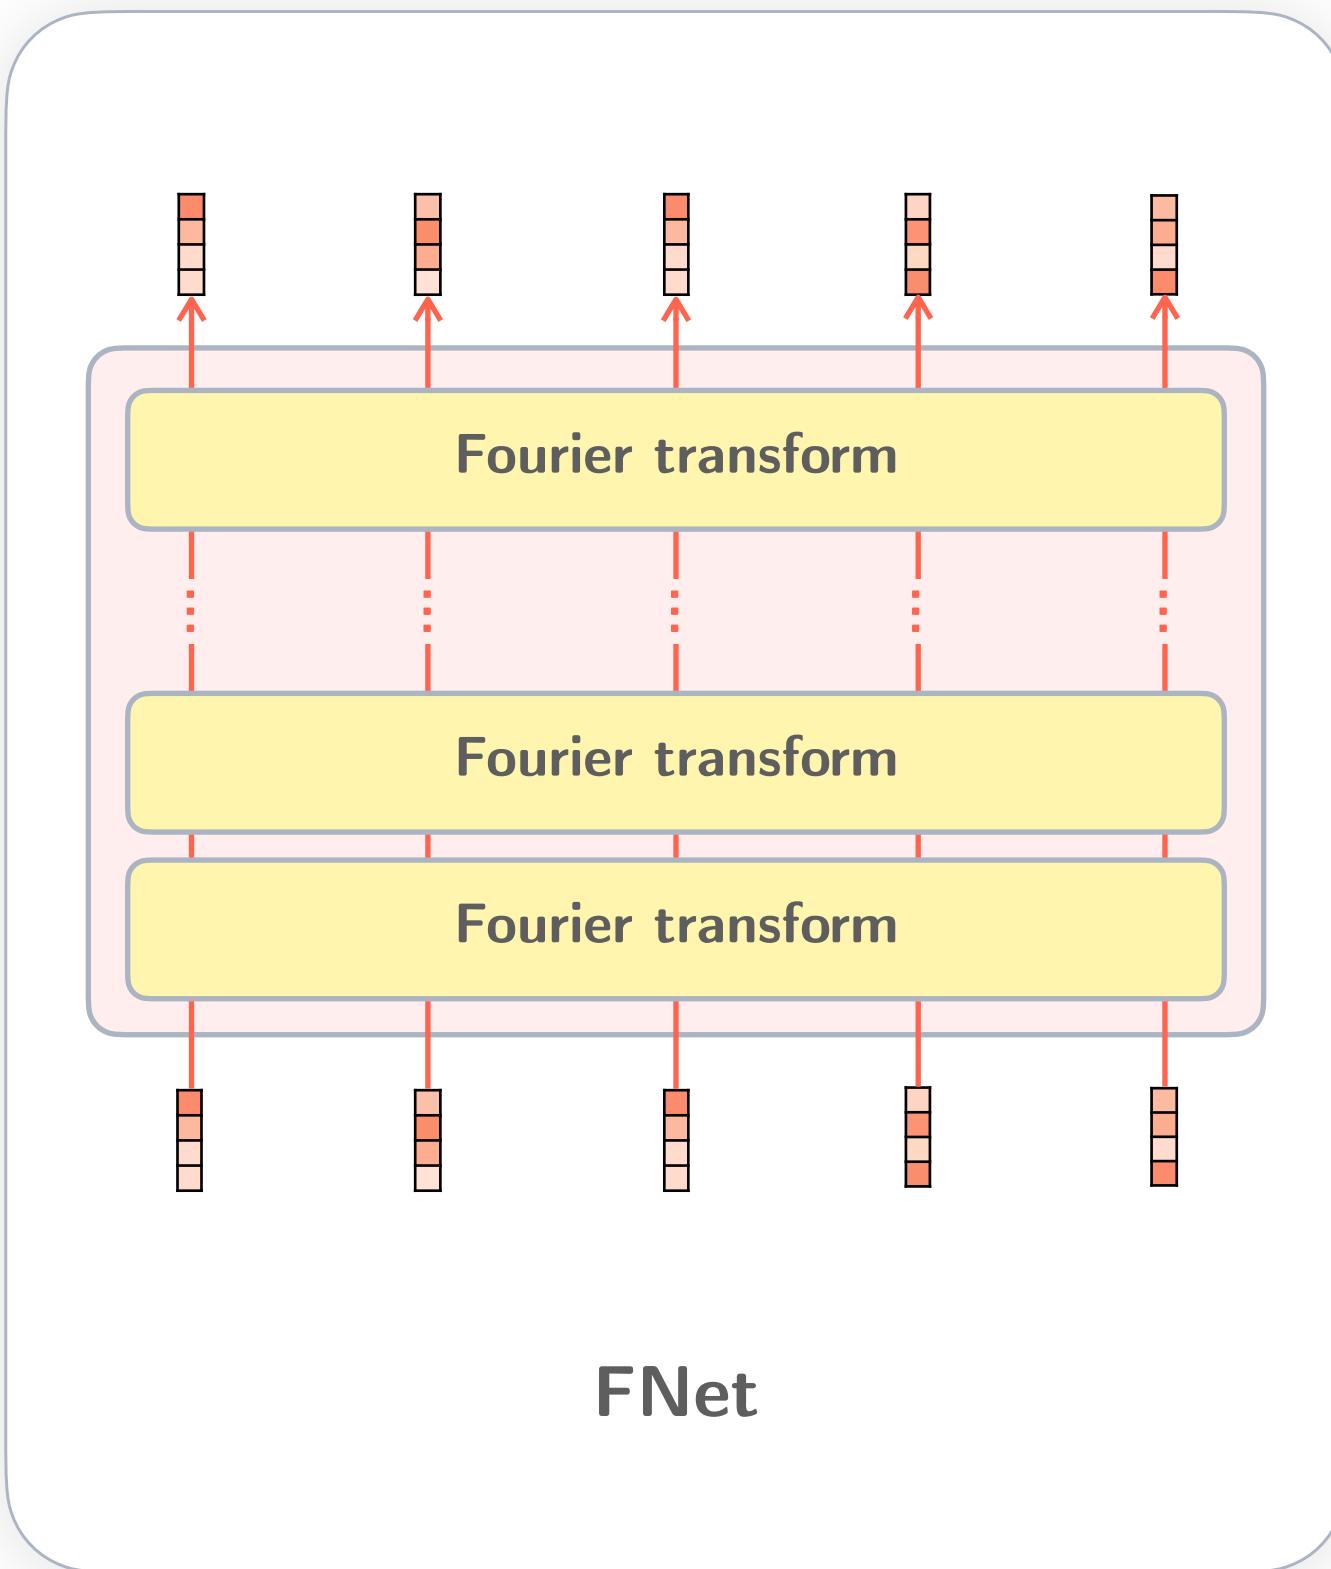
pre-trained convolutions

Any non-recurrent sequence-processing neural network where the sequence representation gets refined with each layer by letting individual time steps learn about all the other time steps?

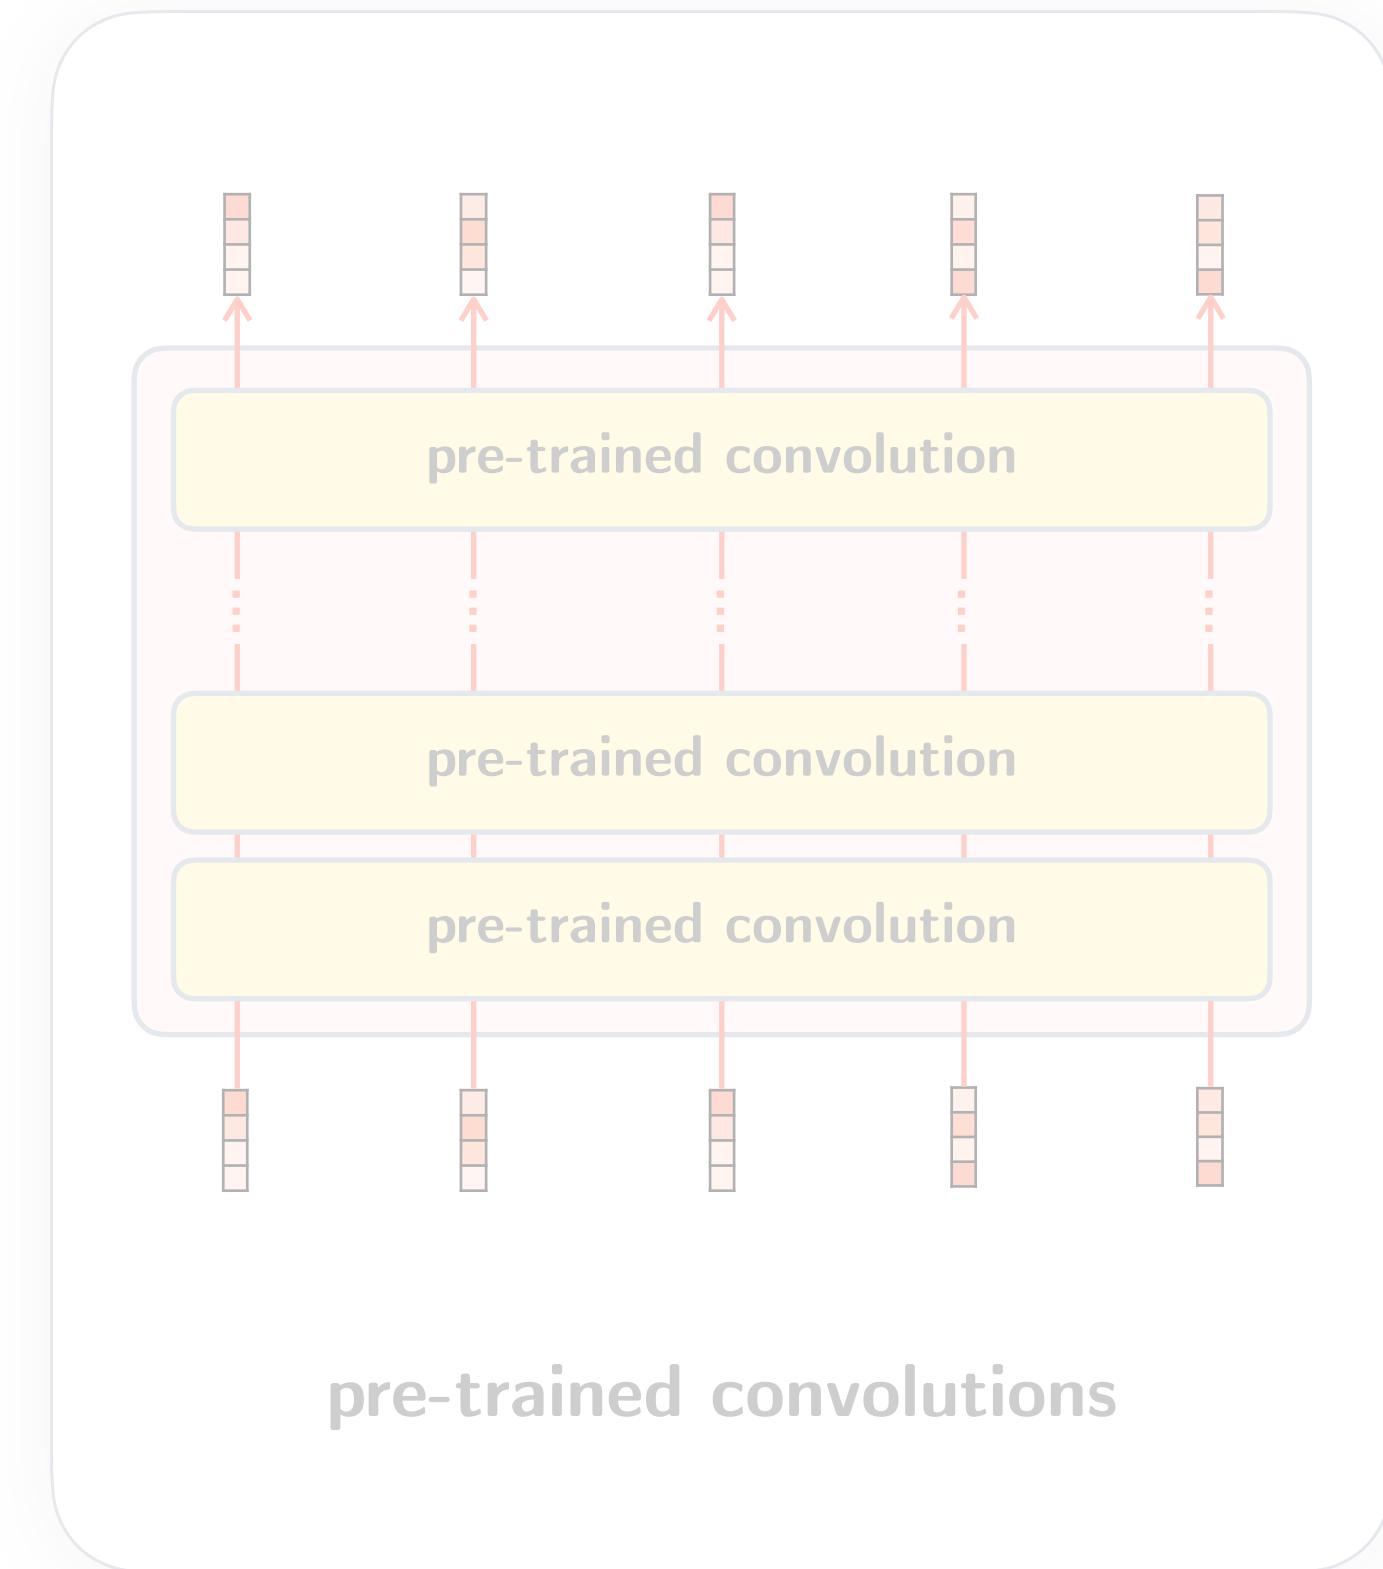




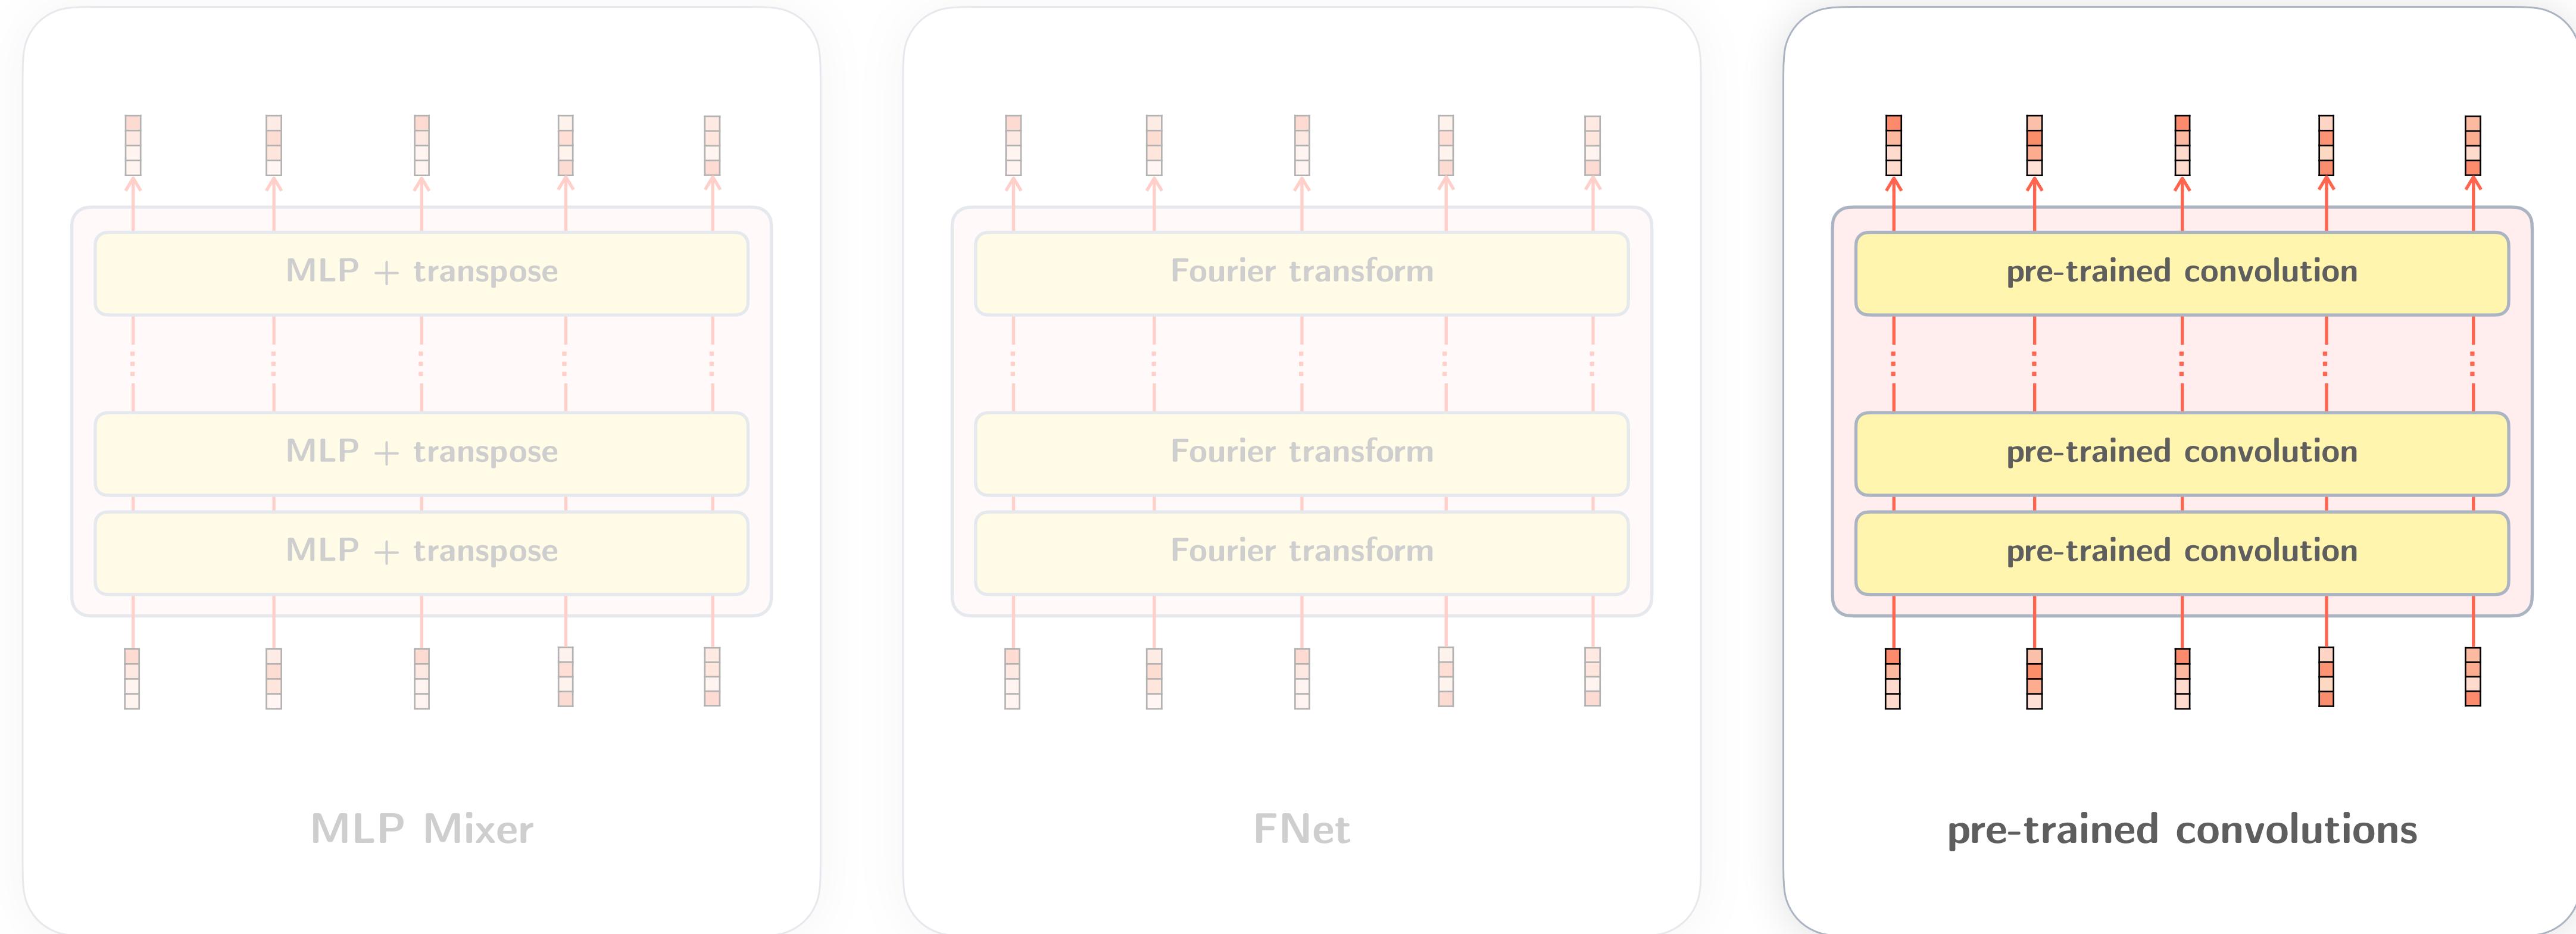
MLP Mixer



FNet

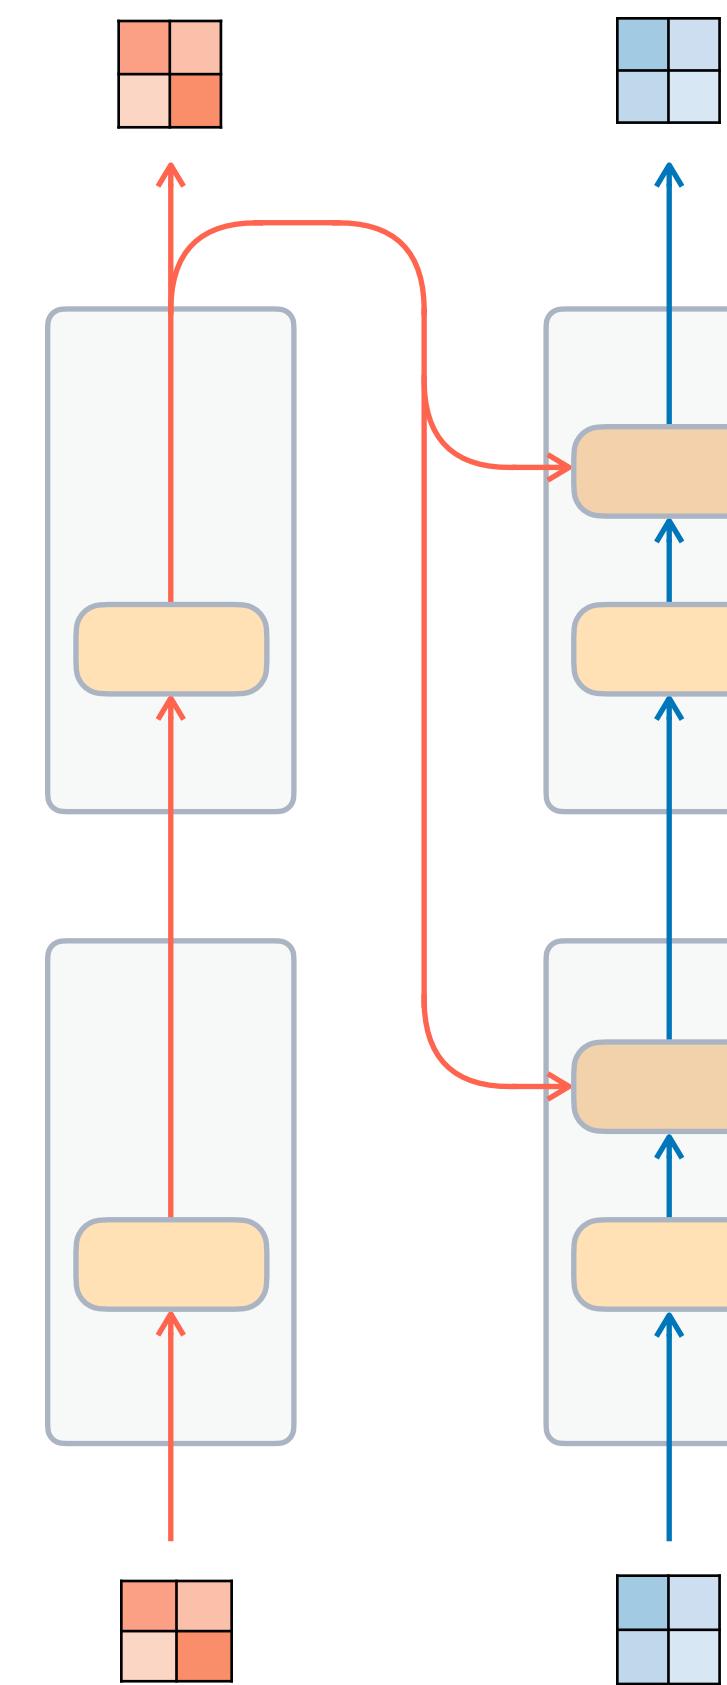


pre-trained convolutions





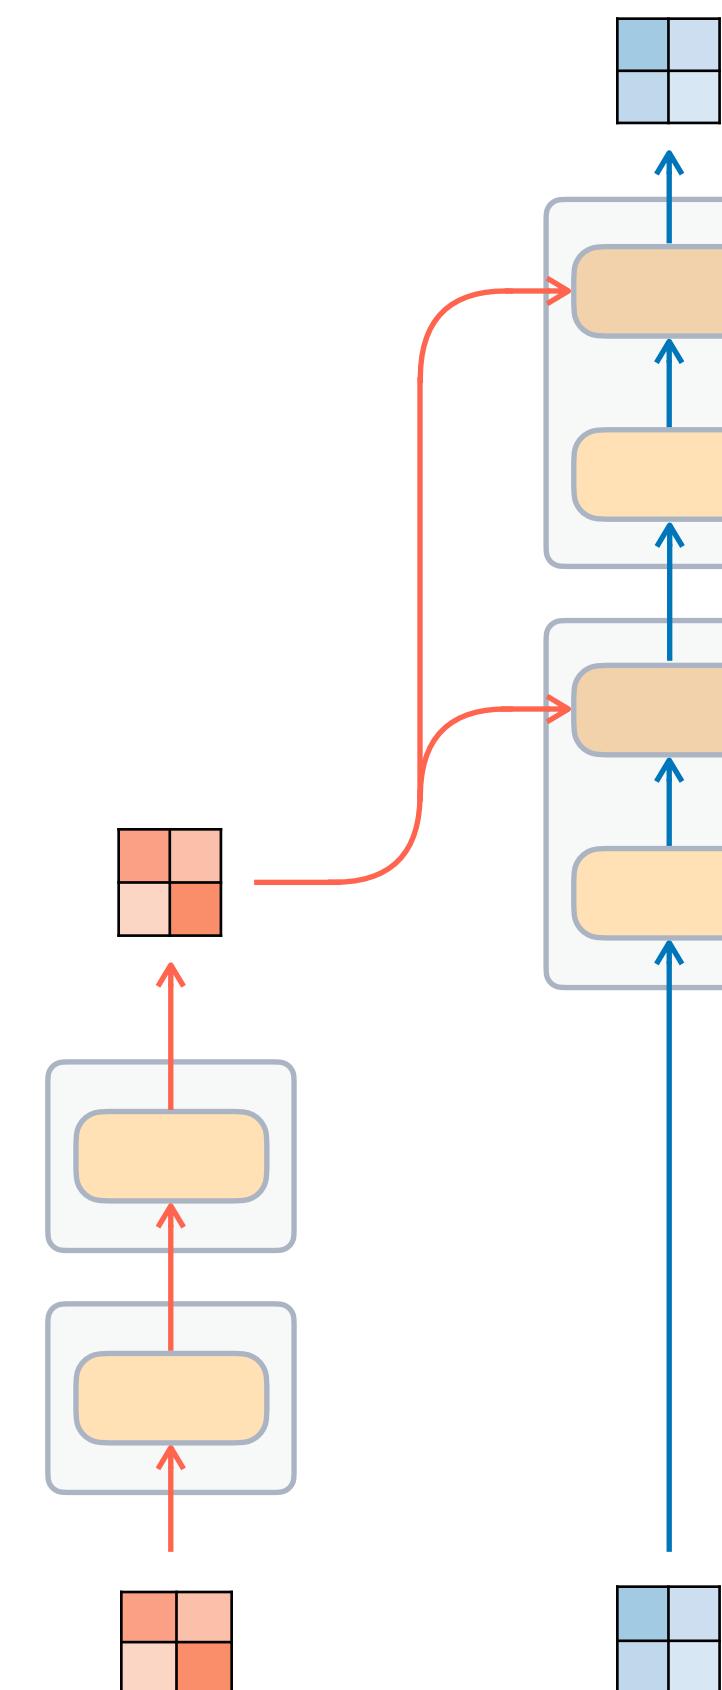
Multi-track transformers



Encoder–Decoder
is not 2-tracked

First, the encoder fully encodes the input into a representation, then, the decoder decodes the output by attending to this representation.

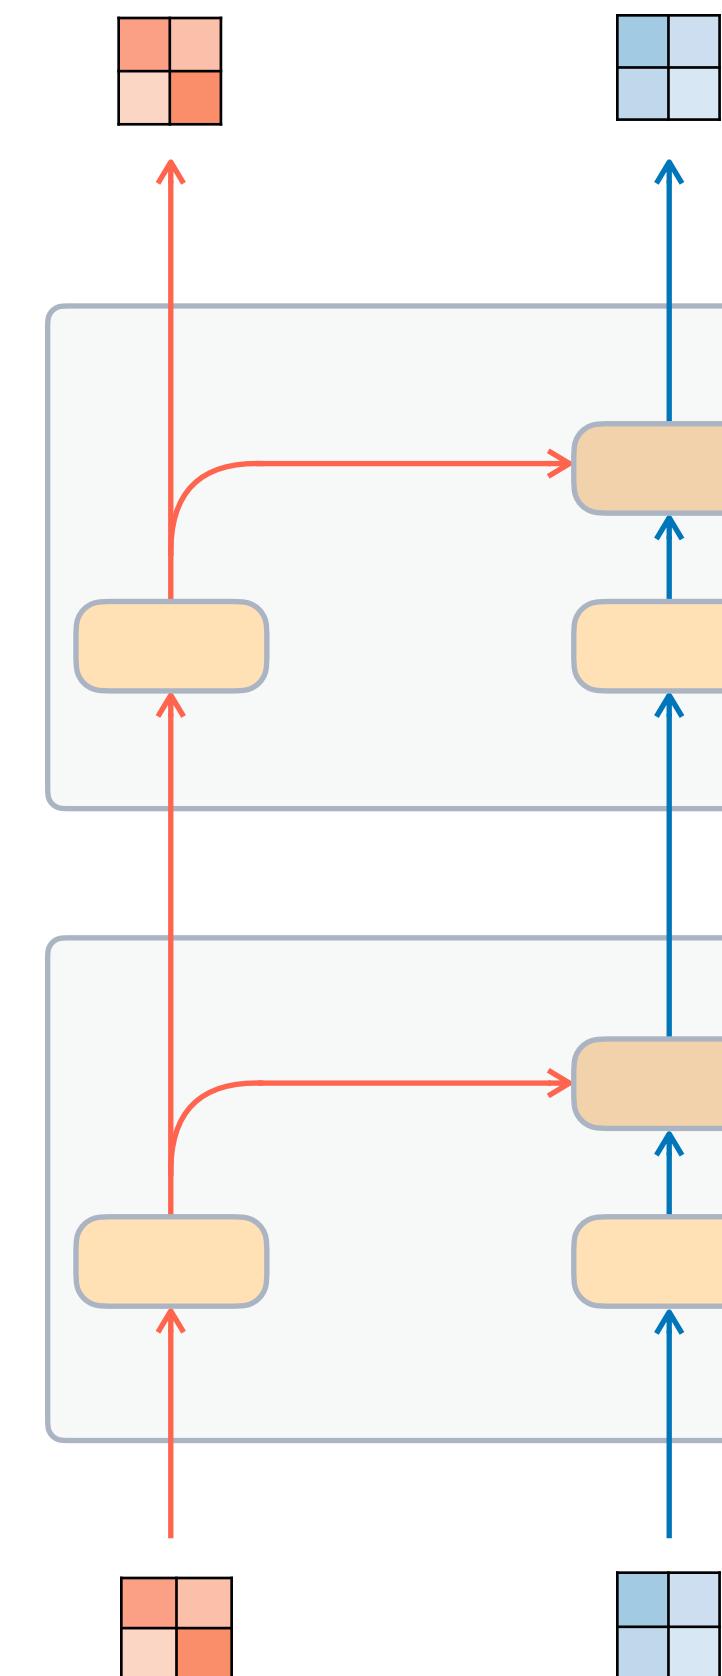
In other words, first, fully run the encoder, *then*, fully run the decoder.



Encoder–Decoder
is not 2-tracked

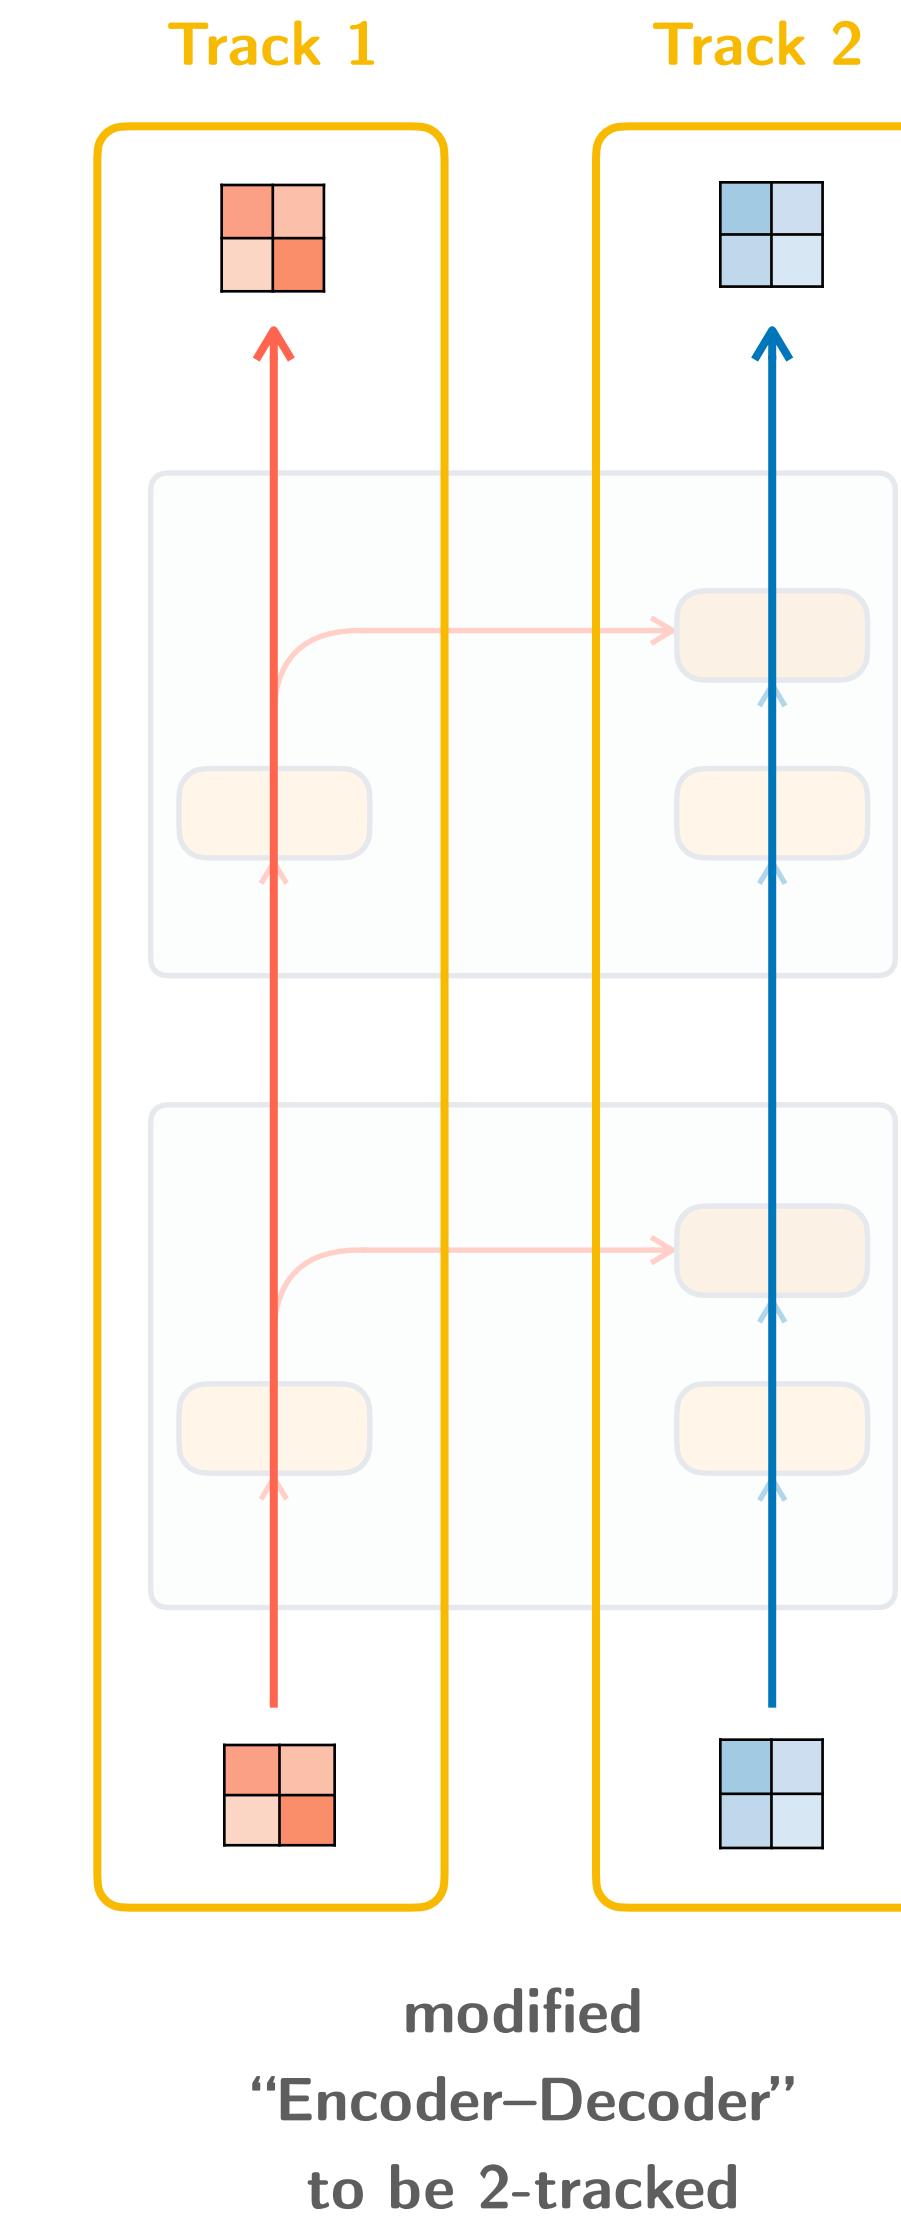
First, the encoder fully encodes the input into a representation, then, the decoder decodes the output by attending to this representation.

In other words, first, fully run the encoder, *then*, fully run the decoder.



modified
“Encoder–Decoder”
to be 2-tracked

In a transformer with multiple tracks,
the tracks communicate in each “block”,
e.g. via cross-attention or biased self-attention.

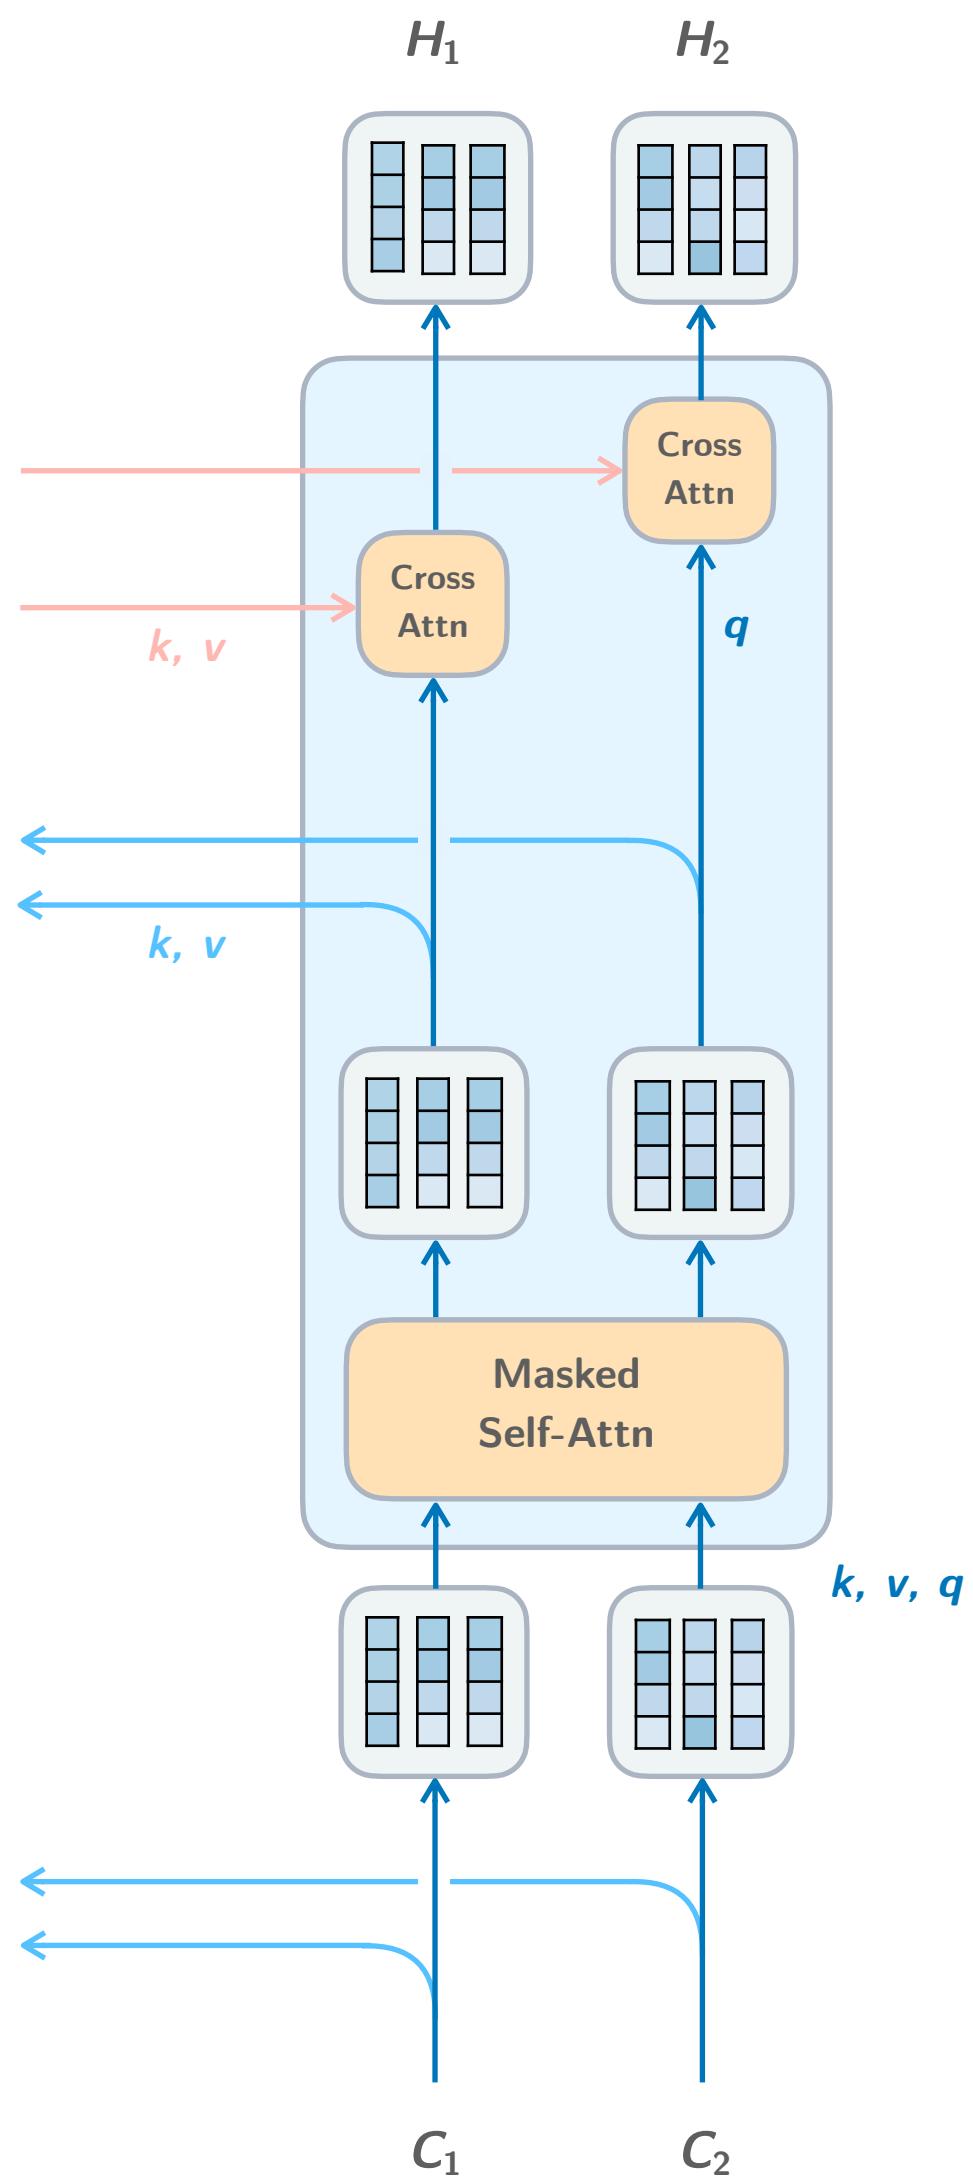


In a transformer with multiple tracks,
the tracks communicate in each “block”,
e.g. via cross-attention or biased self-attention.

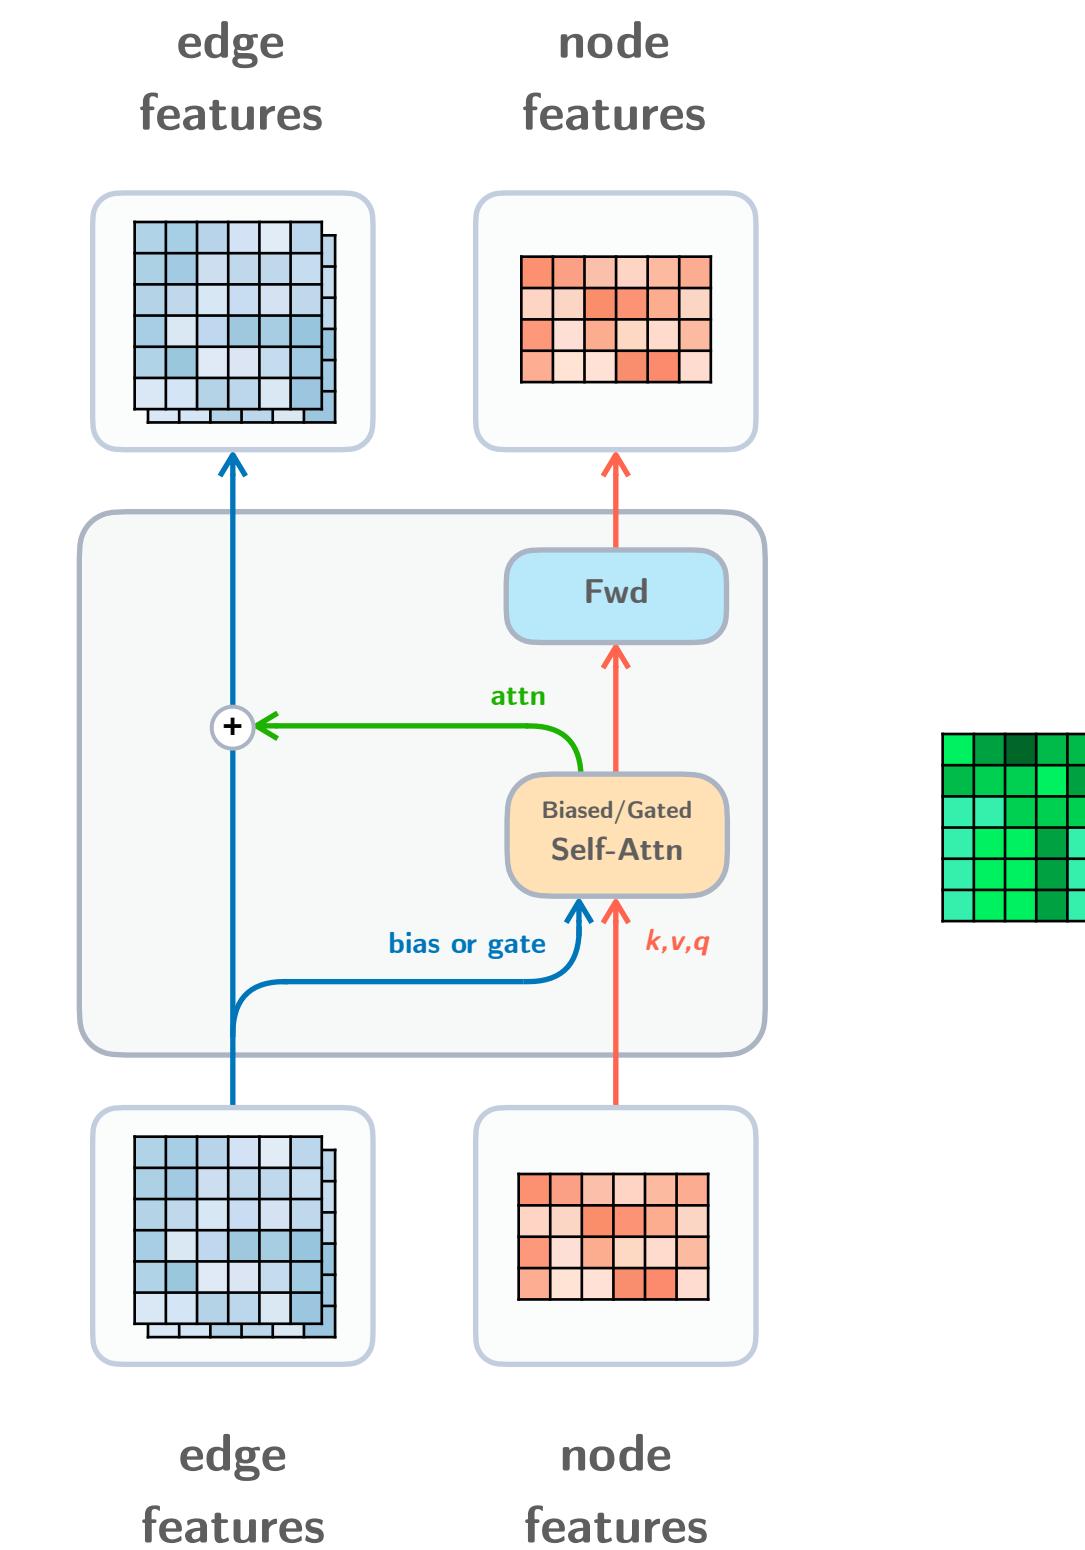
Interpreted as a N -track transformer

Each track maintains and refines a representation of one chunk.

Each track works the same.

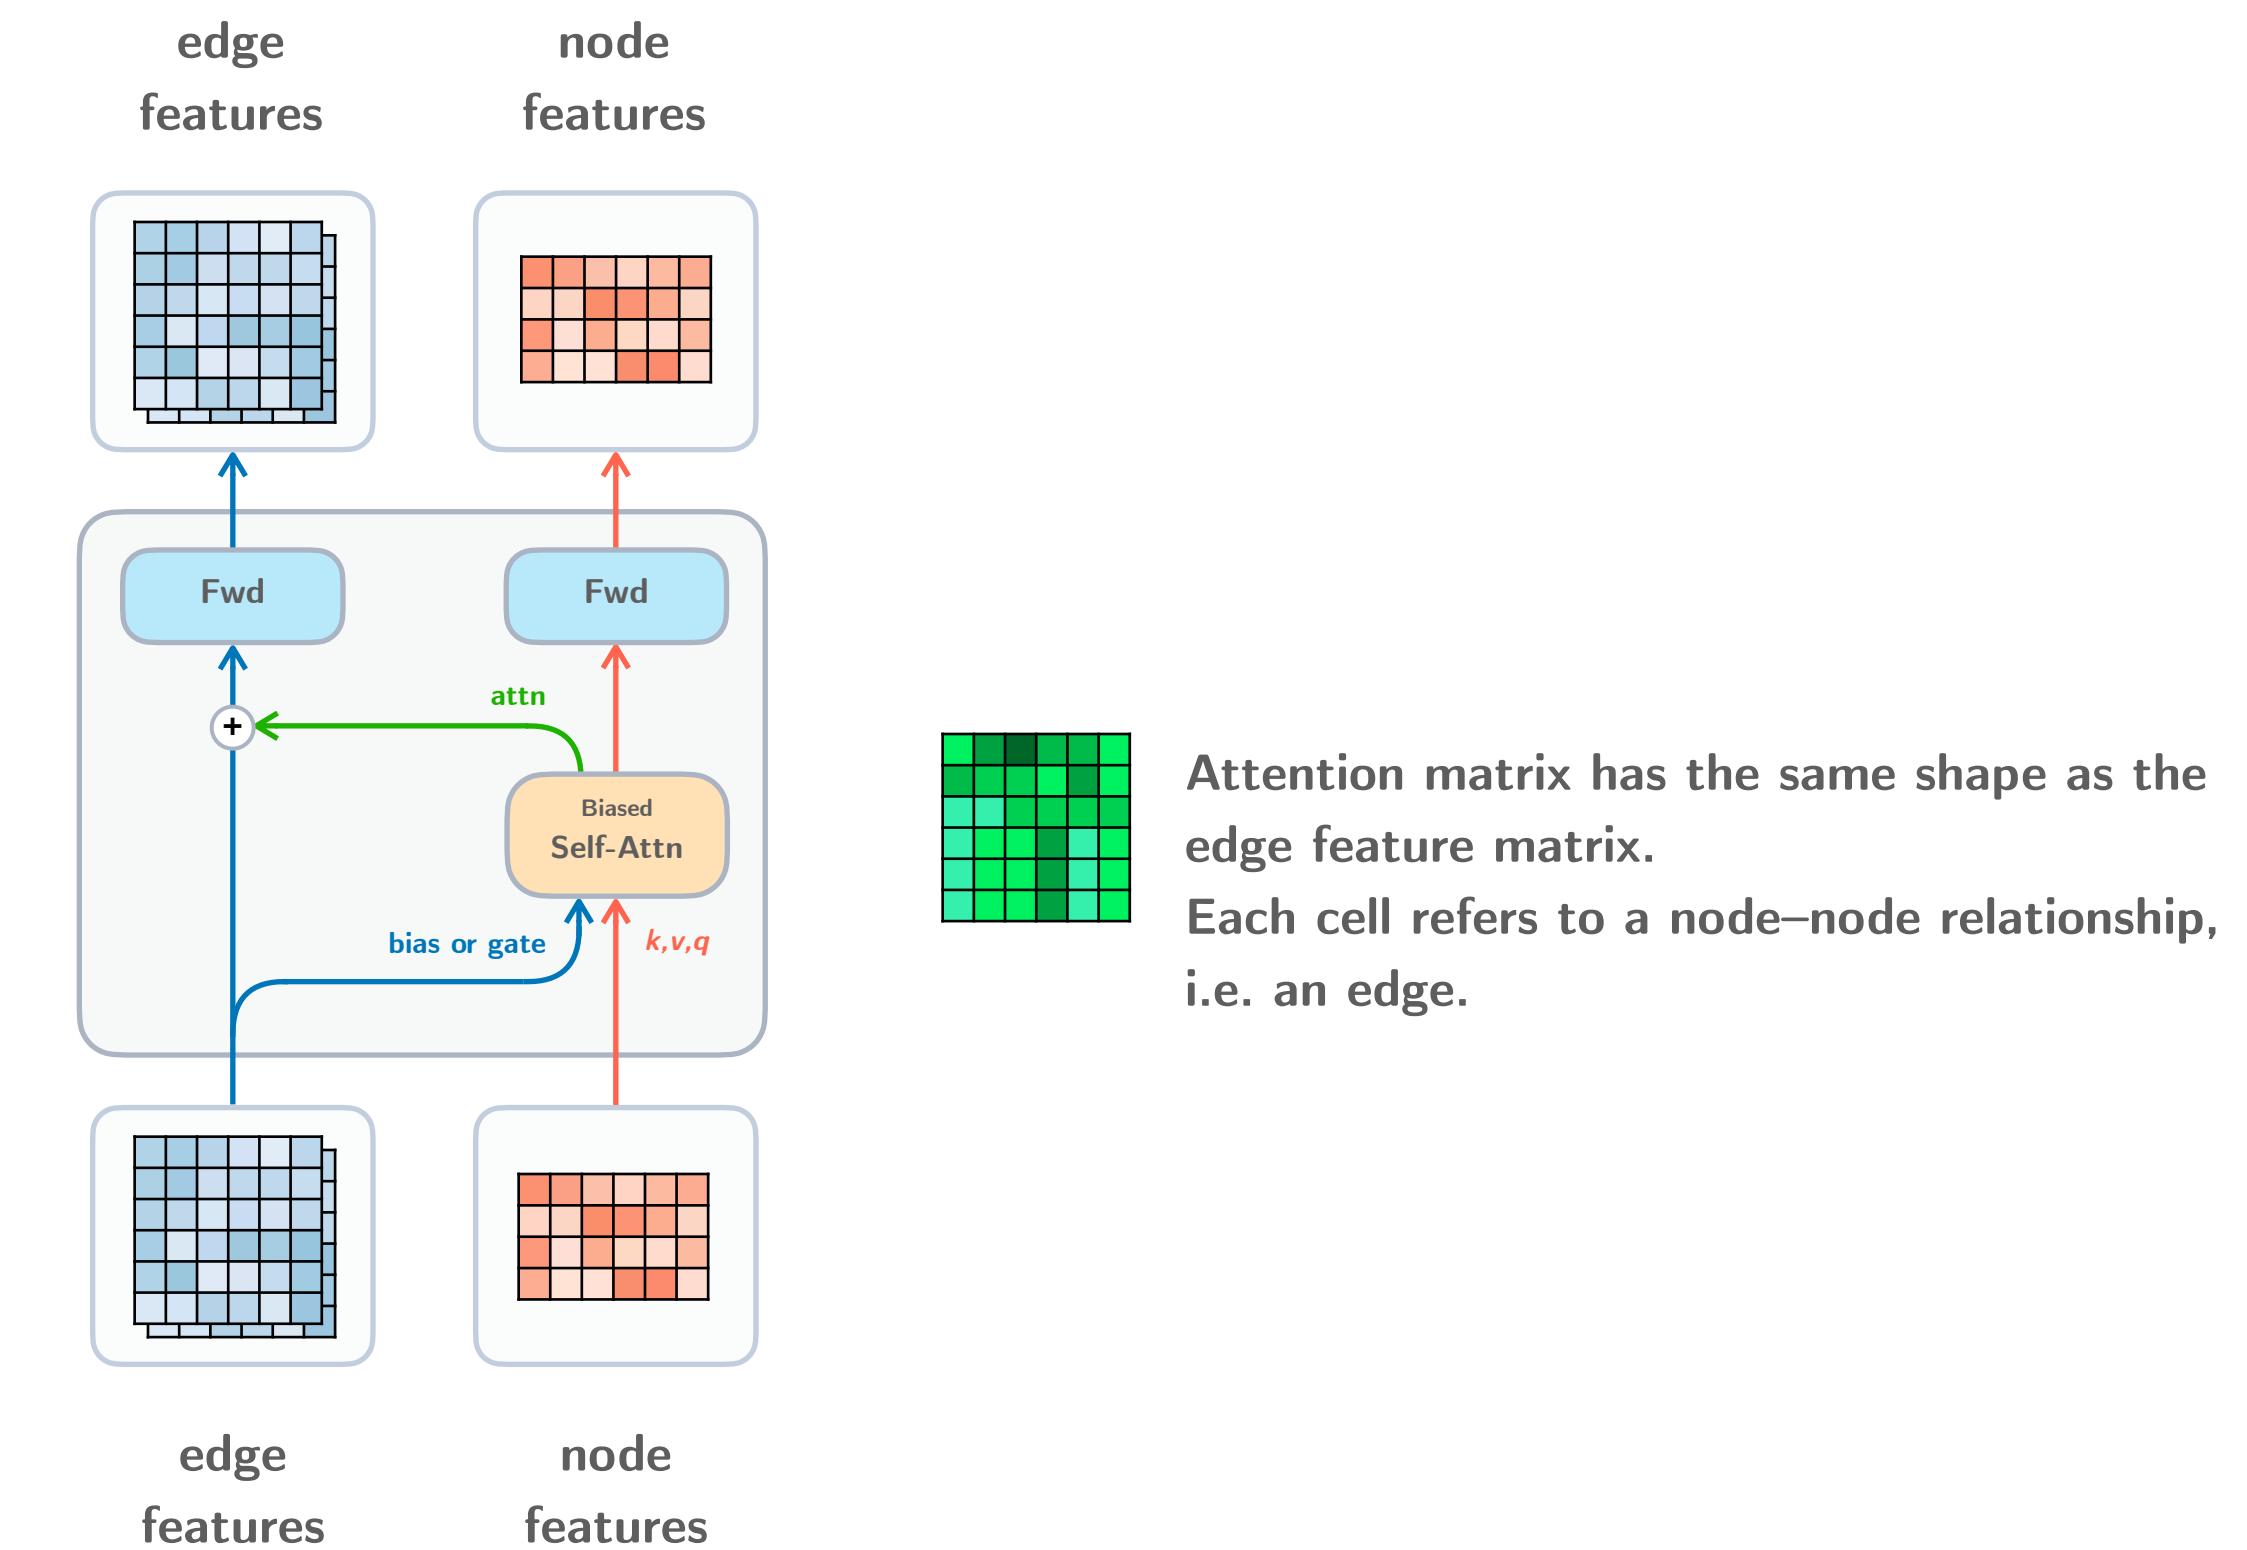


Interpreted as a 2-track transformer
to maintain and refine *node* and *edge* representations
Different tracks operate on different representations.



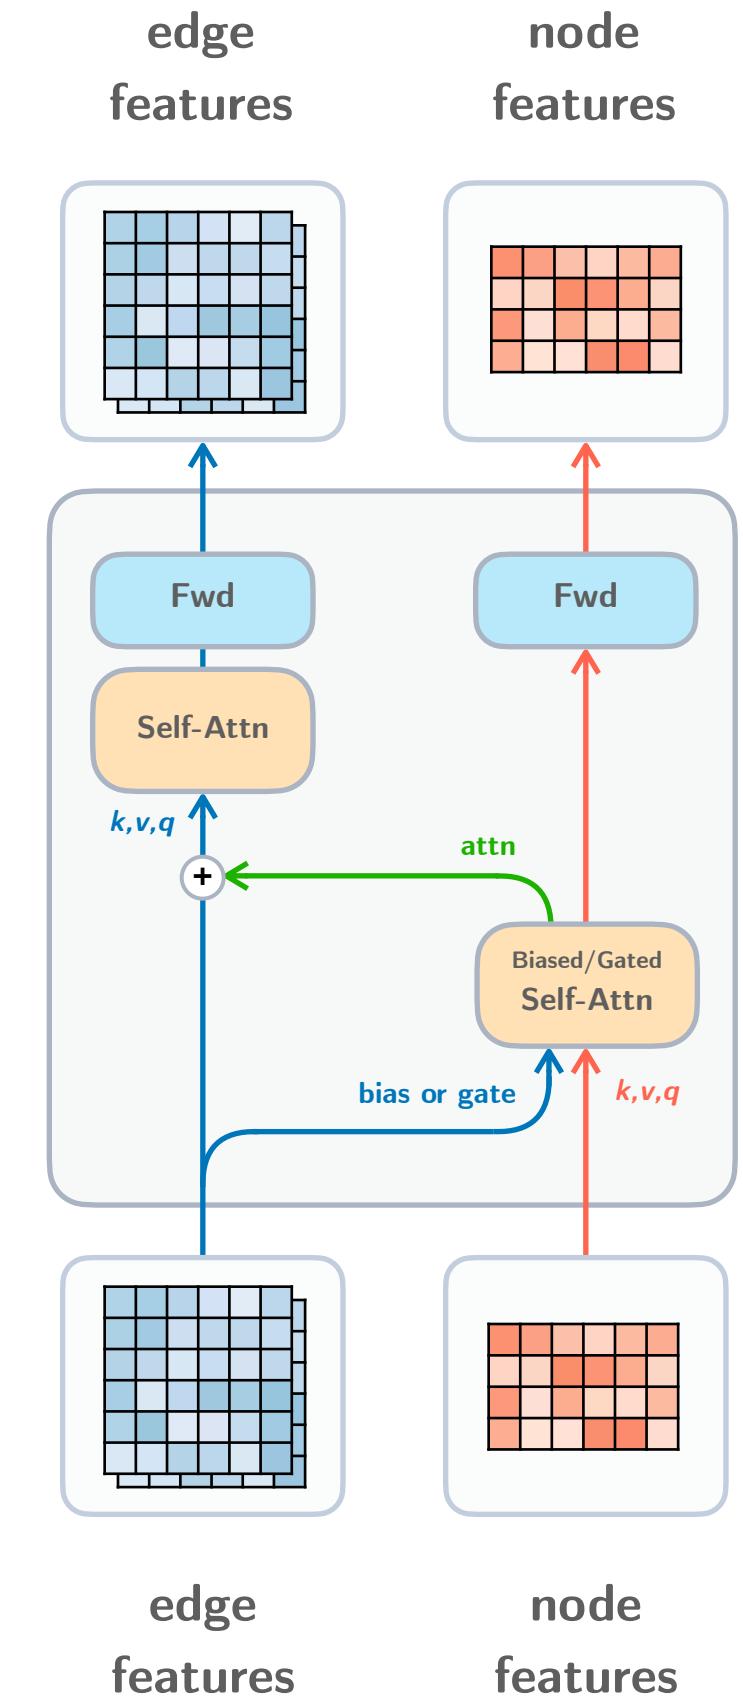
Attention matrix has the same shape as the edge feature matrix.
Each cell refers to a node–node relationship, i.e. an edge.

Interpreted as a 2-track transformer
to maintain and refine *node* and *edge* representations
Different tracks operate on different representations.

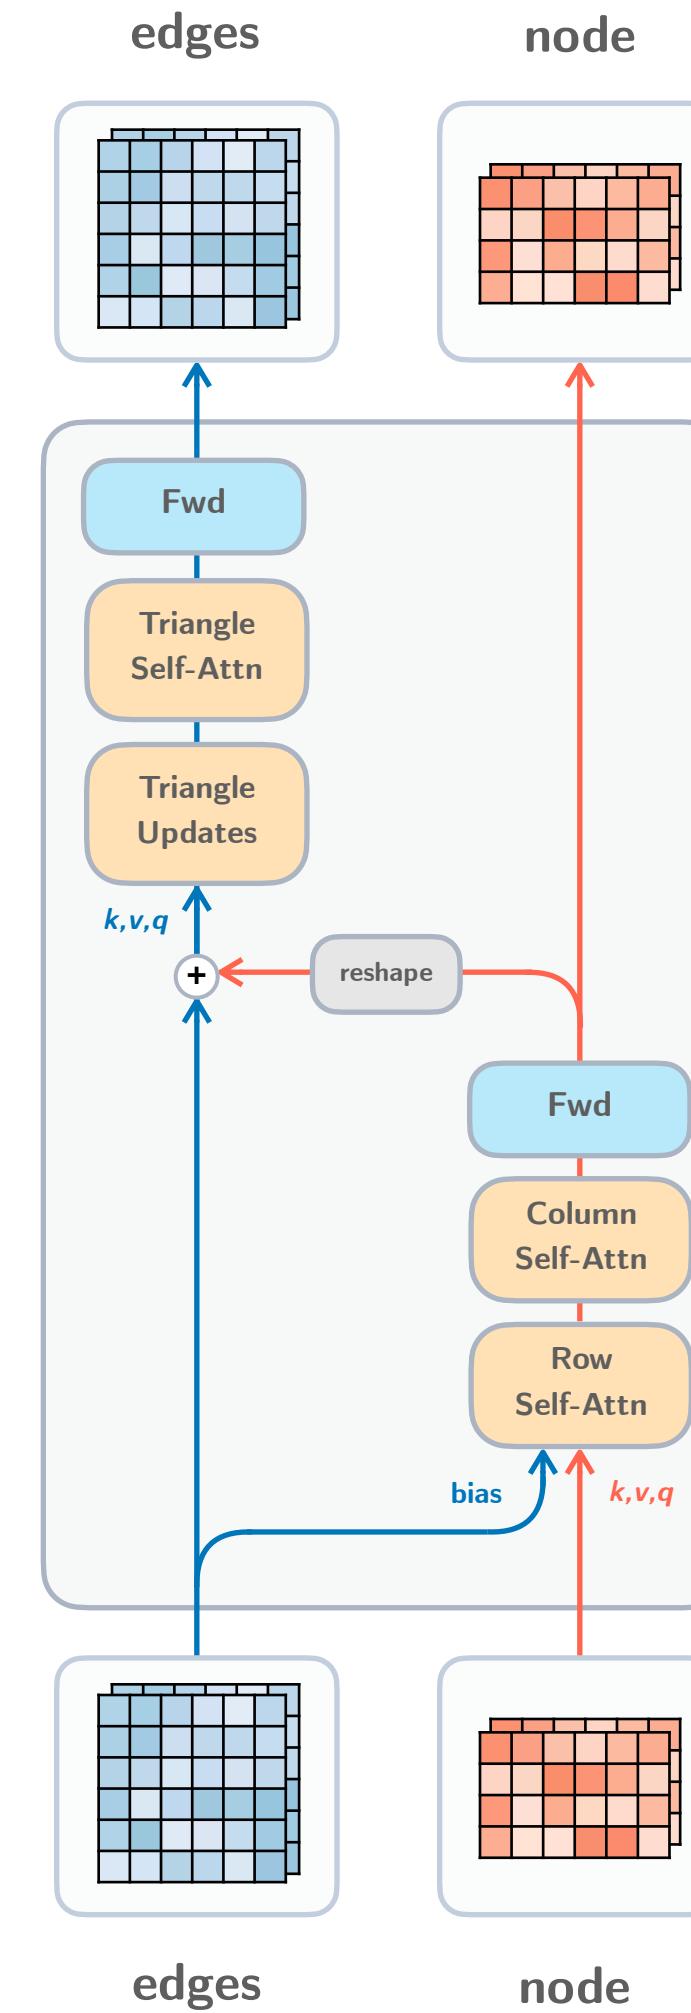


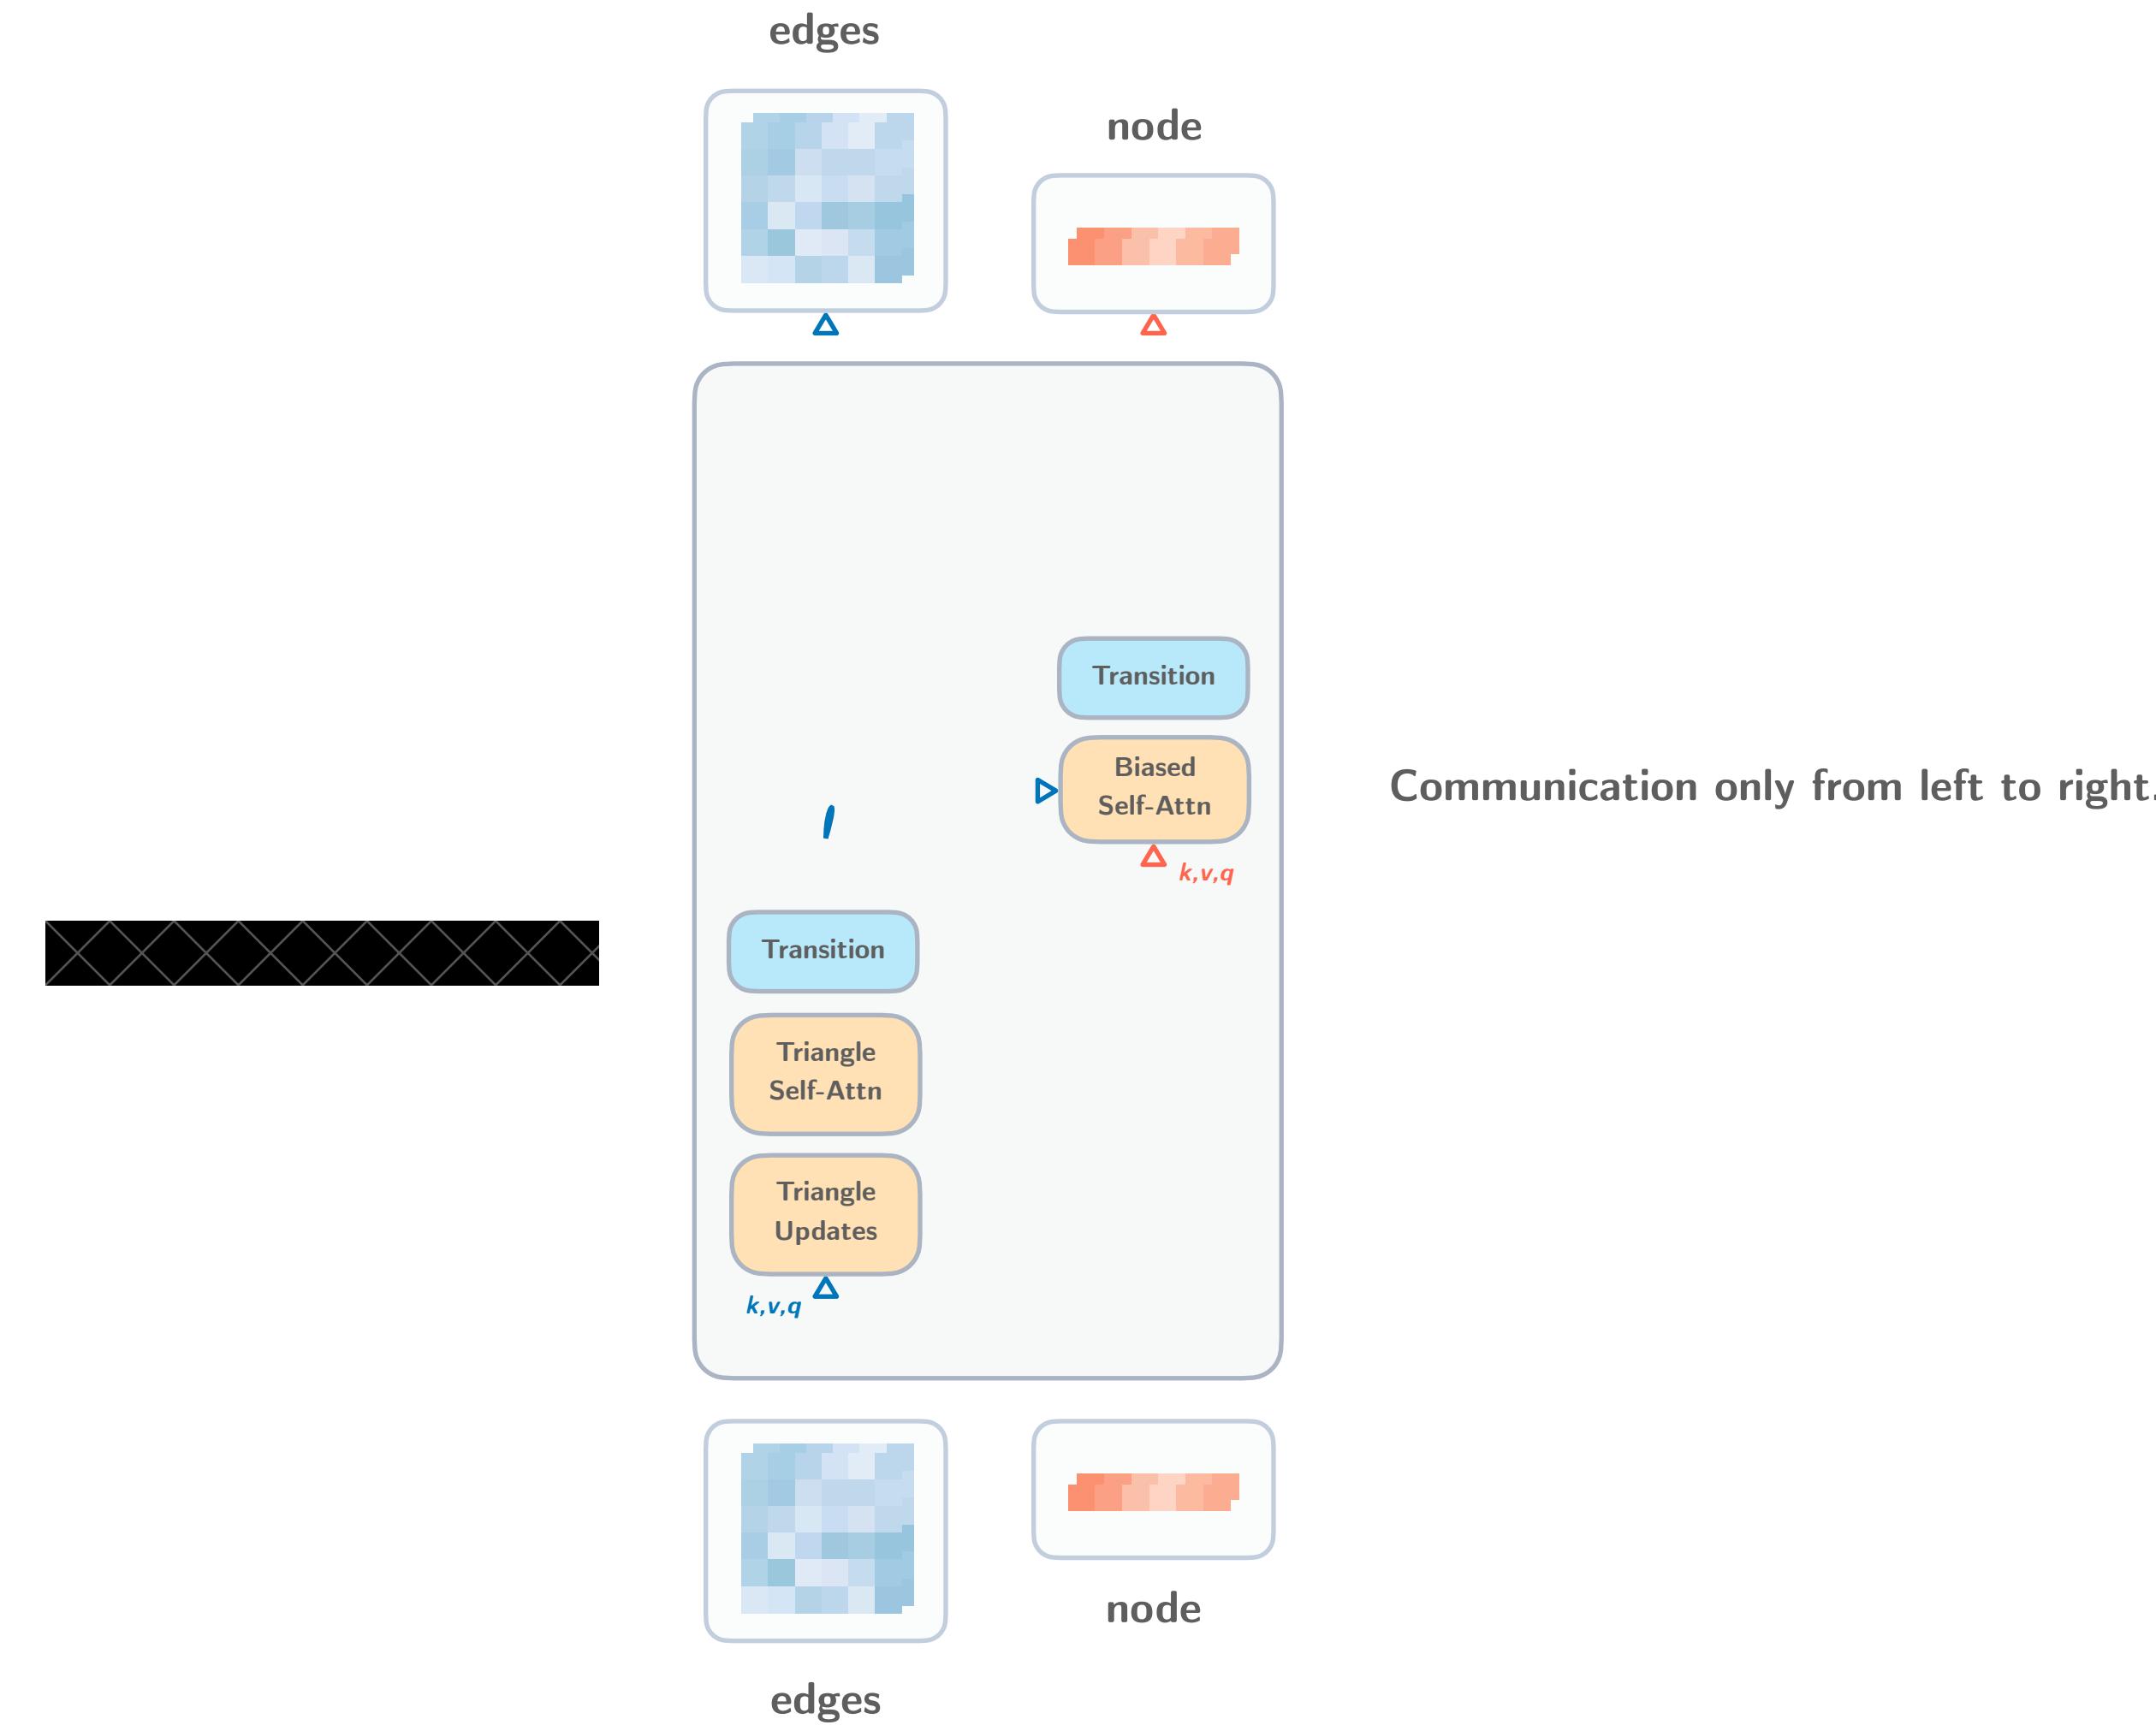
Interpreted as a 2-track transformer
to maintain and refine *node* and *edge* representations
Different tracks operate on different representations.

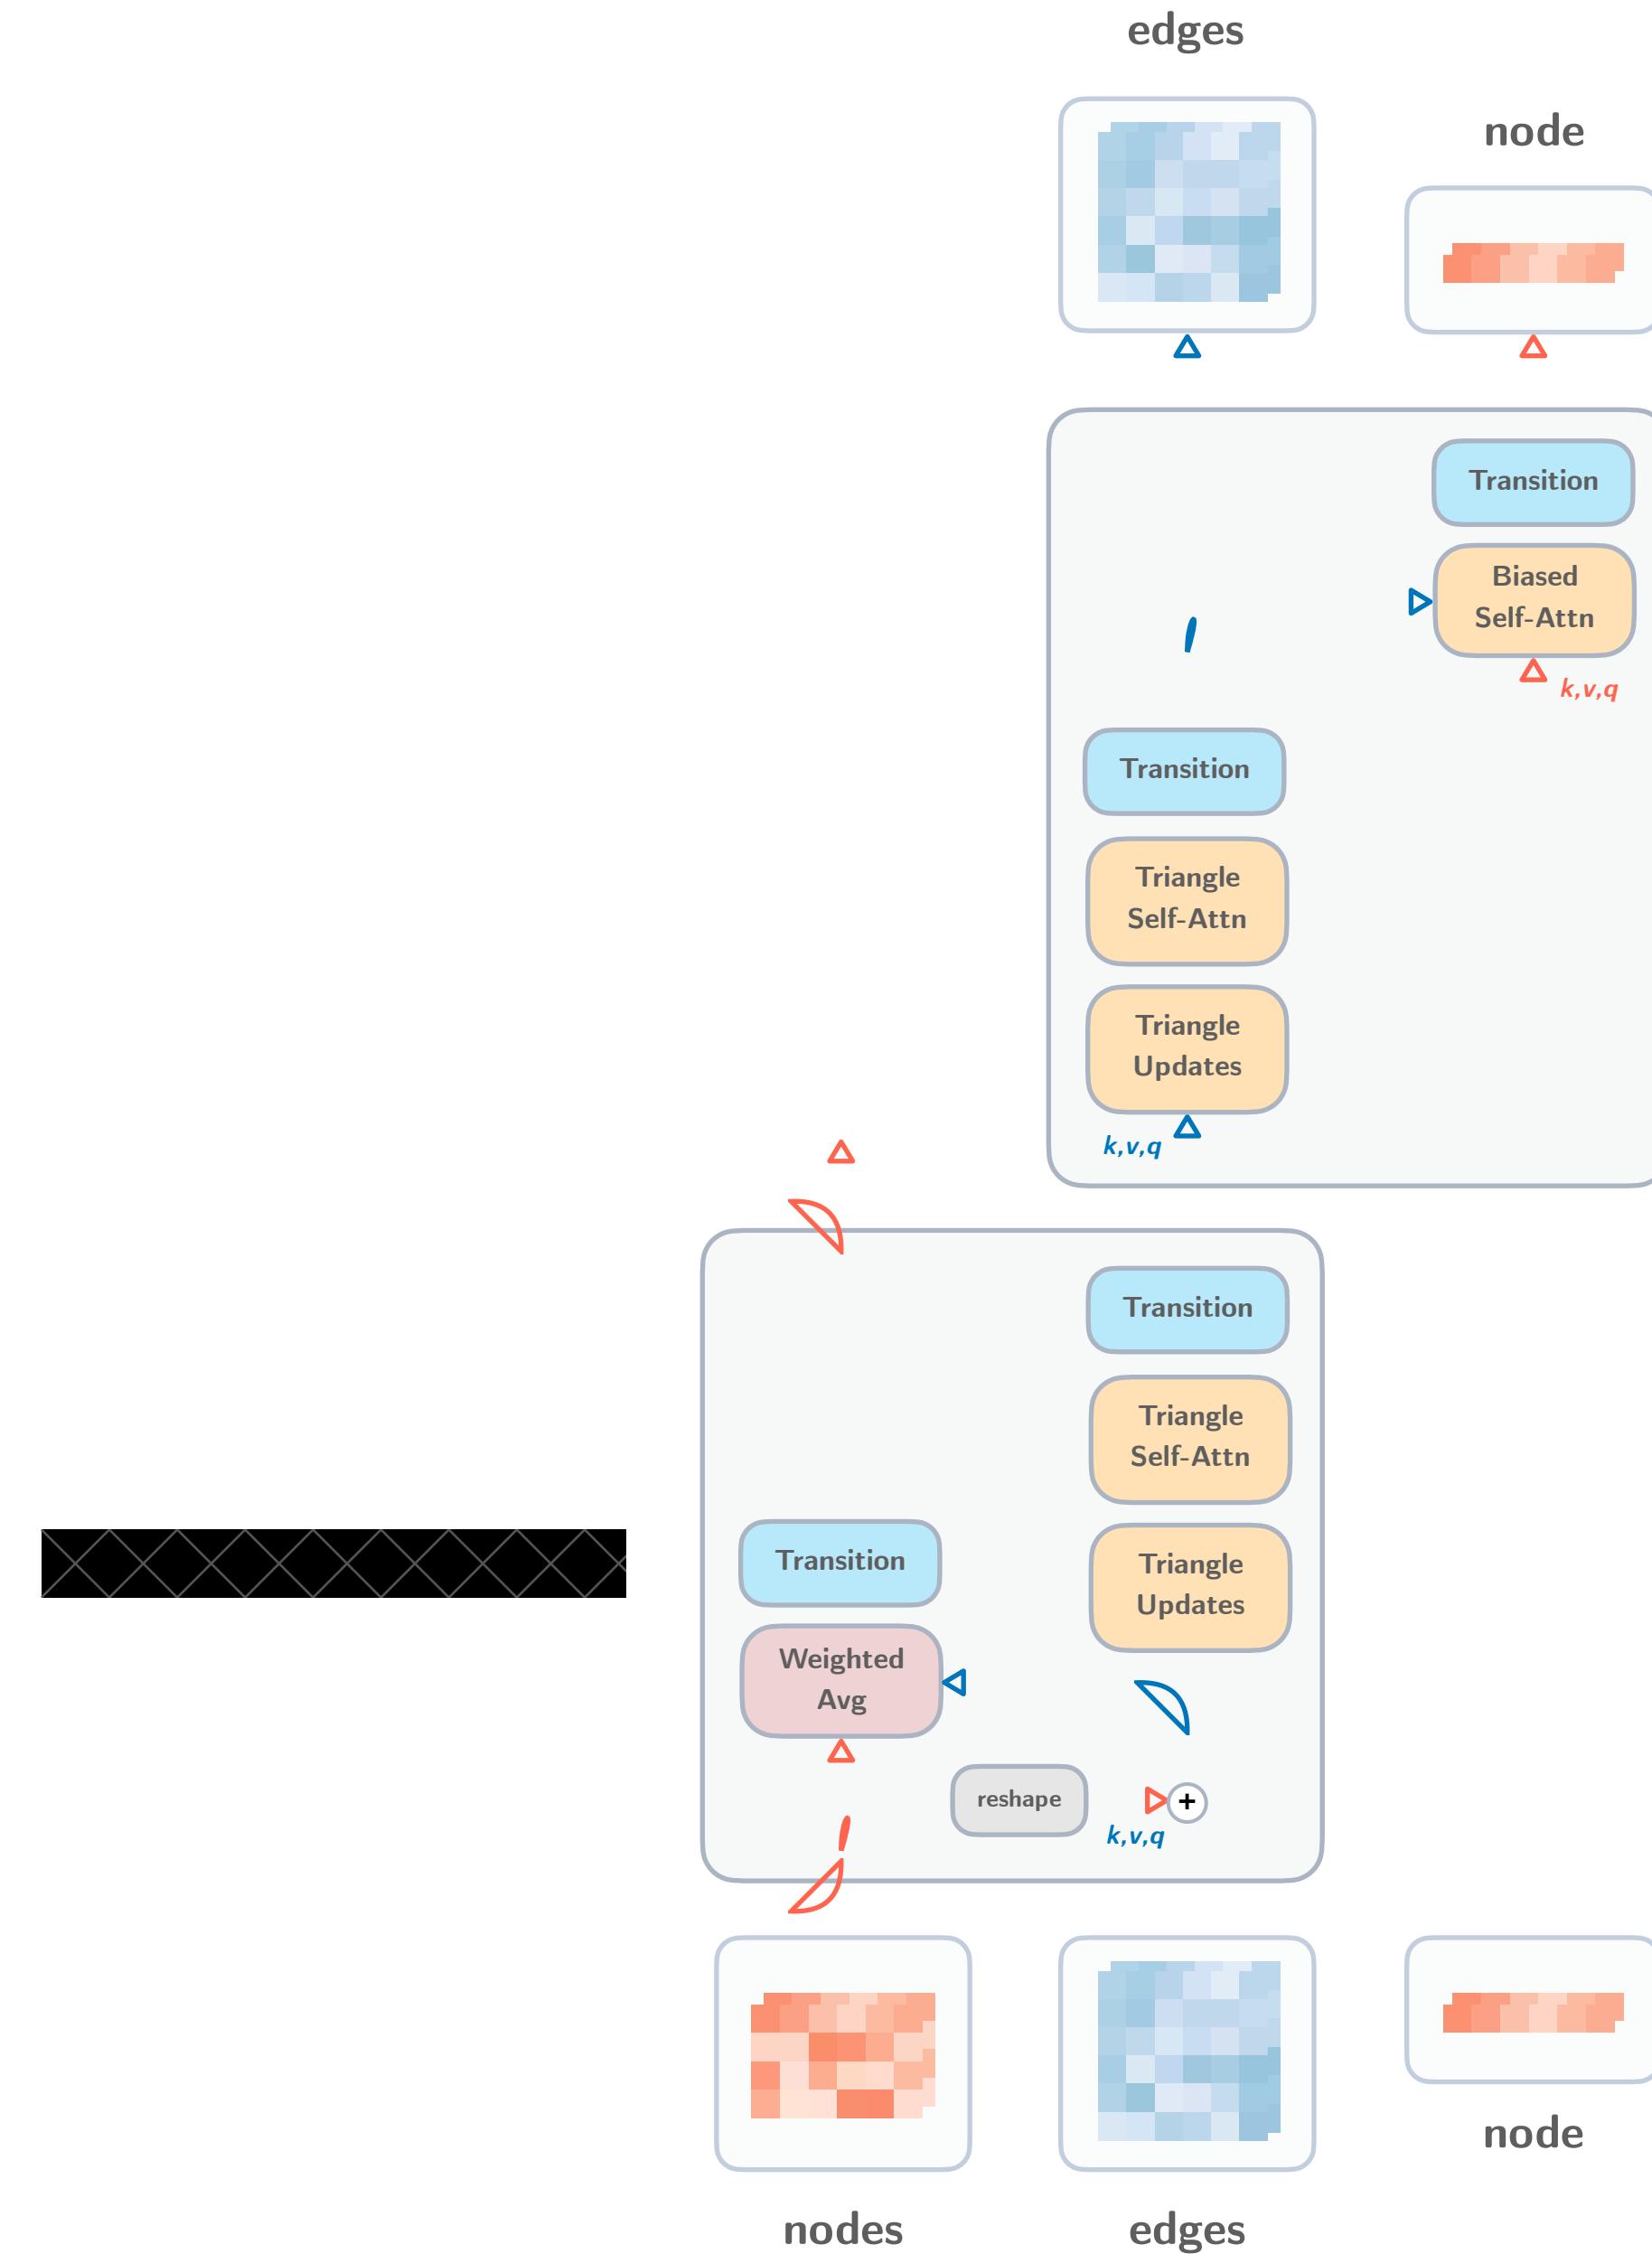
You could also do self-attention to the
edge features directly.

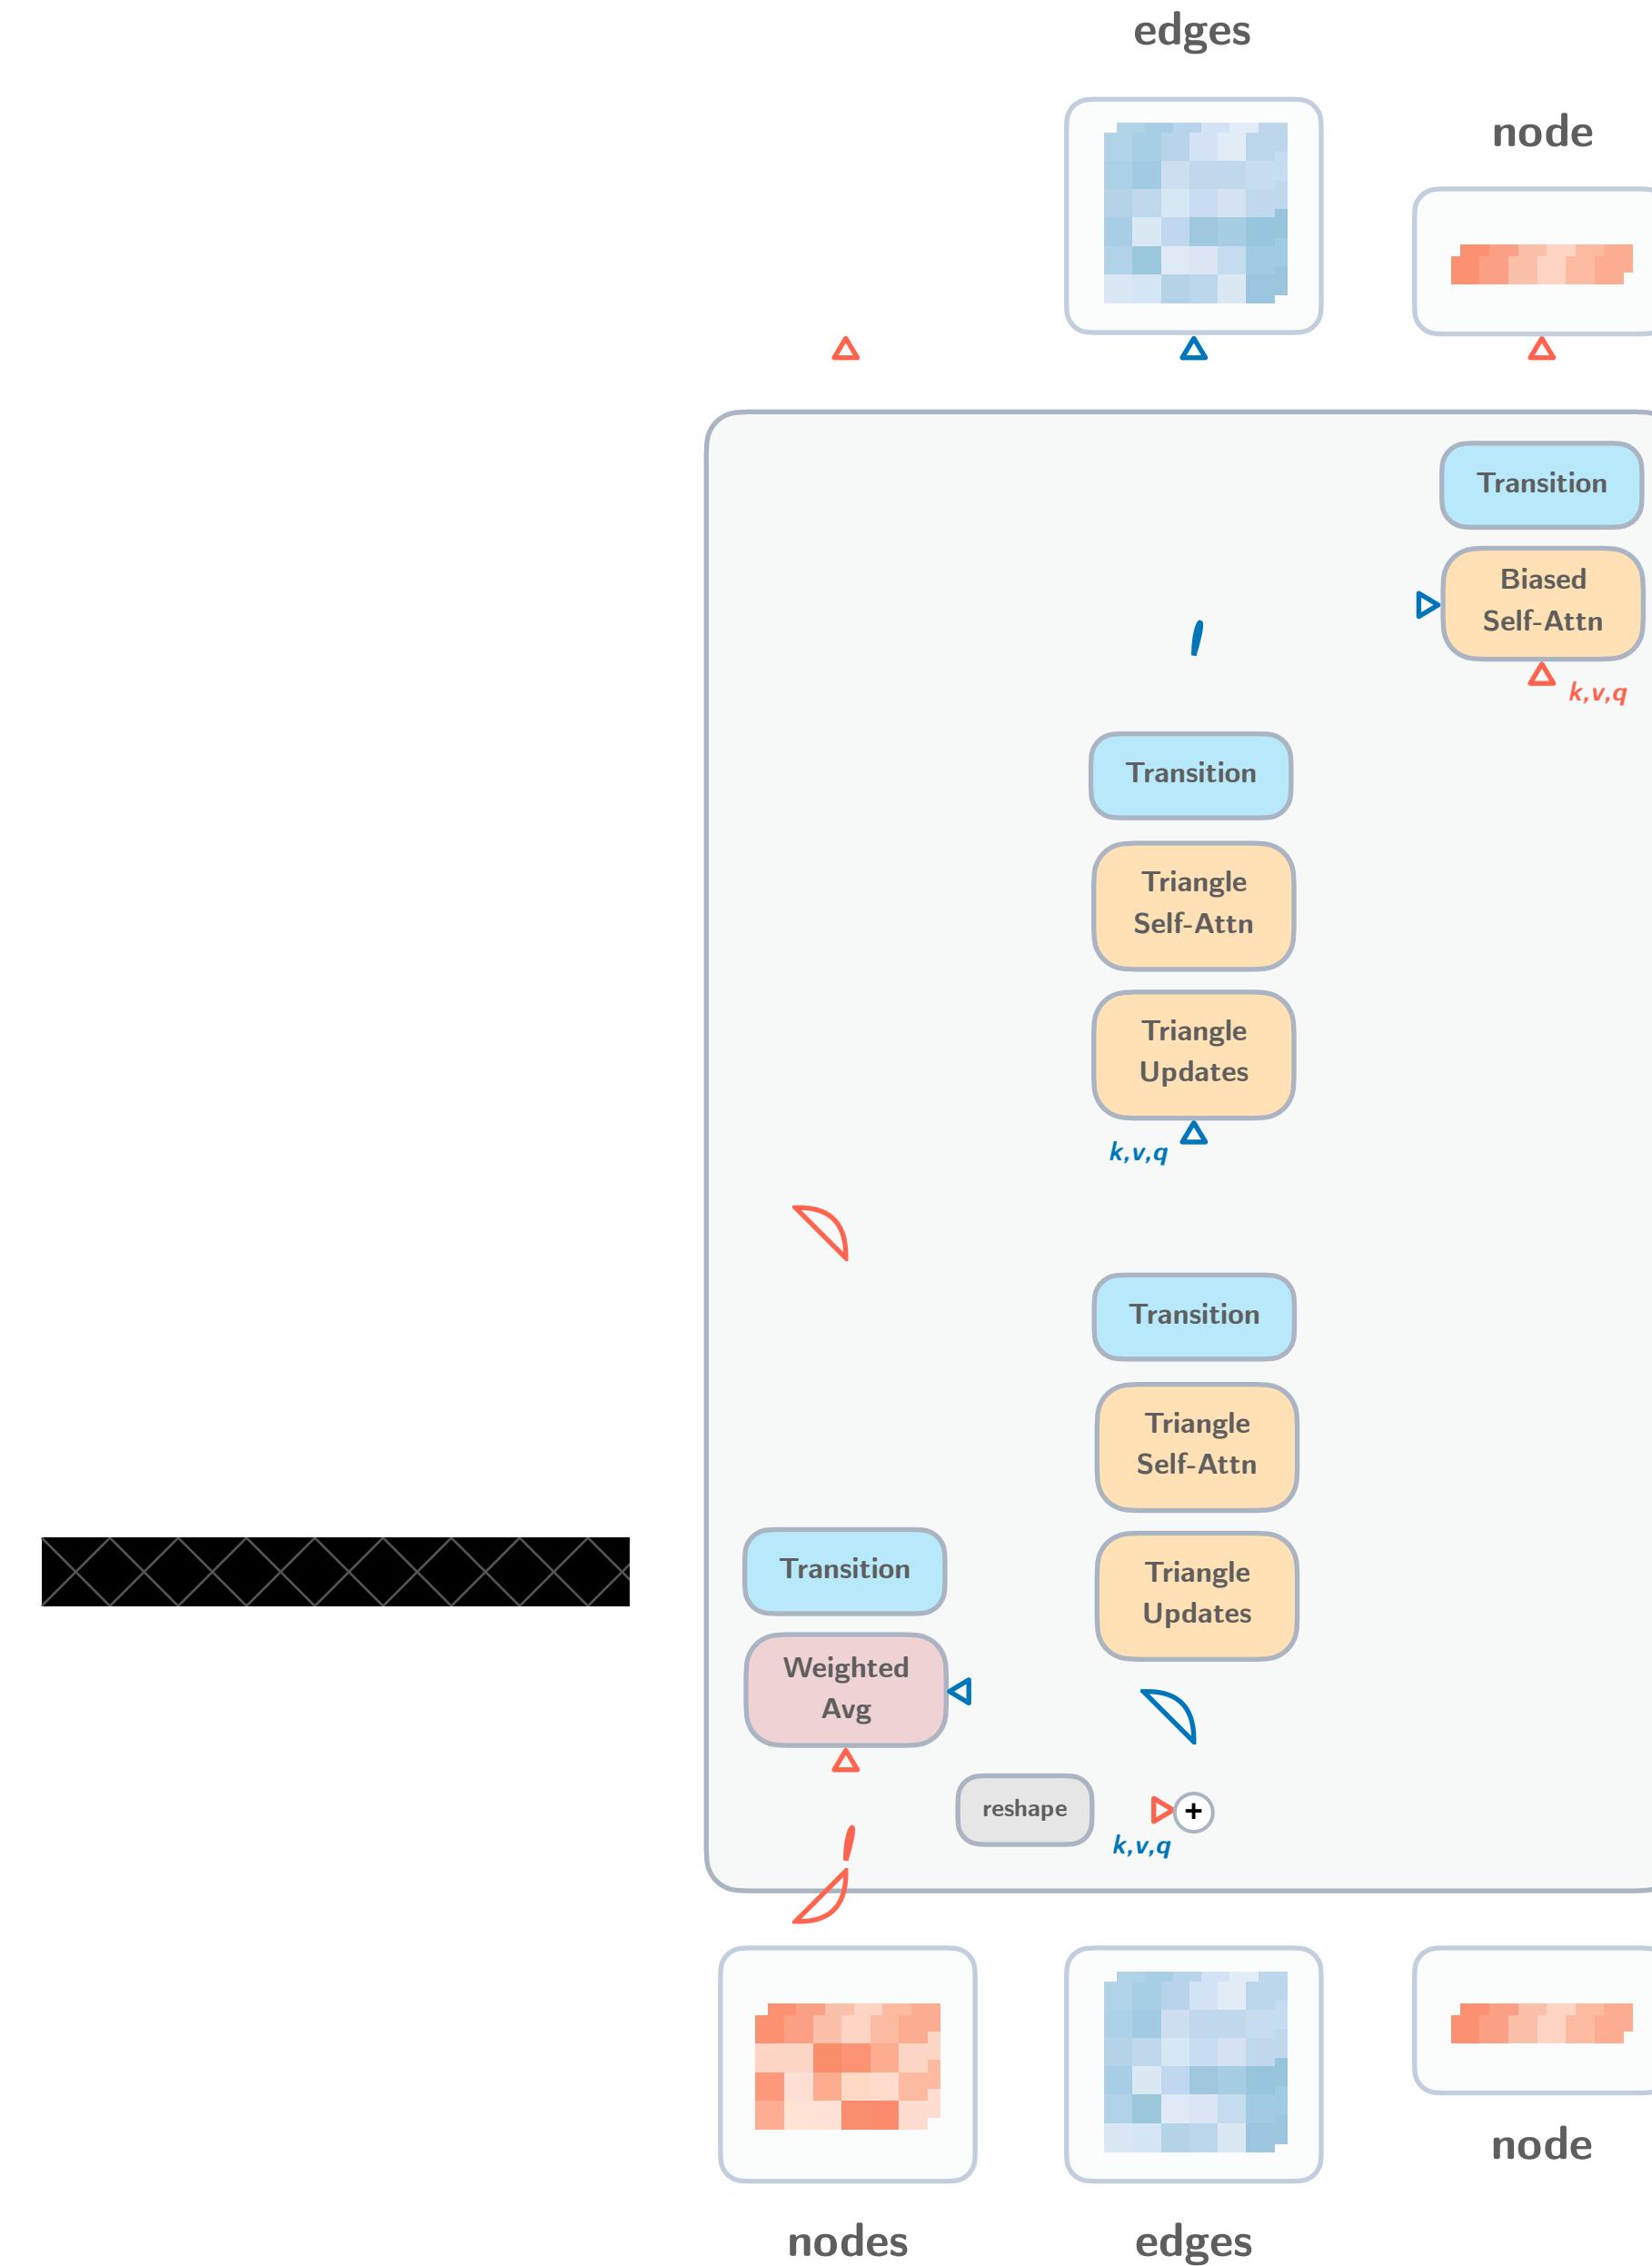


Attention matrix has the same shape as the
edge feature matrix.
Each cell refers to a node–node relationship,
i.e. an edge.



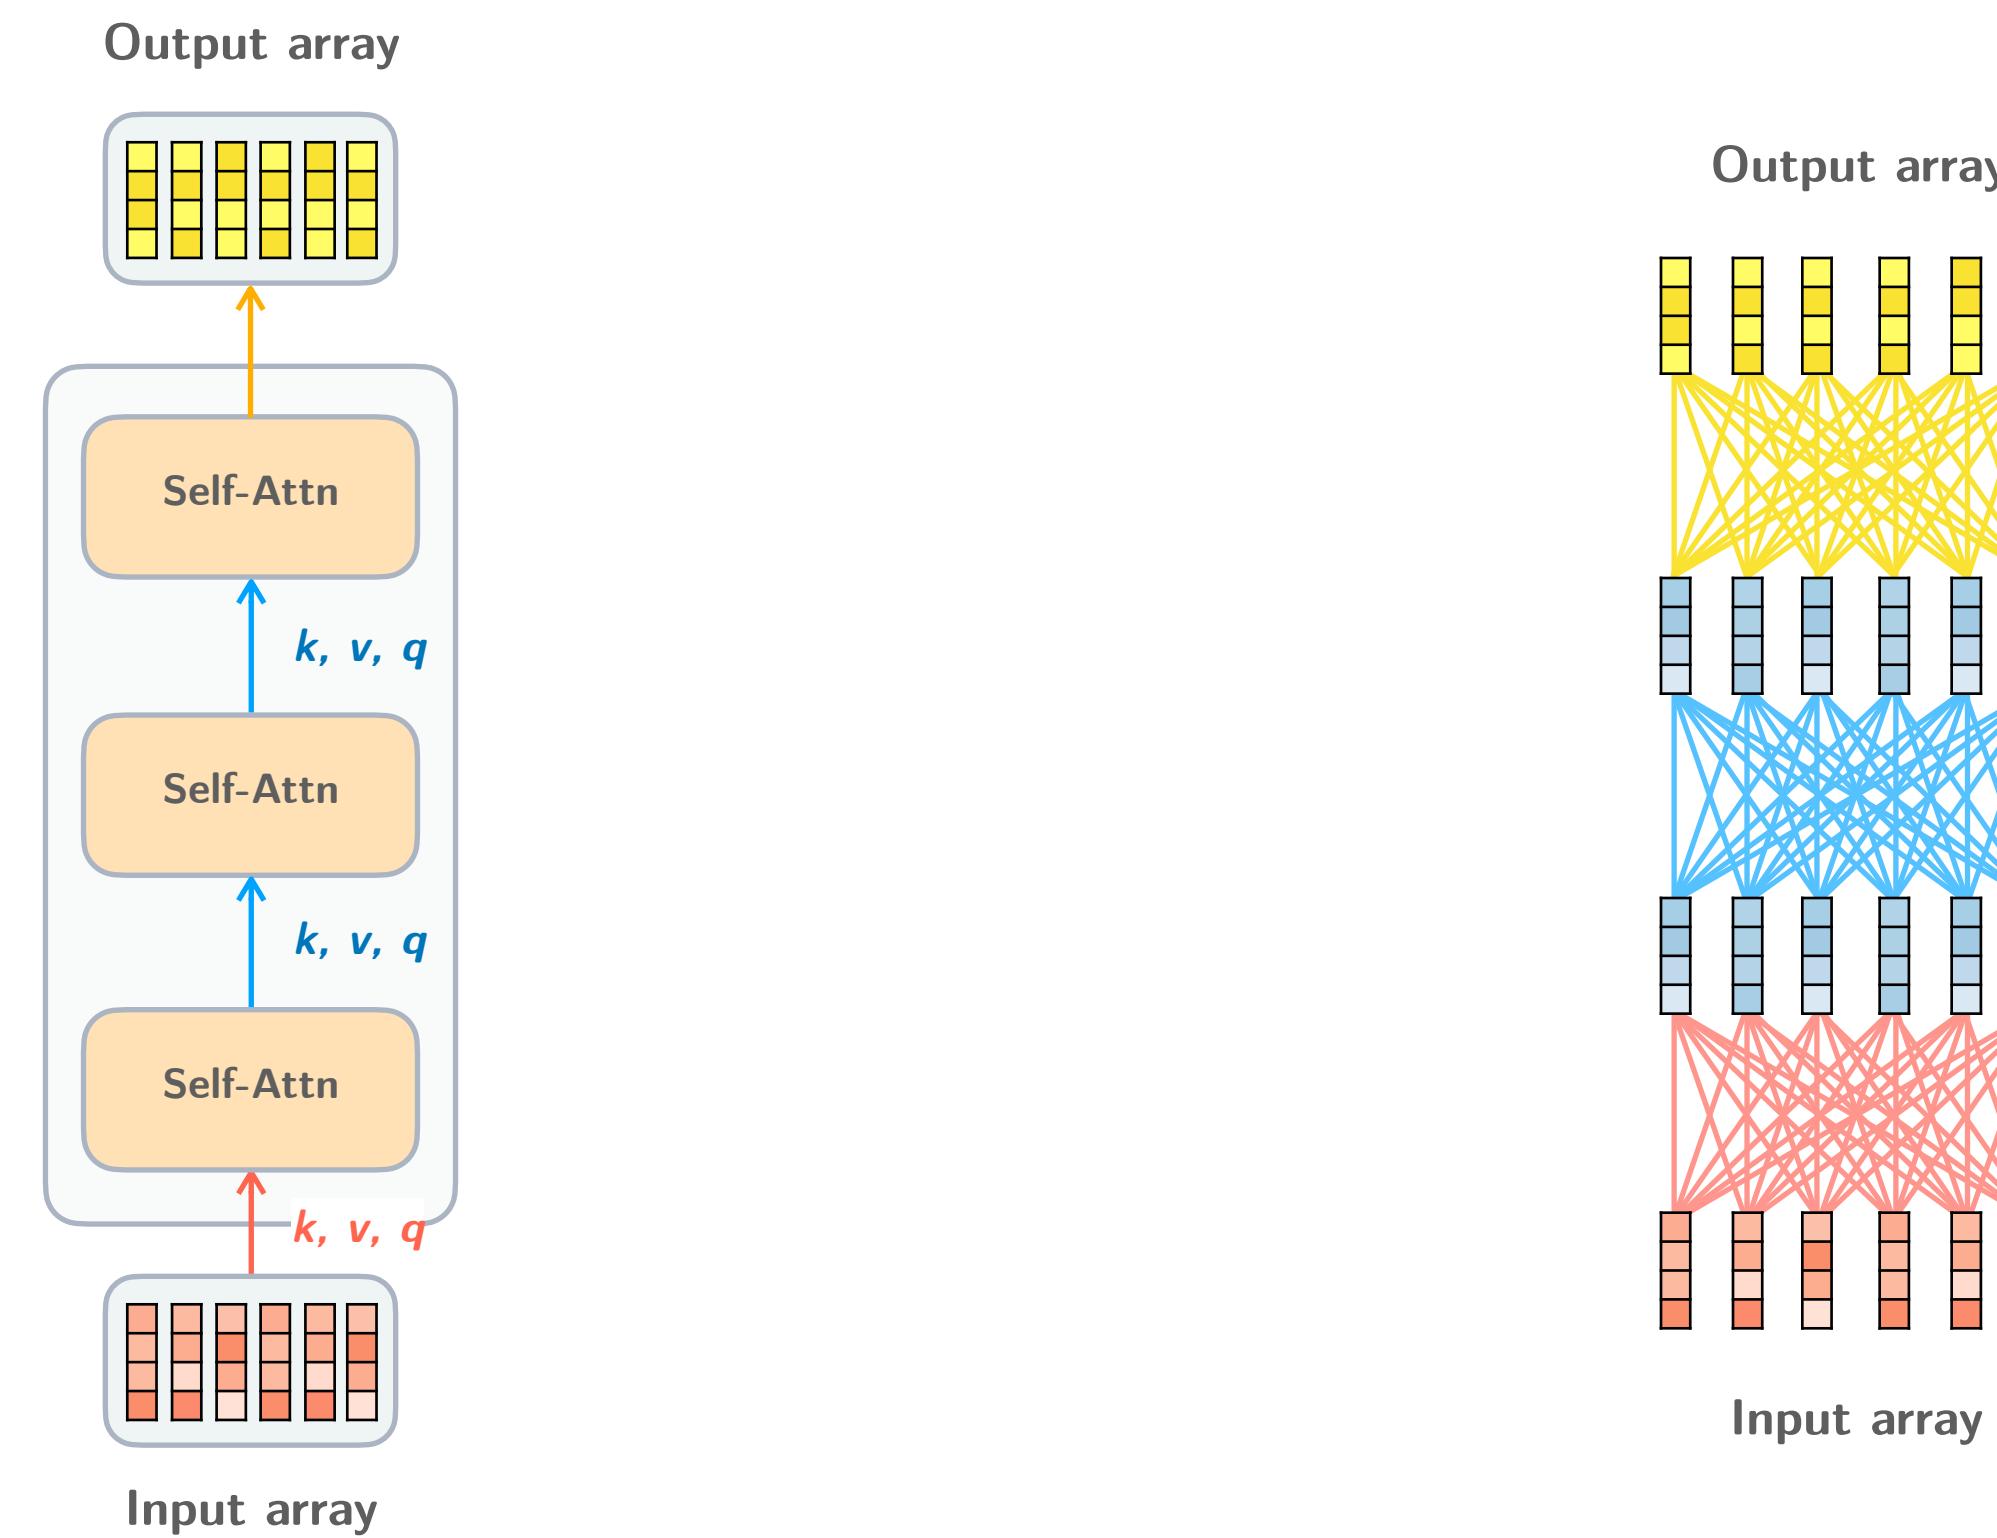








Perceiver (Latent Transformer)

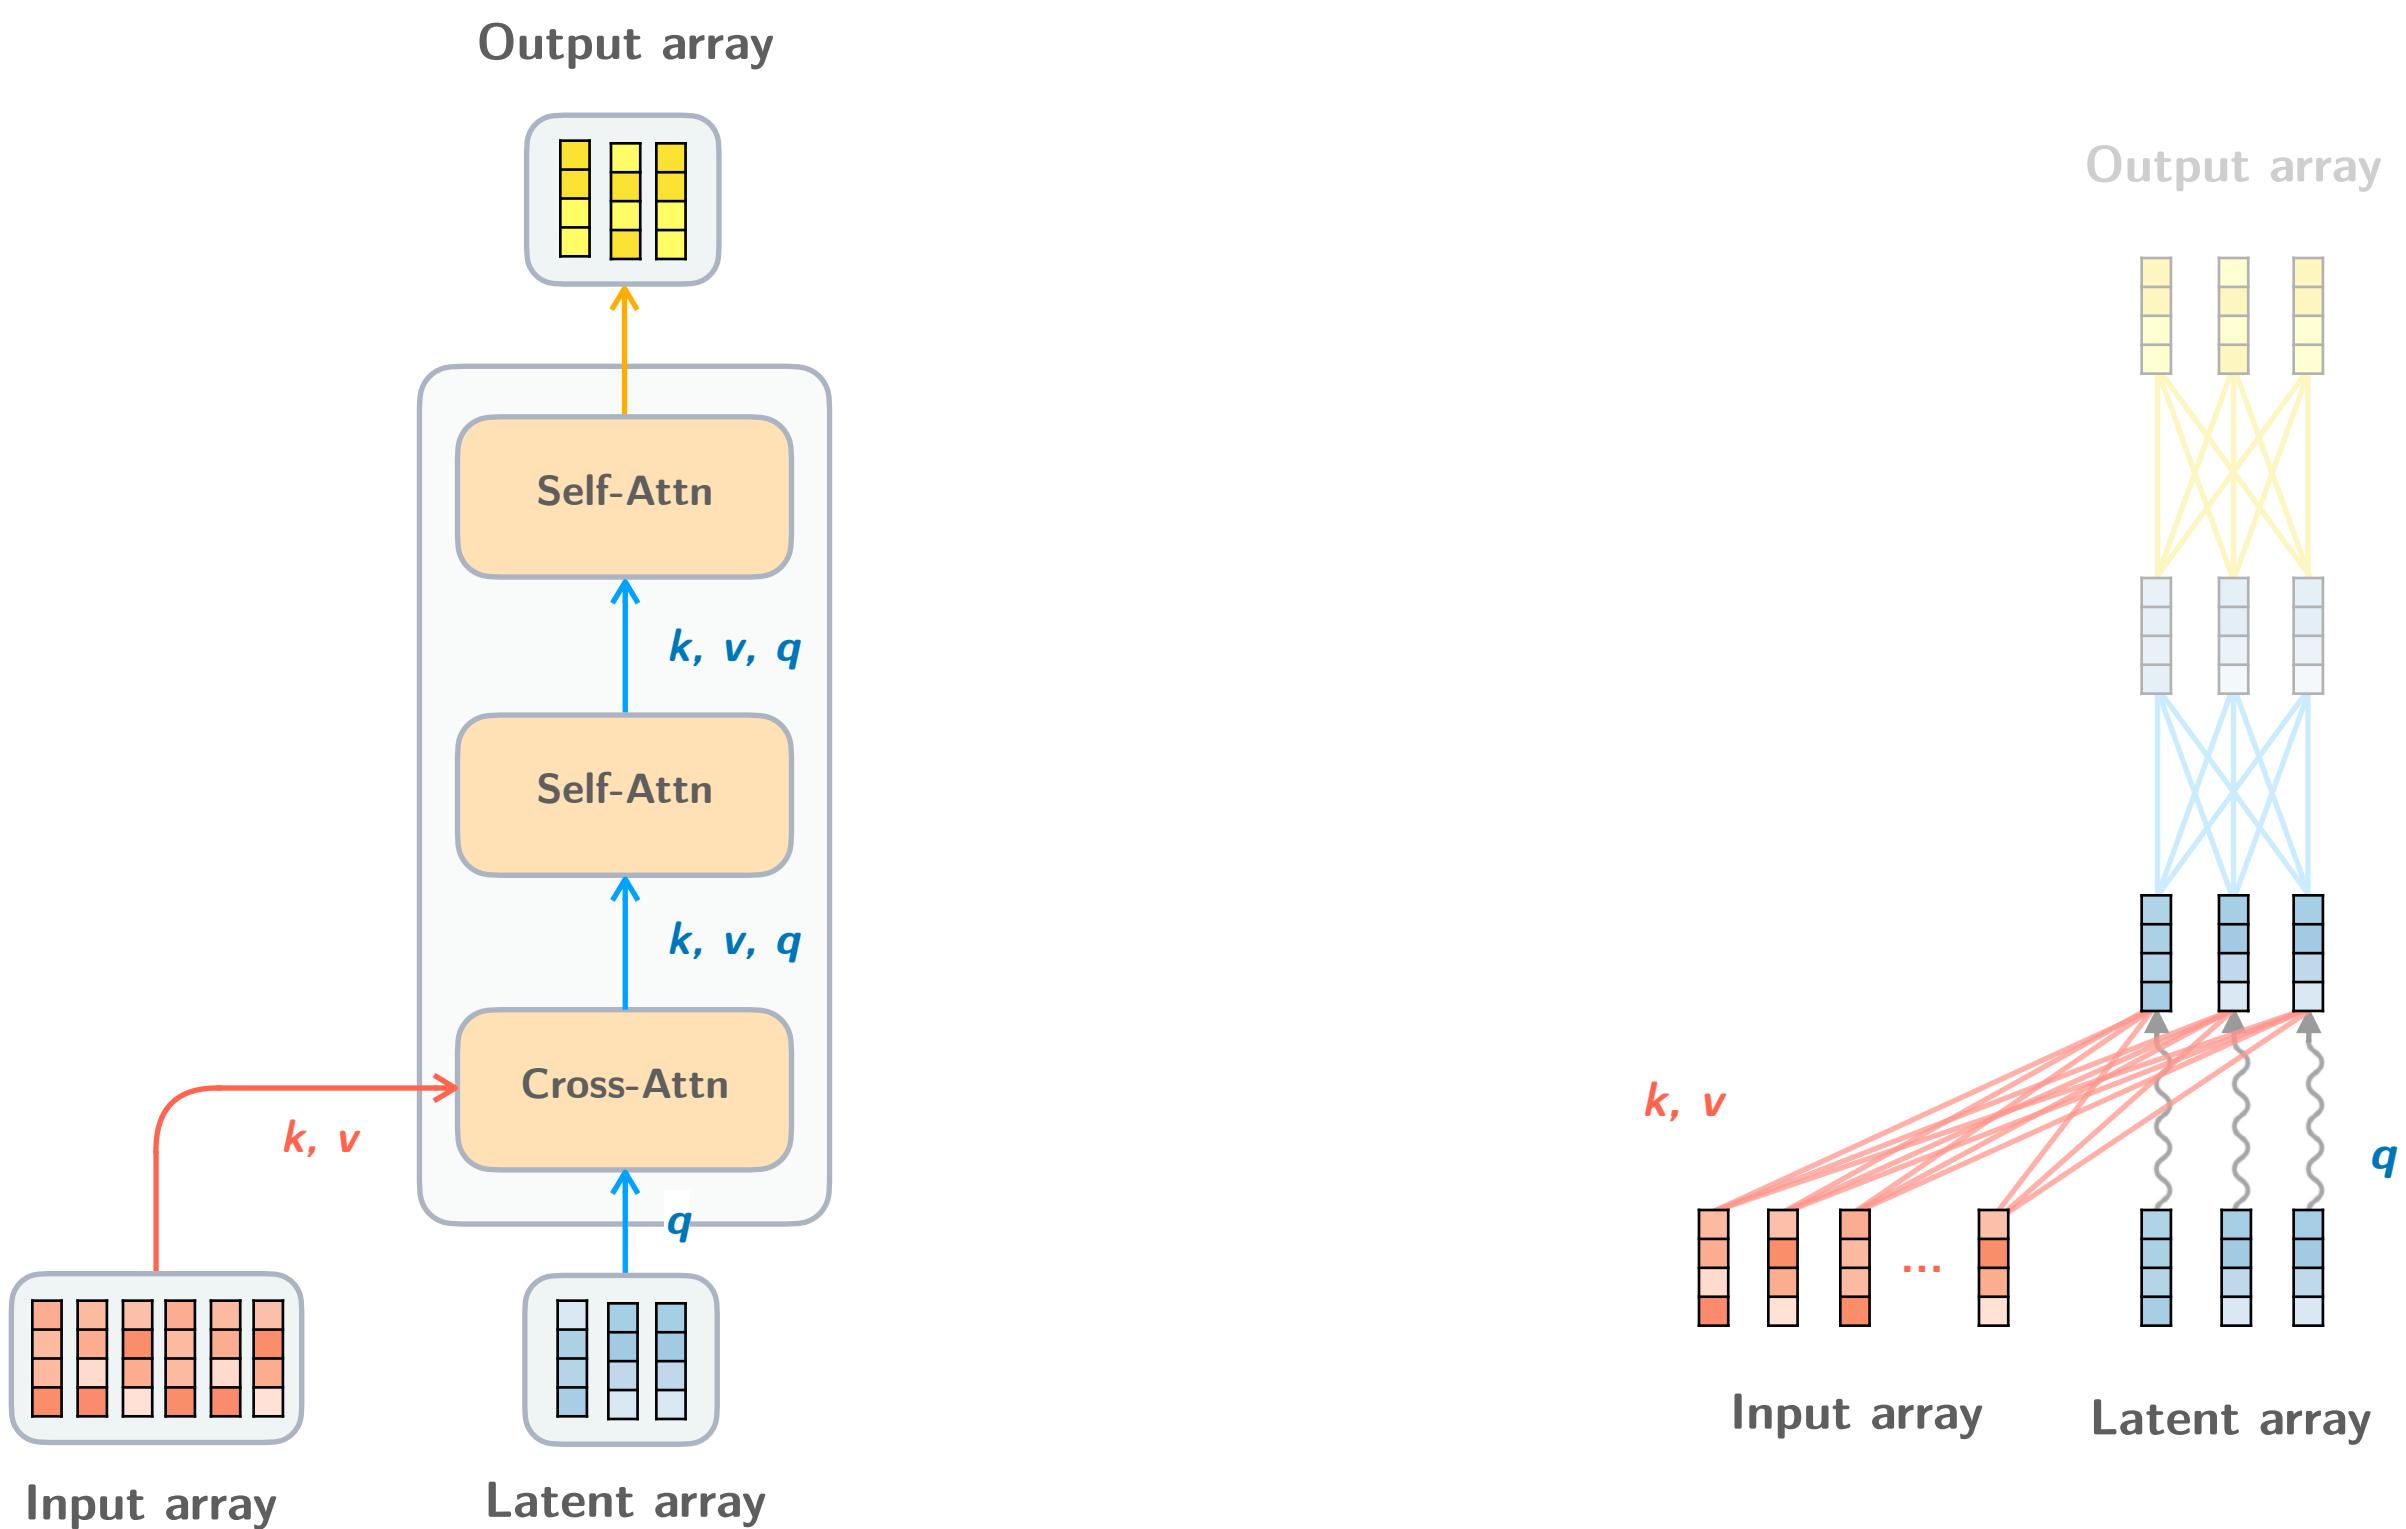


“Latent transformer”

It's a very general-purpose architecture.

Weights are optionally shared across layers.

If so, it's basically a *vertical RNN*,
i.e. recurrence in depth rather than time.

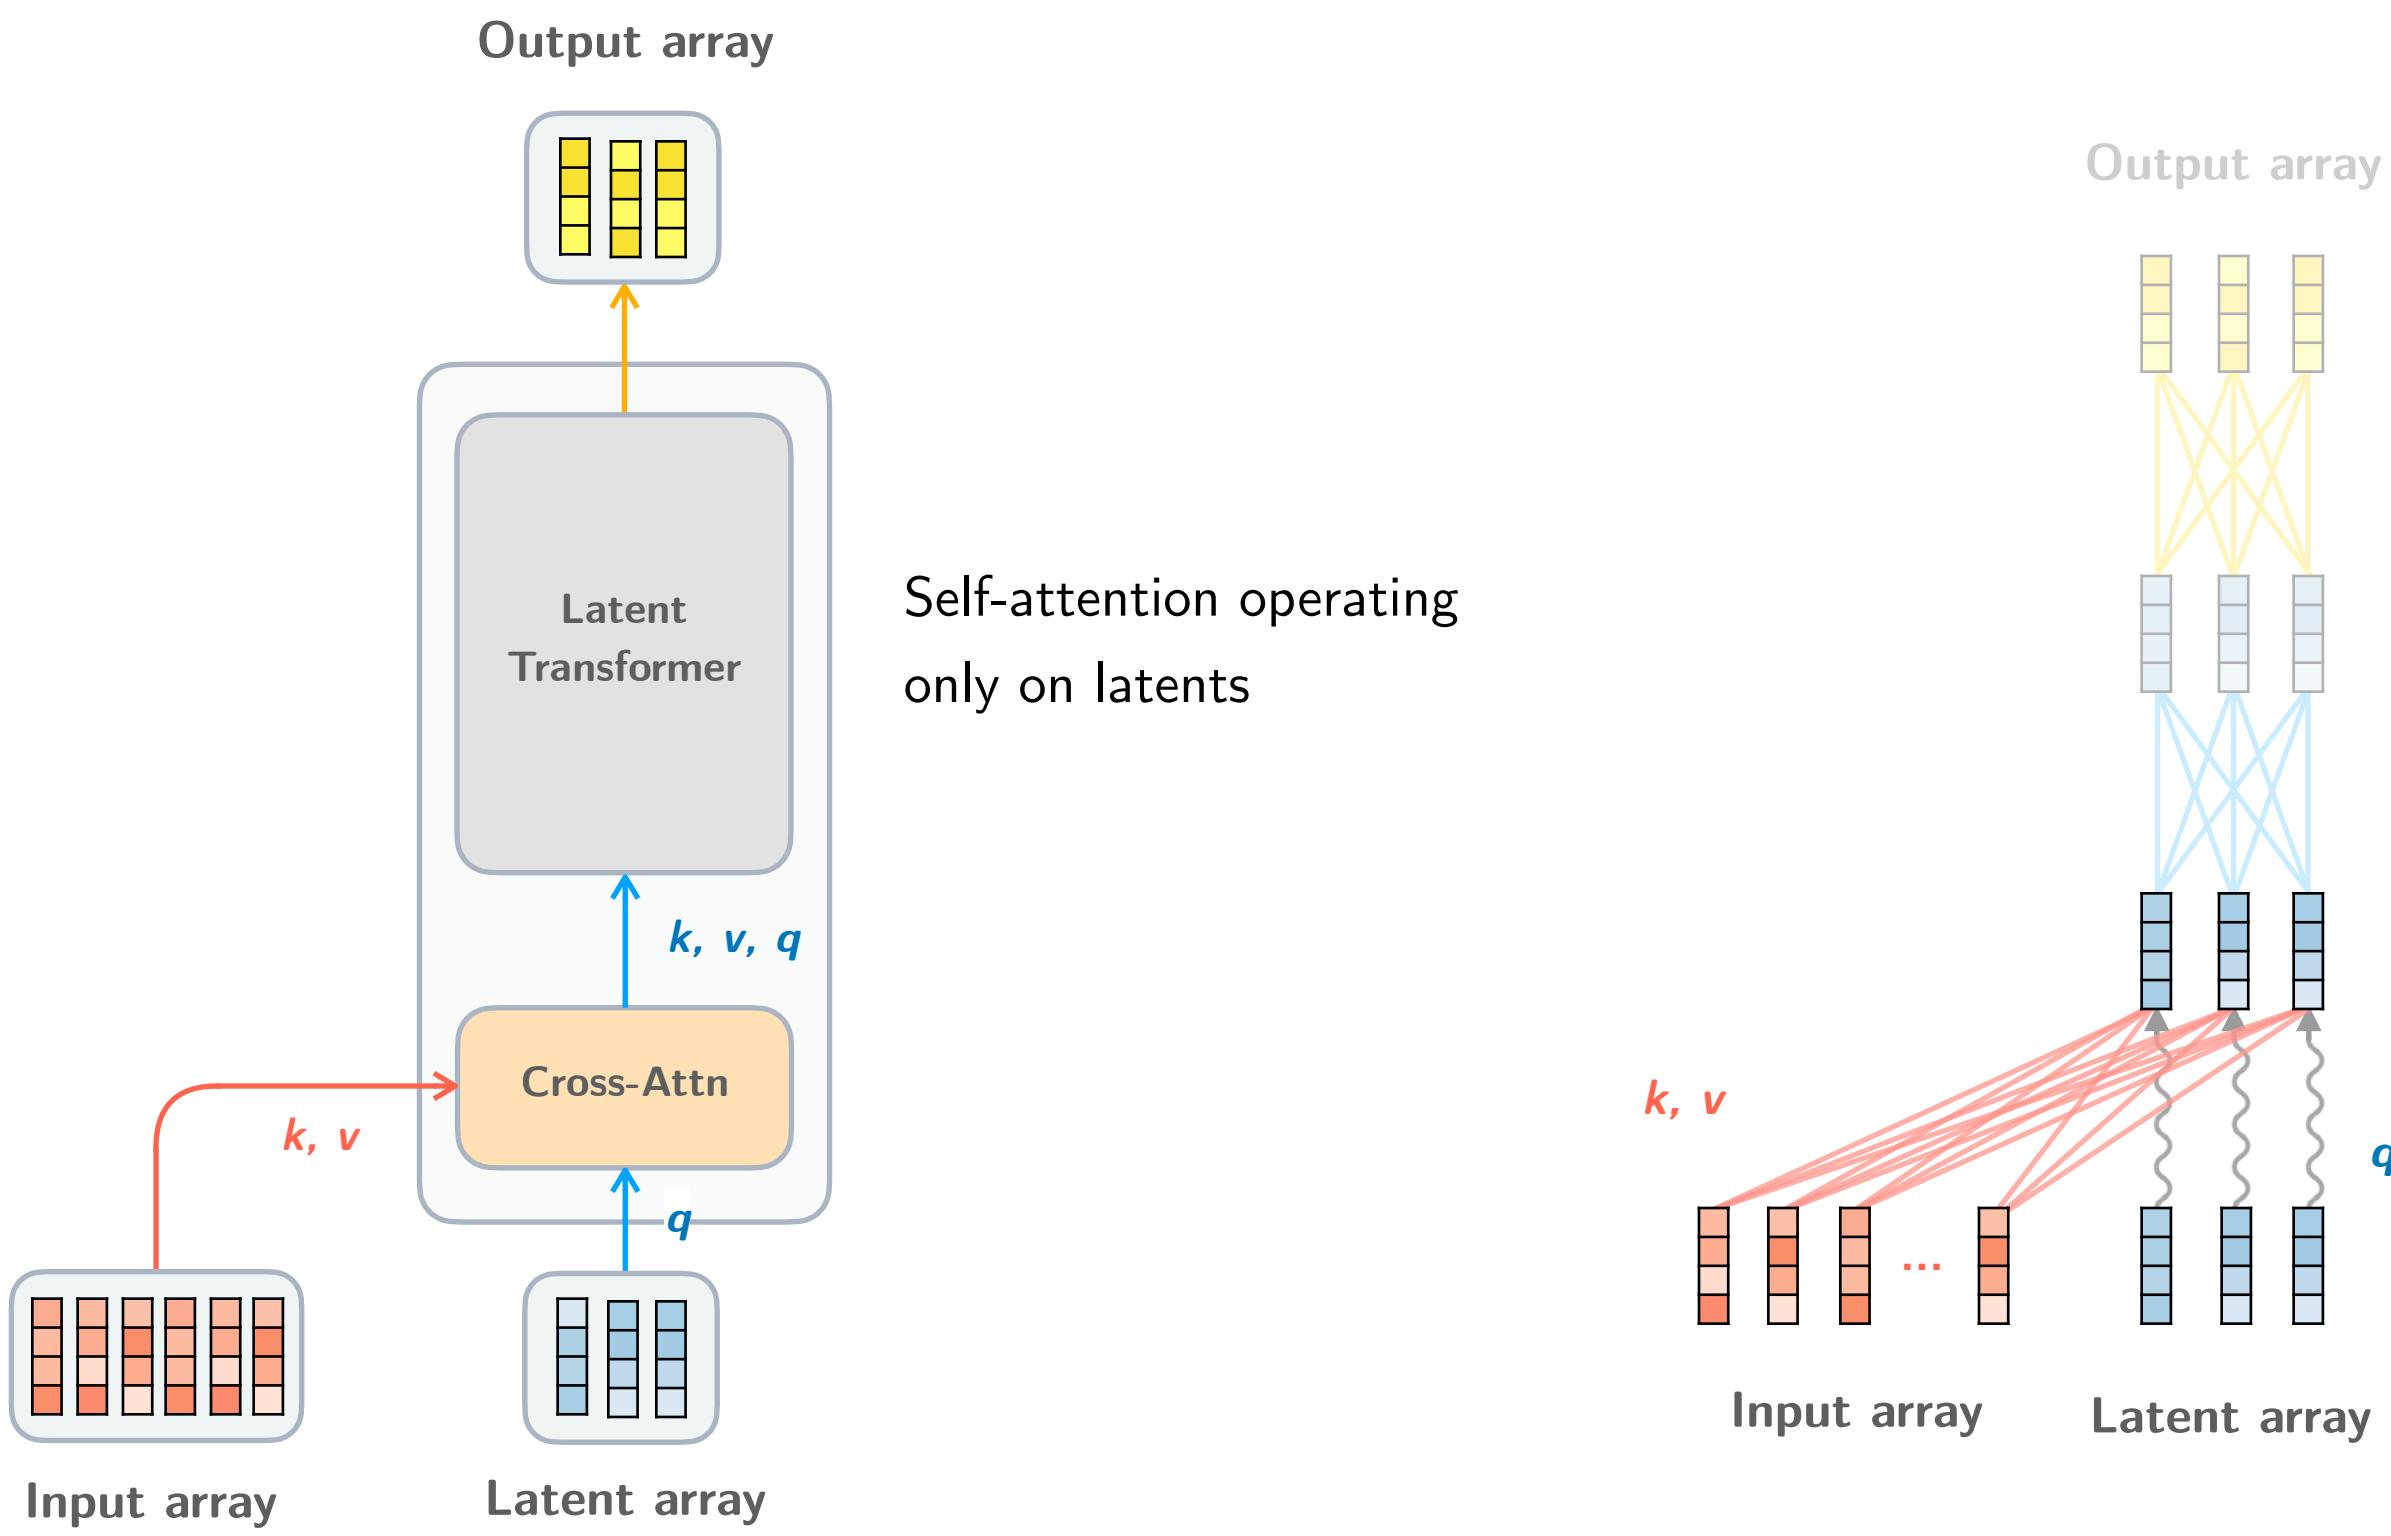


“Latent transformer”

It's a very general-purpose architecture.

Weights are optionally shared across layers.

If so, it's basically a *vertical RNN*,
i.e. recurrence in depth rather than time.

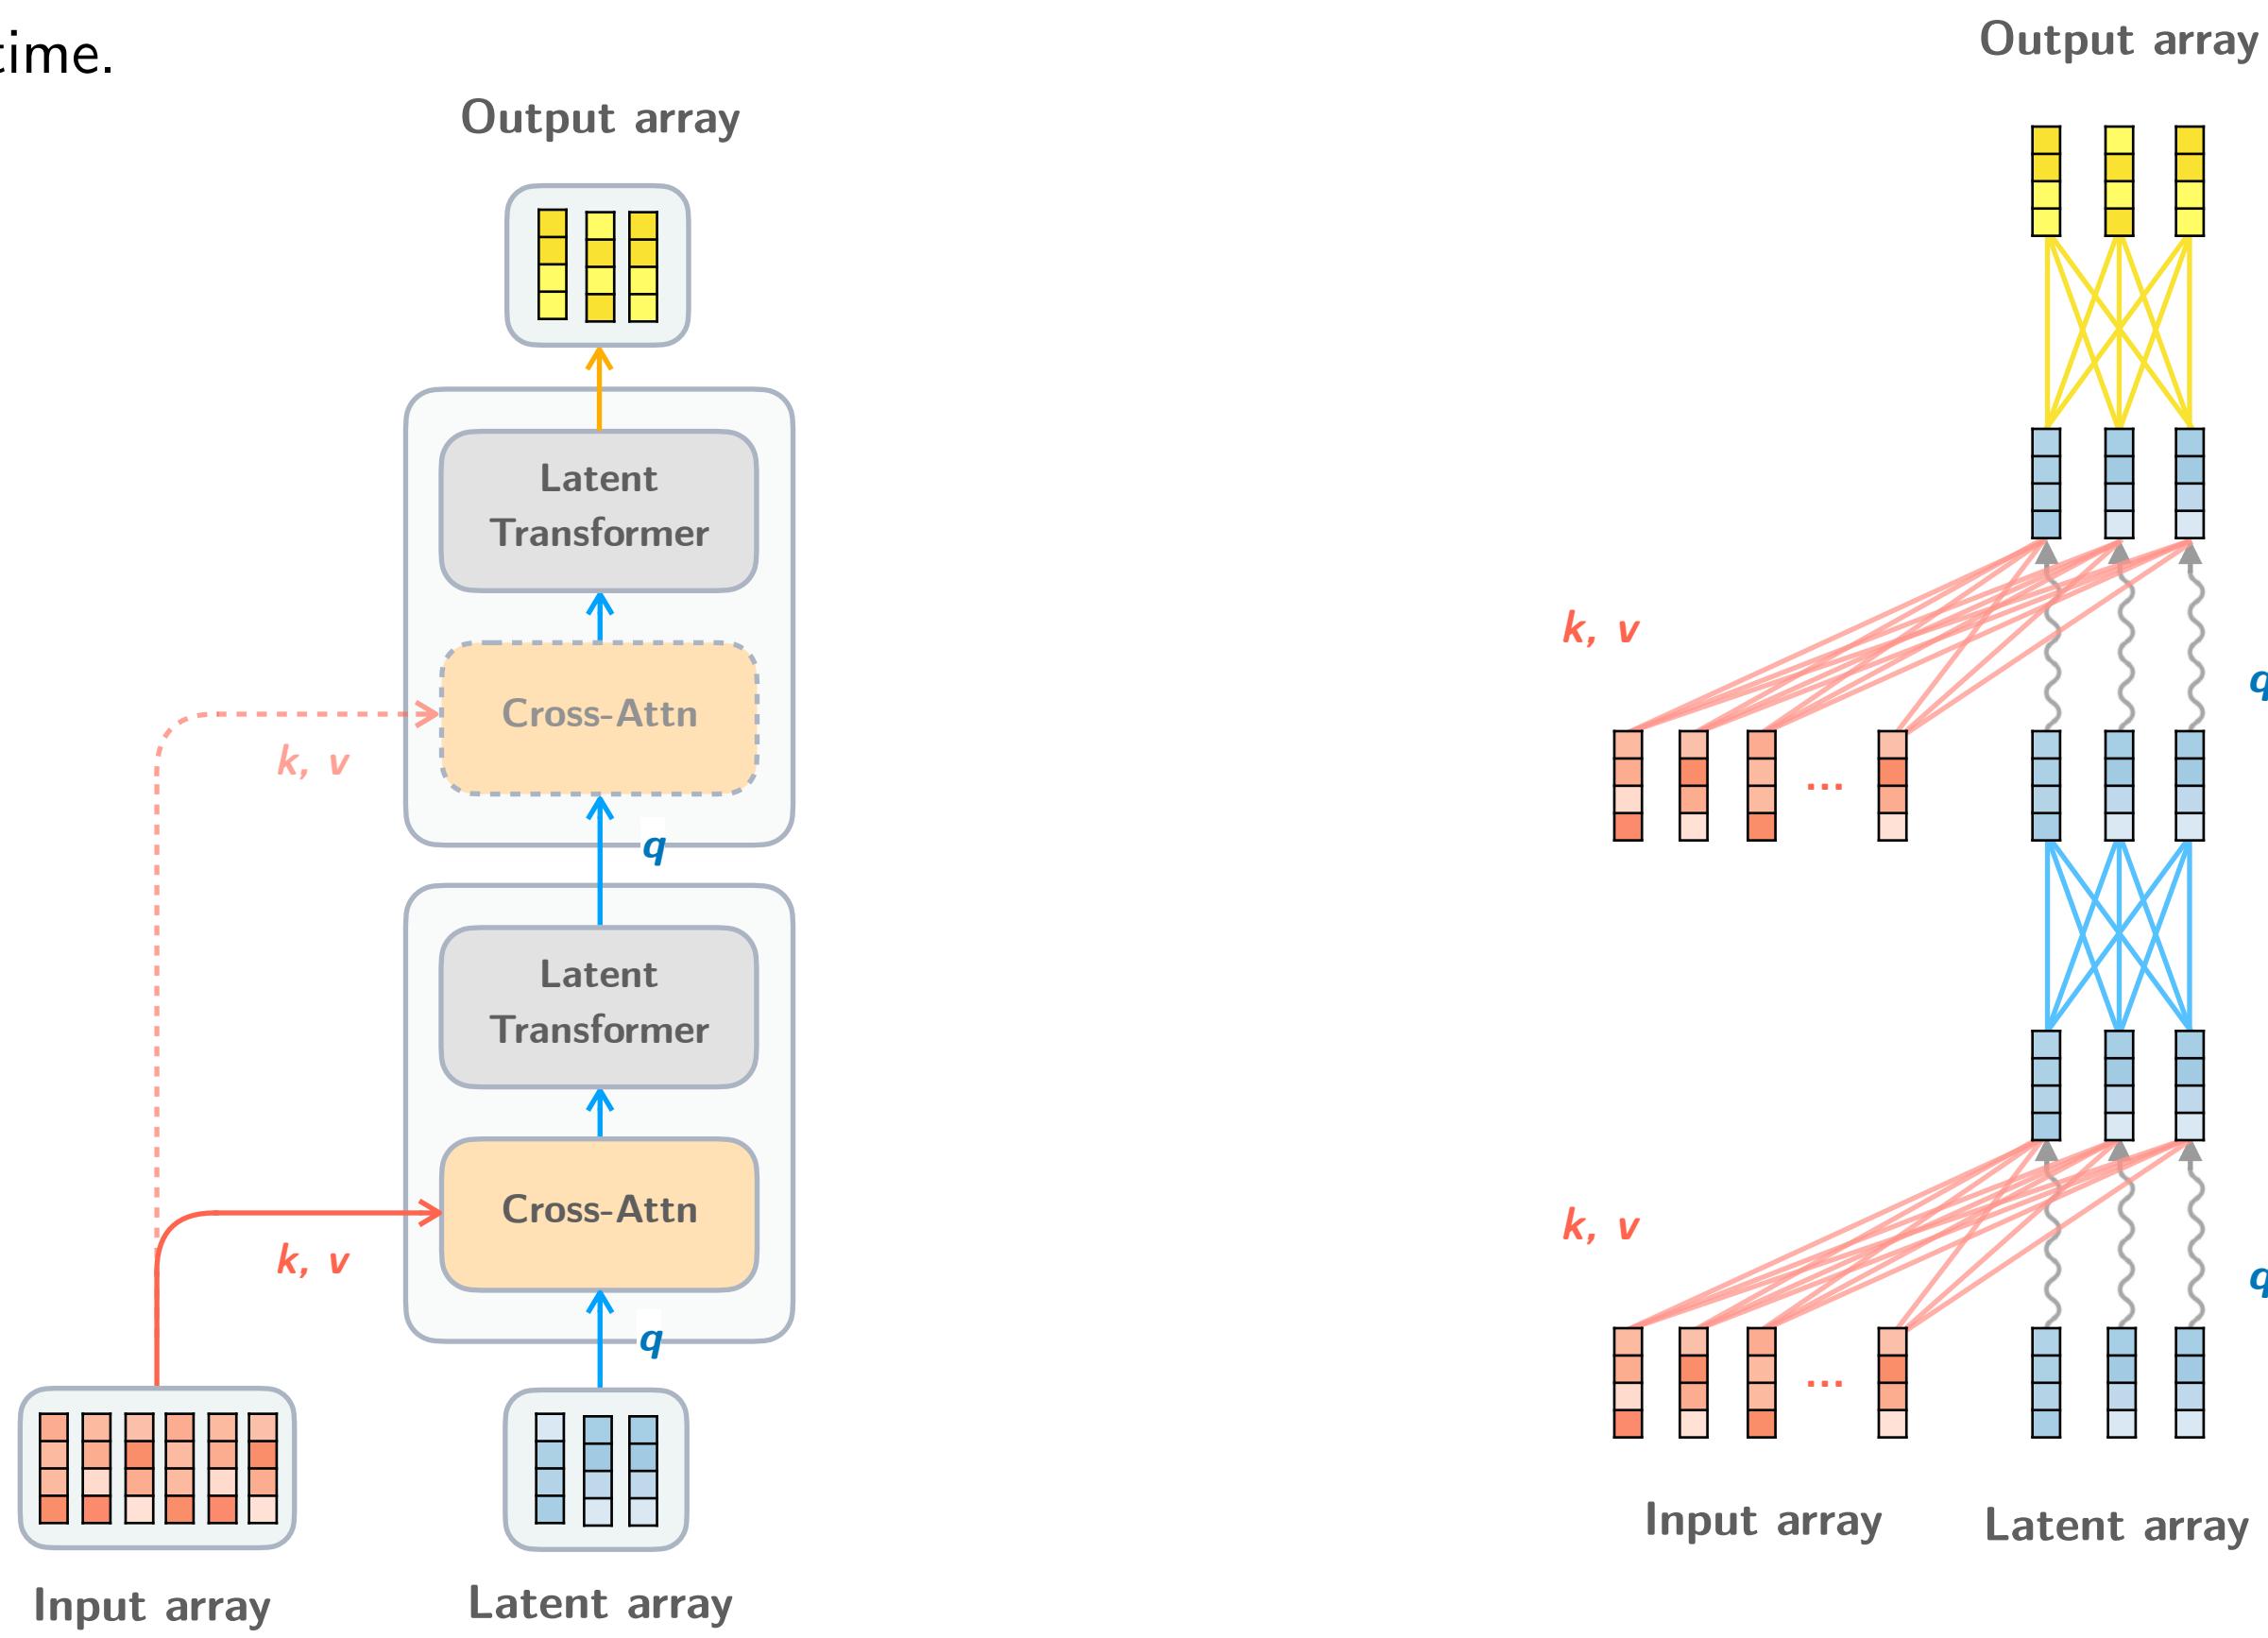


“Latent transformer”

It's a very general-purpose architecture.

Weights are optionally shared across layers.

If so, it's basically a *vertical RNN*,
i.e. recurrence in depth rather than time.



Applied to autoregressive modeling.

Each latent is associated with an output.

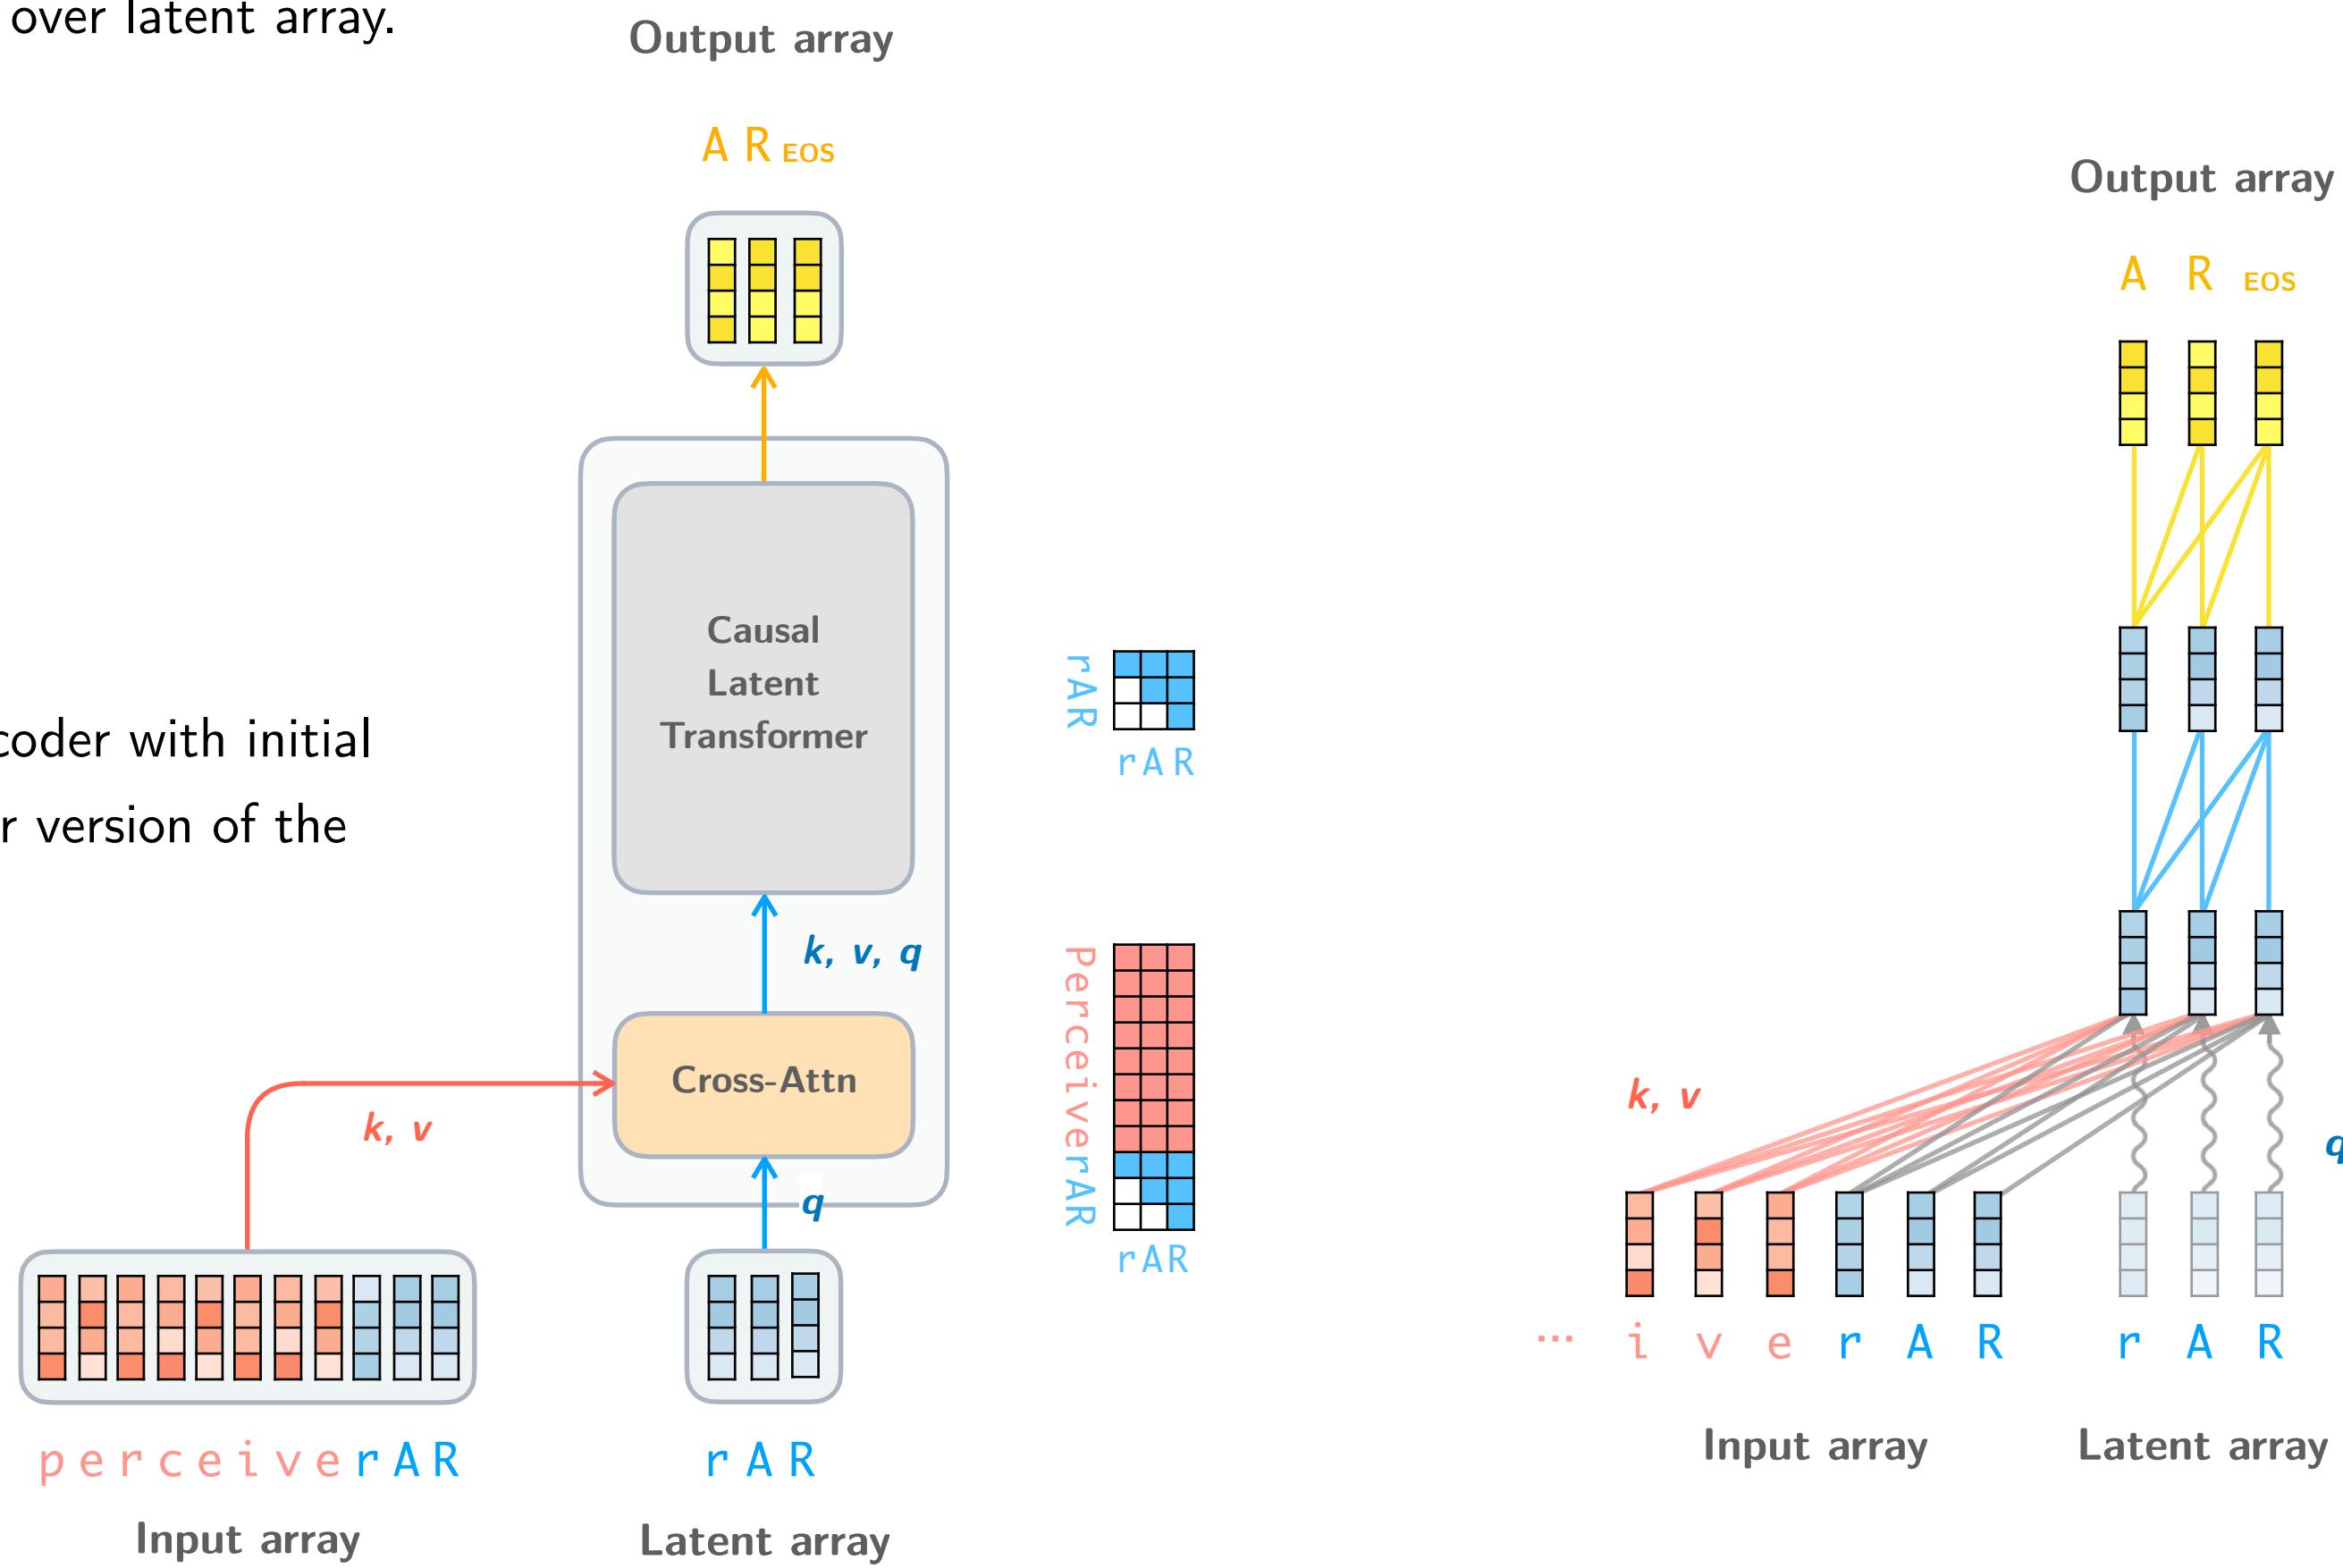
We can still vary the number of latents as before.

Latent self-attention is also causal now.

Input array can be much larger than the latent array.

We only do deep self-attention over latent array.

Basically just a transformer decoder with initial cross-attention layer to a longer version of the autoregressive array?



Applied to autoregressive modeling.

Each latent is associated with an output.

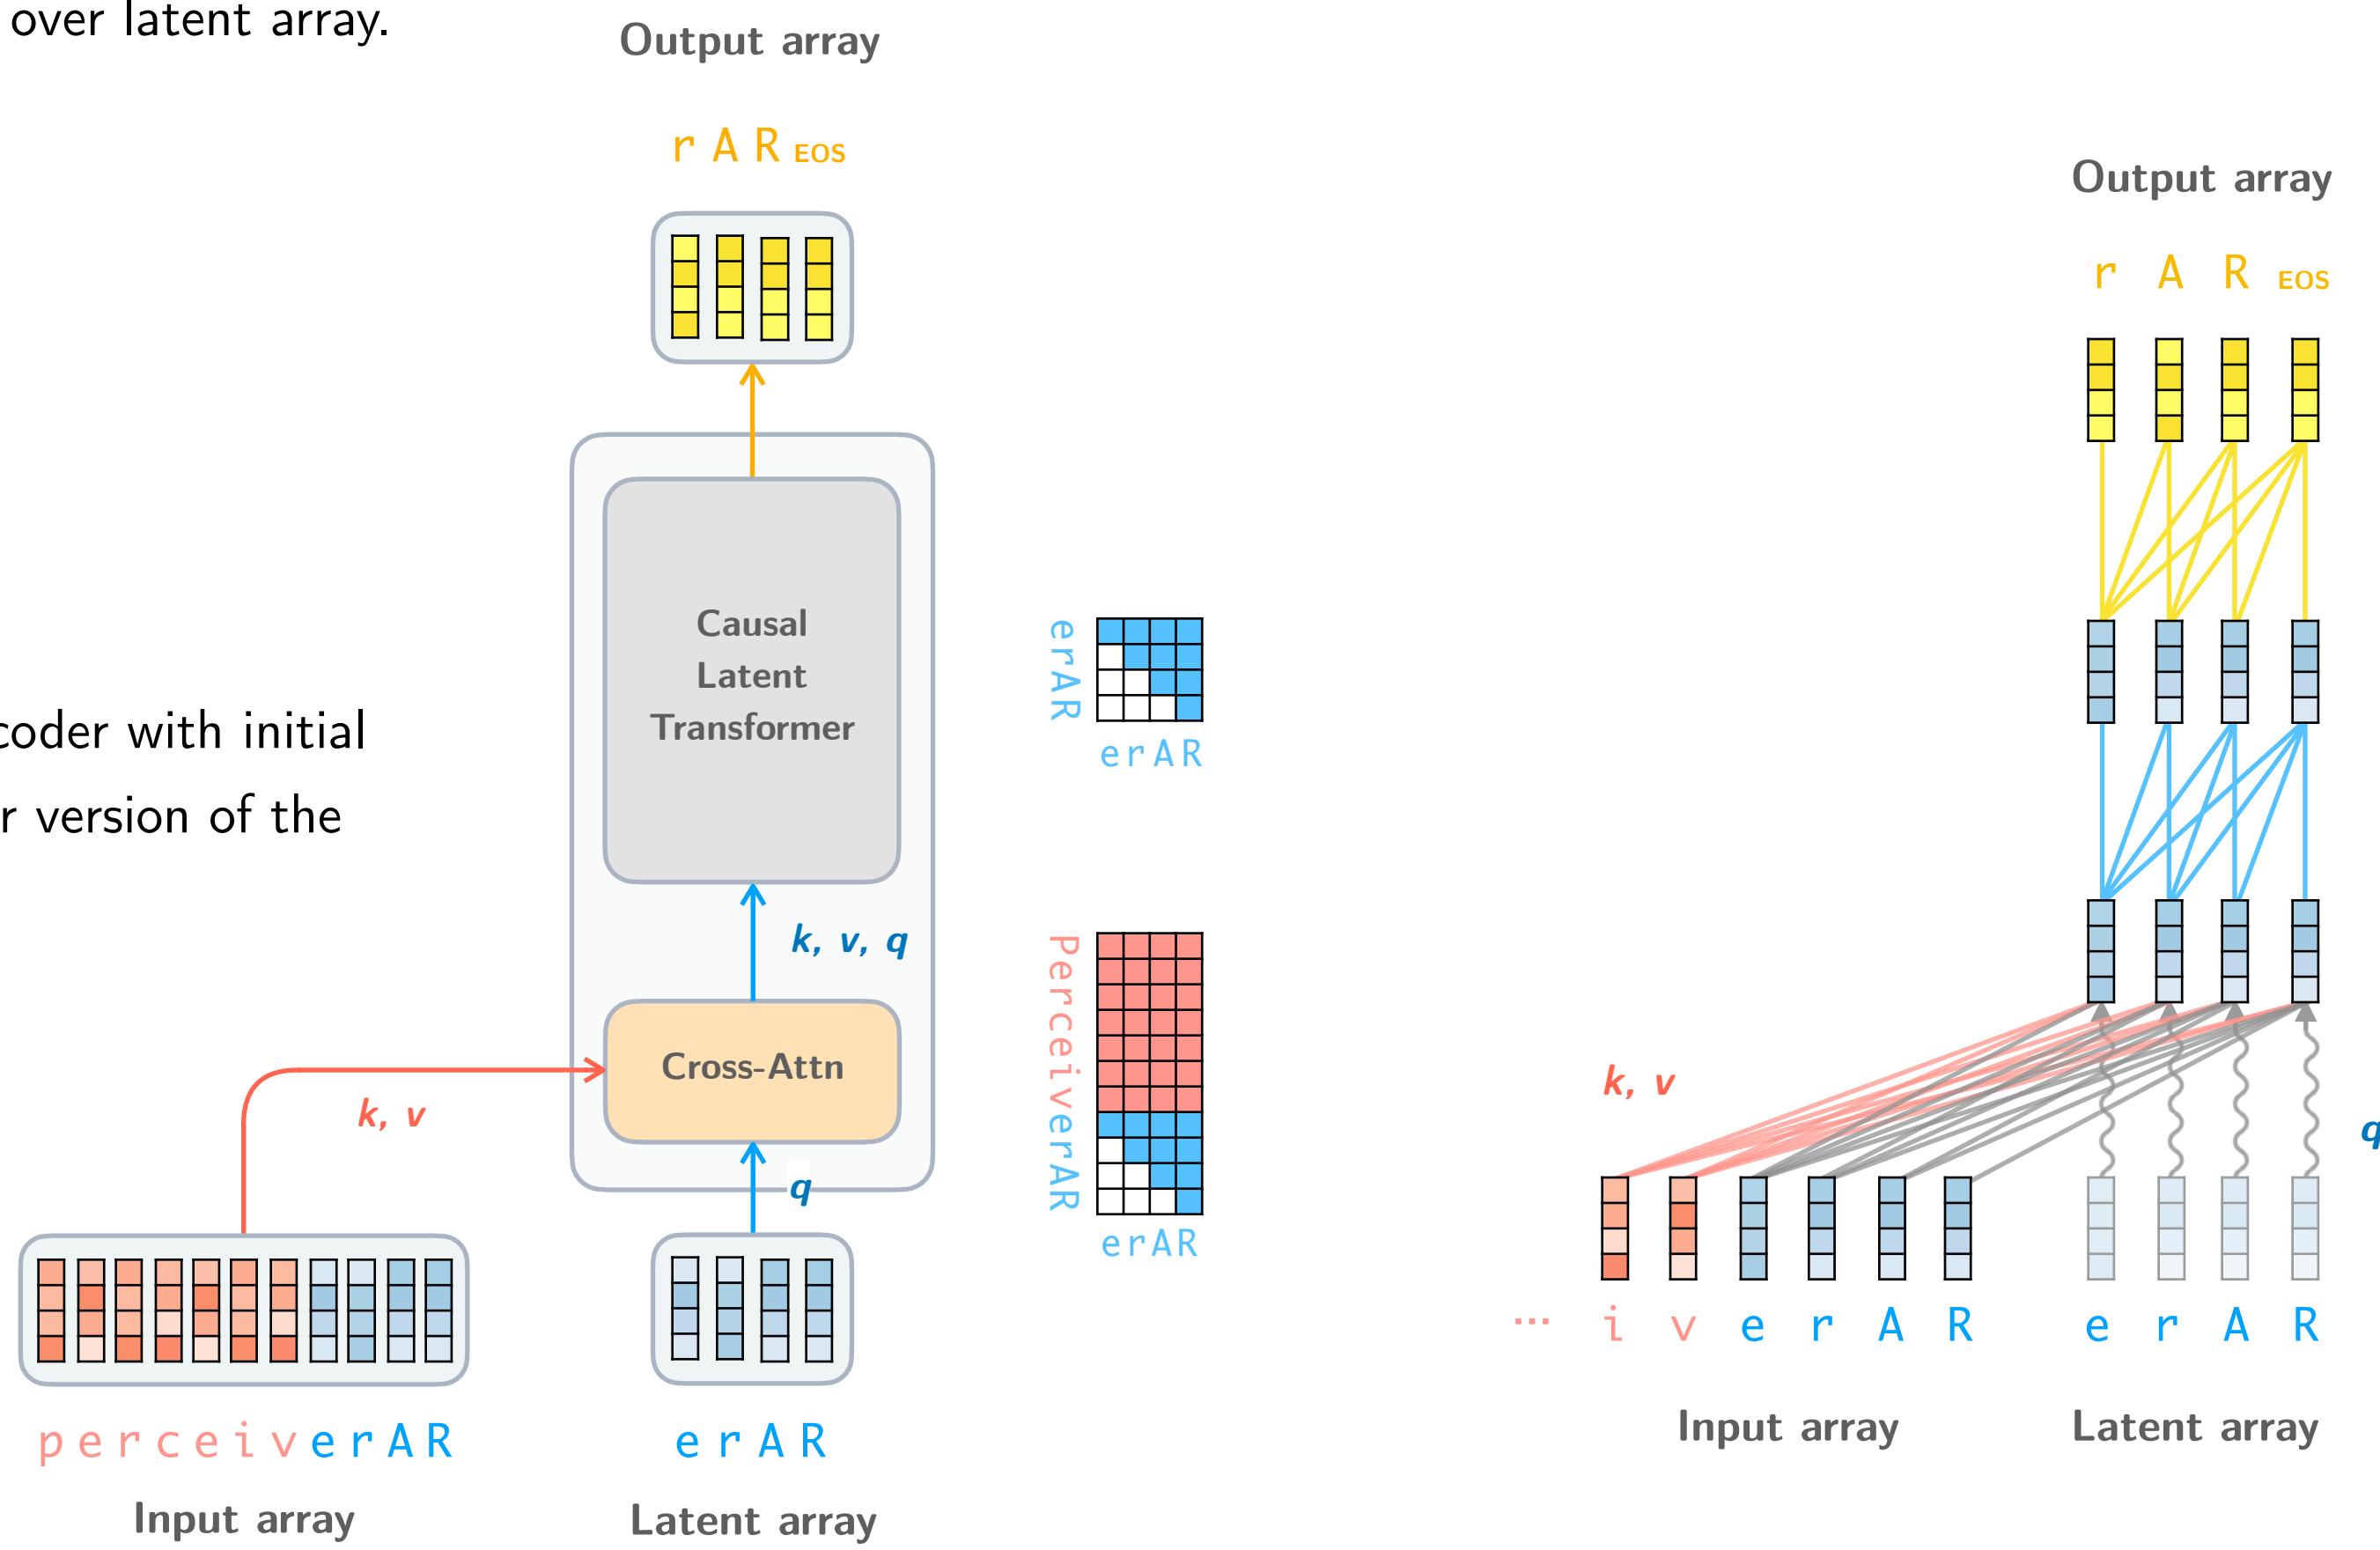
We can still vary the number of latents as before.

Latent self-attention is also causal now.

Input array can be much larger than the latent array.

We only do deep self-attention over latent array.

Basically just a transformer decoder with initial cross-attention layer to a longer version of the autoregressive array?



Applied to autoregressive modeling.

Each latent is associated with an output.

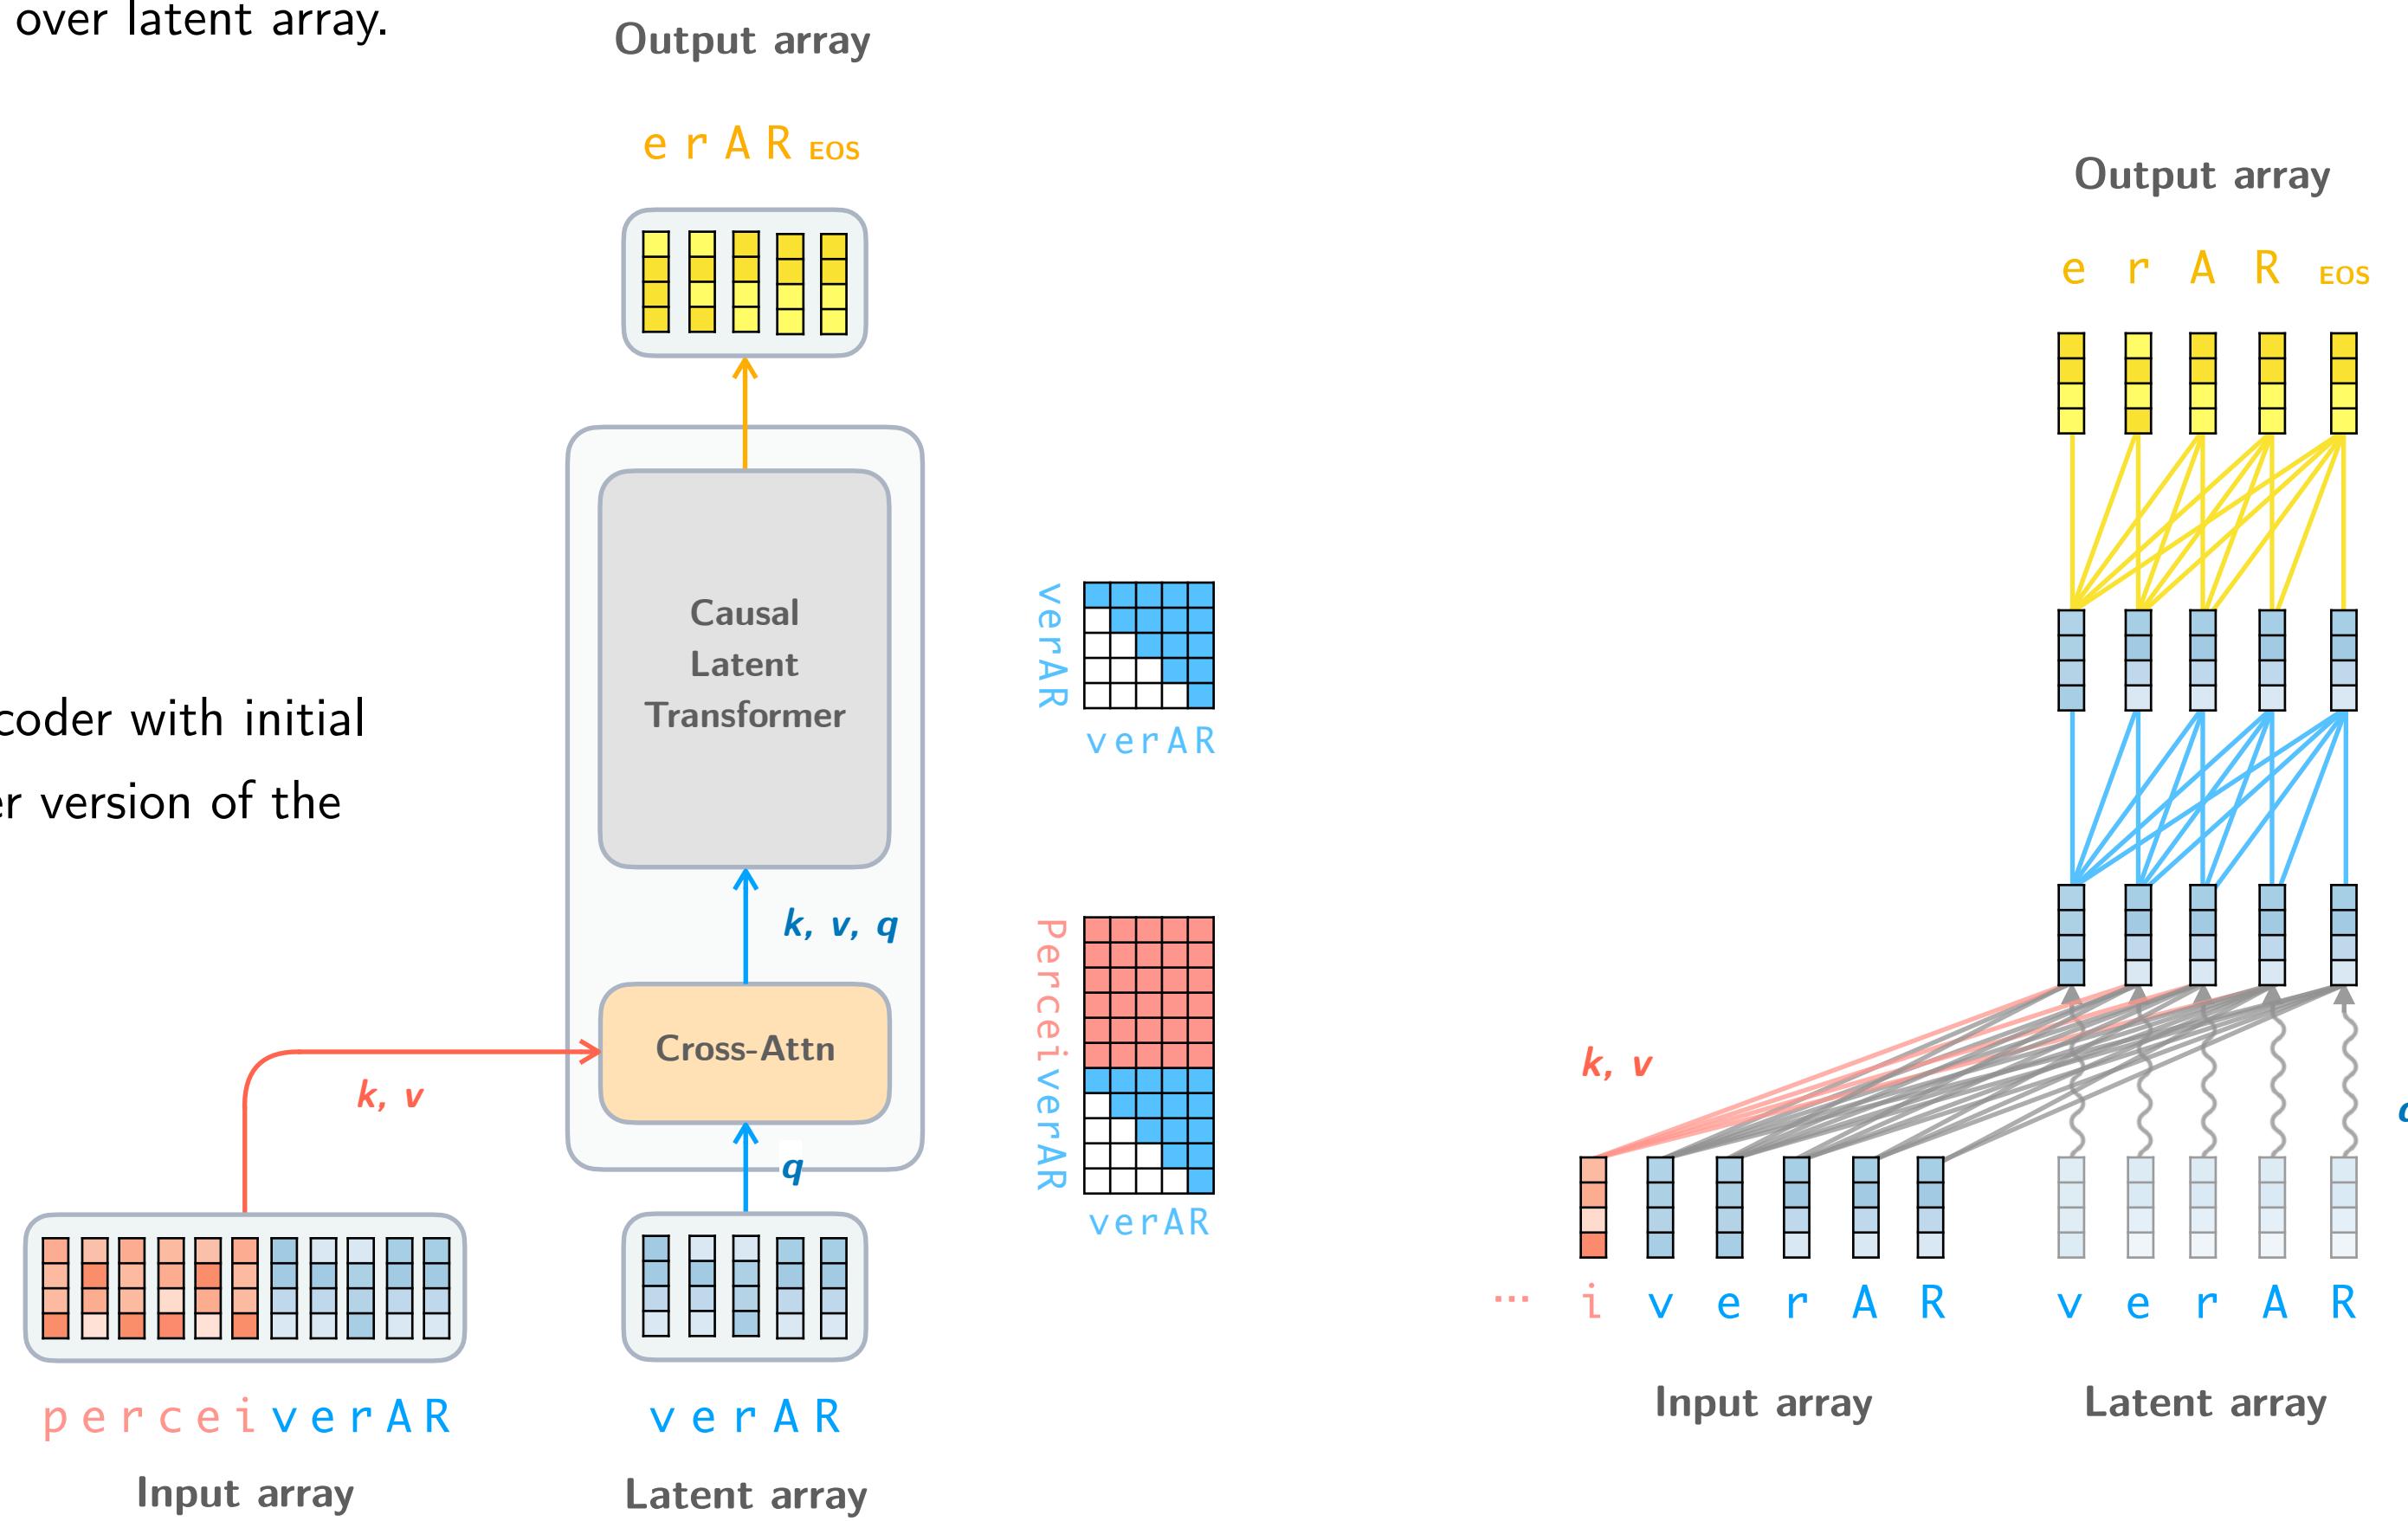
We can still vary the number of latents as before.

Latent self-attention is also causal now.

Input array can be much larger than the latent array.

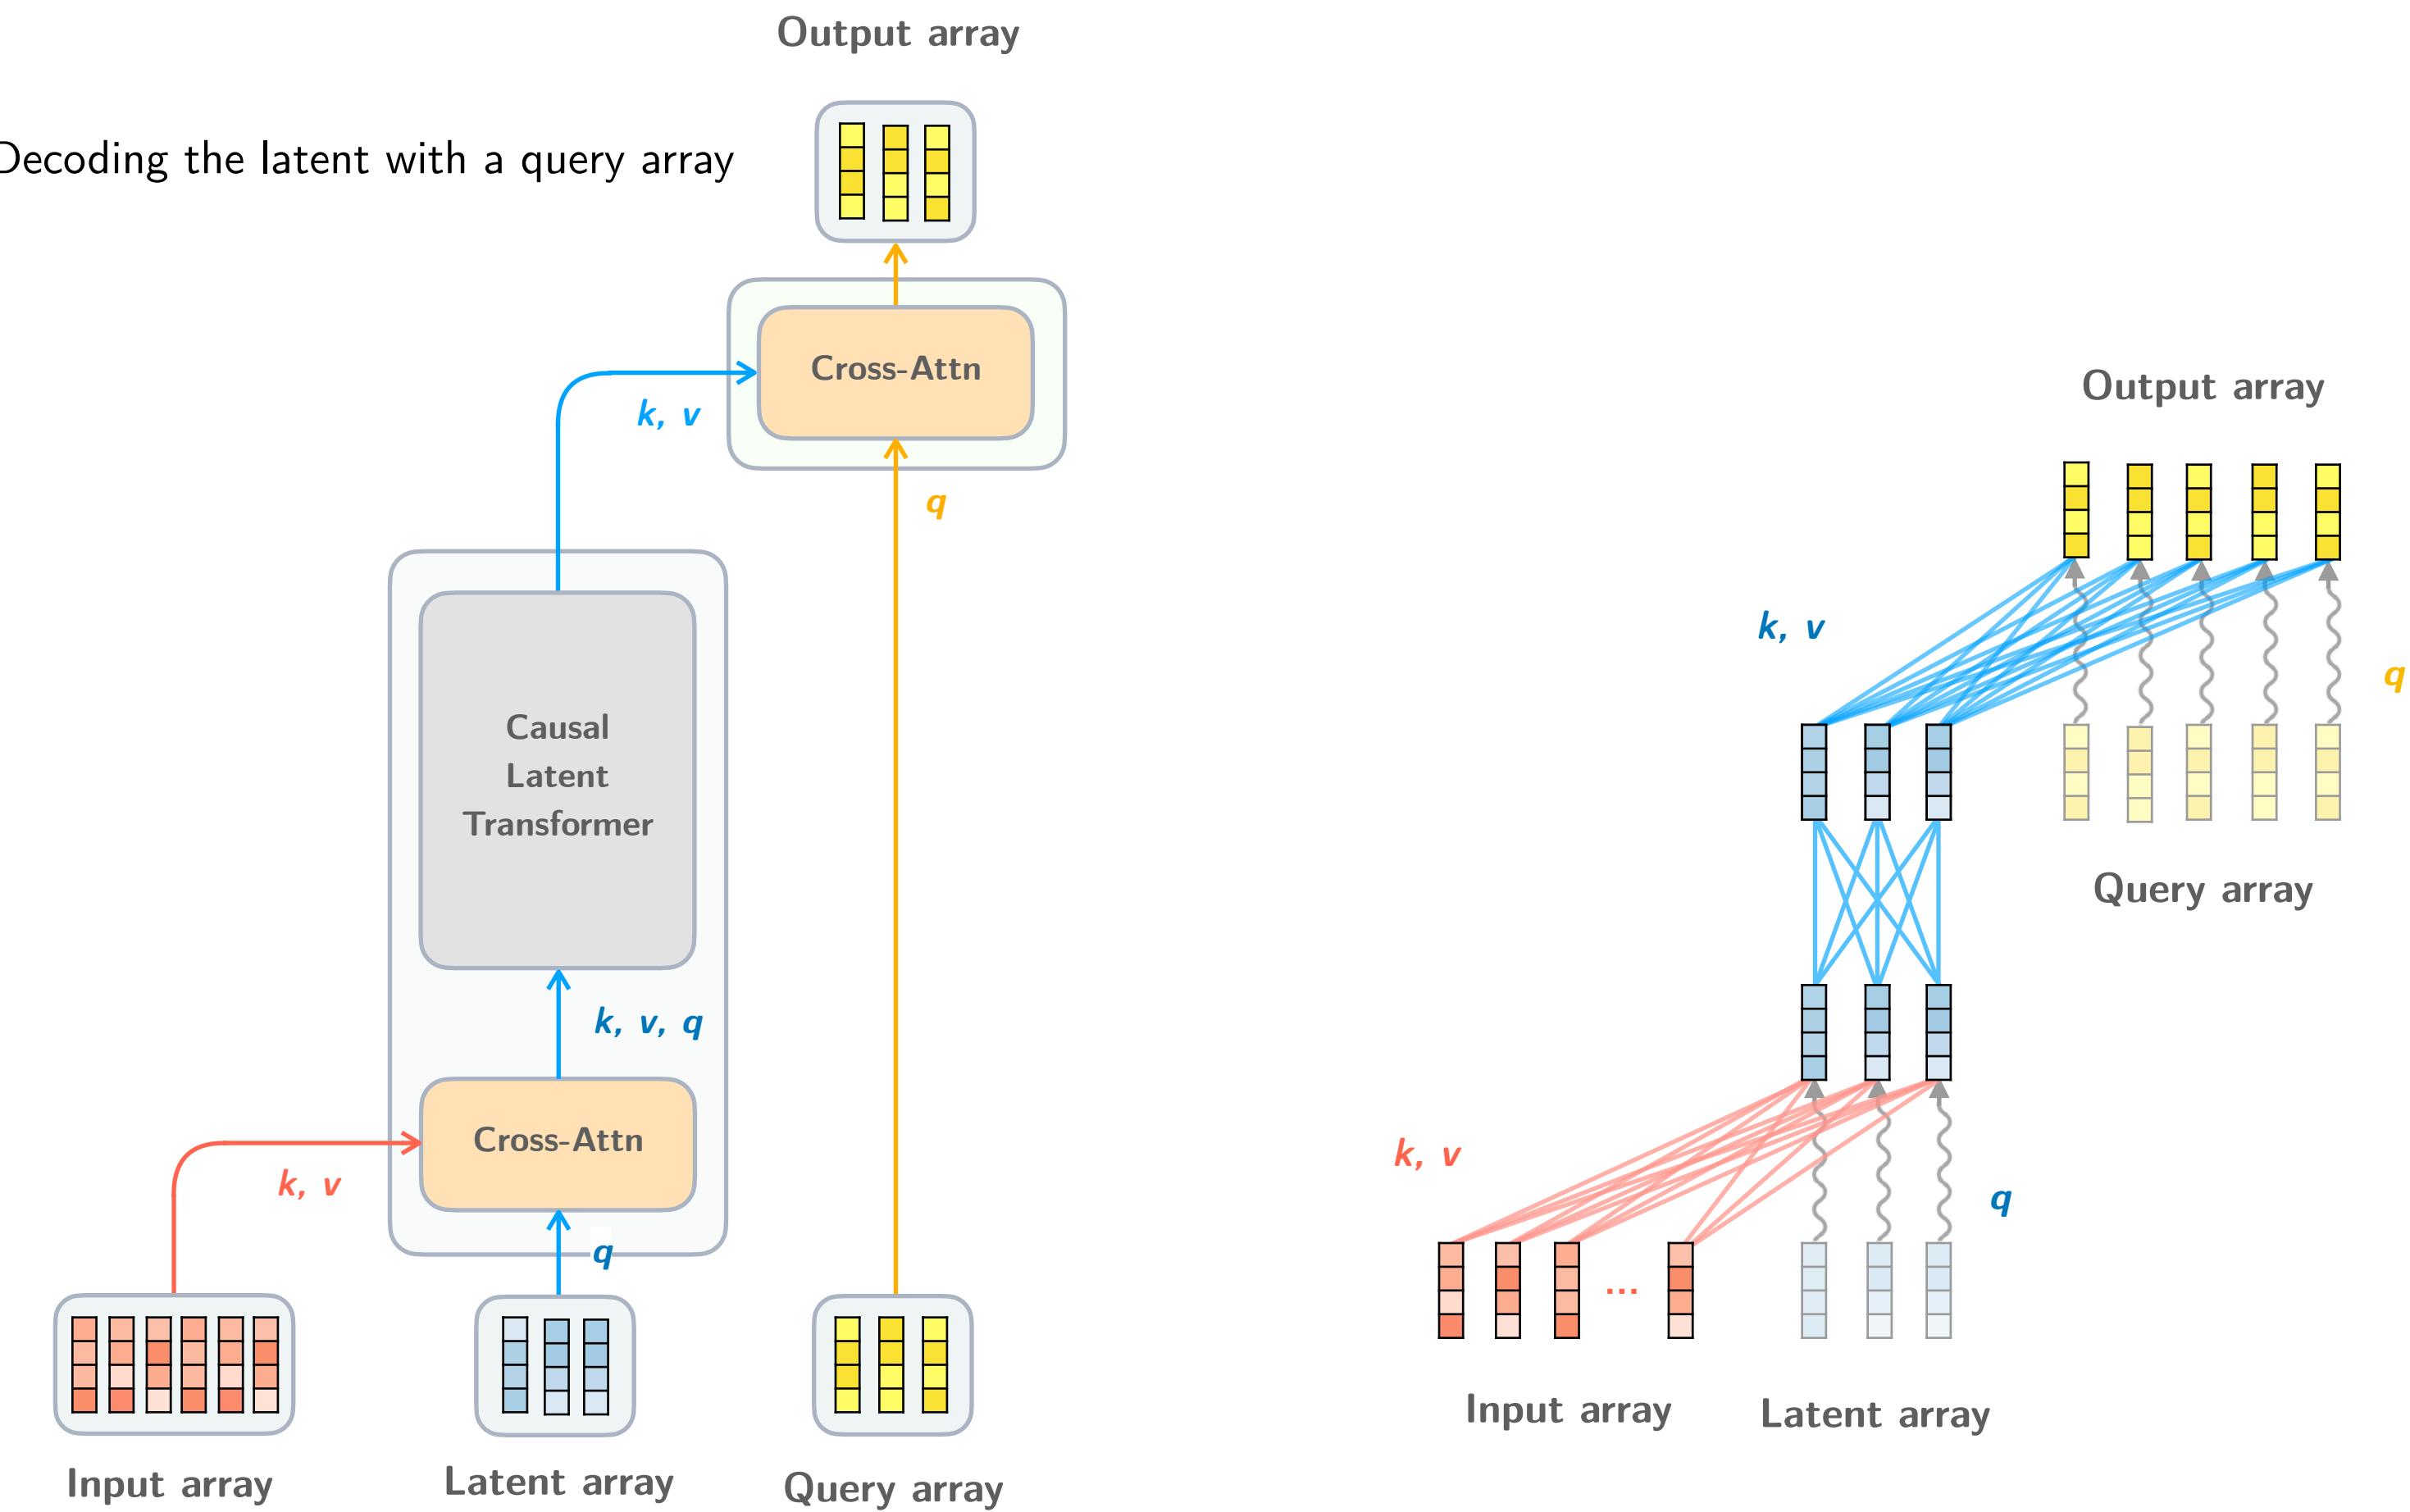
We only do deep self-attention over latent array.

Basically just a transformer decoder with initial cross-attention layer to a longer version of the autoregressive array?



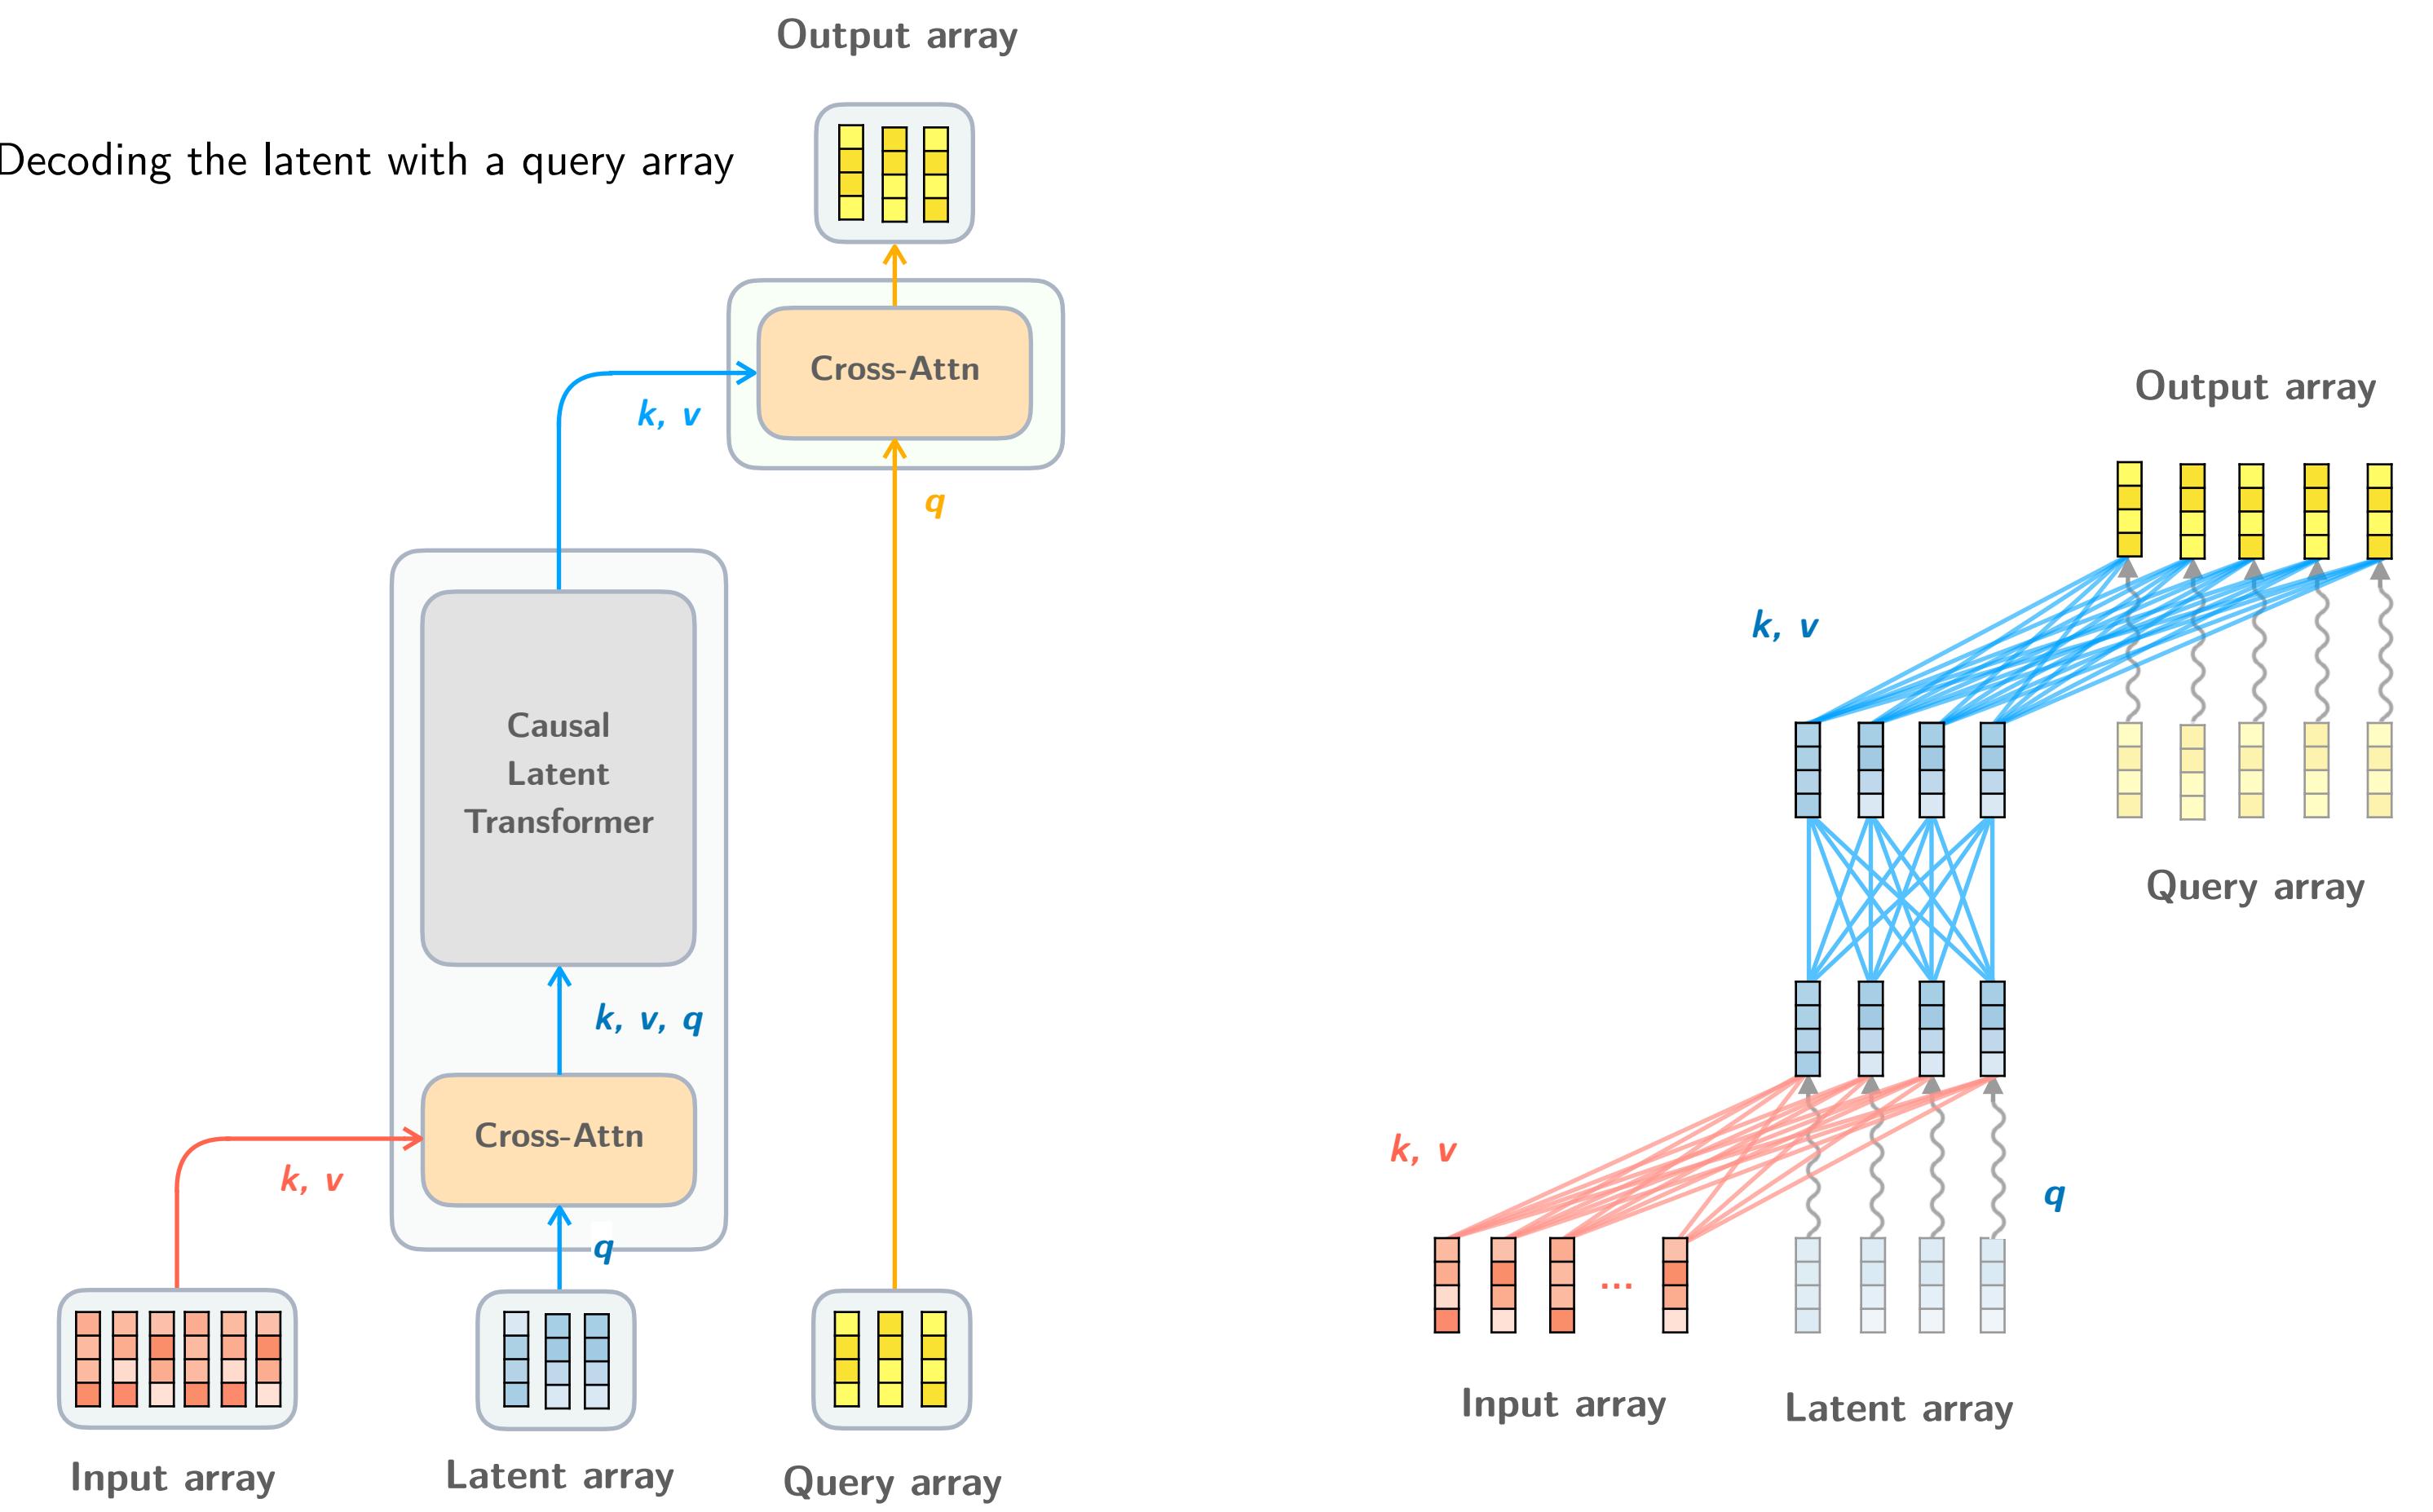
Even more general-purpose.

Better for outputs more complicated than, say,
simple classification or regression.



Even more general-purpose.

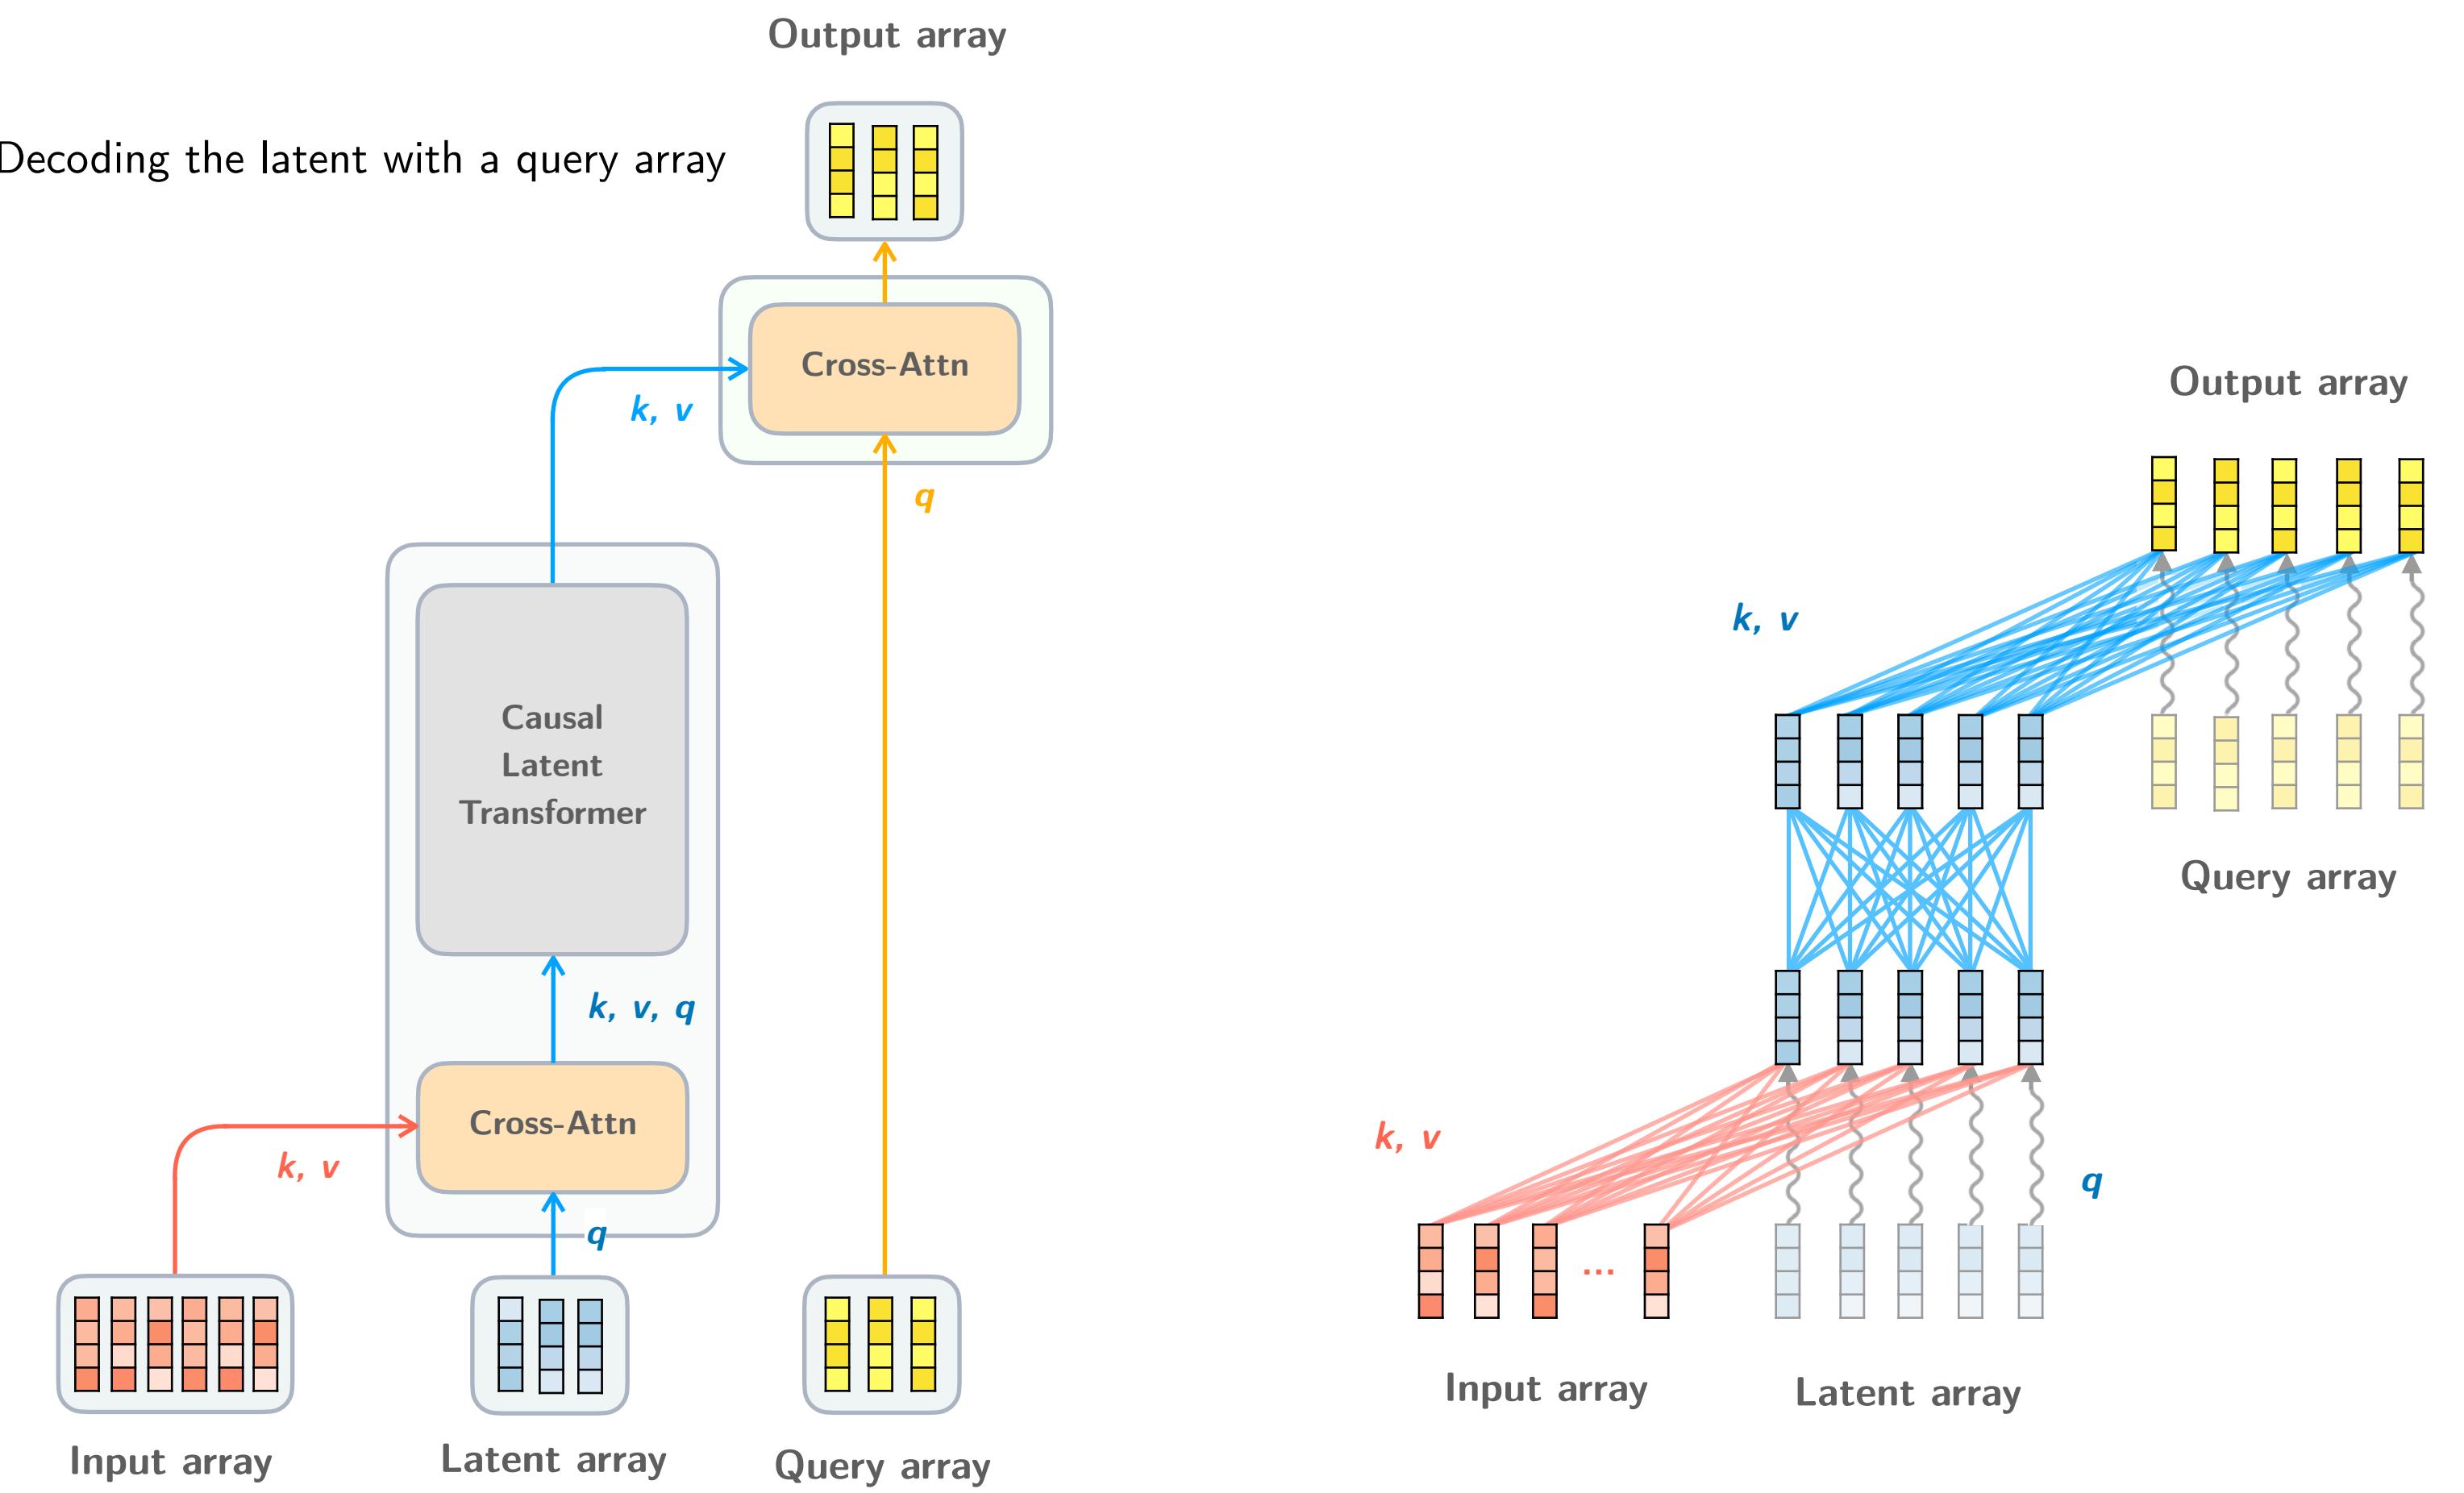
Better for outputs more complicated than, say,
simple classification or regression.



Increasing the “latent-array bandwidth”
(Not the same as increasing the bandwidth of each token,
which is the dimensionality of the latents)

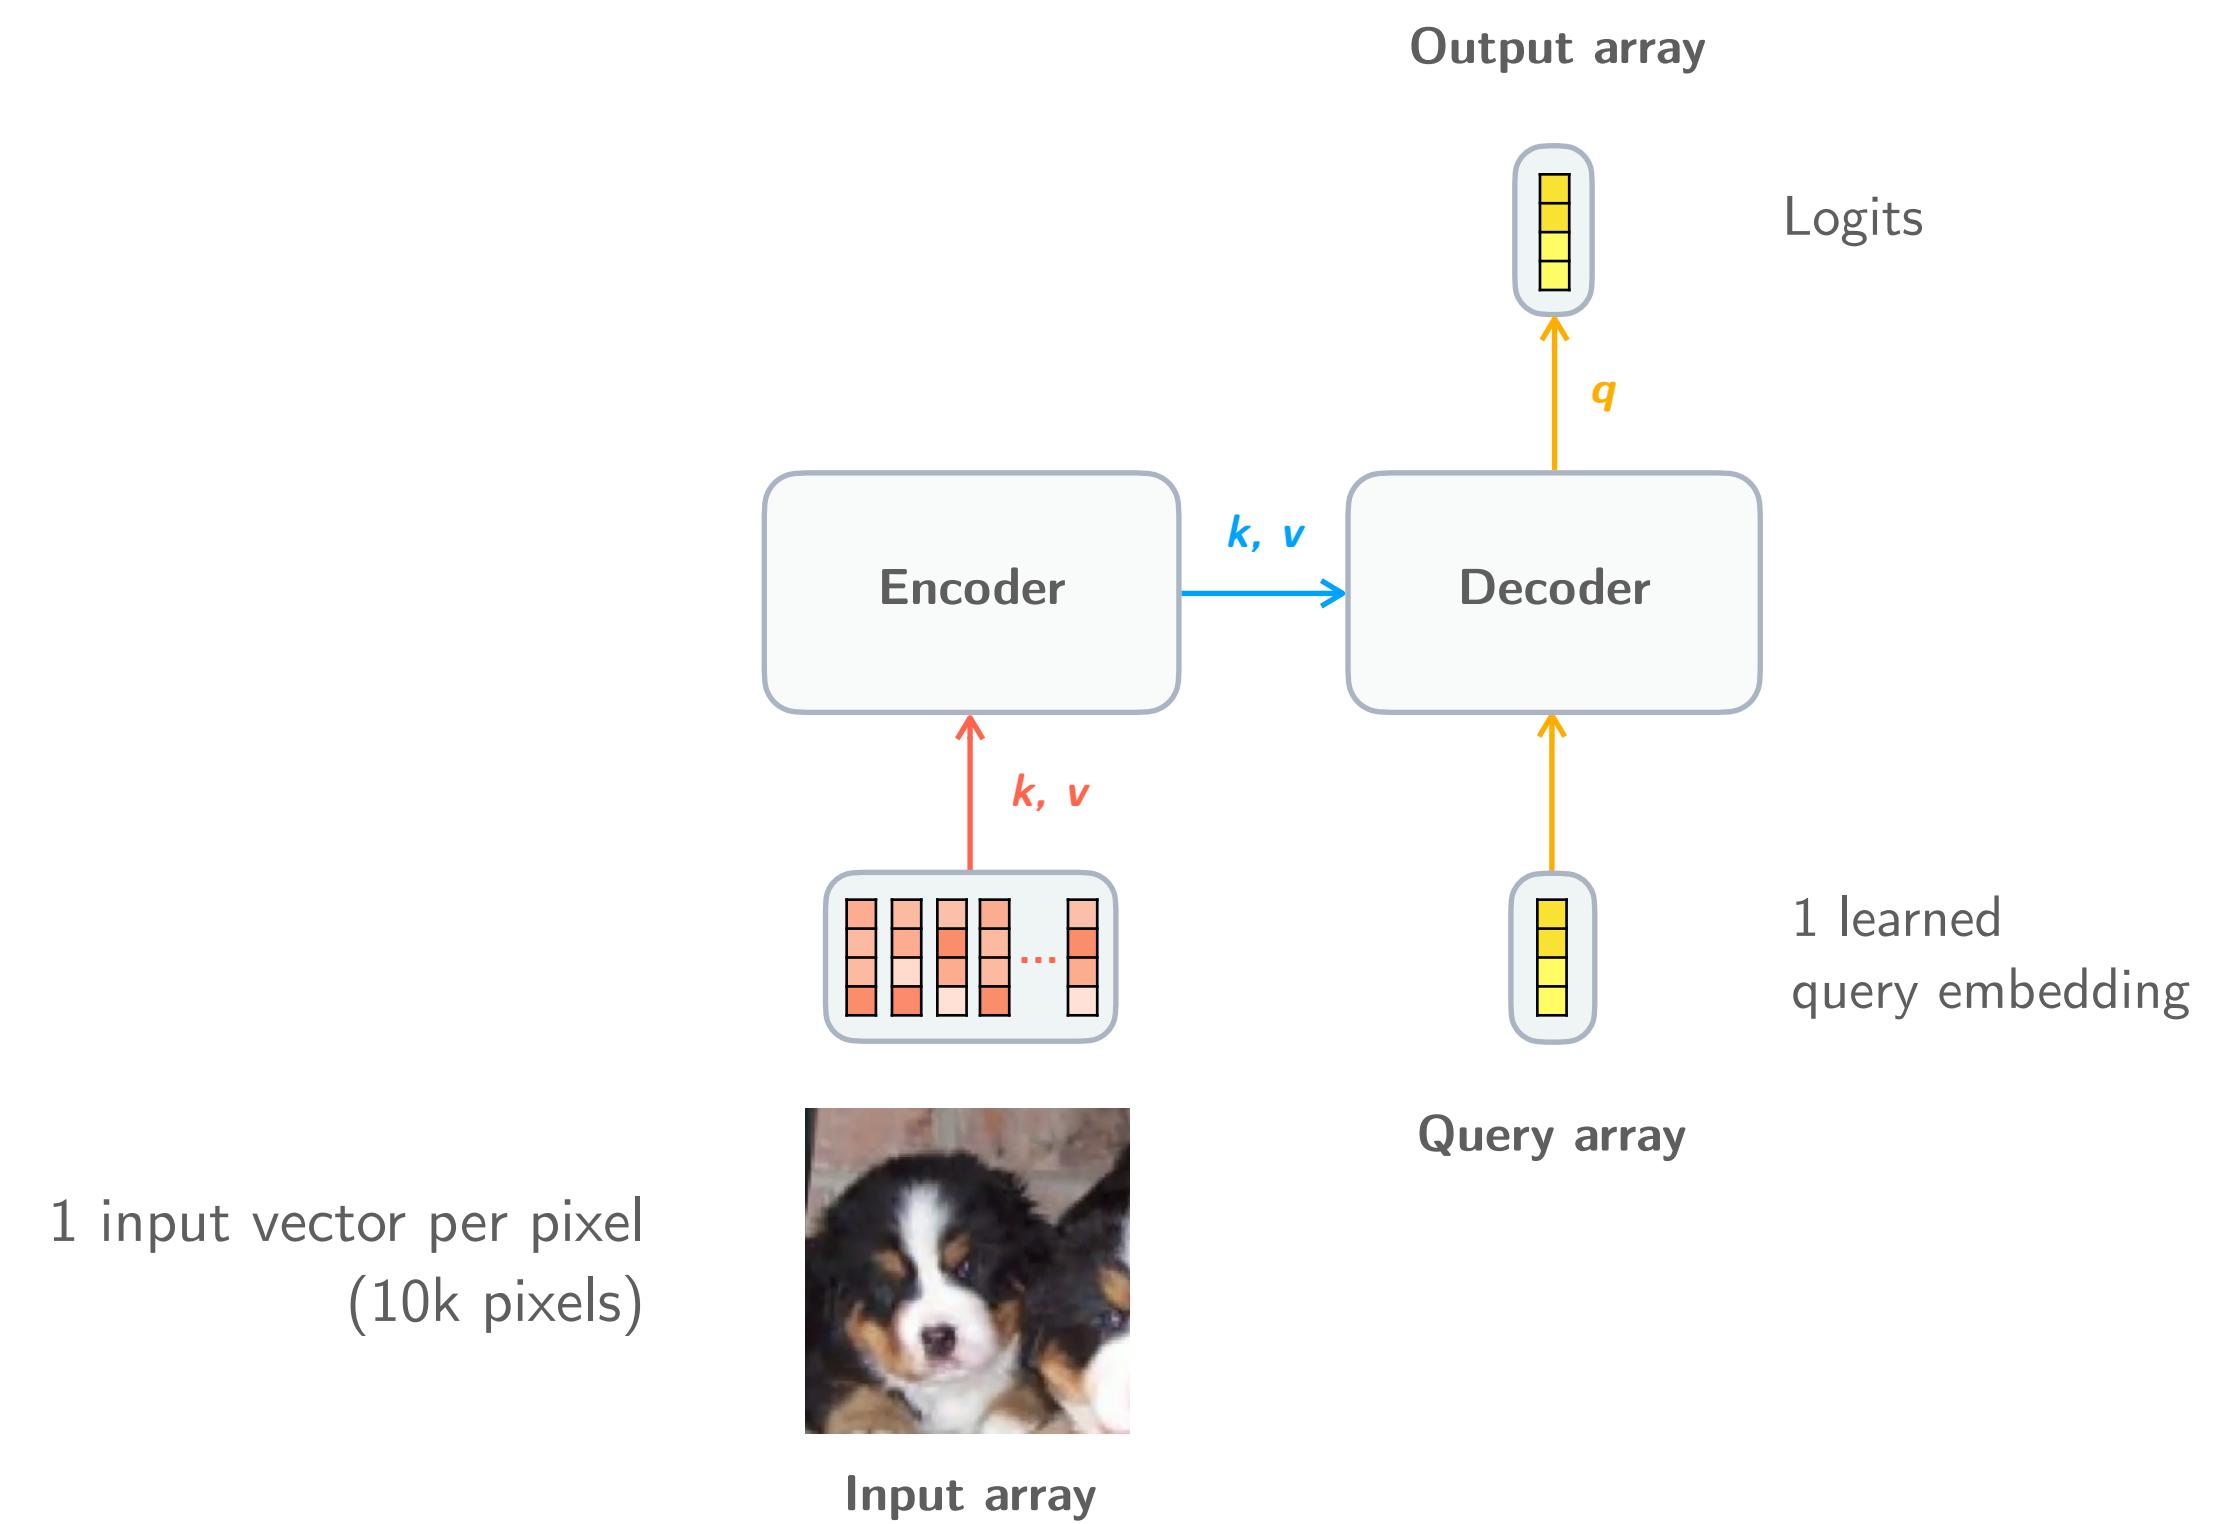
Even more general-purpose.

Better for outputs more complicated than, say,
simple classification or regression.



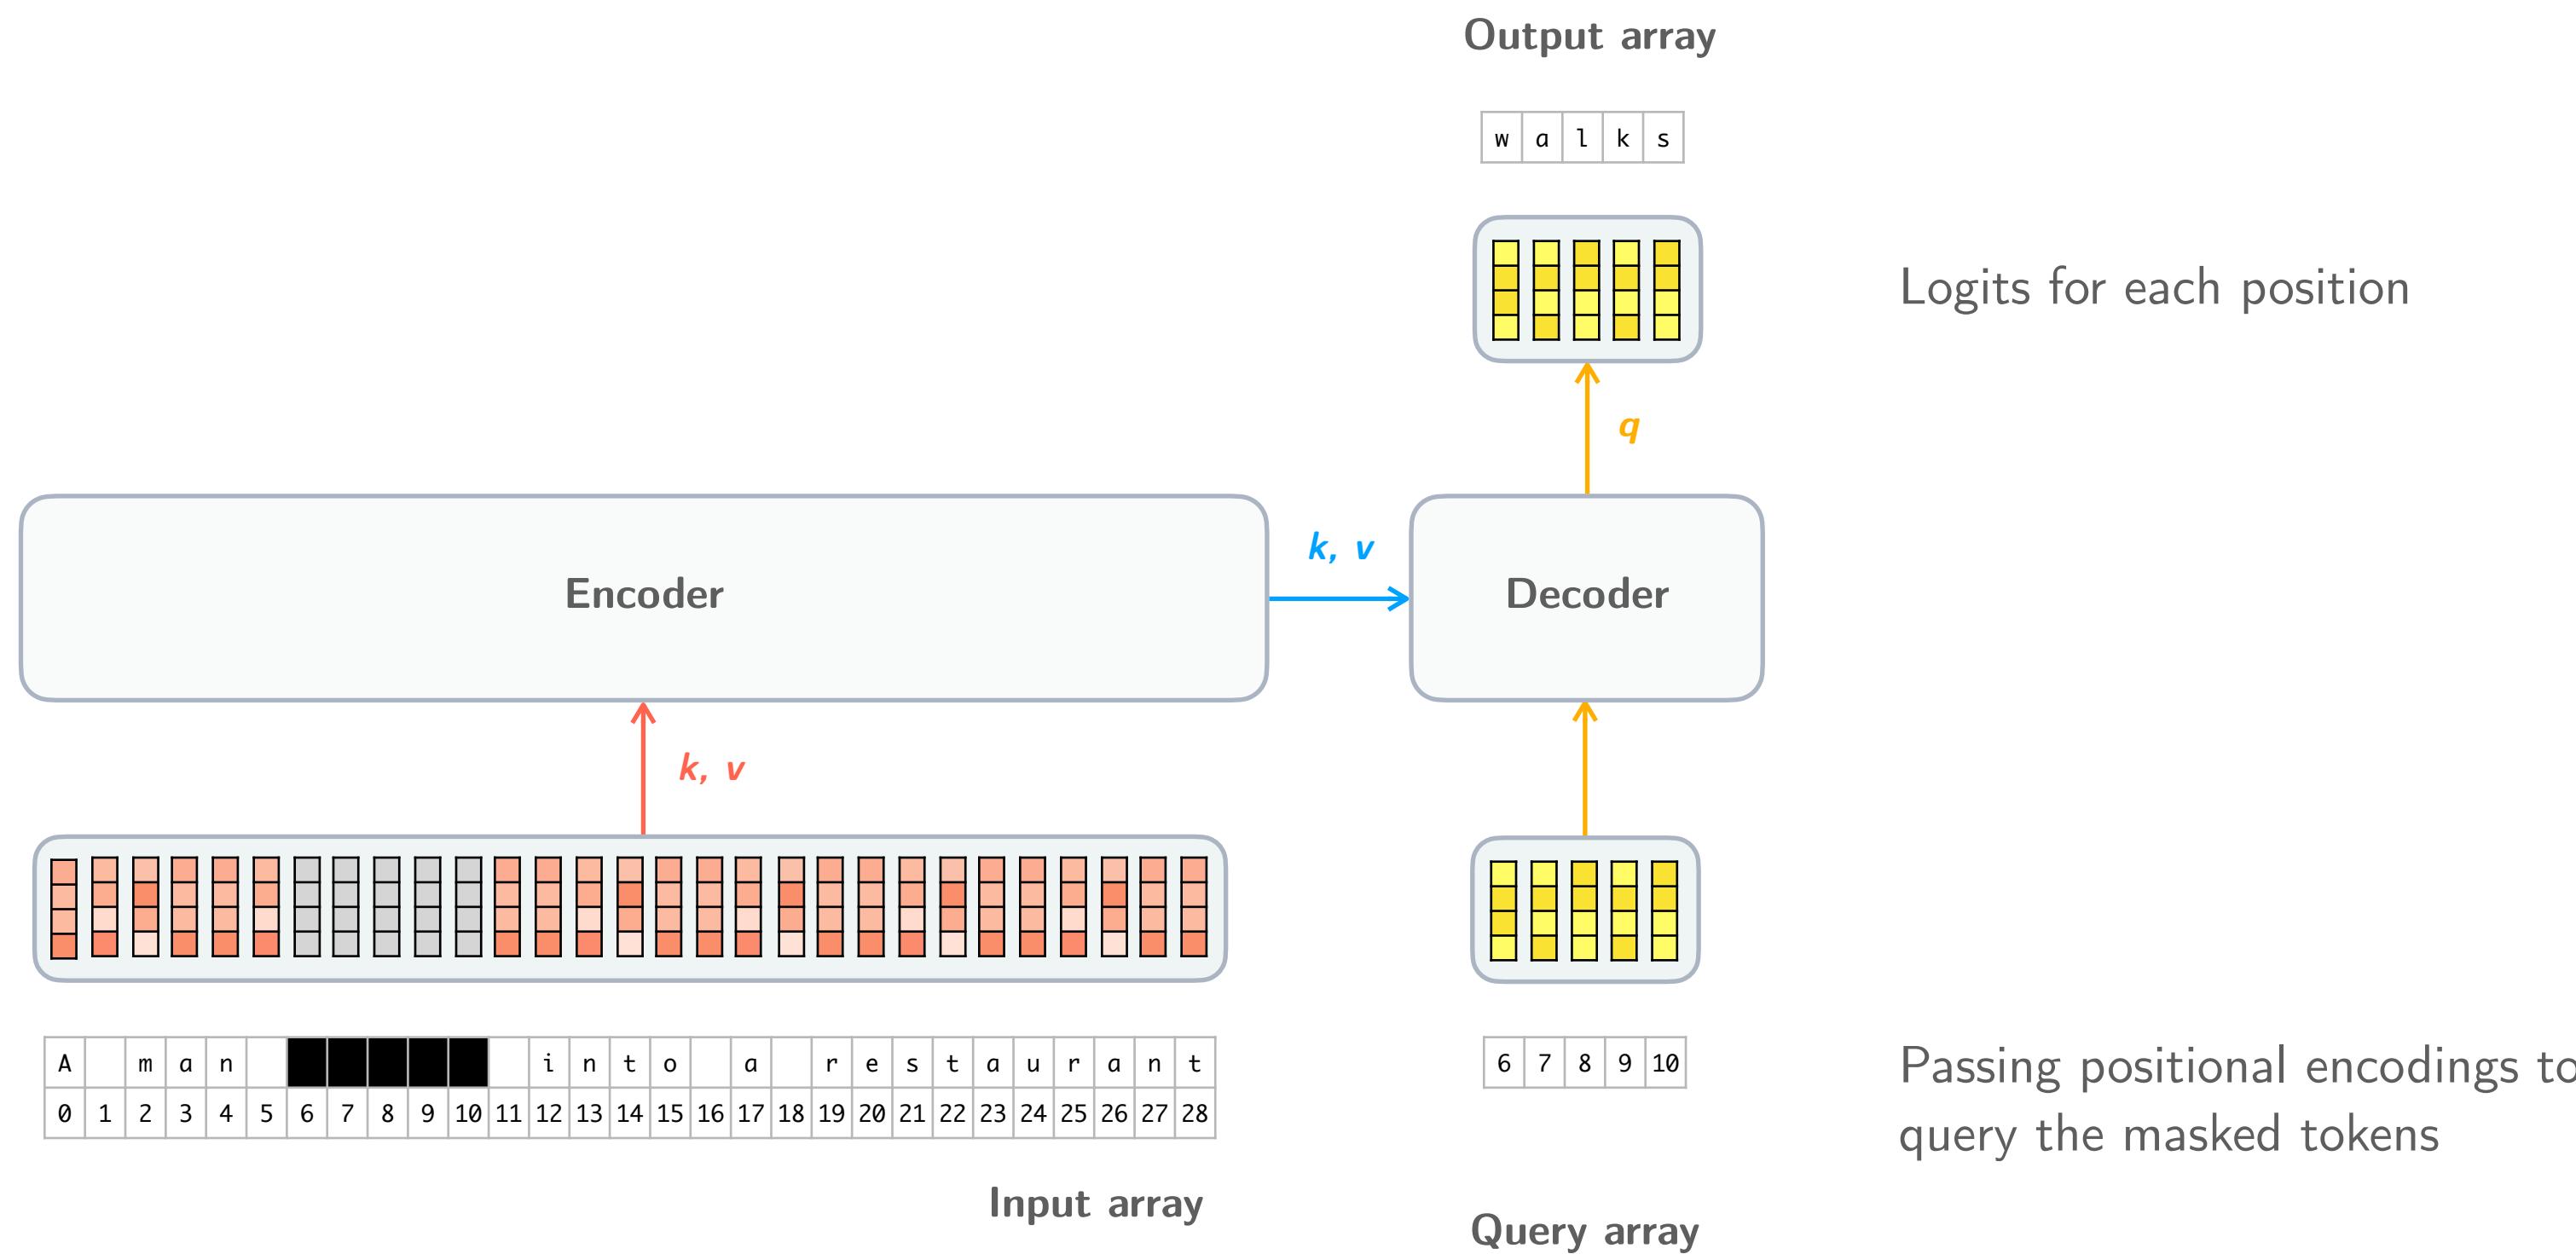
Increasing the “latent-array bandwidth”
(Not the same as increasing the bandwidth of each token,
which is the dimensionality of the latents)

Example: Classification



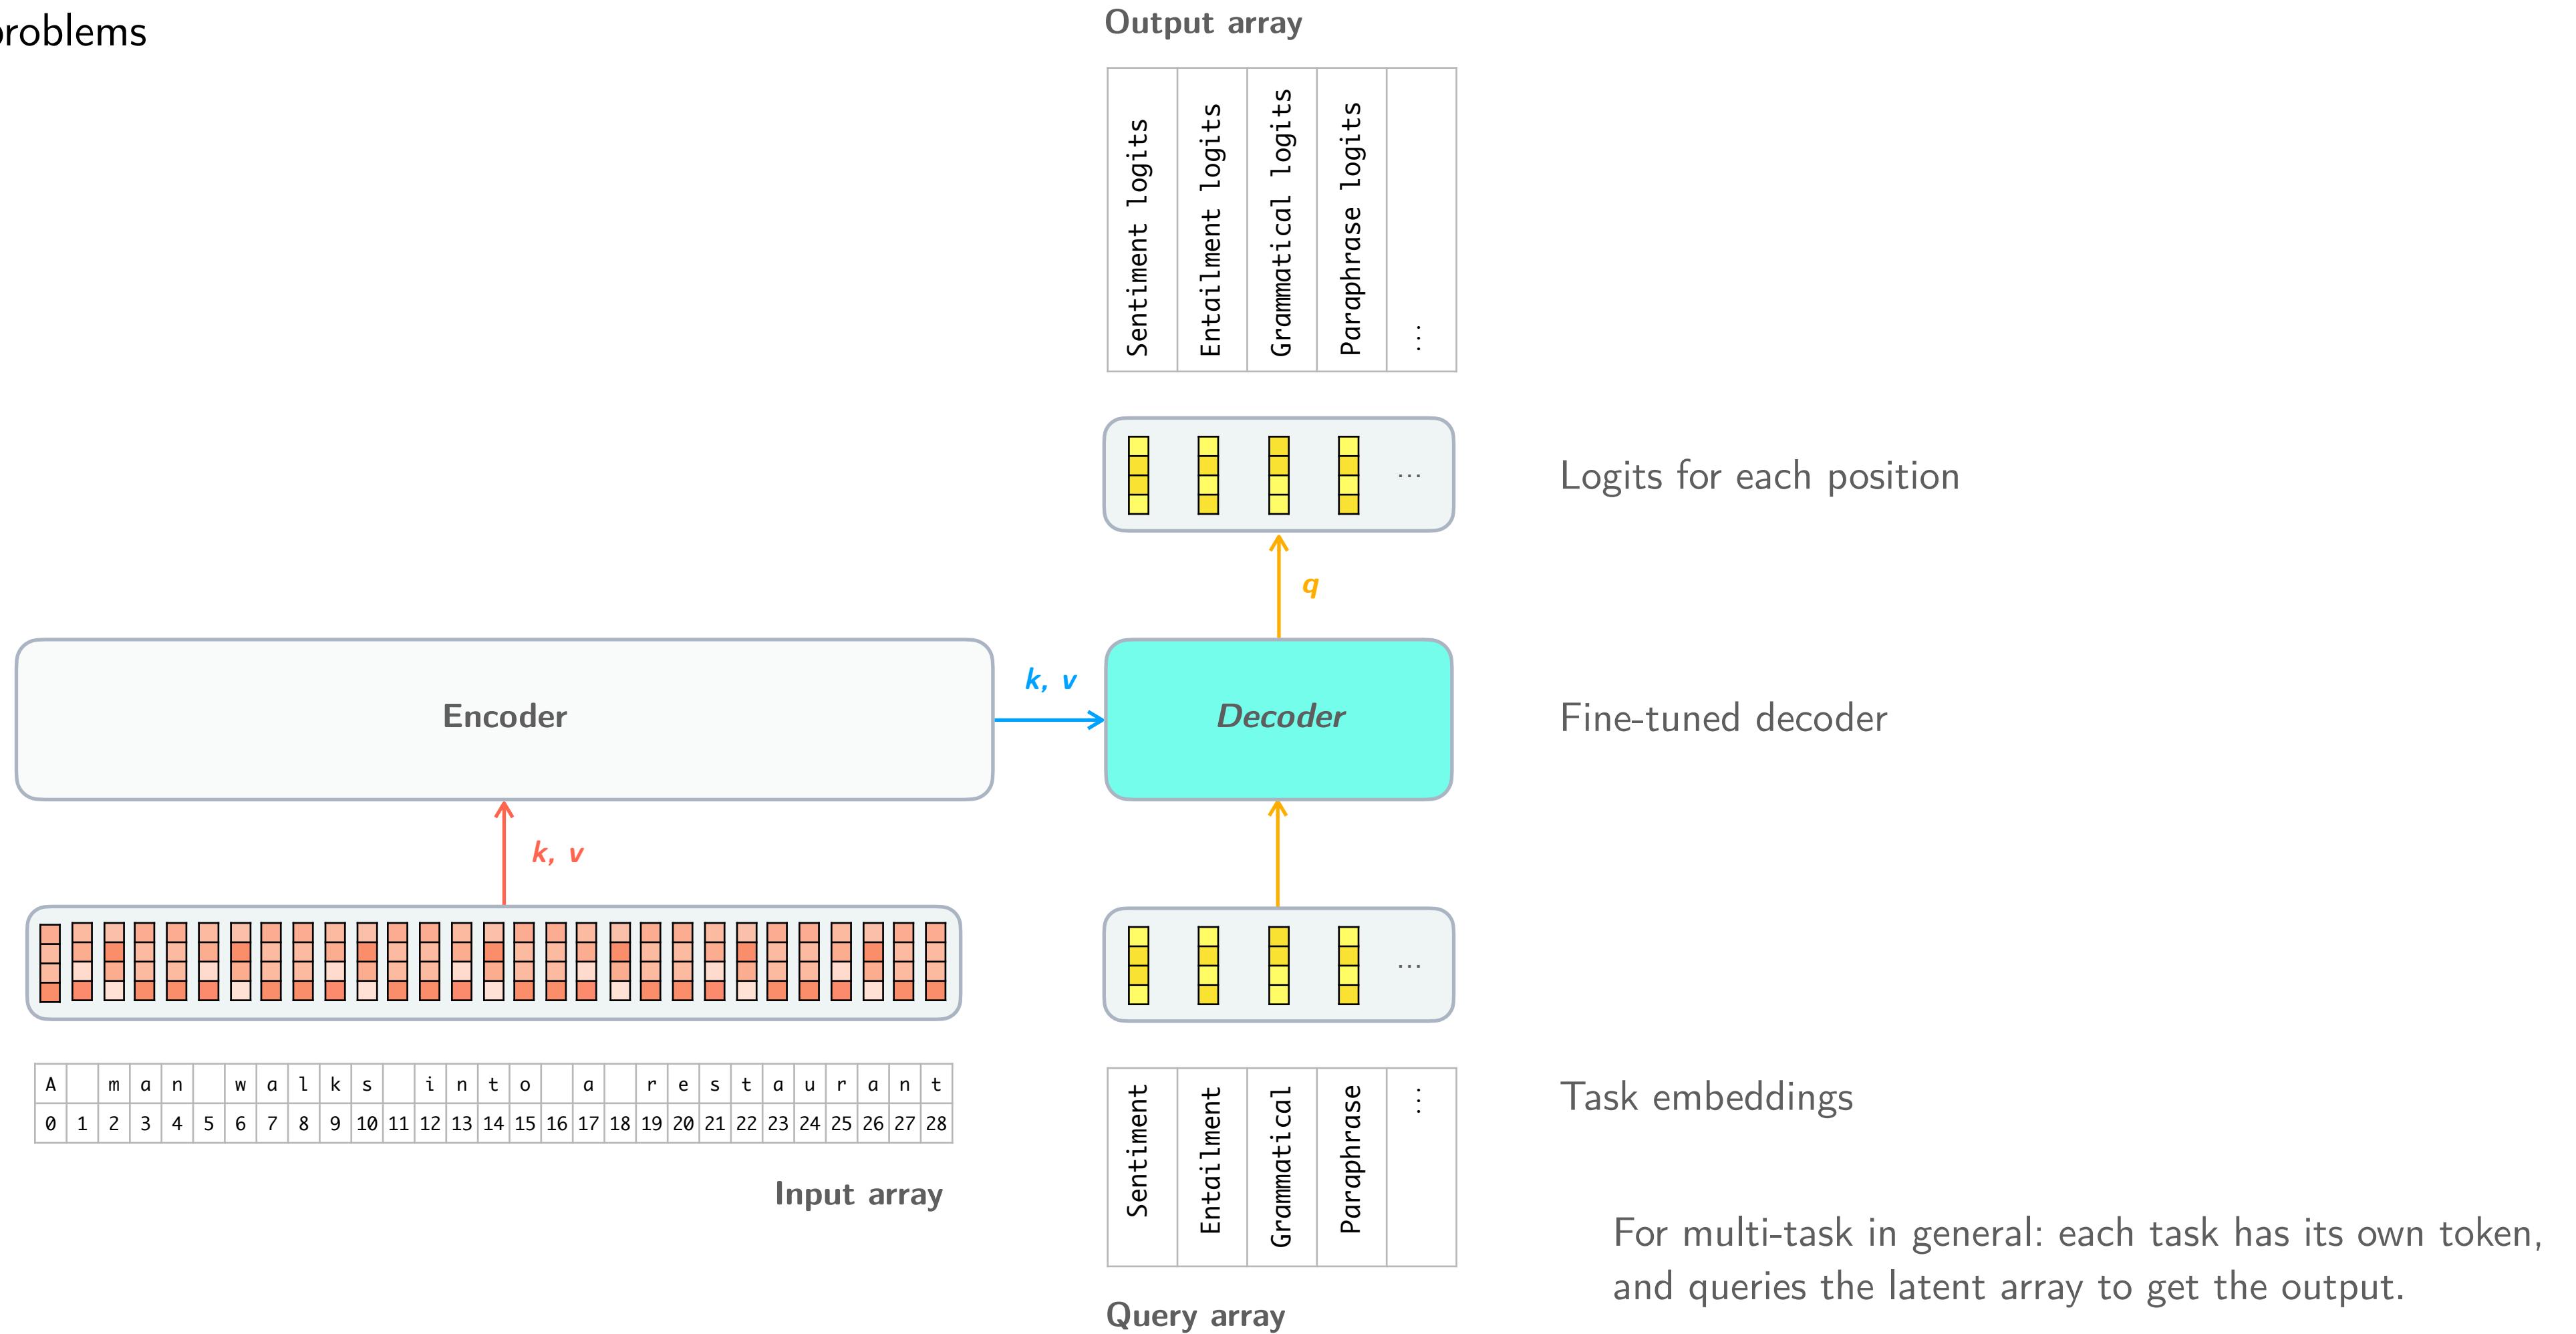
Example: Masked Language Modeling

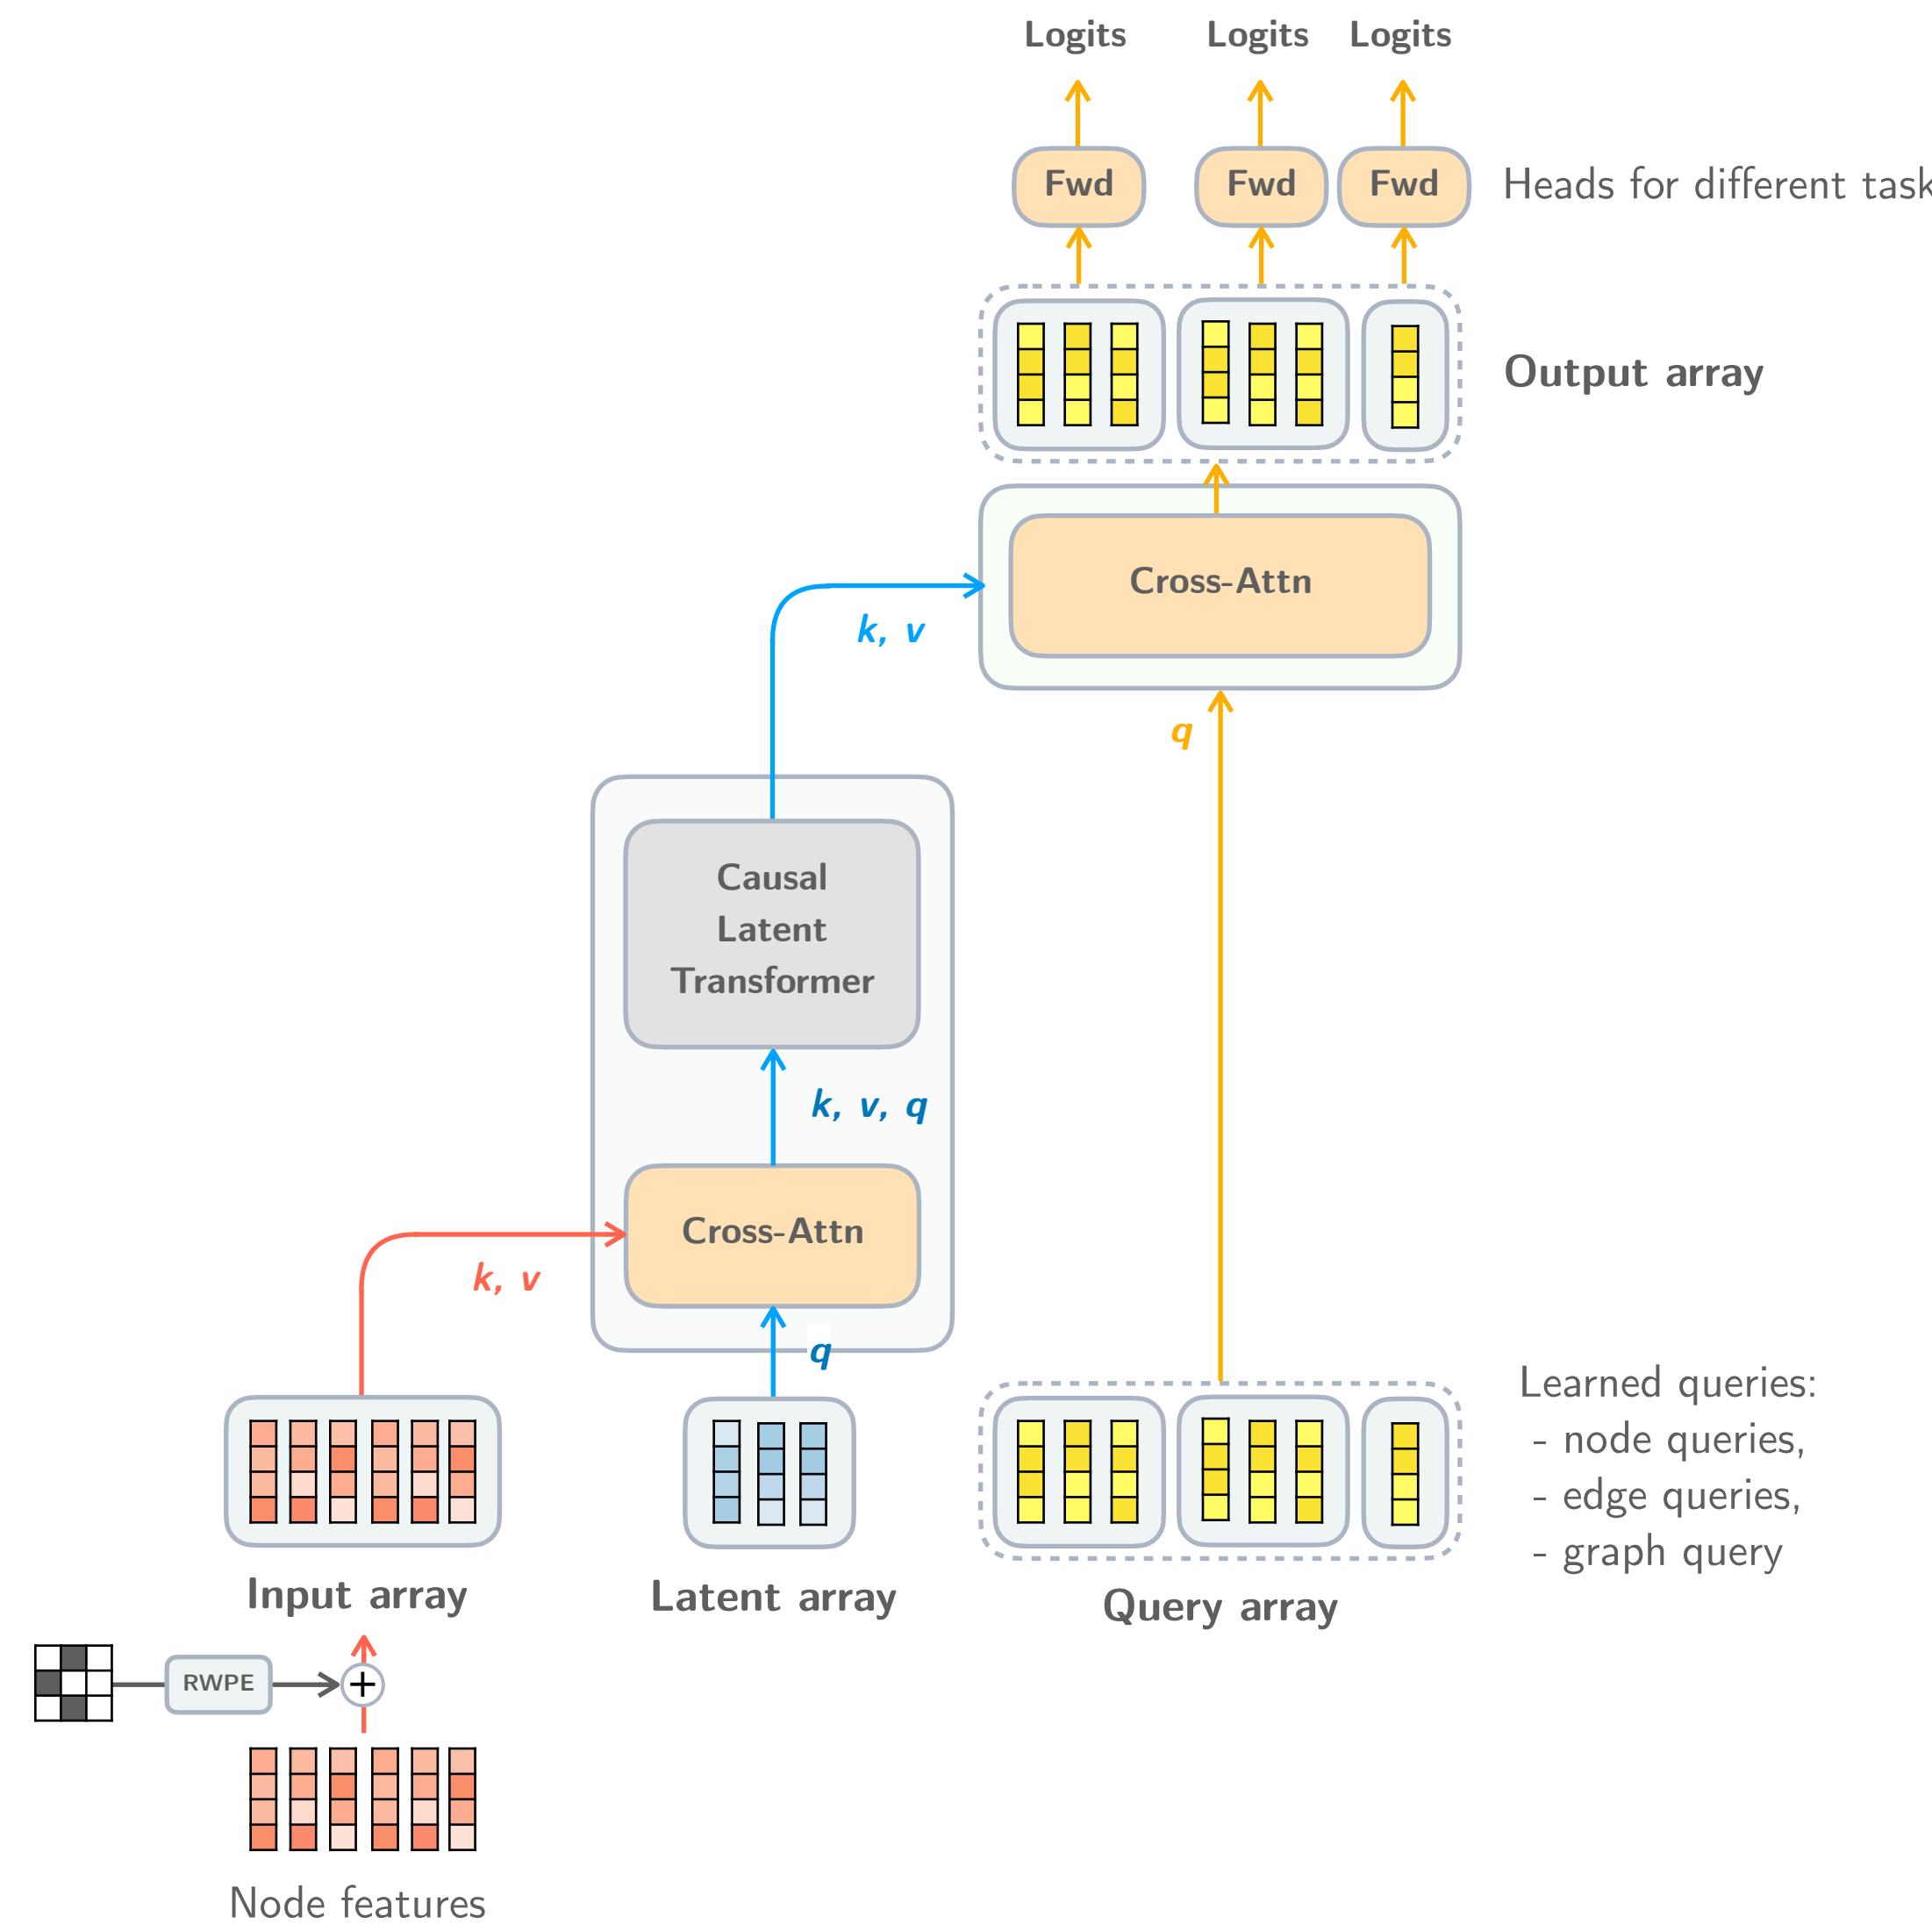
Similar for other spatial or sequence-structured arrays.



Example: Masked Language Modeling → GLUE fine-tuning

Similar for other multi-task problems





Thanks for your *attention*.